Empirical Study of Off-Policy Policy Evaluation for Reinforcement Learning

Cameron Voloshin 1 Hoang M. Le 1 Nan Jiang 2 Yisong Yue 1

Abstract

The disparate experimental conditions in recent off-policy policy evaluation (OPE) literature make it difficult both for practitioners to choose a reliable estimator for their application domain, as well as for researchers to identify fruitful research directions. In this work, we present the first detailed empirical study of a broad suite of OPE methods. Based on thousands of experiments and empirical analysis, we offer a summarized set of guidelines to advance the understanding of OPE performance in practice, and suggest directions for future research. Along the way, our empirical findings challenge several commonly held beliefs about which class of approaches tends to perform well. Our accompanying software implementation serves as a first comprehensive benchmark for OPE.

1. Introduction

We focus on understanding the relative performance of existing methods for off-policy policy evaluation (OPE), which is the problem of estimating the value of a target policy using only pre-collected historical data generated by another policy. The earliest OPE methods rely on classical importance sampling to handle the distribution mismatch between the target and behavior policies (Precup et al., 2000). Many advanced OPE methods have since been proposed for both contextual bandits (Dudík et al., 2011; Bottou et al., 2013; Swaminathan et al., 2017; Wang et al., 2017; Li et al., 2015; Ma et al.) and reinforcement learning settings (Jiang & Li, 2016; Dudík et al., 2011; Farajtabar et al., 2018; Liu et al., 2018; Xie et al., 2019). These new developments reflect practical interests in deploying reinforcement learning to safety-critical situations (Li et al., 2011; Wiering, 2000; Bottou et al., 2013; Bang & Robins, 2005), and the increasing importance of off-policy learning and counterfactual reasoning more broadly (Degris et al., 2012; Thomas et al., 2017; Munos et al., 2016; Le et al., 2019; Liu et al., 2019; Nie et al., 2019). OPE is also closely related to the problem of dynamic treatment regimes in the causal inference literature (Murphy et al., 2001).

Empirical validations have long contributed to the scientific understanding and advancement of machine learning techniques (Chapelle & Li, 2011; Caruana et al., 2008; Caruana & Niculescu-Mizil, 2006). Recently, many have called for careful examination of empirical findings of contemporary deep learning and deep reinforcement learning efforts (Henderson et al., 2018; Locatello et al., 2019). As OPE is central to real-world applications of reinforcement learning, an in-depth empirical understanding is critical to ensure usefulness and accelerate progress. While many recent methods are built on sound mathematical principles, a practitioner is often faced with a non-trivial task of selecting the most appropriate estimator for their application. A notable gap in current literature is a comprehensive empirical understanding of contemporary methods, due in part to the disparate testing environments and varying experimental conditions among prior work. Consequently, there is little holistic insight into where different methods particularly shine, nor a systematic summary of the challenges one may encounter when in different scenarios. Researchers and practitioners may reasonably deduce the following commonly held impressions from surveying the literature:

- 1. Doubly robust methods are often assumed to outperform direct and importance sampling methods.
- 2. Horizon length is the primary driver of poor performance for OPE estimators.
- 3. Model-based is the go-to direct method, either standalone or as part of a doubly-robust estimator.

The reality, as we will discuss, is much more nuanced. In this work, we take a closer look at recently proposed methods and offer a thorough empirical study of a wide range of estimators. We design various experimental conditions to explore the success and failure modes of different methods. We synthesize general insights to guide practitioners, and suggest directions for future research. Finally, we provide a highly extensive software package that can interface with new experimental environments and methods to run new OPE experiments at scale.¹

¹Caltech ²UIUC. Correspondence to: Cameron Voloshin <cvoloshi@caltech.edu>.

¹see https://github.com/clvoloshin/OPE-tools

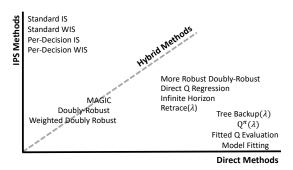


Figure 1. Categorization of OPE methods. Some methods are direct but have IPS influence and thus fit slightly away from the direct methods axis.

2. Preliminaries

As per RL standard, we represent the environment by $\langle X,A,P,R,\gamma\rangle$. X is the state space (or observation space in the non-Markov case), A is the (finite) action space, $P:X\times A\times X\to [0,1]$ is the transition function, $R:X\times A\to \mathbb{R}$ is the reward, and discount factor $\gamma\in(0,1]$. A policy π maps states to a distribution over actions, and $\pi(a|x)$ denotes the probability of choosing $a\in A$ in $x\in X$.

OPE is typically considered in the episodic RL setting. A behavior policy π_b generates a historical data set, $D = \{\tau^i\}_{i=1}^N$, of N trajectories (or episodes), where i indexes over trajectories, and $\tau^i = (x_0^i, a_0^i, r_0^i, \dots, x_{T-1}^i, a_{T-1}^i, r_{T-1}^i)$. The episode length T is frequently assumed to be fixed for notational convenience. In practice, one can pad additional absorbing states to handle variable lengths. Given a desired evaluation policy π_e , the OPE problem is to estimate the value $V(\pi_e)$, defined as:

$$V(\pi_e) = \mathbb{E}_{x \sim d_0} \left[\sum_{t=0}^{T-1} \gamma^t r_t | x_0 = x \right],$$

with $a_t \sim \pi_e(\cdot|x_t)$, $x_{t+1} \sim P(\cdot|x_t, a_t)$, $r_t \sim R(x_t, a_t)$, and d_0 is the initial state distribution.

3. Overview of OPE Methods

OPE methods were historically categorized into importance sampling, direct, and doubly robust methods. This demarcation was first introduced for contextual bandits (Dudík et al., 2011), and later to reinforcement learning (Jiang & Li, 2016). Some recent methods have blurred the boundary of these categories, such as Retrace(λ) (Munos et al., 2016) that uses a product of importance weights of multiple time steps for off-policy Q correction, and MAGIC (Thomas & Brunskill, 2016) that switches between importance weighting and direct methods.

In this paper, we propose to regroup OPE into three similar classes of methods, but with expanded definition for each

category. Figure 1 provides an overview of OPE methods that we consider. The relative positioning of different methods reflects how close they are to being a pure regression-based estimator versus a pure importance sampling-based estimator. Appendix D contains a full description of all methods under consideration.

3.1. Inverse Propensity Scoring (IPS)

Inverse Propensity Scoring (IPS), also called importance sampling, is widely used in statistics (Powell & Swann, 1966; Hammersley & Handscomb, 1964; Horvitz & Thompson, 1952) and RL (Precup et al., 2000). The key idea is to reweight the rewards in the historical data by the importance sampling ratio between π_e and π_b , i.e., how likely a reward is under π_e versus π_b . IPS methods yield consistent and (typically) unbiased estimates; however the product of importance weights can be unstable for long time horizons. The cumulative importance weight between π_e and π_b is written as $\rho^i_{j:j'} = \prod_{t=j}^{\min(j',T-1)} \frac{\pi_e(a^i_t|x^i_t)}{\pi_b(a^i_t|x^i_t)}$ (where $\rho^i_{t:t'} = 1$ for t' < t). Weighted IPS replaces a normalization factor N by $w_{j:j'} = \frac{1}{N} \sum_{i=1}^{N} \rho^i_{j:j'}$. The weighted versions are biased but strongly consistent.

Importance Sampling (IS) takes the form: $\sum_{i=1}^{N} \frac{\rho_{0:T-1}^{i}}{N} \sum_{t=0}^{T-1} \gamma^{t} r_{t}.$ There are three other main IPS variants that we consider: Per-Decision Importance Sampling (PDIS), Weighted Importance Sampling (WIS) and Per-Decision WIS (PDWIS) (see Appendix Table 16 for full definitions). Other variants of IPS exist but are neither consistent nor unbiased (Thomas, 2015). IPS often assumes known π_{b} , which may not be possible – one approach is to estimate π_{b} from data (Hanna et al., 2019), resulting in a potentially biased estimator that can sometimes outperform traditional IPS methods.

3.2. Direct Methods (DM)

The main distinction of direct methods from IPS is the focus on regression-based techniques to (more) directly estimate the value functions of the evaluation policy (Q^{π_e} or V^{π_e}). We consider eight different direct approaches, described completely in appendix D. Similar to policy learning literature, we can view OPE through the lens of model-based vs. model-free approaches².

Model-based. Perhaps the most commonly used DM is model-based (also called approximate model, denoted AM), where the transition dynamics, reward function and termination condition are directly estimated from historical data (Jiang & Li, 2016; Paduraru, 2013). The resulting learned MDP is then used to compute the value of π_e , e.g., by Monte-Carlo policy evaluation.

²the distinction is arguably blurry. We stick with this convention simply for linguistic convenience

			Table 1. Env	rironment pa	arameters			
Environment	Graph	Graph-MC	MC	Pix-MC	Enduro	Graph-POMDP	GW	Pix-GW
Markov?	yes	yes	yes	yes	yes	no	yes	yes
State/Obs	position	position	[pos, vel]	pixels	pixels	position	position	pixels
T	4 or 16	250	250	250	1000	2 or 8	25	25
Stoch Env?	variable	no	no	no	no	no	no	variable
Stoch Rew?	variable	no	no	no	no	no	no	no
Sparse Rew?	variable	terminal	terminal	terminal	dense	terminal	dense	dense
\hat{Q} Func. Class	tabular	tabular	linear/NN	NN	NN	tabular	tabular	NN

Model-free. Estimating the action-value function $\widehat{Q}(\cdot;\theta)$, parameterized by θ , is the focus of several model-free The value estimate is then: $\hat{V}(\pi_e) =$ $\frac{1}{N}\sum_{i=1}^{N}\sum_{a\in A}\pi_e(a|s_0^i)\widehat{Q}(x_0^i,a;\theta)$. A simple example is Fitted Q Evaluation (FQE) (Le et al., 2019), which is a model-free counterpart to AM, and is functionally a policy evaluation counterpart to batch Q learning. FQE learns a sequence of estimators $Q(\cdot, \theta) = \lim_{k \to \infty} Q_k$, where:

$$\widehat{Q}_{k} = \min_{\theta} \frac{1}{N} \sum_{i=1}^{N} \sum_{t=0}^{T-1} (\widehat{Q}_{k-1}(x_{t}^{i}, a_{t}^{i}; \theta) - y_{t}^{i})^{2},$$

$$y_{t}^{i} \equiv r_{t}^{i} + \gamma \mathbb{E}_{\pi_{e}} \widehat{Q}_{k-1}(x_{t+1}^{i}, \cdot; \theta), \quad \widehat{Q}_{0} \equiv 0.$$

Indeed, several model-free methods originated from offpolicy learning settings, but are also natural for OPE. $Q^{\pi}(\lambda)$ (Harutyunyan et al., 2016) can be viewed as a generalization of FQE that looks to the horizon limit to incorporate the long-term value into the backup step. Retrace(λ) (Munos et al., 2016) and Tree-Backup(λ) (Precup et al., 2000) also use full trajectories, but additionally incorporate varying levels of clipped importance weights adjustment. The λ -dependent term mitigates instability in the backup step, and is chosen based on experimental findings of Munos et al. (2016).

Q Regression (Q-Reg) and More Robust Doubly-Robust (MRDR) (Farajtabar et al., 2018) are two recently proposed direct methods that make use of cumulative importance weights in deriving the regression estimate for Q^{π_e} , solved through a quadratic program. MRDR changes the objective of the regression to that of directly minimizing the variance of the Doubly-Robust estimator (see Section 3.3).

Liu et al. (2018) recently proposed a method for the infinite horizon setting (IH). While IH can be viewed as a Rao-Blackwellization of the IS estimator, we include it in the DM category because it essentially solves the Bellman equation for state distributions and requires function approximation, which are more characteristic of DM. IH shifts the focus from importance sampling over action sequences to estimating the importance ratio ω between *state* density distributions induced by π_b and π_e . This ratio replaces all but the final importance weights ρ_{T-1} in the IH estimate, which resembles IS. More recently, several estimators inspired by density ratio estimation idea have been proposed (Nachum et al., 2019; Uehara & Jiang, 2019; Xie et al., 2019) - we will leave evaluation of these new extensions for future work.

3.3. Hybrid Methods (HM)

Hybrid methods subsume doubly robust-like approaches, which combine aspects of both IPS and DM. Standard doubly robust OPE (denoted DR) (Jiang & Li, 2016) is an unbiased estimator that leverages a DM to decrease the variance of the unbiased estimates produced by importance sampling techniques:

$$\sum_{i=1}^{N} \frac{\widehat{V}(x_0^i)}{N} + \frac{1}{N} \sum_{i=1}^{N} \sum_{t=0}^{T-1} \gamma^t \rho_{0:t}^i [r_t^i - \widehat{Q}(x_t^i, a_t^i) + \gamma \widehat{V}(x_{t+1}^i)].$$

Other HMs include Weighted Doubly-Robust (WDR) and MAGIC (see Appendix D). WDR self-normalizes the importance weights (similar to WIS). MAGIC introduces adaptive switching between DR and DM; in particular, one can imagine using DR to estimate the value for part of a trajectory and then using DM for the remainder. Using this idea, MAGIC (Thomas & Brunskill, 2016) finds an optimal linear combination among a set that varies the switch point between WDR and DM. Note that any DM that returns $Q^{\pi_e}(x, a; \theta)$ yields a set of corresponding DR, WDR, and MAGIC estimators. As a result, we consider twentyone hybrid approaches in our experiments.

4. Experiments

Experiment Design Principles. We consider several domain characteristics (simple-complex, deterministicstochastic, sparse-dense rewards, short-long horizon), π_b, π_e pairs (close-far), and data sizes N (small-large), to study OPE performance under varying conditions.

We use two standard RL benchmarks from OpenAI (Brockman et al., 2016): Mountain Car (MC) and Enduro Atari game. As many RL benchmarks are fixed and deterministic, we design 6 additional environments that allow control over various conditions: (i) Graph domain (tabular, varying stochasticity and horizon), (ii) Graph-POMDP (tabular, control for representation), (iii) Graph-MC (simplifying MC to tabular case), (iv) Pixel-MC (study MC in highdimensional setting), (v) Gridworld (tabular, long horizon version) and (vi) Pixel-Gridworld (controlled Gridworld experiments with function approximation).

All together, our benchmark consists of eight environments with characteristics summarized in Table 1. Complete descriptions can be found in Appendix E.

Protocol & Metrics. Each experiment depends on specifying environment and its properties, behavior policy π_b , evaluation policy π_e , and number of trajectories N from rolling-out π_b for historical data. The true on-policy value $V(\pi_e)$ is the Monte-Carlo estimate via 10,000 rollouts of π_e . We repeat each experiment m=10 times with different random seeds. We judge the quality of a method via two metrics:

- Relative mean squared error (*Relative MSE*): $\frac{1}{m} \sum_{i=1}^{m} \frac{(\widehat{V}(\pi_e)_i \frac{1}{m} \sum_{j=1}^{m} V(\pi_e)_j)^2}{(\frac{1}{m} \sum_{j=1}^{m} V(\pi_e)_j)^2}, \text{ which allows a fair comparison across different conditions.}^3$
- Near-top Frequency: For each experimental condition, we include the number of times each OPE estimator is within 10% of the best performing estimator to facilitate aggregate comparison across domains.

Implementation & Hyperparameters. With thirty-three different OPE methods considered, we run thousands of experiments across the above eight domains. Hyperparameters are selected based on publication, code release or author consultation. We maintain a consistent set of hyperparameters for each estimator and each environment across experimental conditions (see hyperparameter choice in appendix Table 26). We create a software package that allows running experiments at scale and easy integration with new domains and techniques for future research. Due to limited space, we will show the results from selected experiment conditions. The complete results, with highlighted best method in each class, are available in the appendix.

5. Results

5.1. What is the best method?

The first important takeaway is that *there is no clear-cut winner*: no single method or method class is consistently the best performer, as multiple environmental factors can influence the accuracy of each estimator. With that caveat in mind, based on the aggregate top performance metrics, we can recommend the following estimators for each method class (See Table 2 and appendix Table 4).

Inverse propensity scoring (IPS). In practice, weighted importance sampling, which is biased, tends to be more accurate and data-efficient than unbiased basic importance

sampling methods. Among the four IPS-based estimators, *PDWIS tends to perform best* (Figure 4 left).

Direct methods (DM). Generally, FQE, $Q^{\pi}(\lambda)$, and IH tend to perform the best among DM (appendix Table 4). FQE tends to be more data efficient and is the best method when data is limited (Figure 5). $Q^{\pi}(\lambda)$ generalizes FQE to multi-step backup, and works particularly well with more data, but is computationally expensive in complex domains. IH is highly competitive in long horizons and with high policy mismatch in a tabular setting (appendix Tables 8, 9). In pixel-based domains, however, choosing a good kernel function for IH is not straightforward, and IH can underperform other DM (appendix Table 12). We provide a numerical comparison among direct methods for tabular (appendix Figure 16) and complex settings (Figure 4 center).

Hybrid methods (**HM**). With the exception of IH, each DM corresponds to three HM: standard doubly robust (DR), weighted doubly robust (WDR), and MAGIC. For each DM, its WDR version often outperforms its DR version. MAGIC can often outperform WDR and DR. However, MAGIC comes with additional hyperparameters, as one needs to specify the set of partial trajectory length to be considered. Unsurprisingly, their performance highly depends on the underlying DM. In our experiments, FQE and $Q^{\pi}(\lambda)$ are typically the most reliable: MAGIC with FQE or MAGIC with $Q^{\pi}(\lambda)$ tend to be among the best hybrid methods (see appendix Figures 22 - 26).

5.2. Key drivers of method accuracy

The main reason for the inconsistent performance of estimators is various environmental factors that are inadequately studied from prior work. These coupled factors often impact accuracy interdependently:

- Representation mismatch: Function approximators with insufficient representation power weaken DM, and so do overly rich ones as they cause overfitting (e.g., tabular classes). These issues do not impact IPS. Severe misspecification favors HM and weakens DM.
- Horizon length: Long horizons hurt all methods, but especially those dependent on importance weights (including IPS, HM and some DM).
- *Policy mismatch:* Large divergence between π_b and π_e hurts all methods, but tends to favor DM in the small data regime relative to HM and IPS. HM will catch up with DM as data size increases.
- Bad estimation of unknown behavior policy: 4 π_b estimation quality depends on the state and action dimensionality, and historical data size. Poor π_b estimates cause HM and IPS to underperform simple DM.

 $^{^3}$ The performance metric in prior OPE work is typically mean squared error MSE= $\frac{1}{m}\sum_{i=1}^m(\widehat{V}(\pi_e)_i-V(\pi_e)_i)^2$

⁴Poor estimation of π_b can also be seen as model misspecification. We distinguish the representation issue of π_b from other representation issues related to DM

	There 2. Froder selection current of Froduction, see definition in section 1 and support in facts 1,					
Class	Recommendation	When to use	Prototypical env.	Near-top Freq.		
Direct	FQE	Stochastic env, severe policy mismatch	Graph, MC, Pix-MC	23.7%		
	$Q(\lambda)$	Compute non-issue, moderate policy mismatch	GW/Pix-GW	15.0%		
	IH	Long horizon, mild policy mismatch, good kernel	Graph-MC	19.0%		
IPS	PDWIS	Short horizon, mild policy mismatch	Graph	4.7%		
Hybrid	MAGIC FQE	Severe model misspecification	Graph-POMDP, Enduro	30.0%		
	MAGIC $Q(\lambda)$	Compute non-issue, severe model misspecification	Graph-POMDP	17.3%		

Table 2. Model Selection Guidelines. (For Near-top Frequency, see definition in Section 4 and support in Table 4)

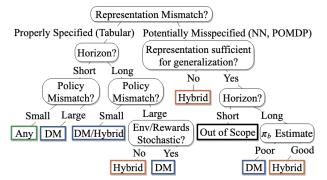


Figure 2. Method Class Selection Decision Tree. Numerical support can be found in Appendix B.1.

Environment / Reward stochasticity: Stochastic environments hurt the data efficiency of all methods, but favor DM over HM and IPS.

We perform a series of controlled experiments to isolate the impact of these factors. Figure 3 shows a typical comparison of the best performing method in each class, under a tabular setting with both short and long horizons, and a large mismatch between π_b and π_e . The particular best method in each class may change depending on the specific conditions. Within each class, a general guideline for method selection is summarized in Table 2. The appendix contains the full empirical results of all experiments.

5.3. A recipe for method selection

Figure 2 summarizes our general guideline for navigating key factors that affect the accuracy of different estimators. To guide the readers through the process, we now dive further into our experimental design to test various factors, and discuss the resulting insights.

Do we potentially have representation mismatch? Representation mismatch comes from two sources: model misspecification and poor generalization. Model misspecification refers to the insufficient representation power of the function class used to approximate either the transition dynamics (AM), value function (other DM), or state distribution density ratio (in IH).

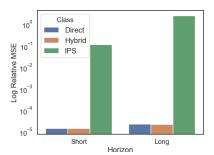
Tabular representation for MDP controls for representation mismatch by ensuring adequate function class capacity, as well as zero inherent Bellman error (left branch, Fig 2). In such case, we may still suffer from poor generalization without sufficient data coverage, which depends on other factors in the domain settings.

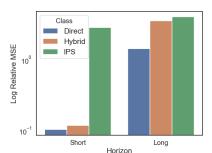
The effect of representation mismatch (right branch, Fig 2) can be understood via two controlled scenarios:

- Misspecified and poor generalization: We expose the impact of this severe mismatch scenario via the Graph POMDP construction, where selected information are omitted from an otherwise equivalent Graph MDP. HM substantially outperform DM in this setting (Figure 3 right versus left).
- Misspecified but good generalization: Function class such as neural networks has powerful generalization ability, but may introduce bias and inherent Bellman error⁵ (Munos & Szepesvári, 2008; Chen & Jiang, 2019) (see linear vs. neural networks comparison for Mountain Car in appendix Fig 13). Still, powerful function approximation makes (biased) DM very competitive with HM, especially under limited data and in complex domains (see pixel-Gridworld in appendix Fig 27-29). However, function approximation bias may cause serious problem for high dimensional and long horizon settings. In the extreme case of Enduro (very long horizon and sparse rewards), all DM fail to convincingly outperform a naïve average of behavior data (appendix Fig 12).

Short horizon vs. Long horizon? It is well-known that IPS methods are sensitive to trajectory length (Li et al., 2015). Long horizon leads to an exponential blow-up of the importance sampling term, and is exacerbated by significant mismatch between π_b and π_e . This issue is inevitable for any unbiased estimator (Jiang & Li, 2016) (a.k.a., the curse of horizon (Liu et al., 2018)). Similar to IPS, DM relying on importance weights also suffer from long horizon (appendix Fig 16), though to a lesser degree. IH aims to bypass the effect of cumulative weighting in long horizons, and indeed performs substantially better than IPS methods in very long horizon domains (Fig 4 left).

⁵defined as $\sup_{g \in \mathcal{F}} \inf_{f \in \mathcal{F}} ||f - \mathbb{T}^{\pi}g||_{d_{\pi}}$, where \mathcal{F} is function class chosen for approximation, and d_{π} is state distribution induced by evaluation policy π





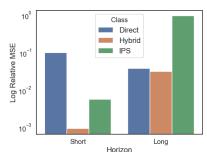


Figure 3. Comparing IPS versus DM versus HM under short and long horizon, large policy mismatch and large data. Left: (Graph domain) Deterministic environment. Center: (Graph domain) Stochastic environment and rewards. Right: (Graph-POMDP) Model misspecification (POMDP). Minimum error per class is shown.

A frequently ignored aspect in previous OPE work is a proper distinction between fixed, finite horizon tasks (IPS focus), infinite horizon tasks (IH focus), and indefinite horizon tasks, where the trajectory length is finite but varies depending on the policy. Many applications should properly belong to the indefinite horizon category.⁶ Applying HM in this setting requires proper padding of the rewards (without altering the value function in the infinite horizon limit) as DR correction typically assumes fixed length trajectories.

How different are behavior and target policies? Similar to IPS, the performance of DM is negatively correlated with the degree of policy mismatch. Figure 5 shows the interplay of increasing policy mismatch and historical data size, on the top DM in the deterministic gridworld. We use $(\sup_{a \in A, x \in X} \frac{\pi_e(a|x)}{\pi_b(a|x)})^T$ as an environment-independent metric of mismatch between the two policies. The performance of the top DM (FQE, $Q^{\pi}(\lambda)$, IH) tend to hold up better than IPS methods when the policy gap increases (appendix Figure 18). FQE and IH are best in the small data regime, and $Q^{\pi}(\lambda)$ performs better as data size increases (Figure 5). Increased policy mismatch weakens the DM that use importance weights (Q-Reg, MRDR, Retrace(λ) and Tree-Backup(λ)).

Do we have a good estimate of the behavior policy? Often the behavior policy may not be known exactly and requires estimation, which can introduce bias and cause HM to underperform DM, especially in low data regime (e.g., pixel gridworld appendix Figure 27-29). Similar phenomenon was observed in the statistics literature (Kang et al., 2007). As the data size increases, HMs regain the advantage as the quality of the π_b estimate improves.

Is the environment stochastic or deterministic? While stochasticity affects all methods by straining the data requirement, HM are more negatively impacted than DM (Figure 3 center, Figure 17). This can be justified by e.g., the variance analysis of DR, which shows that the vari-

ance of the value function with respect to stochastic transitions will be amplified by cumulative importance weights and then contribute to the overall variance of the estimator; see Jiang & Li (2016, Theorem 1) for further details. We empirically observe that DM frequently outperform their DR versions in the small data case (Figure 17). In a stochastic environment and tabular setting, HM do not provide significant edge over DM, even in short horizon case. The gap closes as the data size increases (Figure 17).

5.4. Challenging common wisdom

We close this section by briefly revisiting commonly held beliefs about high-level performance of OPE methods.

Are HM always better than DM? No. Overall, DM are surprisingly competitive with HM. Under high-dimensionality, long horizons, estimated behavior policies, or reward/environment stochasticity, HM can underperform simple DM, sometimes significantly (e.g., see appendix Figure 17).

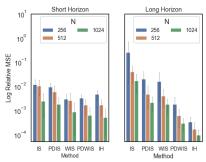
Concretely, HM can perform worse than DM in the following scenarios that we tested:

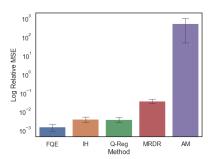
- Tabular with large policy mismatch, or stochastic environments (appendix Figure 17, Table 6, 9).
- Complex domains with long horizon and unknown behavior policy (appendix Figure 27-29, Table 11).

When data is sufficient, or model misspecification is severe, HM do provide consistent improvement over DM.

Is horizon length the most important factor? No. Despite conventional wisdom suggesting IPS methods are most sensitive to horizon length, we find that this is not always the case. Policy divergence $\sup_{a\in A,x\in X}\frac{\pi_e(a|x)}{\pi_b(a|x)}$ can be just as, if not more, meaningful. For comparison, we designed two scenarios with identical mismatch $(\sup_{a\in A,x\in X}\frac{\pi_e(a|x)}{\pi_b(a|x)})^T$ as defined in Section 5.3 (see appendix Tables 14, 15). Starting from a baseline scenario of short horizon and small policy divergence (appendix Table

⁶Applying IH in the indefinite horizon case requires setting up an absorbing state that loops over itself with zero terminal reward.





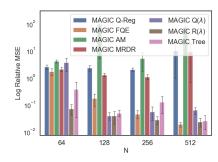


Figure 4. Left: (Graph domain) Comparing IPS (and IH) under short and long horizon. Mild policy mismatch setting. PDWIS is often best among IPS. But IH outperforms in long horizon. Center: (Pixel-MC) Comparing direct methods in high-dimensional, long horizon setting. Relatively large policy mismatch. FQE and IH tend to outperform. AM is significantly worse in complex domains. Retrace(λ), Q(λ) and Tree-Backup(λ) are very computationally expensive and thus excluded. Right: (Pixel Gridworld) Comparing MAGIC with different base DM and different data size. Large policy mismatch, deterministic environment, known π_b .

13), extending horizon length leads to $10\times$ degradation in accuracy, while a comparable increase in policy divergence causes a $100\times$ degradation.

How good is model-based direct method (AM)? AM can be among the worst performing direct methods (appendix Table 4). While AM performs well in tabular setting in the large data case (appendix Figure 16), it tends to perform poorly in high dimensional settings with function approximation (e.g., Figure 4 center). Fitting the transition model P(x'|x,a) is often more prone to small errors than directly approximating Q(x,a). Model fitting errors also compound with long horizons.

5.5. Other Considerations

Hypeparameter selection. As with many machine learning techniques, hyperparameter choice affects the performance of most estimators (except IPS estimators). The situation is more acute for OPE than the online off-policy learning setting, due to the lack of proper validation signal (such as online game score). When using function approximation, direct methods may not have satisfactory convergence, and require setting a reasonable termination threshold hyperparameter. Q-Reg and MRDR require extra care to avoid ill-conditioning, such as tuning with L1 and L2 regularization. Similarly, the various choice of the kernel function for IH and the index set for hybrid method such as MAGIC have large impact on the performance. In general, given the choice among different hybrid (or direct) methods, we recommend opting for simplicity as a guiding principle.

Computational considerations. DM are generally significantly more computationally demanding than IPS. In complex domains, model-free iterative methods can be expensive in training time. Iterative DM that incorporate rollouts until the end of trajectories during training (Retrace(λ), $Q^{\pi}(\lambda)$, Tree-Backup(λ)) are the most computationally de-

manding⁸, requiring an order of T times the number of $\hat{Q}_{k-1}(x,a)$ lookups per gradient step compared to FQE. Model-based method (AM) are expensive at test time when coupled with HM, since rolling-out the learned model is required at every state along the trajectory. HM versions of direct methods require T times more inference steps, which is often fast after training. In difficult tasks such as Atari games, running AM, Retrace(λ), $Q^{\pi}(\lambda)$, Tree-Backup(λ) can be prohibitively expensive. Q-Reg, MRDR are noniterative methods and thus are the fastest to execute among DM. The run-time of IH is dependent on the batch size in building a kernel matrix to compute state similarity. The batch size for IH should be as large as possible, but could significantly slow the training.

Sparsity (non-smoothness) of the rewards: Methods that are dependent on cumulative importance weights are also sensitive to reward sparsity (Figure 19). We recommend normalizing the rewards. As a rough guideline, zero-centering rewards often improve performance of methods that depend on importance weights. This seemingly naïve practice can be actually viewed as a special case of DR using a constant DM component (baseline), and can yield improvements over vanilla IPS (Jiang & Li, 2016).

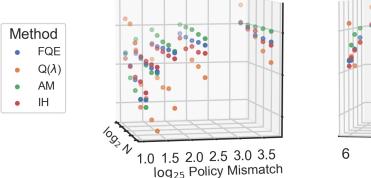
6. Discussion and Future Directions

The most difficult environments break all estimators. Atari games pose significant challenges for contemporary techniques due to long horizon and high state dimensionality. It is possible that substantially more historical data is required for current OPE methods to succeed. However, to overcome computational challenge in complex RL domains, it is important to identify principled ways to stabilize itera-

⁷From correspondence with the authors.

⁸Munos et al. (2016) limits the rolling-out horizon to 16 in Atari domains, but for the policy learning scenario.

⁹Unlike iterative DM (e.g., FQE), model-based method AM does not benefit from stochastic gradient speedup. Parallelizing the rollouts of AM is highly recommended.



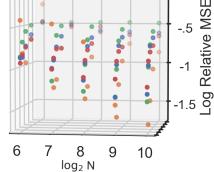


Figure 5. (Gridworld domain) Errors are directly correlated with policy mismatch but inversely correlated with data size. We pick the best direct methods for illustration. The two plots represent the same figure from two different vantage points. See full figures in appendix.

tive methods such as FQE, Retrace(λ), Q(λ) when using function approximation, as convergence is typically not attainable. Some recent progress has been made in stabilizing batch Q-learning in the off-policy learning setting (Fujimoto et al., 2019). It remains to be seen whether similar approach can also benefit DM for OPE.

Lack of short-horizon benchmark in high-dimensional settings. Evaluation of other complex RL tasks with short horizon is currently beyond the scope of our study, due to the lack of a natural benchmark. We refer to prior work on OPE for contextual bandits, which are RL problems with horizon 1 (Dudík et al., 2011). For contextual bandits, it has been shown that while DR is highly competitive, it is sometimes substantially outperformed by DM (Wang et al., 2017). New benchmark tasks should have longer horizon than contextual bandits, but shorter than typical Atari games. We also currently lack natural stochastic environments in high-dimensional RL benchmarks. An example candidate for medium horizon, complex OPE domain is NLP tasks such as dialogue.

Other OPE settings. Below we outline several practically relevant settings that current literature has overlooked:

- Continuous actions. Recent literature on OPE has exclusively focused on finite actions. OPE for continuous action domains will benefit continuous control applications. Currently, continuous action domains will not work with all IPS and HM (see IPS for continuous contextual bandits by Kallus & Zhou (2018)). Among DM, perhaps only FQE may reasonable work with continuous action tasks with some adaptation.
- Missing data coverage. A common assumption in the analysis of OPE is a full support assumption: $\pi_e(a|x) > 0$ implies $\pi_b(a|x) > 0$, which often ensure unbiasedness of estimators (Precup et al., 2000; Liu et al., 2018; Dudík et al., 2011). This assumption may not hold, and is often not verifiable in practice. Practi-

- cally, violation of this assumption requires regularization of unbiased estimators to avoid ill-conditioning (Liu et al., 2018; Farajtabar et al., 2018). One avenue to investigate is to optimize bias-variance tradeoff when the full support is not applicable.
- Confounding variables. Existing OPE research often assumes that the behavior policy chooses actions solely based on the state. This assumption is often violated when the decisions in the historical data are made by humans instead of algorithms, who may base their decisions on variables not recorded in the data, causing confounding effects. Tackling this challenge, possibly using techniques from causal inference (Tennenholtz et al., 2019; Oberst & Sontag, 2019), is an important future direction.

Evaluating new OPE estimators. More recently, several new OPE estimators have been proposed: Nachum et al. (2019); Zhang et al. (2020) further build on the perspective of density ratio estimation from IH; Uehara & Jiang (2019) provides a closely related approach that learns value functions from important ratios; Xie et al. (2019) proposes improvement over standard IPS by estimating marginalized state distribution in an analogous fashion to IH; Kallus & Uehara (2019a;b) analyze double reinforcement learning estimator that makes use of both estimates for Q function and state density ratio. While we have not included these new additions in our analysis, our software implementation is highly modular and can easily accommodate new estimators and environments.

Algorithmic approach to method selection. While we have identified a general guideline for selecting OPE method, often it is not easy to judge whether some decision criteria are satisfied (e.g., quantifying model misspecification, degree of stochasticity, or appropriate data size). As more OPE methods continue to be developed, an important missing piece is a systematic technique for model selection, given a high degree of variability among existing techniques.

References

- Bang, H. and Robins, J. M. Doubly robust estimation in missing data and causal inference models. *Biometrics*, 61(4):962–973, 2005. doi: 10.1111/j.1541-0420.2005. 00377.x.
- Bottou, L., Peters, J., nonero Candela, J. Q., Charles, D. X., Chickering, D. M., Portugaly, E., Ray, D., Simard, P., and Snelson, E. Counterfactual reasoning and learning systems: The example of computational advertising. *Journal of Machine Learning Research (JMLR)*, 14: 3207–3260, 2013.
- Brockman, G., Cheung, V., Pettersson, L., Schneider, J., Schulman, J., Tang, J., and Zaremba, W. Openai gym, 2016.
- Caruana, R. and Niculescu-Mizil, A. An empirical comparison of supervised learning algorithms. In *Proceedings* of the 23rd international conference on Machine learning, pp. 161–168, 2006.
- Caruana, R., Karampatziakis, N., and Yessenalina, A. An empirical evaluation of supervised learning in high dimensions. In *Proceedings of the 25th international conference on Machine learning*, pp. 96–103, 2008.
- Chapelle, O. and Li, L. An empirical evaluation of thompson sampling. In *Advances in neural information processing systems*, pp. 2249–2257, 2011.
- Chen, J. and Jiang, N. Information-theoretic considerations in batch reinforcement learning. In *International Conference on Machine Learning (ICML)*, 2019.
- Degris, T., White, M., and Sutton, R. S. Off-policy actor-critic. 2012.
- Dudík, M., Langford, J., and Li, L. Doubly robust policy evaluation and learning. In *International Conference on Machine Learning (ICML)*, 2011.
- Farajtabar, M., Chow, Y., and Ghavamzadeh, M. More robust doubly robust off-policy evaluation. In *International Conference on Machine Learning (ICML)*, 2018.
- Fujimoto, S., Meger, D., and Precup, D. Off-policy deep reinforcement learning without exploration. In *International Conference on Machine Learning*, pp. 2052–2062, 2019.
- Gauci, J., Conti, E., Liang, Y., Virochsiri, K., Chen, Z., He, Y., Kaden, Z., Narayanan, V., and Ye, X. Horizon: Facebook's open source applied reinforcement learning platform. *arXiv preprint arXiv:1811.00260*, 2018.
- Hammersley, J. M. and Handscomb, D. C. Monte carlo methods. 1964.

- Hanna, J., Niekum, S., and Stone, P. Importance sampling policy evaluation with an estimated behavior policy. In *International Conference on Machine Learning (ICML)*, 2019.
- Harutyunyan, A., Bellemare, M. G., Stepleton, T., and Munos, R. Q(lambda) with off-policy corrections. In *Conference on Algorithmic Learning Theory (ALT)*, 2016.
- Henderson, P., Islam, R., Bachman, P., Pineau, J., Precup, D., and Meger, D. Deep reinforcement learning that matters. In *Thirty-Second AAAI Conference on Artificial In*telligence, 2018.
- Horvitz, D. G. and Thompson, D. J. A generalization of sampling without replacement from a finite universe. *Journal of the American statistical Association*, 47(260): 663–685, 1952.
- Jiang, N. and Li, L. Doubly robust off-policy value evaluation for reinforcement learning. In *International Conference on Machine Learning (ICML)*, 2016.
- Kallus, N. and Uehara, M. Double reinforcement learning for efficient off-policy evaluation in markov decision processes. *arXiv preprint arXiv:1908.08526*, 2019a.
- Kallus, N. and Uehara, M. Efficiently breaking the curse of horizon: Double reinforcement learning in infinite-horizon processes. *arXiv preprint arXiv:1909.05850*, 2019b.
- Kallus, N. and Zhou, A. Policy evaluation and optimization with continuous treatments. In *Proceedings of the 21st International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2018.
- Kang, J. D., Schafer, J. L., et al. Demystifying double robustness: A comparison of alternative strategies for estimating a population mean from incomplete data. *Statistical science*, 22(4):523–539, 2007.
- Le, H. M., Voloshin, C., and Yue, Y. Batch policy learning under constraints. In *International Conference on Machine Learning (ICML)*, 2019.
- Li, L., Chu, W., Langford, J., and Wang, X. Unbiased offline evaluation of contextual-bandit-based news article recommendation algorithms. In *ACM Conference on Web Search and Data Mining (WSDM)*, 2011.
- Li, L., Munos, R., and Szepesvári, C. Toward minimax off-policy value estimation. In *Conference on Artificial Intelligence and Statistics (AISTATS)*, 2015.
- Liu, Q., Li, L., Tang, Z., and Zhou, D. Breaking the curse of horizon: Infinite-horizon off-policy estimation. In *Neural Information Processing Systems (NeurIPS)*, 2018.

- Liu, Y., Swaminathan, A., Agarwal, A., and Brunskill, E. Off-policy policy gradient with state distribution correction. arXiv preprint arXiv:1904.08473, 2019.
- Locatello, F., Bauer, S., Lucic, M., Raetsch, G., Gelly, S., Schölkopf, B., and Bachem, O. Challenging common assumptions in the unsupervised learning of disentangled representations. In *International Conference on Machine Learning*, pp. 4114–4124, 2019.
- Ma, J., Zhao, Z., Yi, X., Yang, J., Chen, M., Tang, J., Hong, L., and Chi, E. H. Off-policy learning in two-stage recommender systems.
- Munos, R. and Szepesvári, C. Finite-time bounds for fitted value iteration. *Journal of Machine Learning Research* (*JMLR*), 9(May):815–857, 2008.
- Munos, R., Stepleton, T., Harutyunyan, A., and Bellemare,
 M. Safe and efficient off-policy reinforcement learning.
 In *Neural Information Processing Systems (NeurIPS)*.
 2016.
- Murphy, S. A., van der Laan, M. J., Robins, J. M., and Group, C. P. P. R. Marginal mean models for dynamic regimes. *Journal of the American Statistical Association*, 96(456):1410–1423, 2001.
- Nachum, O., Chow, Y., Dai, B., and Li, L. Dualdice: Behavior-agnostic estimation of discounted stationary distribution corrections. In *Advances in Neural Information Processing Systems*, pp. 2315–2325, 2019.
- Nie, X., Brunskill, E., and Wager, S. Learning when-to-treat policies. *arXiv preprint arXiv:1905.09751*, 2019.
- Oberst, M. and Sontag, D. Counterfactual off-policy evaluation with gumbel-max structural causal models. In *International Conference on Machine Learning*, pp. 4881–4890, 2019.
- Paduraru, C. *Off-policy evaluation in Markov decision processes*. PhD thesis, McGill University Libraries, 2013.
- Powell, M. J. and Swann, J. Weighted uniform samplinga monte carlo technique for reducing variance. *IMA Journal of Applied Mathematics*, 2(3):228–236, 1966.
- Precup, D., Sutton, R. S., and Singh, S. P. Eligibility traces for off-policy policy evaluation. In *International Conference on Machine Learning (ICML)*, 2000.
- Sherman, J. and Morrison, W. J. Adjustment of an inverse matrix corresponding to a change in one element of a given matrix. *Ann. Math. Statist.*, 21(1):124–127, 03 1950. doi: 10.1214/aoms/1177729893.
- Sutton, R. S. and Barto, A. G. *Reinforcement learning: An introduction*. MIT press, 2018.

- Swaminathan, A., Krishnamurthy, A., Agarwal, A., Dudik, M., Langford, J., Jose, D., and Zitouni, I. Off-policy evaluation for slate recommendation. In *Neural Infor*mation Processing Systems (NeurIPS). 2017.
- Tennenholtz, G., Mannor, S., and Shalit, U. Off-policy evaluation in partially observable environments. *arXiv* preprint arXiv:1909.03739, 2019.
- Thomas, P. Reinforcement learning: An introduction, 2015.
- Thomas, P. and Brunskill, E. Data-efficient off-policy policy evaluation for reinforcement learning. In *International Conference on Machine Learning (ICML)*, 2016.
- Thomas, P. S., Theocharous, G., Ghavamzadeh, M., Durugkar, I., and Brunskill, E. Predictive off-policy policy evaluation for nonstationary decision problems, with applications to digital marketing. In *AAAI Conference on Innovative Applications (IAA)*, 2017.
- Uehara, M. and Jiang, N. Minimax weight and q-function learning for off-policy evaluation. *arXiv* preprint *arXiv*:1910.12809, 2019.
- Wang, Y.-X., Agarwal, A., and Dudík, M. Optimal and adaptive off-policy evaluation in contextual bandits. In *International Conference on Machine Learning (ICML)*, 2017.
- Wiering, M. Multi-agent reinforcement learning for traffic light control. In *International Conference on Machine Learning (ICML)*, 2000.
- Xie, T., Ma, Y., and Wang, Y.-X. Towards optimal off-policy evaluation for reinforcement learning with marginalized importance sampling. In *Advances in Neural Information Processing Systems*, pp. 9665–9675, 2019.
- Zhang, S., Liu, B., and Whiteson, S. Gradientdice: Rethinking generalized offline estimation of stationary values. *arXiv preprint arXiv:2001.11113*, 2020.

Empirical Study of Off-Policy Policy Evaluation for Reinforcement Learning

C	ontei	nts				E.1.4	Mountain Car (MC)	18
1	Intr	oduction	1			E.1.5	Pixel-based Mountain Car (Pix-MC)	19
1	11111	oduction	1			E.1.6	Enduro	19
2	Prel	iminaries	2			E.1.7	Gridworld (GW)	19
2	0	on a CODE Made de	•			E.1.8	Pixel-Gridworld (Pixel-GW)	20
3		rview of OPE Methods	2 2	F	Exp	erimen	tal Setup	20
	3.1	Inverse Propensity Scoring (IPS)	2		F.1		iption of the policies	20
	3.2	Direct Methods (DM)	3		F.2		eration of Experiments	
	3.3	Tryona Methods (Thvi)	3			F.2.1	Graph	
4	Exp	eriments	3			F.2.2	Graph-POMDP	20
5	Resu	ulte	4			F.2.3	Gridworld	20
3	5.1	What is the best method?	4			F.2.4	Pixel-Gridworld (Pix-GW)	20
	5.2	Key drivers of method accuracy	4			F.2.5	Graph-MC	20
	5.3	A recipe for method selection	5			F.2.6	Mountain Car (MC)	20
	5.4	Challenging common wisdom	6			F.2.7	Pixel-Mountain Car (Pix-MC)	20
	5.5	Other Considerations	7			F.2.8	Enduro	20
					F.3	Repres	sentation and Function Class	20
6	Disc	ussion and Future Directions	7		F.4	Choice	e of hyperparameters	22
A	Glos	ssary of Terms	12	G	Add	itional	Supporting Figures	24
В	Ran	king of Methods	13	Н	Tabl	les of R	esults, per Environment	30
	B.1	Decision Tree Support	13		H.1	Detail	ed Results for Graph	31
C	C	nlam antaur. Fallalaus Baskum	16		H.2	Detail	ed Results for Graph-POMDP	95
C	Sup	plementary Folklore Backup	16		H.3		ed Results for Graph Mountain Car	110
D	Met	hods	17		H.4		h-MC)	
	D.1	Inverse Propensity Scoring (IPS) Methods .	17		H.5		ed Results for Pixel-Based Mountain	123
	D.2	Hybrid Methods	17		11.5		Pix-MC)	133
	D.3	Direct Methods (DM)	17		H.6	Detail	ed Results for Gridworld	134
		D.3.1 Model-Based	17		H.7	Detail	ed Results for Pixel Gridworld	141
		D.3.2 Model-Free	17		H.8	Detail	ed Results for Enduro	154
E	Envi	ironments	18					
	E.1	Environment Descriptions	18					
		E.1.1 Graph	18					
		E.1.2 Graph-POMDP	18					
		E.1.3 Graph Mountain Car (Graph-MC) .	18					

A. Glossary of Terms

See Table 3 for a description of the terms used in this paper.

Table 3. Glossary of terms

	Table 3. Glossary of terms
Acronym	Term
OPE	Off-Policy Policy Evaluation
X	State Space
A	Action Space
P	Transition Function
R	Reward Function
γ	Discount Factor
d_0	Initial State Distribution
D	Dataset
au	Trajectory/Episode
T	Horizon/Episode Length
N	Number of episodes in D
π_b	Behavior Policy
π_e	Evaluation Policy
V	Value, ex: $V(\pi_e)$
Q	Action-Value, ex: $Q(\pi_e, a)$
$ ho^i_{j:j'}$	Cumulative Importance Weight, $\prod_{t=j}^{\min(j',T-1)} \frac{\pi_e(a_t^i x_t^i)}{\pi_b(a_t^i x_t^i)}$. If $j>j'$ then default is $\rho=1$
IPS	Inverse Propensity Scoring
DM	Direct Method
HM	Hybrid Method
IS	Importance Sampling
PDIS	Per-Decision Importance Sampling
WIS	Weighted Importance Sampling
PDWIS	Per-Decision Weighted Importance Sampling
PDWIS	Per-Decision Weighted Importance Sampling
FQE	Fitted Q Evaluation (Le et al., 2019)
IH	Infinite Horizon (Liu et al., 2018)
Q-Reg	Q Regression (Farajtabar et al., 2018)
MRDR	More Robust Doubly Robst (Farajtabar et al., 2018)
AM	Approximate Model (Model Based)
$Q(\lambda)$	$Q^{\pi}(\lambda)$ (Harutyunyan et al., 2016)
$R(\lambda)$	Retrace(λ) (Munos et al., 2016)
Tree	Tree-Backup(λ) (Precup et al., 2000)
DR	Doubly-Robust (Jiang & Li, 2016; Dudík et al., 2011)
WDR	Weighted Doubly-Robust (Dudík et al., 2011)
MAGIC	Model And Guided Importance Sampling Combining (Estimator) (Thomas & Brunskill, 2016)
Graph	Graph Environment
Graph-MC	Graph Mountain Car Environment
MC	Mountain Car Environment
Pix-MC	Pixel-Based Mountain Car Environment
Enduro	Enduro Environment
Graph-POMDP	Graph-POMDP Environment
GW	Gridworld Environment
Pix-GW	Pixel-Based Gridworld Environment

B. Ranking of Methods

A method that is within 10% of the method with the lowest Relative MSE is counted as a top method, called Near-top Frequency, and then we aggregate across all experiments. See Table 4 for a sorted list of how often the methods appear within 10% of the best method.

Table 4. Fraction of time among the top estimators across all experiments

Method	Near-top Frequency
MAGIC FQE	0.300211
DM FQE	0.236786
IH	0.190275
WDR FQE	0.177590
MAGIC $Q^{\pi}(\lambda)$	0.173362
WDR $Q^{\pi}(\lambda)$	0.173362
$DM Q^{\pi}(\lambda)$	0.150106
$DR Q^{\pi}(\lambda)$	0.135307
WDR $R(\lambda)$	0.133192
DR FQE	0.128964
MAGIC $R(\lambda)$	0.107822
WDR Tree	0.105708
DR $R(\lambda)$	0.105708
DM $R(\lambda)$	0.097252
DM Tree	0.084567
MAGIC Tree	0.076110
DR Tree	0.073996
DR MRDR	0.073996
WDR Q-Reg	0.071882
DM AM	0.065539
IS	0.063425
WDR MRDR	0.054968
PDWIS	0.046512
DR Q-Reg	0.044397
MAGIC AM	0.038055
MAGIC MRDR	0.033827
DM MRDR	0.033827
PDIS	0.033827
MAGIC Q-Reg	0.027484
WIS	0.025370
NAIVE	0.025370
DM Q-Reg	0.019027
DR AM	0.012685
WDR AM	0.006342

B.1. Decision Tree Support

Tables 5-12 provide a numerical support for the decision tree in the main paper (Figure 2). Each table refers to a child node in the decision tree, ordered from left to right, respectively. For example, Table 5 refers to the leftmost child node (propery specified, short horizon, small policy mismatch) while Table 12 refers to the right-most child node (misspecified, good representation, long horizon, good π_b estimate).

Table 5. Near-top Frequency among the properly specified, short horizon, small policy mismatch experiments

	DM		Hybrid	ı
	DIRECT	DR	WDR	MAGIC
AM	4.7%	4.7%	3.1%	4.7%
Q-Reg	0.0%	4.7%	6.2%	4.7%
MRDR	7.8%	14.1%	7.8%	7.8%
FQE	40.6%	23.4%	21.9%	34.4%
$R(\lambda)$	17.2%	20.3%	20.3%	14.1%
$Q^{\hat{\pi}}(\lambda)$	21.9%	18.8%	18.8%	17.2%
TREE	15.6%	12.5%	12.5%	14.1%
IH	17.2%	-	-	-

	IPS				
	STANDARD PER-DECI				
IS	4.7%	4.7%			
WIS	3.1%	3.1%			
NAIVE	1.6%	-			

Table 6. Near-top Frequency among the properly specified, short horizon, large policy mismatch experiments

	DM	•	Hybrid	,
	DIRECT	DR	WDR	MAGIC
AM	20.3%	1.6%	0.0%	7.8%
Q-Reg	1.6%	1.6%	3.1%	1.6%
MRDR	3.1%	1.6%	6.2%	1.6%
FQE	35.9%	14.1%	17.2%	37.5 %
$R(\lambda)$	23.4%	14.1%	20.3%	23.4%
$Q^{\hat{\pi}}(\lambda)$	15.6%	15.6%	14.1%	20.3%
TREE	21.9%	12.5%	18.8%	21.9%
IH	29.7%	-	-	-

	IPS			
	STANDARD	PER-DECISION		
IS	0.0%	0.0%		
WIS	0.0%	1.6%		
NAIVE	3.1%	-		

Table 7. Near-top Frequency among the properly specified, long horizon, small policy mismatch experiments

	DM		Hybrid	ı
	DIRECT	DR	WDR	MAGIC
AM	6.9%	0.0%	0.0%	5.6%
Q-Reg	0.0%	1.4%	1.4%	1.4%
MRDR	1.4%	0.0%	1.4%	2.8%
FQE	50.0%	22.2%	23.6%	50.0%
$R(\lambda)$	13.9%	12.5%	11.1%	9.7%
$Q^{\hat{\pi}}(\lambda)$	20.8%	18.1%	18.1%	18.1%
TREE	2.8%	1.4%	0.0%	2.8%
IH	29.2%	-	-	-

	IPS				
	STANDARD PER-DECISION				
IS	0.0%	0.0%			
WIS	0.0%	0.0%			
NAIVE	5.6%	-			

Table 8. Near-top Frequency among the properly specified, long horizon, large policy mismatch, deterministic env/rew experiments

	DM		Hybrid)
	DIRECT	DR	WDR	MAGIC
AM	3.5%	3.5%	1.8%	1.8%
Q-Reg	3.5%	1.8%	0.0%	0.0%
MRDR	3.5%	1.8%	0.0%	0.0%
FQE	15.8%	17.5%	29.8%	28.1%
$R(\lambda)$	1.8%	3.5%	0.0%	0.0%
$Q^{\hat{\pi}}(\lambda)$	22.8%	15.8%	38.6%	24.6%
TREE	3.5%	3.5%	1.8%	1.8%
IH	21.1%	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	5.3%	3.5%		
WIS	0.0%	8.8%		
NAIVE	0.0%	-		

Table 9. Near-top Frequency among the properly specified, long horizon, large policy mismatch, stochastic env/rew experiments

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	14.6%	0.0%	0.0%	8.3%
Q-Reg	4.2%	2.1%	0.0%	2.1%
MRDR	4.2%	2.1%	0.0%	0.0%
FQE	31.2%	2.1%	0.0%	25.0%
$R(\lambda)$	4.2%	6.2%	0.0%	0.0%
$Q^{\pi}(\lambda)$	2.1%	0.0%	0.0%	2.1%
TREE	4.2%	6.2%	0.0%	0.0%
IH	41.7%	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	25.0%	4.2%		
WIS	0.0%	0.0%		
NAIVE	2.1%	-		

Table 10. Near-top Frequency among the potentially misspecified, insufficient representation experiments

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	-	-	-	-	
Q-REG	3.9%	13.7%	25.5%	6.9%	
MRDR	0.0%	18.6%	15.7%	5.9%	
FQE	0.0%	5.9%	13.7%	24.5%	
$R(\lambda)$	-	-	-	-	
$Q^{\hat{\pi}}(\lambda)$	-	-	-	-	
TREE	-	-	-	-	
IH	6.9%	-	-	-	

	IPS			
	STANDARD PER-DECISIO			
IS	10.8%	8.8%		
WIS	9.8%	13.7%		
NAIVE	3.9%	-		

Table 11. Near-top Frequency among the potentially misspecified, sufficient representation, poor π_b estimate experiments

DM	Hybrid		D
DIRECT	DR	WDR	MAGIC
0.0%	0.0%	0.0%	0.0%
0.0%	0.0%	3.3%	0.0%
13.3%	6.7%	0.0%	0.0%
0.0%	3.3%	6.7%	10.0%
16.7%	0.0%	6.7%	20.0%
6.7%	0.0%	0.0%	3.3%
20.0%	0.0%	6.7%	6.7%
0.0%	-	-	-
	0.0% 0.0% 13.3% 0.0% 16.7% 6.7% 20.0%	DIRECT DR 0.0% 0.0% 0.0% 0.0% 13.3% 6.7% 0.0% 3.3% 16.7% 0.0% 6.7% 0.0% 20.0% 0.0%	DIRECT DR WDR 0.0% 0.0% 0.0% 0.0% 0.0% 3.3% 13.3% 6.7% 0.0% 0.0% 3.3% 6.7% 16.7% 0.0% 6.7% 6.7% 0.0% 0.0% 20.0% 0.0% 6.7%

	IPS			
	STANDARD PER-DECISION			
IS	3.3%	0.0%		
WIS	0.0%	0.0%		
NAIVE	0.0%	-		

Table 12. Near-top Frequency among the potentially misspecified, sufficient representation, good π_b estimate experiments

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	0.0%	0.0%	0.0%	2.8%	
Q-Reg	0.0%	0.0%	0.0%	0.0%	
MRDR	0.0%	5.6%	0.0%	5.6%	
FQE	8.3%	8.3%	25.0%	11.1%	
$R(\lambda)$	2.8%	8.3%	8.3%	19.4%	
$Q^{\pi}(\lambda)$	5.6%	5.6%	8.3%	0.0%	
TREE	5.6%	8.3%	16.7%	5.6%	
IH	0.0%	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	0.0%	0.0%		
WIS	0.0%	0.0%		
NAIVE	0.0%	-		

C. Supplementary Folklore Backup

The following tables represent the numerical support for how horizon and policy difference affect the performance of the OPE estimators when policy mismatch is held constant. Notice that the policy mismatch for table 14 and 15 are identical: $\left(\frac{.124573...}{.1}\right)^{100}=\left(\frac{.9}{.1}\right)^{10}$. What we see here is that despite identical policy mismatch, the longer horizon does not impact the error as much (compared to the baseline, Table 13) as moving π_e to .9, far from .1 and keeping the horizon the same.

Table 13. Graph, relative MSE. $T=10, N=50, \pi_b(a=0)=0.1, \pi_e(a=0)=0.1246.$ Dense rewards. Baseline.

	DM	Hybrid		
	DIRECT	DR	WDR	MAGIC
AM	1.9E-3	4.9E-3	5.0E-3	3.4E-3
Q-Reg	2.4E-3	4.3E-3	4.2E-3	4.5E-3
MRDR	5.8E-3	8.9E-3	9.4E-3	9.2E-3
FQE	1.8E-3	1.8E-3	1.8E-3	1.8E-3
$R(\lambda)$	1.8E-3	1.8E-3	1.8E-3	1.8E-3
$Q^{\hat{\pi}}(\lambda)$	1.8E-3	1.8E-3	1.8E-3	1.8E-3
TREE	1.8E-3	1.8E-3	1.8E-3	1.8E-3
IH	1.6E-3	-	-	-

	IPS				
	STANDARD PER-DECISIO				
IS	5.6E-4	8.4E-4			
WIS	1.4E-3	1.4E-3			
NAIVE	6.1E-3	-			

Table 14. Graph, relative MSE. $T=100, N=50, \pi_b(a=0)=0.1, \pi_e(a=0)=0.1246$. Dense rewards. Increasing horizon compared to baseline, fixed π_e .

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.6E-2	5.9E-2	5.9E-2	5.3E-2
Q-Reg	3.4E-3	1.1E-1	1.2E-1	9.2E-2
MRDR	1.1E-2	2.5E-1	2.9E-1	3.1E-1
FQE	6.0E-2	6.0E-2	6.0E-2	6.0E-2
$R(\lambda)$	6.0E-2	6.0E-2	6.0E-2	6.0E-2
$Q^{\hat{\pi}}(\lambda)$	6.0E-2	6.0E-2	6.0E-2	6.0E-2
TREE	3.4E-1	7.0E-3	1.6E-3	2.3E-3
IH	4.7E-4	-	-	-

	IPS		
	STANDARD PER-DECISIO		
IS	1.7E-2	2.5E-3	
WIS	9.5E-4	4.9E-4	
NAIVE	5.4E-3	-	

Table 15. Graph, relative MSE. $T=10, N=50, \pi_b(a=0)=0.1, \pi_e(a=0)=0.9$. Dense rewards. Increasing π_e compared to baseline, fixed horizon.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.6E-1	6.7E-1	6.6E-1	6.6E-1
Q-Reg	5.4E-1	6.3E-1	1.3E0	9.3E-1
MRDR	5.4E-1	7.3E-1	2.0E0	2.0E0
FQE	6.6E-1	6.6E-1	6.6E-1	6.6E-1
$R(\lambda)$	6.7E-1	6.6E-1	9.3E-1	1.0E0
$Q^{\hat{\pi}}(\lambda)$	6.6E-1	6.6E-1	6.6E-1	6.6E-1
TREE	6.7E-1	6.6E-1	9.4E-1	1.0E0
IH	1.4E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	5.4E-1	
WIS	2.0E0	9.7E-1	
NAIVE	4.0E0	-	

D. Methods

Below we include a description of each of the methods we tested. Let $\tilde{T} = T - 1$.

D.1. Inverse Propensity Scoring (IPS) Methods

Table 16. IPS methods. (Dudík et al., 2011; Jiang & Li, 2016)

	Standard	PER-DECISION
IS WIS	$\sum_{i=1}^{N} \frac{\rho_{0;\tilde{T}}^{i}}{N} \sum_{t=0}^{\tilde{T}} \gamma^{t} r_{t} \\ \sum_{i=1}^{N} \frac{\rho_{0;\tilde{T}}^{i}}{w_{0;\tilde{T}}} \sum_{t=0}^{\tilde{T}} \gamma^{t} r_{t}$	$\sum_{i=1}^{N} \sum_{t=0}^{\tilde{T}} \gamma^{t} \frac{\rho_{0:t}^{i}}{N} r_{t} \\ \sum_{i=1}^{N} \sum_{t=0}^{\tilde{T}} \gamma^{t} \frac{\rho_{0:t}^{i}}{w_{0:t}} r_{t}$

Table 16 shows the calculation for the four traditional IPS estimators: V_{IS} , V_{PDIS} , V_{WIS} , V_{PDWIS} . In addition, we include the following method as well since it is a Rao-Blackwellization (Liu et al., 2018) of the IPS estimators:

D.2. Hybrid Methods

Hybrid rely on being supplied an action-value function \widehat{Q} , an estimate of Q, from which one can also yield $\widehat{V}(x) = \sum_{a \in A} \pi(a|x) \widehat{Q}(x,a)$. Doubly-Robust (DR): (Thomas & Brunskill, 2016; Jiang & Li, 2016)

$$\begin{split} V_{DR} &= \frac{1}{N} \sum_{i=1}^{N} \widehat{V}(x_{0}^{i}) + \\ &\frac{1}{N} \sum_{i=1}^{N} \sum_{t=0}^{\infty} \gamma^{t} \rho_{0:t}^{i} [r_{t}^{i} - \widehat{Q}(x_{t}^{i}, a_{t}^{i}) + \gamma \widehat{V}(x_{t+1}^{i})] \end{split}$$

Weighted Doubly-Robust (WDR): (Thomas & Brunskill, 2016)

$$V_{WDR} = \frac{1}{N} \sum_{i=1}^{N} \hat{V}(x_0^i) + \sum_{i=1}^{N} \sum_{t=0}^{\infty} \gamma^t \frac{\rho_{0:t}^i}{w_{0:t}} [r_t^i - \hat{Q}(x_t^i, a_t^i) + \gamma \hat{V}(x_{t+1}^i)]$$

MAGIC: (Thomas & Brunskill, 2016) Given $g_J = \{g^i | i \in J \subseteq \mathbb{N} \cup \{-1\}\}$ where

$$g^{j}(D) = \sum_{i=1}^{N} \sum_{t=0}^{j} \gamma^{t} \frac{\rho_{0:t}^{i}}{w_{0:t}} r_{t}^{i} + \sum_{i=1}^{N} \gamma^{j+1} \frac{\rho_{0:t}^{i}}{w_{0:t}} \widehat{V}(x_{j+1}^{i}) - \sum_{i=1}^{N} \sum_{t=0}^{j} \gamma^{t} (\frac{\rho_{0:t}^{i}}{w_{0:t}} \widehat{Q}(x_{t}^{i}, a_{t}^{i}) - \frac{\rho_{0:\tilde{T}}^{i}}{w_{0:\tilde{T}}} \widehat{V}(x_{t}^{i})),$$

then define $dist(y, Z) = \min_{z \in Z} |y - z|$ and

$$\widehat{b}_n(j) = dist(g_j^J(D), CI(g^{\infty}(D), 0.5))$$

$$\widehat{\Omega}_n(i, j) = Cov(g_i^J(D), g_j^J(D))$$

then, for a |J|-simplex $\Delta^{|J|}$ we can calculate

$$\widehat{x}^* \in \arg\min_{x \in \Delta^{|J|}} x^T [\widehat{\Omega}_n + \widehat{b}\widehat{b}^T] x$$

which, finally, yields

$$V_{MAGIC} = (\widehat{x}^*)^T g_J.$$

MAGIC can be thought of as a weighted average of different blends of the DM and Hybrid. In particular, for some $i \in J$, g^i represents estimating the first i steps of $V(\pi_e)$ according to DR (or WDR) and then estimating the remaining steps via \widehat{Q} . Hence, V_{MAGIC} finds the most appropriate set of weights which trades off between using a direct method and a Hybrid.

D.3. Direct Methods (DM)

D.3.1. MODEL-BASED

Approximate Model (AM): (Jiang & Li, 2016) An approach to model-based value estimation is to directly fit the transition dynamics $P(x_{t+1}|x_t,a_t)$, reward $R(x_t,a_t)$, and terminal condition $P(x_{t+1} \in X_{terminal}|x_t,a_t)$ of the MDP using some for of maximum likelihood or function approximation. This yields a simulation environment from which one can extract the value of a policy using an average over rollouts. Thus, $V(\pi) = \mathbb{E}[\sum_{t=1}^T \gamma^t r(x_t,a_t)|x_0 = x, a_0 = \pi(x_0)]$ where the expectation is over initial conditions $x \sim d_0$ and the transition dynamics of the simulator.

D.3.2. MODEL-FREE

Every estimator in this section will approximate Q with $\widehat{Q}(\cdot;\theta)$, parametrized by some θ . From \widehat{Q} the OPE estimate we seek is

$$V = \frac{1}{N} \sum_{i=1}^{N} \sum_{a \in A} \pi_e(a|s) \widehat{Q}(s_0^i, a; \theta)$$

Note that $\mathbb{E}_{\pi_e}Q(x_{t+1},\cdot) = \sum_{a\in A} \pi_e(a|x_{t+1})Q(x_{t+1},a)$.

Direct Model Regression (Q-Reg): (Farajtabar et al., 2018)

$$\begin{split} \widehat{Q}(\cdot,\theta) &= \min_{\theta} \frac{1}{N} \sum_{i=1}^{N} \sum_{t=0}^{\tilde{T}} \gamma^{t} \rho_{0:t}^{i} \left(R_{t:\tilde{T}}^{i} - \widehat{Q}(x_{t}^{i}, a_{t}^{i}; \theta) \right)^{2} \\ R_{t:\tilde{T}}^{i} &= \sum_{t'=1}^{\tilde{T}} \gamma^{t'-t} \rho_{(t+1):t'}^{i} r_{t'}^{i} \end{split}$$

<u>Fitted Q Evaluation (FQE)</u>: (Le et al., 2019) $\widehat{Q}(\cdot, \theta) = \lim_{k \to \infty} \widehat{Q}_k$ where

$$\widehat{Q}_k = \min_{\theta} \frac{1}{N} \sum_{i=1}^{N} \sum_{t=0}^{\tilde{T}} (\widehat{Q}_{k-1}(x_t^i, a_t^i; \theta) - y_t^i)^2$$

$$y_t^i \equiv r_t^i + \gamma \mathbb{E}_{\pi_e} \widehat{Q}_{k-1}(x_{t+1}^i, \cdot; \theta)$$

Retrace(λ) (R(λ)), Tree-Backup (Tree), $Q^{\pi}(\lambda)$: (Munos et al., 2016; Precup et al., 2000; Harutyunyan et al., 2016)

$$\widehat{Q}(\cdot,\theta) = \lim_{k \to \infty} \widehat{Q}_k$$
 where
$$\widehat{Q}_k(x,a;\theta) = \widehat{Q}_{k-1}(x,a;\theta) + \\ \mathbb{E}_{\pi_b}[\sum_{t>0} \gamma^t \prod_{s=1}^t c_s y_t | x_0 = x, a_0 = a]$$

and

$$y_t = r^t + \gamma \mathbb{E}_{\pi_e} \widehat{Q}_{k-1}(x_{t+1}, \cdot; \theta) - \widehat{Q}_{k-1}(x_t, a_t; \theta)$$

$$c_s = \begin{cases} \lambda \min(1, \frac{\pi_e(a_s | x_s)}{\pi_b(a_s | x_s)}) & R(\lambda) \\ \lambda \pi_e(a_s | x_s) & Tree \\ \lambda & Q^{\pi}(\lambda) \end{cases}$$

More Robust Doubly-Robust (MRDR): (Farajtabar et al., 2018) Given

$$\Omega_{\pi_b}(x) = diag[1/\pi_b(a|x)]_{a \in A} - ee^T$$

$$e = [1, \dots, 1]^T$$

$$R_{t:\tilde{T}}^i = \sum_{j=t}^{\tilde{T}} \gamma^{j-t} \rho_{(t+1):j}^i r(x_j^i, a_j^i)$$

and

$$q_{\theta}(x, a, r) = diag[\pi_{e}(a'|x)]_{a' \in A}[\widehat{Q}(x, a'; \theta)]_{a' \in A} - r[\mathbf{1}\{a' = a\}]_{a' \in A}$$

where 1 is the indicator function, then

$$\widehat{Q}(\cdot, \theta) = \min_{\theta} \frac{1}{N} \sum_{i=1}^{N} \sum_{t=0}^{\tilde{T}} \gamma^{2t} (\rho_{0:\tilde{T}}^{i})^{2} \times \rho_{t}^{i} q_{\theta}(x_{t}^{i}, a_{t}^{i}, R_{t:\tilde{T}}^{i})^{T} \Omega_{\pi_{b}}(x_{t}^{i}) q_{\theta}(x_{t}^{i}, a_{t}^{i}, R_{t:\tilde{T}}^{i})$$

State Density Ratio Estimation (IH): (Liu et al., 2018)

$$V_{IH} = \sum_{i=1}^{N} \sum_{t=0}^{\tilde{T}} \frac{\gamma^{t} \omega(s_{t}^{i}) \rho_{t:t} r_{t}^{i}}{\sum_{i'=0}^{N} \sum_{t'=1}^{\tilde{T}} \gamma^{t'} \omega(s_{t'}^{i'}) \rho_{t':t'}}$$

$$\omega(s_{t}^{i}) = \lim_{t \to \infty} \frac{\sum_{t=0}^{T} \gamma^{t} d_{\pi_{e}}(s_{t}^{i})}{\sum_{t=0}^{T} \gamma^{t} d_{\pi_{b}}(s_{t}^{i})}$$

where π_b is assumed to be a fixed data-generating policy, and d_{π} is the distribution of states when executing π from $s_0 \sim d_0$. The details for how to find ω can be found in Algorithm 1 and 2 of (Liu et al., 2018).

E. Environments

For every environment, we initialize the environment with a fixed horizon length T. If the agent reaches a goal before T or if the episode is not over by step T, it will transition to an environment-dependent absorbing state where it will stay until time T. For a high level description of the environment features, see Table 1.

E.1. Environment Descriptions

E.1.1. GRAPH

Figure 6 shows a visualization of the Toy-Graph environment. The graph is initialized with horizon T and with absorbing state $x_{abs} = 2T$. In each episode, the agent starts at a single starting state $x_0 = 0$ and has two actions, a = 0and a = 1. At each time step t < T, the agent can enter state $x_{t+1} = 2t+1$ by taking action a = 0, or $x_{t+1} = 2t+2$ by taking action a = 1. If the environment is stochastic, we simulate noisy transitions by allowing the agent to slip into $x_{t+1} = 2t + 2$ instead of $x_{t+1} = 2t + 1$ and vice-versa with probability .25. At the final time t = T, the agent always enters the terminal state x_{abs} . The reward is +1 if the agent transitions to an odd state, otherwise is -1. If the environment provides sparse rewards, then r = +1 if x_{T-1} is odd, r = -1 if x_{T-1} is even, otherwise r = 0. Similarly to deterministic rewards, if the environment's rewards are stochastic, then the reward is $r \sim N(1,1)$ if the agent transitions to an odd state, otherwise $r \sim N(-1, 1)$. If the rewards are sparse and stochastic then $r \sim N(1,1)$ if x_{T-1} is odd, otherwise $r \sim N(-1, 1)$ and r = 0 otherwise.

E.1.2. GRAPH-POMDP

Figure 10 shows a visualization of the Graph-POMDP environment. The underlying state structure of Graph-POMDP is exactly the Graph environment. However, the states are grouped together based on a choice of Graph-POMDP horizon length, H. This parameter groups states into H observable states. The agent only is able to observe among these states, and not the underlying MDP structure. Model-Fail (Thomas & Brunskill, 2016) is a special case of this environment when H = T = 2.

E.1.3. GRAPH MOUNTAIN CAR (GRAPH-MC)

Figure 7 shows a visualization of the Toy-MC environment. This environment is a 1-D graph-based simplification of Mountain Car. The agent starts at $x_0=0$, the center of the valley and can go left or right. There are 21 total states, 10 to the left of the starting position and 11 to the right of the starting position, and a terminal absorbing state $x_{abs}=22$. The agent receives a reward of r=-1 at every timestep. The reward becomes zero if the agent reaches the goal, which is state x=+11. If the agent reaches x=-10 and continues left then the agent remains in x=-10. If the agent does not reach state x=+11 by step x=-10 then the episode terminates and the agent transitions to the absorbing state.

E.1.4. MOUNTAIN CAR (MC)

We use the OpenAI version of Mountain Car (Brockman et al., 2016; Sutton & Barto, 2018) with a few simplifying

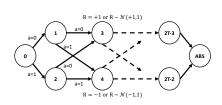


Figure 6. Graph Environment

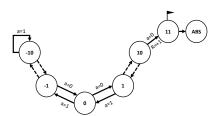


Figure 7. Graph-MC Environment



Figure 8. MC Environment, pixel-version. The non-pixel version involves representing the state of the car as the position and velocity.



Figure 9. Enduro Environment

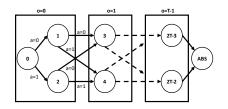


Figure 10. Graph-POMDP Environment. Model-Fail (Thomas & Brunskill, 2016) is a special case of this environment where T=2. We also extend the environment to arbitrary horizon which makes it a semi-mdp.

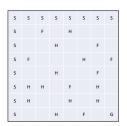


Figure 11. Gridworld environment. Blank spaces indicate areas of a small negative reward, S indicates the starting states, F indicates a field of slightly less negative reward, H indicates a hole of severe penalty, G indicates the goal of positive reward.

modifications. The car starts in a valley and has to go back and forth to gain enough momentum to scale the mountain and reach the end goal. The state space is given by the position and velocity of the car. At each time step, the car has the following options: accelerate backwards, forwards or do nothing. The reward is r=-1 for every time step until the car reaches the goal. While the original trajectory length is capped at 200, we decrease the effective length by applying every action a_t five times before observing x_{t+1} . Furthermore, we modify the random initial position from being uniformly between [-.6, -.4] to being one of $\{-.6, -.5, -.4\}$, with no velocity. The environment is initialized with a horizon T and absorbing state $x_{abs} = [.5, 0]$, position at .5 and no velocity.

E.1.5. PIXEL-BASED MOUNTAIN CAR (PIX-MC)

This environment is identical to Mountain Car except the state space has been modified from position and velocity to a pixel based representation of a ball, representing a car, rolling on a hill, see Figure 8. Each frame f_t is a 80×120 image of the ball on the mountain. One cannot deduce velocity from a single frame, so we represent the state as $x_t = \{f_{t-1}, f_t\}$ where $f_{-1} = f_0$, the initial state. Every-

thing else is identical between the pixel-based version and the position-velocity version described earlier.

E.1.6. ENDURO

We use OpenAI's implementation of Enduro-v0, an Atari 2600 racing game. We downsample the image to a grayscale of size (84,84). We apply every action one time and we represent the state as $x_t = \{f_{t-3}, f_{t-2}, f_{t-1}, f_t\}$ where $f_i = f_0$, the initial state, for i < 0. See Figure 9 for a visualization.

E.1.7. GRIDWORLD (GW)

Figure 11 shows a visualization of the Gridworld environment. The agent starts at a state in the first row or column (denoted S in the figure), and proceeds through the grid by taking actions, given by the four cardinal directions, for T=25 timesteps. An agent remains in the same state if it chooses an action which would take it out of the environment. If the agent reaches the goal state G, in the bottom right corner of the environment, it transitions to a terminal state x=64 for the remainder of the trajectory and receives a reward of +1. In the grid, there is a field

(denoted F) which gives the agent a reward of -.005 and holes (denoted H) which give -.5. The remaining states give a reward of -.01.

E.1.8. PIXEL-GRIDWORLD (PIXEL-GW)

This environment is identical to Gridworld except the state space has been modified from position to a pixel based representation of the position: 1 for the agent's location, 0 otherwise. We use the same policies as in the Gridworld case.

F. Experimental Setup

F.1. Description of the policies

Graph, Graph-POMDP and Graph-MC use static policies with some probability of going left and another probability of going right, ex: $\pi(a=0)=p, \pi(a=1)=1-p$, independent of state. We vary p in our experiments.

GW, Pix-GW, MC, Pixel-MC, and Enduro all use an ϵ -Greedy policy. In other words, we train a policy Q^* (using value iteration or DDQN) and then vary the deviation away from the policy. Hence $\epsilon - Greedy(Q^*)$ implies we follow a mixed policy $\pi = \arg\max_a Q^*(x,a)$ with probability $1 - \epsilon$ and uniform with probability ϵ . We vary ϵ in our experiments.

F.2. Enumeration of Experiments

F.2.1. GRAPH

See Table 18 for a description of the parameters of the experiment we ran in the Graph Environment. The experiments are the Cartesian product of the table.

F.2.2. GRAPH-POMDP

See Table 19 for a description of the parameters of the experiment we ran in the Graph-POMDP Environment. The experiments are the Cartesian product of the table.

F.2.3. GRIDWORLD

See Table 20 for a description of the parameters of the experiment we ran in the Gridworld Environment. The experiments are the Cartesian product of the table.

F.2.4. PIXEL-GRIDWORLD (PIX-GW)

See Table 21 for a description of the parameters of the experiment we ran in the Pix-GW Environment. The experiments are the Cartesian product of the table.

F.2.5. GRAPH-MC

See Table 22 for a description of the parameters of the experiment we ran in the TMC Environment. The experi-

ments are the Cartesian product of the table.

F.2.6. MOUNTAIN CAR (MC)

See Table 23 for a description of the parameters of the experiment we ran in the MC Environment. The experiments are the Cartesian product of the table.

F.2.7. PIXEL-MOUNTAIN CAR (PIX-MC)

See Table 24 for a description of the parameters of the experiment we ran in the Pix-MC Environment. The experiments are the Cartesian product of the table.

F.2.8. ENDURO

See Table 25 for a description of the parameters of the experiment we ran in the Enduro Environment. The experiments are the Cartesian product of the table.

F.3. Representation and Function Class

For the simpler environments, we use a tabular representation for all the methods. AM amounts to solving for the transition dynamics, rewards, terminal state, etc. through maximum likelihood. FQE, Retrace(λ), $Q^{\pi}(\lambda)$, and Tree-Backup are all implemented through dynamics programming with Q tables. MRDR and Q-Reg used the Sherman Morrison (Sherman & Morrison, 1950) method to solve the weighted-least square problem, using a basis which spans a table.

In the cases where we needed function approximation, we did not directly fit the dynamics for AM; instead, we fit on the difference in states T(x'-x|x,a), which is common practice.

For the MC environment, we ran experiments with both a linear and NN function class. In both cases, the representation of the state was not changed and remained [position, velocity]. The NN architecture was dense with [16,8,4,2] as the layers. The layers had relu activations (except the last, with a linear activation) and were all initialized with truncated normal centered at 0 with a standard deviation of 0.1.

For the pixel-based environments (MC, Enduro), we use a convolutional NN. The architechure is a layer of size 8 with filter (7,7) and stride 3, followed by maxpooling and a layer of size 16 with filter (3,3) and stride 1, followed by max pooling, flattening and a dense layer of size 256. The final layer is a dense layer with the size of the action space, with a linear activation. The layers had elu activations and were all initialized with truncated normal centered at 0 with a standard deviation of 0.1. The layers also have kernel L2 regularizers with weight 1e-6.

When using NNs for the IH method, we used the radial-

Table 17	Environment parameters	- Full de	scription

Environment	Graph	Graph-MC	MC	Pix-MC	Enduro	Graph-POMDP	GW	Pix-GW
Is MDP? State desc. T Stoch Env?	yes	yes	yes	yes	yes	no	yes	yes
	position	position	[pos, vel]	pixels	pixels	position	position	pixels
	4 or 16	250	250	250	1000	2 or 8	25	25
	variable	no	no	no	no	no	no	variable
Stoch Rew?	variable	no	no	no	no	no	no	no
Sparse Rew?	variable	terminal	terminal	terminal	dense	terminal	dense	dense
\hat{Q} Func. Class	tabular	tabular	linear/NN	NN	NN	tabular	tabular	NN
Initial state Absorb. state Frame height Frame skip	0 2T 1	0 22 1	variable [.5,0]	variable [.5,0] 2 5	gray img zero img 4	0 2T 1	variable 64 1	variable zero img 1

Table 18. Graph parameters

	Parameters
${\gamma}$.98
N	$2^{3:11}$
T	$\{4, 16\}$
$\pi_b(a=0)$	$\{.2,.6\}$
$\pi_e(a=0)$.8
Stochastic Env	{True, False}
Stochastic Rew	{True, False}
Sparse Rew	{True, False}
Seed	$\{10 \text{ of random}(0:2^{16})\}$
ModelType	Tabular
Regress π_b	False

Table 19. Graph-POMDP parameters

	Parameters
γ	.98
N	$2^{8:11}$
(T,H)	$\{(2,2),(16,6)\}$
$\pi_b(a=0)$	$\{.2,.6\}$
$\pi_e(a=0)$.8
Stochastic Env	{True, False}
Stochastic Rew	{True, False}
Sparse Rew	{True, False}
Seed	$\{10 \text{ of random}(0:2^{16})\}$
ModelType	Tabular
Regress π_b	False

Table 20. Gridworld parameters

	Parameters
γ	.98
N	$2^{6:11}$
T	25
ϵ – Greedy, π_b	$\{.2, .4, .6, .8, 1.\}$
ϵ – Greedy, π_e	.1
Stochastic Env	False
Stochastic Rew	False
Sparse Rew	False
Seed	$\{10 \text{ of random}(0:2^{16})\}$
ModelType	Tabular
Regress π_b	True

Table 21. Pix-GW parameters

	Parameters
γ	.96
N	$2^{6:9}$
T	25
ϵ – Greedy, π_b	$\{.2, .4, .6, .8, 1.\}$
ϵ – Greedy, π_e	.1
Stochastic Env	{True, False}
Stochastic Rew	False
Sparse Rew	False
Seed	$\{10 \text{ of random}(0:2^{16})\}$
ModelType	NN
Regress π_b	{True, False}

Table 22. Graph-MC parameters

	Parameters
γ	.99
N	$2^{7:11}$
T	250
$(\pi_b(a=0), \pi_e(a=0))$	$\{(.45, .45), (.6, .6), (.45.6), (.6, .45), (.8, .2), (.2, .8)\}$
Stochastic Env	False
Stochastic Rew	False
Sparse Rew	False
Seed	$\{10 \text{ of random}(0:2^{16})\}$
ModelType	Tabular
Regress π_b	False

Table 23. MC parameters

	Parameters
γ	.99
N	$2^{7:10}$
T	250
ϵ – Greedy, (π_b, π_e)	$\{(.1,0),(1,0)\ (1,.1),(.1,1)\}$
Stochastic Env	False
Stochastic Rew	False
Sparse Rew	False
Seed	$\{10 \text{ of random}(0:2^{16})\}$
ModelType	{Tabular, NN}
Regress π_b	False

Table 24. Pix-MC parameters

	Parameters
γ	.97
N	512
T	500
ϵ – Greedy, (π_b, π_e)	$\{(.25,0),(.1,0) \ (.25,.1)\}$
Stochastic Env	False
Stochastic Rew	False
Sparse Rew	False
Seed	$\{10 \text{ of random}(0:2^{16})\}$
ModelType	{Tabular, NN}
Regress π_b	False

Table 25. Enduro parameters

	Parameters
γ	.9999
N	512
T	500
ϵ – Greedy, (π_b, π_e)	$\{(.25,0),(.1,0) \ (.25,.1)\}$
Stochastic Env	False
Stochastic Rew	False
Sparse Rew	False
Seed	$\{10 \text{ of random}(0:2^{16})\}$
ModelType	{Tabular, NN}
Regress π_b	False

basis function and a shallow dense network for the kernel and density estimate respectively.

F.4. Choice of hyperparameters

Many methods require selection of convergence criteria, regularization parameters, batch sizes, and a whole host of other hyperparameters. Often there is a trade-off between computational cost and the accuracy of the method. Hyperparameter search is not feasible in OPE since there is no proper validation (like game score in learning). See Table 26 for a list of hyperparameters that were chosen for the experiments.

Table 26. Hyperparameters for each model by Environment

Method	Parameter	Graph	TMC	MC	Pix-MC	Enduro	Graph-POMDP	GW	Pix-GW
	Max Traj Len	T	T	50	50	-	T	T	T
	NN Fit Epochs	-	-	100	100	-	-	-	100
434	NN Batchsize	-	-	32	32	-	-	-	25
AM	NN Train size	-	-	.8	.8	-	-	-	.8
	NN Val size	-	-	.2	.2	-	-	-	.2
	NN Early Stop delta	-	-	1e-4	1e-4	-	-	-	1e-4
	Omega regul.	1	1	-	-	-	1	1	-
	NN Fit Epochs	-	-	80	80	80	-	-	80
Q-Reg	NN Batchsize	-	-	32	32	32	-	-	32
Q-Reg	NN Train size	-	-	.8	.8	.8	-	-	.8
	NN Val size	-	-	.2	.2	.2	-	-	.2
	NN Early Stop delta	-	-	1e-4	1e-4	1e-4	-	-	1e-4
	Convergence ϵ	1e-5	1e-5	1e-4	1e-4	1e-4	1e-5	4e-4	1e-4
	Max Iter	-	-	160	160	600	-	50	80
FQE	NN Batchsize	-	-	32	32	32	-	-	32
rQE	Optimizer Clipnorm	-	-	1.	1.	1.	-	-	1.
	Quad. prog. regular.	1e-3	1e-3	-	_	-	1e-3	1e-3	-
IH	NN Fit Epochs	-	-	10001	10001	10001	-	-	1001
	NN Batchsize	-	-	1024	128	128	-	-	128
	Omega regul.	1	1	-	-	-	1	1	-
	NN Fit Epochs	-	-	80	80	80	-	-	80
MRDR	NN Batchsize	-	-	1024	1024	1024	-	-	32
MKDK	NN Train size	-	-	.8	.8	.8	-	-	.8
	NN Val size	-	-	.2	.2	.2	-	-	.2
	NN Early Stop delta	-	-	1e-4	1e-4	1e-4	-	-	1e-4
	λ	.9	.9	.9	-	-	.9	.9	.9
	Convergence ϵ	1e-3	2e-3	1e-3	-	-	1e-3	2e-3	1e-3
	Max Iter	500	500	-	-	-	500	50	-
$R(\lambda)$	NN Fit Epochs	-	-	80	-	-	-	-	80
	NN Batchsize	-	-	4	-	-	-	-	4
	NN Train Size	-	-	.03	-	-	-	-	.03
	NN ClipNorm		-	1.	-	-		-	1.
	λ	.9	.9	.9	-	-	.9	.9	.9
	Convergence ϵ	1e-3	2e-3	1e-3	-	-	1e-3	2e-3	1e-3
	Max Iter	500	500	-	-	-	500	50	-
$Q^{\pi}(\lambda)$	NN Fit Epochs	-	-	80	-	-	-	-	80
	NN Batchsize	-	-	4	-	-	-	-	4
	NN Train Size	-	-	.03	-	-	-	-	.03
	NN ClipNorm	-	-	1.	-			-	1.
	λ	.9	.9	.9	-	-	.9	.9	.9
	Convergence ϵ	1e-3	2e-3	1e-3	-	-	1e-3	2e-3	1e-3
	Max Iter	500	500	-	-	-	500	50	-
Tree	NN Fit Epochs	-	-	80	-	-	-	-	80
	NN Batchsize	-	-	4	-	-	-	-	4
	NN Train Size	-	-	.03	-	-	-	-	.03
	NN ClipNorm	-	-	1.	-	-	-	-	1.

G. Additional Supporting Figures

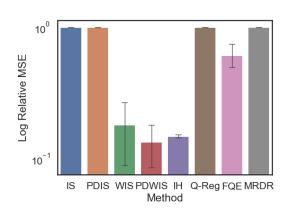


Figure 12. Enduro DM vs IPS. π_b is a policy that deviates uniformly from a trained policy 25% of the time, π_e is a policy trained with DDQN. IH has relatively low error mainly due to tracking the simple average, since the kernel function did not learn useful density ratio. The computational time required to calculate the multi-step rollouts of AM, Retrace(λ), $Q^{\pi}(\lambda)$, Tree-Backup(λ) exceeded our compute budget and were thus excluded.

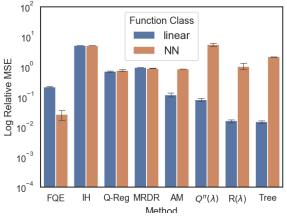


Figure 13. MC comparison. N=256. π_b is a uniform random policy, π_e is a policy trained with DDQN

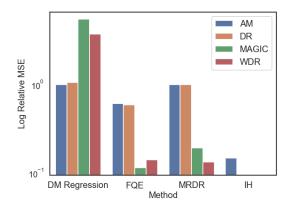


Figure 14. Enduro DM vs HM. π_b is a policy that deviates uniformly from a trained policy 25% of the time, π_e is a policy trained with DDQN.

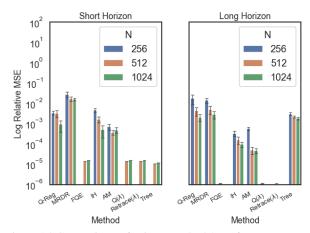


Figure 15. Comparison of Direct methods' performance across horizon and number of trajectories in the Toy-Graph environment. Small policy mismatch under a deterministic environment.

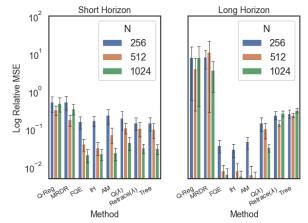


Figure 16. (Graph domain) Comparing DMs across horizon length and number of trajectories. Large policy mismatch and a stochastic environment setting.

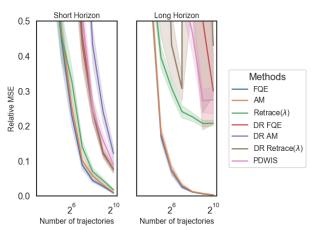


Figure 17. Comparing DM to DR in a stochastic environment with large policy mismatch. (Graph)

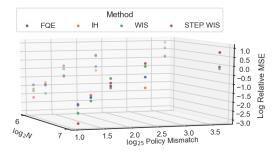


Figure 18. Comparison between FQE, IH and WIS in a low data regime. For low policy mismatch, IPS is competitive to DM in low data, but as the policy mismatch grows, the top DM outperform. Experiments ran in the Gridworld Environment.

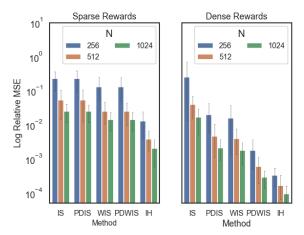


Figure 19. Comparison between IPS methods and IH with dense vs sparse rewards. Per-Decision IPS methods see substantial improvement when the rewards are dense. Experiments ran in the Toy-Graph environment with $\pi(a=0)=.6, \pi_e(a=0)=.8$ See Tables 224, 225, 226, 128, 129, 130

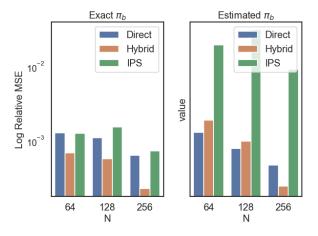


Figure 20. Exact vs Estimated π_b . Exact $\pi_b = .2$ —Greedy(optimal), $\pi_e = .1$ —Greedy(optimal). Min error per class. (Pixel Gridworld, deterministic)

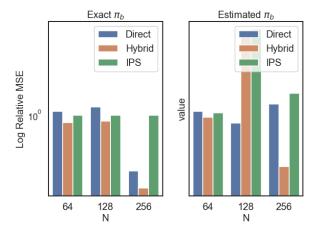


Figure 21. Exact vs Estimated π_b . Exact π_b =uniform, π_e = .1—Greedy(optimal). Min error per class. (Pixel Gridworld, deterministic)

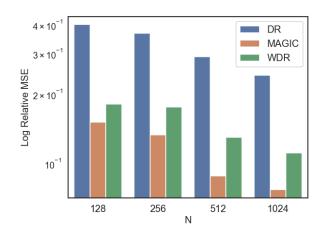


Figure 22. Hybrid Method comparison. $\pi_b(a=0)=.2, \pi_e(a=0)=.8.$ Min error per class. (Graph-MC)

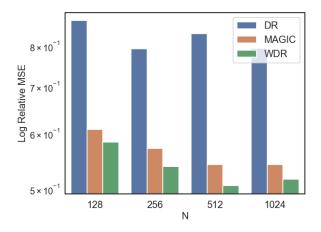


Figure 23. Hybrid Method comparison. $\pi_b(a=0)=.8, \pi_e(a=0)=.2$. Min error per class. (Graph-MC)

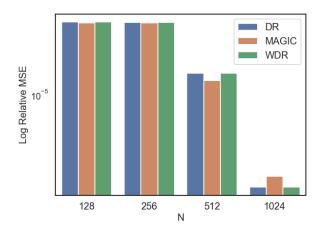


Figure 24. Hybrid Method comparison. $\pi_b(a=0)=.6, \pi_e(a=0)=.6.$ Min error per class. (Graph-MC)

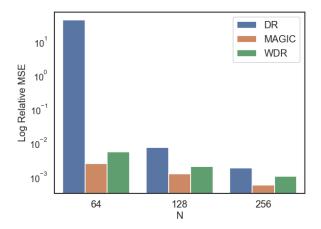


Figure 25. Hybrid Method comparison. Exact $\pi_b = .2$ —Greedy(optimal), $\pi_e = .1$ —Greedy(optimal). Min error per class. (Pixel Gridworld)

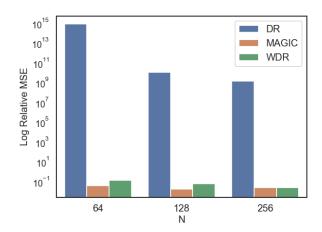


Figure 26. Hybrid Method comparison. $\pi_b = .8$ —Greedy(optimal), $\pi_e = .1$ —Greedy(optimal). Min error per class. (Pixel Gridworld)

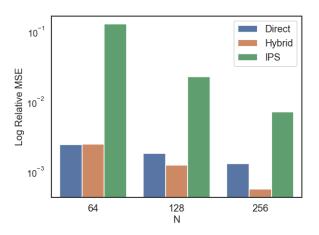


Figure 27. Class comparison with unknown π_b . At first, HM underperform DM because π_b is more difficult to calculate leading to imprecise importance sampling estimates. Exact $\pi_b = .2$ —Greedy(optimal), $\pi_e = .1$ —Greedy(optimal). Min error per class. (Pixel Gridworld, stochastic env with .2 slippage)

Empirical Study of Off-Policy Policy Evaluation for Reinforcement Learning

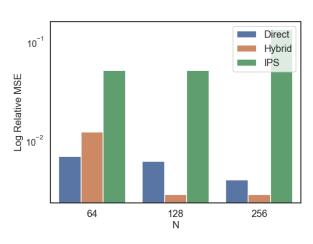


Figure 28. Class comparison with unknown π_b . At first, HM underperform DM because π_b is more difficult to calculate leading to imprecise importance sampling estimates. Exact $\pi_b = .6$ -Greedy(optimal), $\pi_e = .1$ -Greedy(optimal). Min error per class. (Pixel Gridworld, stochastic env with .2 slippage)

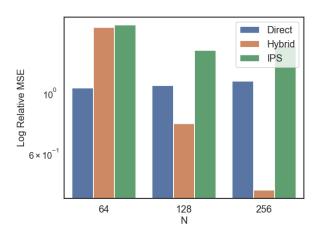


Figure 29. Class comparison with unknown π_b . At first, HM underperform DM because π_b is more difficult to calculate leading to imprecise importance sampling estimates. Exact π_b =uniform, $\pi_e = .1$ -Greedy(optimal). Min error per class. (Pixel Gridworld, stochastic env with .2 slippage)

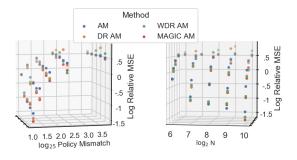


Figure 30. AM Direct vs Hybrid comparison for AM. (Gridworld)

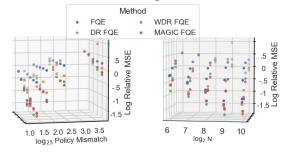


Figure 31. FQE Direct vs Hybrid comparison. (Gridworld)

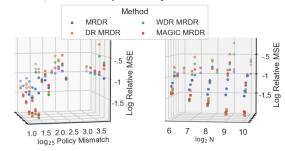


Figure 32. MRDR Direct vs Hybrid comparison. (Gridworld)

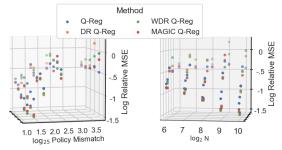


Figure 33. Q-Reg Direct vs Hybrid comparison. (Gridworld)

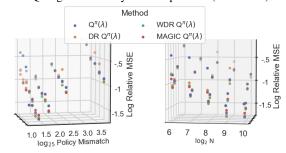


Figure 34. $Q^{\pi}(\lambda)$ Direct vs Hybrid comparison. (Gridworld)

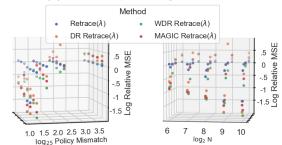


Figure 35. Retrace(λ) Direct vs Hybrid comparison. (Gridworld)

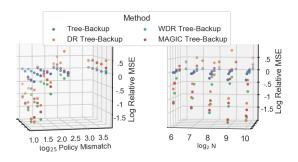


Figure 36. Tree-Backup Direct vs Hybrid comparison. (Gridworld)

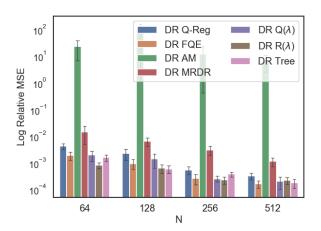


Figure 37. DR comparison with $\pi_b = .2$ —Greedy(optimal), $\pi_e = 1$.—Greedy(optimal). (Pixel Gridworld)

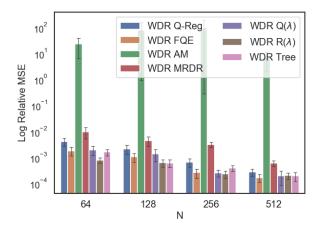


Figure 38. WDR comparison with $\pi_b = .2$ —Greedy(optimal), $\pi_e = 1$.—Greedy(optimal). (Pixel Gridworld)

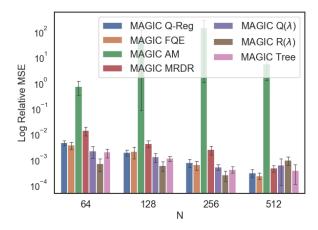


Figure 39. MAGIC comparison with $\pi_b = .2$ —Greedy(optimal), $\pi_e = 1$.—Greedy(optimal). (Pixel Gridworld)

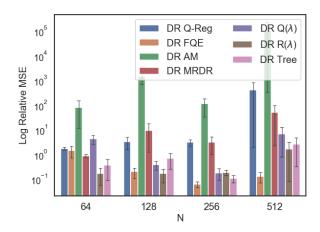


Figure 40. DR comparison with $\pi_b = .8$ —Greedy(optimal), $\pi_e = 1$.—Greedy(optimal). (Pixel Gridworld)

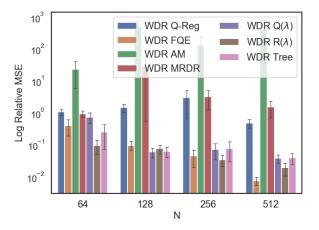


Figure 41. WDR comparison with $\pi_b=.8-$ Greedy(optimal), $\pi_e=1.-$ Greedy(optimal). (Pixel Gridworld)

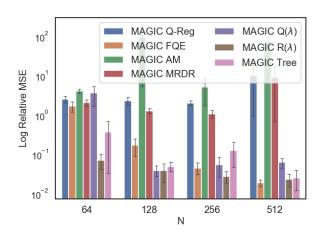


Figure 42. MAGIC comparison with $\pi_b=.8-$ Greedy(optimal), $\pi_e=1.-$ Greedy(optimal). (Pixel Gridworld)

H. Tables of Results, per Environment

H.1. Detailed Results for Graph

Table 27. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.9E-1	5.3E-1	6.4E-1	4.9E-1
Q-Reg	2.0E0	5.1E-1	5.7E-1	1.9E0
MRDR	1.7E0	1.7E0	7.0E-1	9.0E-1
FQE	4.8E-1	4.8E-1	4.8E-1	4.8E-1
$R(\lambda)$	4.8E-1	4.8E-1	4.8E-1	4.8E-1
$Q^{\pi}(\lambda)$	4.8E-1	4.8E-1	4.8E-1	4.8E-1
TREE	4.8E-1	4.8E-1	4.8E-1	4.8E-1
IH	2.9E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	2.4E0	3.3E0		
WIS	1.2E0	7.5E-1		
NAIVE	3.6E0	-		

Table 28. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.3E-1	6.7E-1	6.8E-1	4.9E-1
Q-Reg	4.3E0	7.0E0	5.9E-1	9.2E-1
MRDR	3.4E0	1.4E1	7.1E-1	2.9E0
FQE	3.9E-1	3.9E-1	3.9E-1	3.9E-1
$R(\lambda)$	3.9E-1	3.9E-1	3.9E-1	3.9E-1
$Q^{\hat{\pi}}(\lambda)$	3.9E-1	3.9E-1	3.9E-1	3.9E-1
TREE	3.9E-1	3.9E-1	4.0E-1	4.0E-1
IH	4.8E-1	-	-	-

	IPS				
	STANDARD PER-DECISION				
IS	6.7E1	7.9E0			
WIS	1.2E0	7.1E-1			
NAIVE	3.9E0	-			

Table 29. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.7E-1	2.8E-1	2.7E-1	1.8E-1
Q-Reg	5.0E-1	2.4E-1	3.7E-1	3.6E-1
MRDR	7.6E-1	3.1E-1	5.0E-1	3.1E-1
FQE	1.5E-1	1.5E-1	1.5E-1	1.5E-1
$R(\lambda)$	1.5E-1	1.5E-1	1.5E-1	1.5E-1
$Q^{\hat{\pi}}(\lambda)$	1.5E-1	1.5E-1	1.5E-1	1.5E-1
TREE	1.5E-1	1.5E-1	1.5E-1	1.5E-1
IH	3.9E-2	-	-	-

	IPS				
	STANDARD PER-DECISION				
IS	5.6E-1	3.8E-1			
WIS	5.0E-1	1.9E-1			
NAIVE	3.8E0	-			

Table 30. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.8E-2	1.1E-1	1.2E-1	3.6E-2
Q-REG	4.6E-1	1.2E-1	4.5E-2	2.4E-1
MRDR	3.4E-1	3.2E-1	1.2E-1	3.2E-1
FQE	3.4E-2	3.4E-2	3.4E-2	3.4E-2
$R(\lambda)$	3.4E-2	3.4E-2	3.4E-2	3.4E-2
$Q^{\hat{\pi}}(\lambda)$	3.4E-2	3.4E-2	3.4E-2	3.4E-2
TREE	3.4E-2	3.4E-2	3.4E-2	3.4E-2
IH	4.6E-2	-	-	-

	IPS			
	STANDARD	PER-DECISION		
IS	3.9E0	5.4E-1		
WIS	6.8E-1	2.0E-1		
NAIVE	4.0E0	-		

Table 31. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.5E-3	6.1E-2	5.2E-2	5.8E-3
Q-Reg	4.3E-1	9.7E-2	9.9E-3	1.4E-1
MRDR	3.9E-1	2.5E-1	6.9E-2	1.3E-1
FQE	1.2E-5	1.2E-5	1.2E-5	1.2E-5
$R(\lambda)$	1.2E-5	1.2E-5	9.0E-6	1.2E-5
$Q^{\hat{\pi}}(\lambda)$	1.2E-5	1.2E-5	1.2E-5	1.2E-5
TREE	1.4E-5	1.2E-5	1.1E-5	1.4E-5
IH	2.5E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	4.7E0	5.9E-1		
WIS	2.2E-1	4.5E-2		
NAIVE	3.9E0	-		

Table 32. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.3E-3	2.6E-2	2.2E-2	7.4E-3
Q-REG	6.9E-2	7.8E-3	1.4E-3	5.2E-2
MRDR	8.6E-2	6.9E-2	1.1E-1	2.7E-2
FQE	1.4E-8	1.4E-8	1.4E-8	1.4E-8
$R(\lambda)$	2.5E-8	2.5E-8	6.7E-7	3.3E-8
$Q^{\hat{\pi}}(\lambda)$	1.4E-8	1.4E-8	1.4E-8	1.5E-8
TREE	2.6E-8	4.9E-8	2.4E-6	1.7E-8
IH	7.5E-3	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	5.5E-1	9.7E-2		
WIS	2.4E-1	5.4E-2		
NAIVE	4.1E0	-		

Table 33. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.8E-4	2.5E-2	2.9E-2	7.5E-3
Q-Reg	4.1E-2	1.3E-3	4.6E-4	2.6E-2
MRDR	4.6E-2	2.0E-2	2.4E-2	3.3E-2
FQE	7.0E-6	7.0E-6	7.0E-6	7.0E-6
$R(\lambda)$	7.0E-6	7.0E-6	8.0E-6	7.0E-6
$Q^{\pi}(\lambda)$	7.0E-6	7.0E-6	7.0E-6	7.0E-6
TREE	5.0E-6	7.0E-6	9.0E-6	5.0E-6
IH	2.2E-3	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	5.0E-1	5.5E-2	
WIS	1.5E-1	1.5E-2	
NAIVE	4.0E0	-	

Table 34. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.7E-4	6.5E-3	6.2E-3	4.8E-3
Q-Reg	4.8E-2	6.2E-4	3.7E-4	3.6E-3
MRDR	2.9E-2	4.6E-3	2.9E-2	3.0E-2
FQE	4.4E-5	4.4E-5	4.4E-5	4.4E-5
$R(\lambda)$	4.4E-5	4.4E-5	4.3E-5	4.4E-5
$Q^{\hat{\pi}}(\lambda)$	4.4E-5	4.4E-5	4.4E-5	4.4E-5
TREE	4.0E-5	4.4E-5	4.3E-5	4.1E-5
IH	1.8E-3	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	3.6E-1	3.7E-2		
WIS	5.2E-2	1.2E-2		
NAIVE	4.0E0	-		

Table 35. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.3E-1	7.9E-1	6.5E-1	6.1E-1
Q-Reg	7.4E-1	1.0E0	1.8E0	7.7E-1
MRDR	6.4E-1	1.0E0	8.6E-1	6.8E-1
FQE	5.6E-1	5.8E-1	5.7E-1	5.6E-1
$R(\lambda)$	5.5E-1	6.0E-1	5.7E-1	5.5E-1
$Q^{\hat{\pi}}(\lambda)$	5.5E-1	8.8E-1	5.4E-1	5.5E-1
TREE	5.5E-1	5.9E-1	5.7E-1	5.5E-1
IH	1.3E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.1E0	8.0E-1	
WIS	1.5E0	1.2E0	
NAIVE	3.5E0	-	

Table 36. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.5E-1	7.8E-1	6.9E-1	7.5E-1
Q-Reg	6.6E-1	5.1E-1	5.0E-1	6.8E-1
MRDR	6.5E-1	5.7E-1	4.7E-1	7.8E-1
FQE	6.0E-1	8.0E-1	6.4E-1	6.0E-1
$R(\lambda)$	6.1E-1	7.3E-1	6.4E-1	6.1E-1
$Q^{\hat{\pi}}(\lambda)$	6.2E-1	7.1E-1	6.5E-1	6.2E-1
TREE	6.1E-1	7.4E-1	6.5E-1	6.1E-1
IH	6.8E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	7.9E-1	6.7E-1		
WIS	1.7E0	8.6E-1		
NAIVE	4.2E0	-		

Table 37. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.5E-1	4.7E-1	3.6E-1	2.5E-1
Q-Reg	3.3E-1	4.5E-1	3.0E-1	3.3E-1
MRDR	4.0E-1	2.4E-1	2.7E-1	4.6E-1
FQE	2.2E-1	2.7E-1	2.6E-1	2.2E-1
$R(\lambda)$	2.2E-1	2.7E-1	2.9E-1	2.3E-1
$Q^{\pi}(\lambda)$	2.7E-1	2.8E-1	2.6E-1	2.8E-1
TREE	2.2E-1	2.8E-1	2.9E-1	2.3E-1
IH	1.6E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	7.7E-1	2.7E-1		
WIS	9.9E-1	3.9E-1		
NAIVE	3.7E0	-		

Table 38. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.1E-1	3.1E-1	2.0E-1	2.0E-1
Q-Reg	7.7E-1	2.5E-1	2.4E-1	3.1E-1
MRDR	6.1E-1	2.2E-1	1.6E-1	3.0E-1
FQE	2.0E-1	2.3E-1	2.5E-1	2.0E-1
$R(\lambda)$	2.2E-1	2.1E-1	2.4E-1	2.2E-1
$Q^{\hat{\pi}}(\lambda)$	2.0E-1	1.7E-1	2.0E-1	2.0E-1
TREE	2.2E-1	2.1E-1	2.4E-1	2.2E-1
IH	1.5E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	9.0E0	5.8E-1	
WIS	6.6E-1	1.7E-1	
NAIVE	3.8E0	-	

Table 39. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.5E-2	5.3E-1	2.7E-1	9.8E-2
Q-REG	8.0E-1	1.4E-1	1.3E-1	6.7E-1
MRDR	3.8E-1	1.4E-1	1.4E-1	5.6E-1
FQE	2.8E-2	3.5E-1	1.4E-1	2.9E-2
$R(\lambda)$	6.3E-2	2.7E-1	1.3E-1	6.3E-2
$Q^{\hat{\pi}}(\lambda)$	1.0E-1	3.5E-1	1.5E-1	1.0E-1
TREE	5.6E-2	2.7E-1	1.3E-1	5.6E-2
IH	2.7E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	4.0E0	7.0E-1		
WIS	5.7E-1	1.4E-1		
NAIVE	3.8E0	-		

Table 40. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.8E-2	1.4E-1	7.0E-2	3.6E-2
Q-Reg	3.3E-1	6.8E-2	6.8E-2	2.6E-1
MRDR	2.4E-1	5.3E-2	6.0E-2	2.3E-1
FQE	1.7E-2	2.3E-1	7.1E-2	1.8E-2
$R(\lambda)$	2.6E-2	1.7E-1	6.8E-2	2.6E-2
$Q^{\hat{\pi}}(\lambda)$	4.0E-2	2.3E-1	7.6E-2	4.1E-2
TRÈE	2.4E-2	1.7E-1	6.8E-2	2.4E-2
IH	1.1E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	2.9E0	2.9E-1		
WIS	2.9E-1	5.4E-2		
NAIVE	3.9E0	-		

Table 41. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	9.2E-3	1.4E-1	1.6E-1	2.4E-2	
Q-Reg	8.6E-2	1.3E-1	1.1E-1	1.2E-1	
MRDR	1.0E-1	1.1E-1	1.6E-1	1.7E-1	
FQE	8.3E-3	7.1E-2	6.2E-2	8.3E-3	
$R(\lambda)$	1.2E-2	7.2E-2	6.4E-2	1.2E-2	
$Q^{\hat{\pi}}(\lambda)$	1.3E-2	7.7E-2	6.7E-2	1.3E-2	
TREE	1.1E-2	7.1E-2	6.4E-2	1.1E-2	
IH	1.5E-2	-	-	-	

	IPS					
	STANDARD	STANDARD PER-DECISION				
IS	3.1E-1	9.1E-2				
WIS	2.1E-1 7.5E-					
NAIVE	4.1E0	-				

Table 42. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.6E-3	1.0E-1	8.2E-2	3.5E-2
Q-REG	2.9E-2	2.5E-2	2.3E-2	2.8E-2
MRDR	1.8E-2	2.0E-2	2.3E-2	2.0E-2
FQE	8.4E-3	2.7E-2	2.3E-2	1.1E-2
$R(\lambda)$	8.0E-3	2.6E-2	2.3E-2	1.1E-2
$Q^{\hat{\pi}}(\lambda)$	1.0E-2	2.8E-2	2.4E-2	1.0E-2
TREE	7.7E-3	2.6E-2	2.3E-2	1.1E-2
IH	1.3E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.9E-1	2.7E-2	
WIS	8.4E-2	2.3E-2	
NAIVE	4.0E0	-	

Table 43. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.2E-1	4.5E-1	1.2E0	6.0E-1
Q-Reg	3.9E-1	6.9E-1	2.0E0	4.5E-1
MRDR	4.4E-1	9.3E-1	2.1E0	3.9E-1
FQE	2.7E-1	2.8E-1	2.5E-1	2.7E-1
$R(\lambda)$	2.8E-1	2.9E-1	2.5E-1	2.8E-1
$Q^{\hat{\pi}}(\lambda)$	3.9E-1	3.3E-1	3.6E-1	3.9E-1
TREE	2.8E-1	2.8E-1	2.5E-1	2.8E-1
IH	6.2E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.4E0	4.7E-1	
WIS	2.7E0	1.4E0	
NAIVE	4.0E0	-	

Table 44. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	8.9E-1	8.7E0	7.6E-1	8.1E-1	
Q-Reg	3.2E1	1.2E1	5.8E-1	3.2E1	
MRDR	1.8E1	4.5E1	1.1E0	1.8E1	
FQE	8.1E-1	1.3E0	7.4E-1	8.1E-1	
$R(\lambda)$	7.2E-1	1.8E0	8.3E-1	6.9E-1	
$Q^{\hat{\pi}}(\lambda)$	1.6E0	1.9E0	1.0E0	1.6E0	
TREE	7.4E-1	1.8E0	8.3E-1	7.2E-1	
IH	1.1E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	2.8E2	3.0E1		
WIS	2.3E0	6.8E-1		
NAIVE	3.9E0	-		

Table 45. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.2E-1	2.2E0	9.8E-1	5.6E-1
Q-Reg	6.5E0	1.4E1	4.2E-1	2.7E0
MRDR	3.8E0	2.2E1	4.6E-1	3.9E0
FQE	4.9E-1	5.8E-1	4.1E-1	4.9E-1
$R(\lambda)$	3.8E-1	4.1E-1	3.7E-1	4.0E-1
$Q^{\pi}(\lambda)$	5.3E-1	8.5E-1	4.7E-1	5.1E-1
TREE	3.9E-1	4.2E-1	3.7E-1	4.0E-1
IH	5.2E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	7.4E1	1.1E1		
WIS	1.6E0	5.7E-1		
NAIVE	3.5E0	-		

Table 46. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.0E-1	2.0E0	6.1E-1	2.9E-1
Q-Reg	1.0E0	6.4E-1	5.6E-1	1.0E0
MRDR	8.2E-1	7.7E-1	7.7E-1	7.8E-1
FQE	2.0E-1	9.7E-1	4.2E-1	2.0E-1
$R(\lambda)$	2.5E-1	1.0E0	4.7E-1	2.4E-1
$Q^{\pi}(\lambda)$	4.4E-1	8.4E-1	4.5E-1	4.4E-1
TREE	2.4E-1	1.0E0	4.7E-1	2.4E-1
IH	1.4E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.3E0	9.3E-1	
WIS	2.0E0	7.8E-1	
NAIVE	4.2E0	-	

Table 47. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.1E-1	8.8E-1	5.4E-1	1.0E-1
Q-Reg	1.4E0	6.8E-1	3.5E-1	9.4E-1
MRDR	6.5E-1	4.3E-1	2.3E-1	1.7E0
FQE	8.8E-2	9.1E-1	4.9E-1	8.8E-2
$R(\lambda)$	9.8E-2	8.1E-1	4.7E-1	1.0E-1
$Q^{\hat{\pi}}(\lambda)$	7.2E-2	9.6E-1	5.2E-1	5.3E-2
TREE	9.9E-2	8.1E-1	4.7E-1	1.1E-1
IH	1.8E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	5.3E0	1.1E0	
WIS	8.2E-1	5.1E-1	
NAIVE	4.0E0	-	

Table 48. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.5E-2	2.8E-1	2.7E-1	1.2E-1
Q-Reg	2.8E-1	1.8E-1	2.0E-1	8.1E-2
MRDR	2.6E-1	1.0E-1	1.4E-1	3.0E-1
FQE	6.6E-2	1.9E-1	2.2E-1	6.8E-2
$R(\lambda)$	1.2E-1	2.0E-1	2.1E-1	1.1E-1
$Q^{\hat{\pi}}(\lambda)$	1.1E-1	1.8E-1	2.1E-1	1.1E-1
TREE	1.1E-1	2.0E-1	2.1E-1	1.1E-1
IH	5.7E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	7.2E-1	2.3E-1	
WIS	5.8E-1	2.1E-1	
NAIVE	4.3E0	-	

Table 49. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.4E-2	1.3E-1	1.2E-1	4.3E-2
Q-Reg	5.5E-2	1.2E-1	1.1E-1	4.3E-2
MRDR	4.3E-2	7.5E-2	1.1E-1	1.1E-1
FQE	1.3E-2	8.7E-2	9.8E-2	1.3E-2
$R(\lambda)$	2.2E-2	9.2E-2	9.9E-2	2.5E-2
$Q^{\hat{\pi}}(\lambda)$	1.8E-2	8.9E-2	1.0E-1	1.6E-2
TREE	2.2E-2	9.2E-2	9.8E-2	2.4E-2
IH	1.5E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	3.2E-1	6.9E-2	
WIS	1.4E-1	6.7E-2	
NAIVE	4.0E0	-	

Table 50. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.1E-2	8.9E-2	9.7E-2	2.4E-2	
Q-REG	6.7E-2	6.6E-2	6.1E-2	3.4E-2	
MRDR	5.4E-2	7.3E-2	1.0E-1	7.3E-2	
FQE	9.2E-3	6.2E-2	5.8E-2	1.0E-2	
$R(\lambda)$	1.7E-2	6.4E-2	5.9E-2	1.5E-2	
$Q^{\hat{\pi}}(\lambda)$	3.0E-2	6.4E-2	5.8E-2	1.8E-2	
TREE	1.6E-2	6.4E-2	5.9E-2	1.5E-2	
IH	1.8E-2	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	2.6E-1	8.4E-2	
WIS	2.0E-1	4.3E-2	
NAIVE	3.9E0	-	

Table 51. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic environment. Stochastic rewards. Dense rewards.

DM Hybrid DIRECT DR WDR MAGIC 1.0E0 3.4E0 2.8E0 1.5E0 AMQ-REG 8.4E1 9.4E0 1.6E0 8.4E1 MRDR 4.1E1 4.3E1 2.7E0 4.1E1 FQE 9.0E-1 2.5E0 9.8E-1 9.0E-1 $R(\lambda)$ 8.9E-1 1.2E0 8.2E-1 8.9E-1 $Q^{\hat{\pi}}(\hat{\lambda})$ 1.1E0 1.1E0 9.6E-1 2.1E0 TREE 9.0E-1 1.2E0 8.3E-1 9.0E-1 ΙH 1.7E0

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	1.6E1	5.9E1			
WIS	3.9E0	2.5E0			
NAIVE	6.3E0	-			

Table 52. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	2.3E0	7.8E0	3.3E0	2.2E0
Q-Reg	5.9E0	2.0E0	2.3E0	5.0E0
MRDR	5.3E0	1.5E0	2.4E0	4.6E0
FQE	2.3E0	4.0E0	2.7E0	2.3E0
$R(\lambda)$	2.0E0	2.4E0	2.2E0	2.0E0
$Q^{\pi}(\lambda)$	2.9E0	2.6E0	1.9E0	2.9E0
TREE	2.0E0	2.5E0	2.3E0	2.0E0
IH	2.3E0	-	-	-

		IPS		
	STANDARD	PER-DECISION		
IS	1.1E1	1.0E1		
WIS	5.8E0	4.9E0		
NAIVE	5.8E0	-		

Table 53. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	7.3E-1	8.9E-1	1.2E0	6.8E-1	
Q-Reg	3.4E0	1.8E0	2.5E0	3.4E0	
MRDR	2.1E0	9.4E-1	1.6E0	3.1E0	
FQE	7.1E-1	4.3E0	1.8E0	7.1E-1	
$R(\lambda)$	9.7E-1	3.1E0	1.9E0	9.7E-1	
$Q^{\hat{\pi}}(\hat{\lambda})$	2.6E0	5.4E0	2.4E0	2.6E0	
TREE	8.8E-1	3.1E0	1.8E0	8.8E-1	
IH	1.6E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.5E0	3.4E0		
WIS	3.5E0	2.6E0		
NAIVE	4.4E0	-		

Table 54. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.9E-1	2.6E0	1.9E0	4.8E-1
Q-Reg	3.2E0	3.2E0	1.6E0	3.2E0
MRDR	2.9E0	4.3E0	2.5E0	2.0E0
FQE	4.7E-1	8.3E-1	7.4E-1	4.8E-1
$R(\lambda)$	5.3E-1	8.8E-1	7.8E-1	5.4E-1
$Q^{\hat{\pi}}(\lambda)$	3.1E-1	8.9E-1	6.2E-1	3.4E-1
TREE	5.2E-1	8.6E-1	7.8E-1	5.3E-1
IH	3.8E-1	-	-	-

	IPS		
	STANDARD	Per-Decision	
IS	5.8E1	4.0E0	
WIS	4.6E0	1.4E0	
NAIVE	3.5E0	-	

Table 55. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.3E-1	3.0E0	1.9E0	1.2E0	
Q-Reg	2.5E0	1.1E0	7.3E-1	2.3E0	
MRDR	2.5E0	1.8E0	9.4E-1	3.1E0	
FQE	3.7E-1	1.3E0	9.3E-1	3.8E-1	
$R(\lambda)$	3.8E-1	1.4E0	9.4E-1	3.7E-1	
$Q^{\hat{\pi}}(\lambda)$	4.4E-1	1.2E0	8.4E-1	2.8E-1	
TREE	3.7E-1	1.4E0	9.5E-1	3.7E-1	
IH	3.5E-1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.6E1	3.3E0		
WIS	1.8E0	1.0E0		
NAIVE	3.9E0	-		

Table 56. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.4E-1	1.2E0	9.8E-1	1.9E-1
Q-Reg	5.3E-1	2.6E-1	2.6E-1	3.5E-1
MRDR	5.4E-1	3.7E-1	3.3E-1	3.4E-1
FQE	1.6E-1	7.0E-1	4.6E-1	1.5E-1
$R(\lambda)$	1.5E-1	5.0E-1	3.6E-1	1.6E-1
$Q^{\pi}(\lambda)$	2.0E-1	7.5E-1	4.8E-1	2.2E-1
TREE	1.5E-1	5.0E-1	3.6E-1	1.6E-1
IH	1.7E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.5E0	6.6E-1		
WIS	5.2E-1	3.4E-1		
NAIVE	4.6E0	-		

Table 57. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.3E-2	5.1E-1	4.0E-1	1.3E-1
Q-Reg	3.3E-1	1.8E-1	1.7E-1	1.8E-1
MRDR	1.8E-1	7.8E-2	6.6E-2	2.2E-1
FQE	4.1E-2	2.0E-1	1.9E-1	4.3E-2
$R(\lambda)$	1.1E-1	2.1E-1	2.0E-1	9.9E-2
$Q^{\hat{\pi}}(\lambda)$	1.1E-1	2.0E-1	1.9E-1	8.1E-2
TREE	1.0E-1	2.1E-1	2.0E-1	9.5E-2
IH	3.3E-2	-	-	-

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	3.7E-1	2.4E-1			
WIS	1.5E-1	2.0E-1			
NAIVE	3.9E0	-			

Table 58. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM	Hybrid		
	DIRECT	DR	WDR	MAGIC
AM	2.4E-2	3.0E-1	3.4E-1	8.5E-2
Q-Reg	4.9E-1	2.4E-1	2.6E-1	2.3E-1
MRDR	3.5E-1	3.4E-1	2.9E-1	4.1E-1
FQE	2.1E-2	2.1E-1	2.2E-1	2.3E-2
$R(\lambda)$	3.2E-2	2.0E-1	2.2E-1	3.0E-2
$Q^{\pi}(\lambda)$	4.5E-2	2.1E-1	2.3E-1	3.7E-2
TREE	3.1E-2	2.0E-1	2.2E-1	3.0E-2
IH	2.3E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.5E0	4.5E-1		
WIS	9.2E-1	3.2E-1		
NAIVE	3.9E0	-		

Table 59. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.2E0	1.2E0	1.6E0	1.2E0	
Q-Reg	2.5E0	2.5E0	2.3E0	2.5E0	
MRDR	2.3E0	4.2E0	1.7E0	2.1E0	
FQE	1.3E0	1.3E0	1.3E0	1.3E0	
$R(\lambda)$	1.3E0	1.3E0	1.3E0	1.3E0	
$Q^{\hat{\pi}}(\lambda)$	1.3E0	1.3E0	1.3E0	1.3E0	
TREE	1.2E0	1.3E0	1.3E0	1.3E0	
IH	1.1E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	3.8E0	3.8E0		
WIS	3.9E0	3.9E0		
NAIVE	3.9E0	-		

Table 60. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.4E-1	1.3E0	8.8E-1	7.5E-1
Q-Reg	2.5E0	1.7E0	1.6E0	1.0E0
MRDR	2.7E0	1.8E0	5.7E-1	1.3E0
FQE	6.8E-1	6.8E-1	6.8E-1	6.8E-1
$R(\lambda)$	6.8E-1	6.8E-1	7.0E-1	6.8E-1
$Q^{\pi}(\lambda)$	6.8E-1	6.8E-1	6.8E-1	6.8E-1
TREE	6.8E-1	6.8E-1	7.2E-1	6.8E-1
IH	7.4E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	3.9E0	3.9E0		
WIS	2.2E0	2.2E0		
NAIVE	3.8E0	-		

Table 61. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.8E-1	1.1E0	7.2E-1	4.8E-1
Q-Reg	1.5E1	2.2E1	2.2E0	1.5E1
MRDR	1.4E1	4.2E1	3.0E0	1.5E1
FQE	4.3E-1	4.3E-1	4.3E-1	4.3E-1
$R(\lambda)$	4.2E-1	4.3E-1	4.3E-1	4.3E-1
$Q^{\hat{\pi}}(\lambda)$	4.3E-1	4.3E-1	4.3E-1	4.3E-1
TRÈE	4.2E-1	4.2E-1	4.3E-1	4.2E-1
IH	4.0E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	3.4E1	3.4E1	
WIS	1.1E0	1.1E0	
NAIVE	3.5E0	-	

Table 62. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.7E-1	4.8E-1	6.0E-1	2.9E-1	
Q-REG	7.2E-1	6.4E-1	4.3E-1	4.8E-1	
MRDR	9.3E-1	5.6E-1	5.9E-1	5.1E-1	
FQE	2.7E-1	2.7E-1	2.7E-1	2.7E-1	
$R(\lambda)$	2.7E-1	2.7E-1	2.7E-1	2.7E-1	
$Q^{\hat{\pi}}(\lambda)$	2.7E-1	2.7E-1	2.7E-1	2.7E-1	
TREE	2.7E-1	2.7E-1	2.7E-1	2.7E-1	
IH	4.1E-1	-	-	-	

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	6.7E-1	6.7E-1			
WIS	1.1E0	1.1E0			
NAIVE	3.9E0	-			

Table 63. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	5.9E-3	1.4E-1	1.6E-1	3.3E-2	
Q-Reg	5.8E0	7.9E-1	1.1E-1	6.7E0	
MRDR	4.2E0	2.2E0	7.5E-1	4.1E0	
FQE	9.0E-6	9.0E-6	9.0E-6	9.0E-6	
$R(\lambda)$	7.0E-6	7.0E-6	1.0E-5	7.0E-6	
$Q^{\hat{\pi}}(\lambda)$	9.0E-6	9.0E-6	9.0E-6	9.0E-6	
TRÈE	1.5E-5	7.0E-6	2.3E-5	1.4E-5	
IH	2.4E-1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	4.8E0	4.8E0		
WIS	1.0E-1	1.0E-1		
NAIVE	4.0E0	-		

Table 64. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.6E-3	4.3E-2	6.0E-2	4.0E-2
Q-Reg	8.3E-1	2.5E0	1.8E0	6.8E0
MRDR	7.9E-1	5.9E-1	6.3E-1	8.0E-1
FQE	3.0E-5	3.0E-5	3.0E-5	3.0E-5
$R(\lambda)$	3.0E-5	3.0E-5	2.1E-5	3.0E-5
$Q^{\hat{\pi}}(\lambda)$	3.0E-5	3.0E-5	3.0E-5	3.0E-5
TREE	4.5E-5	3.2E-5	2.1E-5	4.3E-5
IH	1.0E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	7.8E-1	7.8E-1	
WIS	3.6E-1	3.6E-1	
NAIVE	4.1E0	-	

Table 65. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.1E-3	5.1E-2	4.0E-2	2.0E-2
Q-Reg	4.8E-1	3.9E-2	1.2E-2	1.9E-1
MRDR	3.7E-1	2.7E-1	2.7E-1	3.6E-1
FQE	9.3E-7	9.3E-7	9.3E-7	9.3E-7
$R(\lambda)$	8.3E-7	8.4E-7	3.0E-6	8.8E-7
$Q^{\hat{\pi}}(\lambda)$	9.7E-7	9.3E-7	9.3E-7	9.5E-7
TREE	1.8E-7	9.3E-7	8.5E-6	1.5E-7
IH	5.6E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	4.9E-1	4.9E-1	
WIS	2.8E-1	2.8E-1	
NAIVE	3.9E0	-	

Table 66. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.5E-3	1.1E-2	1.1E-2	6.8E-3
Q-Reg	5.1E-1	1.8E-2	4.3E-3	3.6E-2
MRDR	5.3E-1	1.8E-1	4.7E-1	7.5E-1
FQE	1.5E-5	1.5E-5	1.5E-5	1.5E-5
$R(\lambda)$	1.6E-5	1.6E-5	1.5E-5	1.6E-5
$Q^{\hat{\pi}}(\lambda)$	1.5E-5	1.5E-5	1.5E-5	1.5E-5
TRÈE	2.7E-5	1.7E-5	1.6E-5	2.6E-5
IH	2.1E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	4.5E-1	4.5E-1		
WIS	1.2E-1	1.2E-1		
NAIVE	4.1E0	-		

Table 67. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	5.5E0	1.1E1	5.6E0	5.2E0	
Q-Reg	1.2E1	1.1E1	6.4E0	1.2E1	
MRDR	1.1E1	6.5E0	3.6E0	1.1E1	
FQE	5.7E0	5.0E0	3.6E0	5.7E0	
$R(\lambda)$	5.4E0	5.6E0	3.4E0	5.4E0	
$Q^{\hat{\pi}}(\lambda)$	3.9E0	1.9E1	4.0E0	3.9E0	
TREE	5.5E0	5.4E0	3.4E0	5.5E0	
IH	1.5E1	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	1.4E1	4.0E1	
WIS	1.3E1	1.0E1	
NAIVE	8.6E0	-	

Table 68. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.4E0	4.2E0	5.2E0	3.5E0	
Q-Reg	5.8E0	4.8E0	1.2E1	8.6E0	
MRDR	6.9E0	3.6E0	5.3E0	5.4E0	
FQE	3.4E0	4.5E0	5.2E0	3.4E0	
$R(\lambda)$	3.6E0	3.9E0	4.7E0	3.6E0	
$Q^{\pi}(\lambda)$	5.5E0	2.8E0	4.5E0	5.5E0	
TREE	3.6E0	4.0E0	4.8E0	3.6E0	
IH	7.3E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	3.6E0	5.4E0		
WIS	7.3E0	7.1E0		
NAIVE	4.0E0	-		

Table 69. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM	Hybrid		
	DIRECT	DR	WDR	MAGIC
AM	2.7E0	5.3E0	6.1E0	2.9E0
Q-Reg	3.5E0	3.9E0	3.9E0	3.0E0
MRDR	1.3E0	1.8E0	2.2E0	2.2E0
FQE	2.6E0	2.8E0	4.2E0	2.6E0
$R(\lambda)$	2.6E0	2.7E0	3.6E0	2.6E0
$Q^{\pi}(\lambda)$	3.8E0	2.6E0	2.8E0	3.7E0
TREE	2.6E0	2.7E0	3.6E0	2.6E0
IH	2.3E0	-	-	-

	IPS			
	STANDARD PER-DECISIO			
IS	2.0E0	1.9E0		
WIS	4.7E0	4.9E0		
NAIVE	4.7E0	-		

Table 70. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.3E0	9.0E0	7.9E0	3.1E0	
Q-Reg	5.4E1	7.6E0	5.0E0	5.2E1	
MRDR	2.8E1	1.3E1	2.3E0	2.8E1	
FQE	2.8E0	6.4E0	3.8E0	2.8E0	
$R(\lambda)$	3.5E0	5.1E0	3.8E0	3.5E0	
$Q^{\hat{\pi}}(\lambda)$	4.8E0	6.8E0	3.9E0	4.8E0	
TREE	3.4E0	5.1E0	3.8E0	3.4E0	
IH	1.8E0	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	1.3E1	3.2E1	
WIS	5.8E0	3.8E0	
NAIVE	4.0E0	-	

Table 71. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.5E0	4.2E0	2.5E0	1.5E0	
Q-Reg	3.2E1	1.2E1	2.5E0	1.5E1	
MRDR	1.9E1	2.9E1	3.6E0	1.6E1	
FQE	1.5E0	1.9E0	2.4E0	1.5E0	
$R(\lambda)$	1.4E0	1.9E0	2.3E0	1.5E0	
$Q^{\hat{\pi}}(\lambda)$	2.1E0	1.8E0	2.4E0	2.1E0	
TRÈE	1.4E0	2.0E0	2.3E0	1.4E0	
IH	2.2E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	3.7E1	3.0E1		
WIS	3.1E0	3.9E0		
NAIVE	4.1E0	-		

Table 72. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.0E0	1.2E1	6.1E0	1.5E0	
Q-Reg	9.5E0	7.2E0	5.6E0	1.7E1	
MRDR	6.7E0	4.2E0	4.8E0	9.7E0	
FQE	1.1E0	8.0E0	5.2E0	1.1E0	
$R(\lambda)$	2.1E0	7.5E0	5.2E0	2.1E0	
$Q^{\pi}(\lambda)$	1.6E0	8.4E0	5.4E0	1.6E0	
TREE	2.0E0	7.4E0	5.1E0	2.0E0	
IH	1.3E0	-	-	-	

	IPS			
	STANDARD PER-DECISIO			
IS	1.6E1	8.6E0		
WIS	7.5E0	5.2E0		
NAIVE	3.8E0	-		

Table 73. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.5E-1	2.4E0	2.1E0	3.1E-1	
Q-Reg	1.6E0	1.6E0	1.2E0	1.2E0	
MRDR	1.8E0	1.4E0	1.0E0	1.2E0	
FQE	2.8E-1	1.1E0	1.1E0	2.9E-1	
$R(\lambda)$	4.2E-1	1.2E0	1.1E0	4.3E-1	
$Q^{\hat{\pi}}(\lambda)$	4.9E-1	1.1E0	1.1E0	4.9E-1	
TREE	4.0E-1	1.2E0	1.1E0	4.0E-1	
IH	3.0E-1	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	5.5E-1	1.5E0	
WIS	3.8E-1	9.3E-1	
NAIVE	3.9E0	-	

Table 74. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	5.1E-2	8.5E-1	7.2E-1	1.7E-1	
Q-REG	5.2E-1	7.1E-1	5.3E-1	4.8E-1	
MRDR	3.9E-1	5.3E-1	4.9E-1	8.3E-1	
FQE	6.8E-2	4.1E-1	4.2E-1	6.8E-2	
$R(\lambda)$	8.8E-2	4.3E-1	4.4E-1	8.9E-2	
$Q^{\hat{\pi}}(\lambda)$	7.8E-2	4.1E-1	4.1E-1	7.8E-2	
TREE	8.5E-2	4.3E-1	4.4E-1	8.6E-2	
IH	5.1E-2	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	8.7E-1	5.2E-1	
WIS	1.1E0	4.6E-1	
NAIVE	3.9E0	-	

Table 75. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.8E0	7.7E0	7.0E0	3.7E0	
Q-Reg	6.4E2	1.8E3	1.6E1	6.3E2	
MRDR	5.2E2	3.7E3	2.5E1	5.1E2	
FQE	3.5E0	2.9E0	3.3E0	3.4E0	
$R(\lambda)$	3.4E0	2.8E0	3.1E0	3.3E0	
$Q^{\hat{\pi}}(\lambda)$	5.1E0	2.9E0	2.8E0	4.5E0	
TREE	3.4E0	2.8E0	3.1E0	3.4E0	
IH	2.3E0	-	-	-	

	IPS		
	STANDARD	Per-Decision	
IS	1.2E3	1.2E3	
WIS	9.6E0	9.6E0	
NAIVE	6.0E0	-	

Table 76. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	2.3E0	8.7E0	6.2E0	2.3E0
Q-Reg	1.7E0	4.1E0	2.1E1	2.0E1
MRDR	3.3E0	5.6E0	6.7E1	5.1E1
FQE	1.8E0	3.9E0	1.4E0	1.8E0
$R(\lambda)$	1.5E0	1.4E0	1.5E0	1.5E0
$Q^{\hat{\pi}}(\lambda)$	3.6E0	4.2E0	2.6E0	3.3E0
TRÈE	1.5E0	1.4E0	1.5E0	1.5E0
IH	7.9E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	2.5E0	2.5E0	
WIS	5.7E0	5.7E0	
NAIVE	4.5E0	-	

Table 77. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM	Hybrid		
	DIRECT	DR	WDR	MAGIC
AM	3.0E0	2.6E1	1.1E1	3.5E0
Q-Reg	9.1E0	2.2E1	1.4E1	7.4E0
MRDR	1.5E1	2.4E0	2.2E0	2.9E0
FQE	2.6E0	2.1E1	5.3E0	2.6E0
$R(\lambda)$	3.5E0	1.2E1	4.1E0	3.5E0
$Q^{\pi}(\lambda)$	5.6E0	2.1E1	5.5E0	5.3E0
TREE	3.3E0	1.2E1	4.2E0	3.2E0
IH	5.8E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	8.9E0	8.9E0	
WIS	8.0E0	8.0E0	
NAIVE	3.9E0	-	

Table 78. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.2E-1	1.6E1	8.5E0	3.2E-1	
Q-Reg	2.2E1	4.3E0	2.3E0	2.3E1	
MRDR	1.7E1	9.4E0	4.6E0	1.9E1	
FQE	2.4E-1	5.3E0	2.8E0	2.4E-1	
$R(\lambda)$	9.6E-1	5.4E0	2.8E0	9.6E-1	
$Q^{\hat{\pi}}(\lambda)$	7.6E-1	5.5E0	2.4E0	8.8E-1	
TREE	7.7E-1	5.4E0	2.8E0	7.7E-1	
IH	2.6E-1	-	-	-	

	IPS			
	STANDARD PER-DECISIO			
IS	2.0E1	2.0E1		
WIS	3.0E0	3.0E0		
NAIVE	4.0E0	-		

Table 79. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.2E-1	2.1E0	1.5E0	6.0E-1	
Q-Reg	1.7E1	2.8E0	2.9E0	1.5E1	
MRDR	1.4E1	1.1E1	9.8E0	2.1E1	
FQE	3.6E-1	2.3E0	1.8E0	3.6E-1	
$R(\lambda)$	6.8E-1	2.1E0	1.8E0	6.8E-1	
$Q^{\pi}(\lambda)$	4.5E-1	2.5E0	1.9E0	4.8E-1	
TREE	6.5E-1	2.1E0	1.8E0	6.5E-1	
IH	3.0E-1	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	2.0E1	2.0E1	
WIS	5.1E0	5.1E0	
NAIVE	4.8E0	-	

Table 80. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.6E-1	1.9E0	2.3E0	4.9E-1	
Q-Reg	3.4E-1	7.5E-1	5.7E-1	2.7E-1	
MRDR	4.8E-1	5.3E-1	2.1E0	1.9E0	
FQE	1.4E-1	6.5E-1	5.6E-1	1.3E-1	
$R(\lambda)$	2.7E-1	7.1E-1	5.9E-1	2.8E-1	
$Q^{\hat{\pi}}(\lambda)$	2.5E-1	6.6E-1	5.5E-1	2.2E-1	
TREE	2.7E-1	7.1E-1	5.9E-1	2.8E-1	
IH	2.0E-1	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	3.3E-1	3.3E-1	
WIS	3.4E-1	3.4E-1	
NAIVE	3.9E0	-	

Table 81. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.1E-2	1.4E0	1.3E0	1.5E-1
Q-Reg	1.4E0	4.7E-1	3.7E-1	9.8E-1
MRDR	1.8E0	5.1E-1	9.4E-1	1.8E0
FQE	6.1E-2	3.2E-1	3.1E-1	6.4E-2
$R(\lambda)$	9.8E-2	3.3E-1	3.3E-1	1.0E-1
$Q^{\hat{\pi}}(\lambda)$	2.2E-1	3.3E-1	3.3E-1	1.9E-1
TREE	9.0E-2	3.3E-1	3.3E-1	9.4E-2
IH	1.2E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.2E0	1.2E0	
WIS	1.0E0	1.0E0	
NAIVE	4.1E0	-	

Table 82. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.1E-2	2.7E-1	3.1E-1	4.5E-2
Q-REG	2.6E-1	2.4E-1	2.3E-1	1.8E-1
MRDR	1.1E0	3.1E-1	3.0E-1	6.5E-1
FQE	2.5E-2	2.1E-1	2.0E-1	2.5E-2
$R(\lambda)$	4.0E-2	2.1E-1	2.0E-1	3.9E-2
$Q^{\hat{\pi}}(\lambda)$	4.9E-2	2.2E-1	2.1E-1	2.3E-2
TREE	4.0E-2	2.1E-1	2.0E-1	3.9E-2
IH	7.6E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	3.0E-1	3.0E-1	
WIS	3.4E-1	3.4E-1	
NAIVE	3.9E0	-	

Table 83. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

DM Hybrid DIRECT DR WDR MAGIC AM 1.6E1 8.6E1 5.5E1 1.6E1 Q-REG 3.7E0 6.2E1 3.3E1 3.7E0 MRDR 3.5E0 9.6E1 3.5E0 2.6E1 FQE 1.1E1 1.1E1 2.3E1 1.7E1 $R(\lambda)$ 9.5E0 9.4E0 1.1E1 9.5E0 1.8E1 $Q^{\pi}(\lambda)$ 1.1E1 1.1E1 1.2E1 TREE 9.7E0 9.6E0 1.2E1 9.7E0 ΙH 2.0E1

	IPS		
	STANDARD	PER-DECISION	
IS	6.6E1	1.1E1	
WIS	2.6E1	1.4E1	
NAIVE	1.2E1	-	

Table 84. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	1.6E1	7.1E1	2.4E1	1.4E1
Q-REG	8.0E2	1.8E3	8.5E1	7.9E2
MRDR	7.2E2	4.1E3	1.2E2	6.9E2
FQE	1.3E1	1.7E2	1.8E1	1.3E1
$R(\lambda)$	1.4E1	5.2E1	1.5E1	1.3E1
$Q^{\hat{\pi}}(\lambda)$	2.4E1	1.9E2	1.8E1	2.4E1
TREE	1.3E1	5.3E1	1.5E1	1.3E1
IH	2.0E1	-	-	-

		IPS		
	STANDARD	PER-DECISION		
IS	7.2E3	2.0E3		
WIS	4.2E1	3.2E1		
NAIVE	8.7E0	-		

Table 85. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	8.8E0	3.2E1	1.5E1	8.4E0
Q-Reg	4.0E1	5.6E1	1.8E1	4.0E1
MRDR	3.3E1	7.1E1	2.1E1	2.7E1
FQE	9.6E0	2.0E1	1.1E1	9.6E0
$R(\lambda)$	1.3E1	2.6E1	1.6E1	1.3E1
$Q^{\hat{\pi}}(\lambda)$	1.3E1	2.2E1	1.5E1	1.3E1
TREE	1.3E1	2.6E1	1.6E1	1.3E1
IH	1.5E1	-	-	-

		IPS		
	STANDARD	PER-DECISION		
IS	2.0E1	3.7E1		
WIS	1.6E1	1.2E1		
NAIVE	3.3E0	-		

Table 86. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	6.4E0	3.6E1	3.7E1	7.5E0
Q-Reg	6.4E1	2.3E1	1.0E1	6.6E1
MRDR	4.3E1	4.3E1	6.8E0	5.4E1
FQE	6.4E0	8.7E0	8.6E0	6.4E0
$R(\lambda)$	7.1E0	7.4E0	7.0E0	7.1E0
$Q^{\hat{\pi}}(\hat{\lambda})$	8.0E0	1.2E1	9.8E0	8.1E0
TREE	7.1E0	7.2E0	6.9E0	7.1E0
IH	5.1E0	-	-	-

		IPS		
	STANDARD	PER-DECISION		
IS	9.3E1	8.0E1		
WIS	1.4E1	9.1E0		
NAIVE	6.5E0	-		

Table 87. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	2.8E0	5.3E1	2.0E1	2.8E0
Q-Reg	4.4E1	1.2E1	1.4E1	3.9E1
MRDR	3.5E1	2.1E1	1.6E1	4.6E1
FQE	2.8E0	4.4E1	1.4E1	2.8E0
$R(\lambda)$	5.2E0	2.9E1	1.4E1	4.4E0
$Q^{\hat{\pi}}(\lambda)$	6.4E0	4.2E1	1.5E1	5.3E0
TREE	4.9E0	3.0E1	1.4E1	4.3E0
IH	2.6E0	-	-	-

		IPS	
	STANDARD	PER-DECISION	
IS	1.3E2	5.7E1	
WIS	3.1E1	1.5E1	
NAIVE	5.4E0	-	

Table 88. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	9.1E-1	2.6E0	4.8E0	1.5E0
Q-Reg	5.4E0	2.3E0	2.4E0	4.4E0
MRDR	4.2E0	3.4E0	5.5E0	3.8E0
FQE	1.1E0	2.8E0	3.0E0	1.1E0
$R(\lambda)$	8.0E-1	2.3E0	2.9E0	8.1E-1
$Q^{\hat{\pi}}(\lambda)$	9.1E-1	2.7E0	2.9E0	1.1E0
TREE	8.3E-1	2.2E0	2.9E0	8.3E-1
IH	1.1E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	2.2E0	5.0E0	
WIS	5.6E0	6.4E0	
NAIVE	3.8E0	-	

Table 89. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	6.4E-1	7.3E0	6.0E0	1.1E0
Q-Reg	4.0E0	1.9E0	1.8E0	2.6E0
MRDR	2.8E0	1.8E0	2.5E0	2.5E0
FQE	5.5E-1	2.0E0	1.5E0	5.4E-1
$R(\lambda)$	7.0E-1	1.9E0	1.5E0	6.2E-1
$Q^{\hat{\pi}}(\lambda)$	1.1E0	1.9E0	1.4E0	7.6E-1
TREE	6.7E-1	2.0E0	1.5E0	6.1E-1
IH	8.0E-1	-	-	-

		IPS		
	STANDARD	STANDARD PER-DECISION		
IS	1.3E1	4.1E0		
WIS	7.4E0	2.7E0		
NAIVE	4.0E0	-		

Table 90. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	1.0E-1	6.5E0	4.4E0	3.5E-1
Q-Reg	3.6E0	1.4E0	1.5E0	2.1E0
MRDR	3.0E0	1.8E0	3.0E0	2.4E0
FQE	1.2E-1	2.2E0	1.6E0	1.2E-1
$R(\lambda)$	2.0E-1	2.1E0	1.5E0	1.5E-1
$Q^{\hat{\pi}}(\lambda)$	7.9E-1	2.2E0	1.6E0	4.2E-1
TREE	1.8E-1	2.1E0	1.6E0	1.4E-1
IH	1.7E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.4E1	4.0E0	
WIS	1.2E1	2.9E0	
NAIVE	4.6E0	-	

Table 91. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.5E-1	8.7E-1	1.0E0	9.1E-1
Q-Reg	6.8E-1	9.0E-1	4.8E0	2.2E0
MRDR	7.2E-1	9.8E-1	6.5E0	6.1E0
FQE	8.5E-1	8.5E-1	8.5E-1	8.5E-1
$R(\lambda)$	8.5E-1	8.4E-1	1.4E0	1.3E0
$Q^{\hat{\pi}}(\lambda)$	8.5E-1	8.5E-1	8.5E-1	8.5E-1
TRÈE	8.5E-1	8.3E-1	1.5E0	1.4E0
IH	7.5E-2	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	6.5E-1	
WIS	2.6E0	1.8E0	
NAIVE	4.2E0	-	

Table 92. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.6E-1	7.4E-1	1.0E0	7.1E-1
Q-Reg	4.4E-1	5.0E-1	9.4E-1	8.7E-1
MRDR	5.3E-1	8.7E-1	2.3E0	2.1E0
FQE	6.5E-1	6.5E-1	6.5E-1	6.5E-1
$R(\lambda)$	6.6E-1	6.5E-1	9.4E-1	9.4E-1
$Q^{\hat{\pi}}(\lambda)$	6.5E-1	6.5E-1	6.5E-1	6.5E-1
TRÈE	6.7E-1	6.5E-1	1.0E0	1.0E0
IH	1.1E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	4.5E-1	
WIS	2.1E0	1.2E0	
NAIVE	3.9E0	-	

Table 93. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.7E-1	5.2E-1	4.3E-1	5.6E-1
Q-Reg	5.9E-1	5.0E-1	1.2E0	8.5E-1
MRDR	5.9E-1	8.3E-1	5.2E0	5.3E0
FQE	5.4E-1	5.4E-1	5.4E-1	5.4E-1
$R(\lambda)$	6.1E-1	6.0E-1	6.3E-1	7.2E-1
$Q^{\hat{\pi}}(\lambda)$	5.5E-1	5.4E-1	5.4E-1	5.4E-1
TRÈE	6.4E-1	6.1E-1	6.5E-1	7.4E-1
IH	1.7E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	6.1E-1	
WIS	1.7E0	7.4E-1	
NAIVE	4.1E0	-	

Table 94. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.7E-1	2.5E-1	2.7E-1	1.7E-1
Q-Reg	4.9E-1	9.2E0	1.6E1	4.8E-1
MRDR	4.7E-1	7.4E-1	2.2E0	2.2E0
FQE	1.7E-1	1.7E-1	1.7E-1	1.7E-1
$R(\lambda)$	3.4E-1	3.3E-1	4.3E-1	4.9E-1
$Q^{\hat{\pi}}(\lambda)$	1.7E-1	1.7E-1	1.7E-1	1.7E-1
TRÈE	4.0E-1	3.6E-1	4.6E-1	5.4E-1
IH	5.9E-3	-	-	-

	IPS			
	STANDARD PER-DECISIO			
IS	1.0E0	5.0E-1		
WIS	1.5E0	6.7E-1		
NAIVE	4.0E0	-		

Table 95. Graph, relative MSE. $T = 16, N = 128, \pi_b(a = 0) = 0.2, \pi_e(a = 0) = 0.8$. Dense rewards.

DM Hybrid DIRECT DR WDR MAGIC 4.0E-2 AM2.1E-2 1.5E-1 2.5E-1 Q-REG 4.0E-1 7.1E0 1.7E0 3.8E-1 MRDR 3.5E-1 5.6E-1 6.0E06.0E0 2.0E-2 2.0E-22.0E-2 FQE 2.0E-2 $R(\lambda)$ 2.2E-1 1.9E-1 1.0E-1 2.3E-1 $Q^{\pi}(\lambda)$ 2.0E-2 2.0E-2 2.0E-22.0E-2 TREE 3.0E-1 2.6E-1 1.3E-1 3.0E-1 ΙH 3.2E-3

	IPS			
	STANDARD PER-DECISION			
IS	9.0E-1	3.8E-1		
WIS	9.9E-1	2.7E-1		
NAIVE	4.0E0	-		

Table 96. Graph, relative MSE. $T = 16, N = 256, \pi_b(a = 0) = 0.2, \pi_e(a = 0) = 0.8$. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.3E-4	1.7E1	2.5E-1	5.5E-4
Q-REG	2.5E1	4.0E2	2.3E1	9.9E0
MRDR	1.9E1	1.2E3	1.9E1	2.2E1
FQE	9.9E-8	9.9E-8	9.9E-8	9.9E-8
$R(\lambda)$	9.7E-2	2.7E0	7.1E-2	1.2E-1
$Q^{\hat{\pi}}(\lambda)$	1.0E-7	1.1E-7	9.9E-8	1.0E-7
TRÈE	1.8E-1	1.2E1	9.7E-2	1.9E-1
IH	6.9E-4	-	-	-

	IPS				
	STANDARD PER-DECISION				
IS	2.2E1	2.8E1			
WIS	7.0E-1	2.2E-1			
NAIVE	4.0E0	-			

Table 97. Graph, relative MSE. $T = 16, N = 512, \pi_b(a = 0) = 0.2, \pi_e(a = 0) = 0.8$. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.2E-4	5.9E-1	1.8E-1	6.2E-4
Q-Reg	3.1E0	3.4E1	1.8E1	3.3E0
MRDR	1.6E0	1.4E2	4.1E1	1.0E1
FQE	3.6E-7	3.6E-7	3.6E-7	3.6E-7
$R(\lambda)$	1.0E-1	8.3E-1	8.7E-2	1.3E-1
$Q^{\pi}(\lambda)$	3.6E-7	3.6E-7	3.6E-7	3.6E-7
TREE	1.8E-1	2.2E0	1.3E-1	2.1E-1
IH	5.5E-4	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	4.2E0	2.6E0		
WIS	8.6E-1	2.4E-1		
NAIVE	4.0E0	-		

Table 98. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.5E-5	1.2E-1	5.6E-2	3.8E-4
Q-Reg	2.8E-1	1.4E0	3.0E-1	2.4E-1
MRDR	3.7E-1	6.0E-1	3.9E0	3.9E0
FQE	1.0E-6	1.0E-6	1.0E-6	1.0E-6
$R(\lambda)$	9.3E-2	7.5E-2	5.8E-2	9.3E-2
$Q^{\hat{\pi}}(\lambda)$	1.0E-6	1.0E-6	1.0E-6	1.0E-6
TREE	1.7E-1	1.2E-1	7.6E-2	1.7E-1
IH	8.5E-4	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	9.3E-1	2.7E-1		
WIS	6.9E-1	1.5E-1		
NAIVE	4.0E0	-		

Table 99. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.8E-1	2.0E0	1.2E0	7.8E-1
Q-Reg	9.5E-1	2.5E1	2.2E1	1.1E0
MRDR	1.1E0	6.5E1	2.5E1	2.6E0
FQE	7.6E-1	7.1E-1	9.5E-1	7.6E-1
$R(\lambda)$	7.7E-1	7.6E-1	1.1E0	8.7E-1
$Q^{\hat{\pi}}(\lambda)$	7.6E-1	7.2E-1	7.8E-1	7.6E-1
TREE	7.7E-1	7.0E-1	1.2E0	9.0E-1
IH	2.9E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	9.7E-1	4.8E0		
WIS	1.9E0	1.4E0		
NAIVE	3.7E0	-		

Table 100. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.8E-1	1.5E1	3.2E-1	6.2E-1
Q-Reg	2.1E0	2.0E1	2.0E0	2.8E0
MRDR	7.8E-1	3.9E0	4.1E1	4.2E1
FQE	6.9E-1	8.6E-1	6.8E-1	7.0E-1
$R(\lambda)$	7.4E-1	7.7E-1	9.8E-1	1.0E0
$Q^{\hat{\pi}}(\lambda)$	7.1E-1	1.8E0	6.2E-1	7.1E-1
TRÈE	7.6E-1	7.8E-1	1.0E0	9.9E-1
IH	1.2E-1	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	9.9E-1	6.0E-1		
WIS	1.5E0	9.9E-1		
NAIVE	3.8E0	-		

Table 101. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.7E-1	7.8E-1	1.3E0	5.5E-1
Q-Reg	6.9E-1	1.2E0	4.2E0	1.6E0
MRDR	1.5E0	2.5E0	8.6E0	8.1E0
FQE	5.8E-1	7.1E-1	5.7E-1	5.8E-1
$R(\lambda)$	6.6E-1	6.7E-1	1.0E0	1.0E0
$Q^{\hat{\pi}}(\lambda)$	6.7E-1	1.3E0	5.8E-1	6.7E-1
TREE	6.8E-1	6.8E-1	1.1E0	1.1E0
IH	6.5E-2	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	6.9E-1	
WIS	2.3E0	1.2E0	
NAIVE	4.1E0	-	

Table 102. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.6E-1	2.9E-1	3.3E-1	2.6E-1
Q-Reg	4.9E-1	3.6E-1	8.1E-1	7.4E-1
MRDR	4.8E-1	9.1E-1	3.3E0	3.4E0
FQE	2.6E-1	2.5E-1	2.2E-1	2.4E-1
$R(\lambda)$	4.4E-1	3.9E-1	2.9E-1	5.2E-1
$Q^{\pi}(\lambda)$	2.7E-1	2.8E-1	3.0E-1	2.7E-1
TREE	4.7E-1	4.1E-1	3.3E-1	5.8E-1
IH	5.2E-2	-	-	-

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	1.0E0	5.4E-1			
WIS	1.1E0	3.4E-1			
NAIVE	3.9E0	-			

Table 103. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic rewards. Dense rewards.

DMHybrid DIRECT DR WDR MAGIC6.4E-2 6.1E-1 2.9E-1 6.2E-2 AMQ-REG 4.3E-1 3.1E-1 2.3E0 3.6E-1 MRDR 3.7E-1 5.8E-1 2.7E02.7E0 1.5E-1 FQE 6.5E-2 8.3E-2 6.5E-2 7.1E-2 $R(\lambda)$ 1.9E-1 1.8E-1 1.9E-1 $Q^{\pi}(\lambda)$ 5.8E-2 5.1E-2 1.3E-1 5.8E-2 TREE 2.6E-12.2E-1 8.9E-2 2.6E-1 ΙH 3.0E-2

	IPS		
	STANDARD	PER-DECISION	
IS	9.8E-1	4.1E-1	
WIS	1.3E0	2.1E-1	
NAIVE	4.0E0	-	

Table 104. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.6E-3	7.0E-1	4.2E-1	8.8E-3
Q-Reg	3.5E-1	3.2E0	6.6E-1	3.7E-1
MRDR	3.6E-1	6.4E-1	2.1E0	1.8E0
FQE	8.4E-3	9.4E-2	3.4E-2	8.3E-3
$R(\lambda)$	1.4E-1	1.9E-1	9.0E-2	1.4E-1
$Q^{\hat{\pi}}(\lambda)$	5.7E-2	1.4E-1	6.5E-2	5.7E-2
TREE	2.3E-1	2.9E-1	1.4E-1	2.4E-1
IH	1.2E-2	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	9.7E-1	4.0E-1		
WIS	1.1E0	3.0E-1		
NAIVE	3.9E0	-		

Table 105. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.4E-3	3.2E-1	2.1E-1	7.4E-3
Q-Reg	4.1E-1	8.0E-1	2.0E-1	4.0E-1
MRDR	2.7E-1	5.0E-1	3.4E0	3.4E0
FQE	5.0E-3	4.3E-2	3.0E-2	4.8E-3
$R(\lambda)$	1.1E-1	1.7E-1	1.4E-1	1.3E-1
$Q^{\hat{\pi}}(\lambda)$	1.4E-2	6.2E-2	3.6E-2	1.3E-2
TREE	1.9E-1	2.5E-1	1.8E-1	2.1E-1
IH	6.2E-3	-	-	-

	IPS			
	STANDARD PER-DECISIO			
IS	1.0E0	4.1E-1		
WIS	1.0E0	3.0E-1		
NAIVE	4.0E0	-		

Table 106. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.3E-3	1.5E1	1.7E-1	2.4E-2
Q-Reg	1.5E0	9.5E0	1.7E1	1.4E0
MRDR	9.7E-1	5.0E1	3.0E1	8.0E0
FQE	8.6E-4	3.0E0	2.4E-2	9.8E-4
$R(\lambda)$	1.1E-1	5.8E-1	8.3E-2	9.1E-2
$Q^{\pi}(\lambda)$	7.7E-3	3.1E0	2.1E-2	7.7E-3
TREE	2.0E-1	1.7E0	1.1E-1	2.0E-1
IH	1.5E-3	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	1.4E2	1.8E0		
WIS	8.0E-1	1.8E-1		
NAIVE	4.0E0	-		

Table 107. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.5E-1	1.0E0	1.2E0	7.1E-1
Q-Reg	7.6E-1	8.3E-1	2.5E0	7.3E-1
MRDR	7.9E-1	1.6E0	3.0E0	1.1E0
FQE	7.3E-1	6.8E-1	6.6E-1	7.3E-1
$R(\lambda)$	7.2E-1	6.9E-1	1.1E0	8.7E-1
$Q^{\hat{\pi}}(\lambda)$	6.2E-1	8.5E-1	8.7E-1	6.2E-1
TREE	7.4E-1	6.9E-1	1.3E0	9.6E-1
IH	2.3E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	1.1E0		
WIS	2.7E0	1.9E0		
NAIVE	4.2E0	-		

Table 108. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.3E-1	3.9E-1	1.8E0	5.4E-1
Q-Reg	5.4E-1	4.1E-1	7.6E0	8.5E-1
MRDR	5.0E-1	8.0E-1	9.5E0	8.7E0
FQE	5.1E-1	4.7E-1	7.3E-1	5.1E-1
$R(\lambda)$	5.0E-1	5.1E-1	1.4E0	7.6E-1
$Q^{\hat{\pi}}(\lambda)$	3.9E-1	3.3E-1	3.1E-1	3.9E-1
TREE	5.4E-1	5.6E-1	1.7E0	1.0E0
IH	1.6E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	5.3E-1	
WIS	2.3E0	2.0E0	
NAIVE	4.1E0	-	

Table 109. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.3E-1	2.3E0	1.3E0	3.3E-1
Q-Reg	7.6E-1	5.0E-1	5.9E0	8.5E-1
MRDR	9.9E-1	7.8E-1	2.2E1	1.2E0
FQE	1.9E-1	2.7E-1	2.7E-1	1.9E-1
$R(\lambda)$	3.6E-1	3.6E-1	5.7E-1	6.0E-1
$Q^{\hat{\pi}}(\lambda)$	4.4E-1	4.6E-1	3.2E-1	4.3E-1
TREE	4.1E-1	4.5E-1	6.7E-1	6.8E-1
IH	4.8E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	6.9E-1	
WIS	2.4E0	8.1E-1	
NAIVE	4.2E0	-	

Table 110. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.3E-2	2.3E1	1.6E0	6.2E-2
Q-Reg	4.0E0	2.3E0	3.7E0	3.9E0
MRDR	2.1E0	2.5E1	1.1E1	1.1E1
FQE	5.1E-2	1.8E0	3.2E-1	5.1E-2
$R(\lambda)$	2.0E-1	1.7E-1	2.7E-1	3.2E-1
$Q^{\hat{\pi}}(\lambda)$	3.2E-1	1.9E0	7.7E-1	3.4E-1
TREE	2.5E-1	7.3E-1	2.5E-1	3.2E-1
IH	1.3E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.1E0	3.1E0		
WIS	1.9E0	3.1E-1		
NAIVE	4.2E0	-		

Table 111. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic environment. Dense rewards.

DM Hybrid DIRECT DR WDR MAGIC 2.1E-1 4.4E0 2.6E-2 1.5E0 AM4.0E-1 Q-REG 3.3E-1 3.3E1 1.4E1 MRDR 3.8E0 5.0E0 1.7E1 1.4E1 FQE 1.8E-2 1.4E-1 7.4E-21.7E-2 $R(\lambda)$ 2.7E-1 2.7E-12.0E-1 2.9E-1 $Q^{\dot{\pi}}(\dot{\lambda})$ 5.3E-1 1.6E-1 1.1E-1 1.3E-1 TREE 3.3E-1 2.8E-1 2.3E-1 3.5E-1 ΙH 2.2E-2

	IPS			
	STANDARD PER-DECISION			
IS	9.3E-1	2.6E-1		
WIS	8.3E-1 2.9E-1			
NAIVE	4.0E0	-		

Table 112. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.1E-2	3.9E0	1.1E0	1.0E-2
Q-Reg	2.7E-1	1.3E0	5.4E-2	1.8E-1
MRDR	4.3E-1	1.2E0	8.3E0	8.3E0
FQE	7.4E-3	5.5E-2	9.5E-2	1.4E-2
$R(\lambda)$	1.6E-1	1.3E-1	1.5E-1	1.6E-1
$Q^{\pi}(\lambda)$	1.1E-1	1.3E-1	1.3E-1	1.1E-1
TREE	2.0E-1	1.6E-1	1.9E-1	2.1E-1
IH	6.8E-3	-	_	_

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	2.8E-1	
WIS	1.6E0	2.9E-1	
NAIVE	4.0E0	-	

Table 113. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.1E-3	5.5E-1	1.1E0	8.0E-3
Q-Reg	3.9E-1	2.7E1	1.3E1	3.6E-1
MRDR	9.0E-1	1.1E0	1.3E1	1.2E1
FQE	4.9E-3	5.5E-2	1.3E-1	4.8E-3
$R(\lambda)$	1.3E-1	1.5E-1	1.7E-1	1.6E-1
$Q^{\hat{\pi}}(\lambda)$	9.4E-3	9.5E-2	1.4E-1	8.2E-3
TREE	2.2E-1	2.1E-1	2.1E-1	2.5E-1
IH	5.8E-3	-	-	-

	IPS			
	STANDARD PER-DECISI			
IS	9.9E-1	3.9E-1		
WIS	1.3E0	3.1E-1		
NAIVE	4.0E0	-		

Table 114. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.8E-3	9.5E-1	6.8E-1	3.8E-3	
Q-Reg	3.9E-1	2.0E-1	3.9E-1	5.5E-1	
MRDR	4.0E-1	4.6E-1	2.0E0	2.0E0	
FQE	2.7E-3	1.8E-1	8.4E-2	2.6E-2	
$R(\lambda)$	1.5E-1	1.5E-1	4.9E-2	1.5E-1	
$Q^{\pi}(\lambda)$	1.3E-2	1.9E-1	7.9E-2	2.0E-2	
TREE	2.2E-1	2.2E-1	4.6E-2	2.2E-1	
IH	2.0E-3	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	9.3E-1	3.9E-1	
WIS	6.4E-1	5.4E-2	
NAIVE	4.0E0	-	

Table 115. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

DM Hybrid DIRECT DR WDR MAGIC 1.4E0 1.9E0 8.4E-1 9.2E-1 AM3.0E0 Q-REG 1.3E0 1.8E0 2.0E0 MRDR 1.3E0 8.1E-1 2.7E02.3E0 FQE 9.3E-1 8.0E-1 1.0E0 9.3E-1 $R(\lambda)$ 8.6E-1 7.7E-1 2.0E0 1.3E0 $Q^{\hat{\pi}}(\hat{\lambda})$ 1.4E0 8.7E-1 1.4E0 4.7E0 TREE 8.6E-1 1.3E0 7.8E-1 2.1E0 ΙH 1.2E0

	IPS		
	STANDARD PER-DECIS		
IS	1.0E0	1.6E0	
WIS	3.4E0	3.2E0	
NAIVE	4.5E0	-	

Table 116. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	7.4E-1	1.5E0	2.3E0	1.1E0	
Q-Reg	8.6E-1	2.4E0	4.5E1	7.9E-1	
MRDR	9.6E-1	1.4E0	3.0E1	6.3E0	
FQE	7.7E-1	9.1E-1	1.0E0	7.7E-1	
$R(\lambda)$	8.8E-1	8.8E-1	1.4E0	9.4E-1	
$Q^{\pi}(\lambda)$	1.4E0	1.4E0	1.2E0	1.4E0	
TREE	9.2E-1	9.2E-1	1.5E0	9.6E-1	
IH	4.3E-1	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	9.9E-1	8.3E-1	
WIS	1.6E0	1.3E0	
NAIVE	3.7E0	-	

Table 117. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.2E-1	4.2E1	2.1E0	2.3E-1
Q-Reg	1.4E0	6.8E1	7.5E1	1.1E0
MRDR	1.6E0	1.3E2	7.4E1	4.8E0
FQE	1.9E-1	2.2E-1	4.7E-1	1.9E-1
$R(\lambda)$	4.0E-1	1.3E0	8.9E-1	8.6E-1
$Q^{\hat{\pi}}(\lambda)$	1.8E0	3.1E0	2.8E0	1.8E0
TREE	4.4E-1	3.1E0	9.1E-1	8.6E-1
IH	2.1E-1	-	-	-

	IPS		
	STANDARD PER-DECI		
IS	9.9E-1	3.2E0	
WIS	2.8E0	1.2E0	
NAIVE	4.1E0	-	

Table 118. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.6E-2	2.6E0	5.2E0	9.4E-2
Q-Reg	7.1E-1	4.0E0	6.3E0	2.4E0
MRDR	8.1E-1	1.4E0	5.6E0	5.3E0
FQE	5.5E-2	1.5E-1	3.3E-1	5.5E-2
$R(\lambda)$	1.4E-1	2.5E-1	7.2E-1	1.4E-1
$Q^{\pi}(\lambda)$	1.2E0	9.1E-1	1.1E0	1.2E0
TREE	2.0E-1	3.5E-1	8.5E-1	2.1E-1
IH	6.2E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.2E0	7.7E-1	
WIS	3.9E0	1.3E0	
NAIVE	4.1E0	-	

Table 119. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic environment. Stochastic rewards. Dense rewards.

DM HybridDIRECT DR WDR MAGIC 5.7E-2 5.4E-2 3.1E1 3.4E0 AMQ-REG 1.9E0 1.1E1 4.9E0 1.8E0 MRDR 4.0E1 2.3E1 1.6E1 1.8E1 FQE 5.1E-2 7.0E0 4.8E-1 5.1E-2 $R(\lambda)$ 2.3E-1 1.3E0 5.9E-1 4.9E-1 $Q^{\hat{\pi}}(\hat{\lambda})$ 2.0E-1 5.0E-1 1.7E0 2.1E-1 TREE 3.0E-1 2.9E0 6.6E-1 6.9E-1 IΗ 4.0E-2

	IPS				
	STANDARD PER-DECISION				
IS	1.7E0	3.0E0			
WIS	2.2E0	7.9E-1			
NAIVE	4.1E0	-			

Table 120. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.7E-2	5.3E1	3.4E0	4.9E-2	
Q-Reg	7.9E0	3.2E1	1.9E1	8.4E0	
MRDR	8.3E0	1.9E2	4.9E1	2.0E1	
FQE	3.8E-2	1.2E0	3.1E-1	4.0E-2	
$R(\lambda)$	2.4E-1	2.4E0	4.5E-1	2.4E-1	
$Q^{\hat{\pi}}(\lambda)$	1.5E-1	7.5E-1	3.1E-1	1.4E-1	
TREE	2.6E-1	3.2E0	5.3E-1	2.7E-1	
IH	3.0E-2	_	-	_	

	IPS		
	STANDARD PER-DECISION		
IS	7.9E1	1.0E1	
WIS	1.2E0	6.8E-1	
NAIVE	4.1E0	-	

Table 121. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	6.3E-3	2.6E0	1.4E0	1.1E-1	
Q-Reg	4.0E0	9.8E0	1.5E1	5.1E0	
MRDR	1.1E1	3.2E1	3.1E2	3.0E2	
FQE	7.9E-3	1.5E0	2.4E-1	7.9E-3	
$R(\lambda)$	1.5E-1	7.9E-1	2.8E-1	1.5E-1	
$Q^{\hat{\pi}}(\lambda)$	1.0E-1	3.1E0	2.9E-1	9.3E-2	
TREE	2.4E-1	1.1E0	3.5E-1	2.4E-1	
IH	8.1E-3	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	1.6E0	3.5E0	
WIS	1.5E0	4.0E-1	
NAIVE	4.0E0	-	

Table 122. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	5.3E-3	1.3E2	1.3E0	9.0E-3	
Q-Reg	8.1E0	2.6E0	1.4E1	4.0E0	
MRDR	3.7E0	1.2E2	3.2E1	8.8E0	
FQE	6.4E-3	1.3E1	2.4E-1	6.5E-3	
$R(\lambda)$	2.7E-1	3.0E0	2.5E-1	2.7E-1	
$Q^{\hat{\pi}}(\lambda)$	3.3E-2	2.0E1	2.4E-1	3.4E-2	
TREE	3.2E-1	4.9E0	2.9E-1	4.1E-1	
IH	5.8E-3	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	4.8E2	6.8E0	
WIS	1.3E0	3.2E-1	
NAIVE	4.0E0	-	

Table 123. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.0E0	9.1E-1	9.1E-1	1.0E0	
Q-Reg	9.7E-1	1.2E0	1.2E1	1.0E0	
MRDR	9.7E-1	1.3E0	1.2E1	1.0E0	
FQE	1.0E0	1.0E0	1.0E0	1.0E0	
$R(\lambda)$	1.0E0	9.7E-1	4.3E0	3.4E0	
$Q^{\hat{\pi}}(\lambda)$	1.0E0	1.0E0	1.0E0	1.0E0	
TRÈE	1.0E0	9.7E-1	4.3E0	3.4E0	
IH	1.2E0	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	9.7E-1	9.7E-1	
WIS	4.3E0	4.3E0	
NAIVE	5.3E0	-	

Table 124. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.0E0	1.1E0	6.2E0	1.0E0	
Q-REG	1.0E0	9.9E-1	9.7E0	9.9E-1	
MRDR	1.0E0	1.0E0	6.4E0	3.5E0	
FQE	1.0E0	1.0E0	1.0E0	1.0E0	
$R(\lambda)$	1.0E0	1.0E0	1.5E0	1.1E0	
$Q^{\hat{\pi}}(\lambda)$	1.0E0	1.0E0	1.0E0	1.0E0	
TREE	1.0E0	1.0E0	1.5E0	1.1E0	
IH	8.9E-1	_	-	_	

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	1.0E0	
WIS	1.5E0	1.5E0	
NAIVE	4.2E0	-	

Table 125. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.0E0	2.1E0	2.6E0	1.2E0
Q-Reg	9.9E-1	1.0E0	1.2E1	1.1E0
MRDR	9.8E-1	1.0E0	1.3E1	5.7E0
FQE	9.8E-1	9.8E-1	9.8E-1	9.8E-1
$R(\lambda)$	1.0E0	9.9E-1	1.1E0	1.5E0
$Q^{\pi}(\lambda)$	9.8E-1	9.8E-1	9.8E-1	9.8E-1
TREE	1.0E0	9.9E-1	1.1E0	1.5E0
IH	9.6E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	9.9E-1	9.9E-1	
WIS	1.1E0	1.1E0	
NAIVE	4.3E0	-	

Table 126. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.3E-1	2.2E0	9.4E0	1.2E0
Q-Reg	1.0E0	9.6E-1	7.3E0	1.0E0
MRDR	9.9E-1	9.7E-1	5.2E1	3.5E1
FQE	4.9E-1	4.9E-1	4.9E-1	4.9E-1
$R(\lambda)$	1.0E0	1.0E0	1.8E0	1.8E0
$Q^{\pi}(\lambda)$	4.9E-1	4.9E-1	4.9E-1	4.9E-1
TREE	1.0E0	1.0E0	1.8E0	1.8E0
IH	8.4E-1	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	1.0E0	1.0E0		
WIS	1.8E0	1.8E0		
NAIVE	4.0E0	-		

Table 127. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	9.3E-3	1.8E1	9.1E0	6.1E-2	
Q-Reg	1.0E0	6.3E0	1.4E1	1.3E1	
MRDR	1.1E0	1.2E0	1.4E1	1.4E1	
FQE	2.0E-6	2.0E-6	2.0E-6	2.0E-6	
$R(\lambda)$	1.0E0	1.1E0	3.5E0	2.3E0	
$Q^{\hat{\pi}}(\lambda)$	2.0E-6	2.0E-6	2.0E-6	2.0E-6	
TREE	1.0E0	1.1E0	3.5E0	2.3E0	
IH	7.3E-1	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	1.1E0	1.1E0	
WIS	3.5E0	3.5E0	
NAIVE	4.2E0	-	

Table 128. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.4E-3	2.4E1	7.3E0	3.2E-2	
Q-Reg	9.3E-1	1.2E1	1.2E1	8.4E-1	
MRDR	8.6E-1	4.6E0	1.5E2	1.5E2	
FQE	5.1E-5	5.1E-5	5.1E-5	5.1E-5	
$R(\lambda)$	1.0E0	9.3E-1	2.0E0	2.2E0	
$Q^{\pi}(\lambda)$	5.1E-5	5.1E-5	5.1E-5	5.0E-5	
TREE	1.0E0	9.3E-1	2.0E0	2.2E0	
IH	1.5E-1	-	-	-	

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	9.3E-1	9.3E-1		
WIS	2.0E0	2.0E0		
NAIVE	4.0E0	-		

Table 129. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.6E-3	1.8E1	4.6E0	2.6E-2	
Q-Reg	1.7E1	9.2E2	6.1E2	2.4E1	
MRDR	9.5E0	9.6E2	3.7E2	1.3E2	
FQE	5.0E-6	5.0E-6	5.0E-6	5.0E-6	
$R(\lambda)$	1.0E0	1.6E1	1.9E0	1.6E0	
$Q^{\hat{\pi}}(\lambda)$	5.0E-6	1.1E-5	5.0E-6	5.0E-6	
TREE	1.0E0	1.6E1	1.9E0	1.6E0	
IH	8.4E-2	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	1.6E1	1.6E1	
WIS	1.9E0 1.9 E		
NAIVE	3.9E0	-	

Table 130. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.8E-3	2.3E3	2.7E0	3.2E-3	
Q-Reg	1.2E3	2.2E3	2.5E1	1.3E3	
MRDR	1.8E4	2.5E4	1.4E2	9.6E2	
FQE	2.4E-5	2.4E-5	2.4E-5	2.4E-5	
$R(\lambda)$	1.0E0	1.1E3	2.5E0	1.0E0	
$Q^{\pi}(\lambda)$	2.4E-5	2.3E-5	2.4E-5	2.4E-5	
TREE	1.0E0	1.1E3	2.5E0	1.0E0	
IH	2.6E-2	-	-	-	

		IPS			
	STANDARD	STANDARD PER-DECISION			
IS	1.1E3	1.1E3			
WIS	2.5E0	2.5E0			
NAIVE	3.9E0	-			

Table 131. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.7E0	1.3E1	2.5E1	6.0E0
Q-Reg	7.7E2	6.4E3	1.9E4	7.7E2
MRDR	8.2E2	2.8E4	1.7E4	3.9E2
FQE	6.8E0	1.1E1	1.9E1	6.9E0
$R(\lambda)$	6.9E0	3.2E2	3.1E1	8.5E0
$Q^{\hat{\pi}}(\lambda)$	5.8E0	4.2E1	1.9E1	5.8E0
TRÈE	6.9E0	7.1E2	3.8E1	9.5E0
IH	9.9E1	-	-	-

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	5.6E3	8.5E2			
WIS	8.5E1	6.2E1			
NAIVE	1.9E1	-			

Table 132. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	3.8E0	1.2E1	2.7E1	3.4E0
Q-Reg	4.9E2	1.5E3	3.6E2	4.9E2
MRDR	3.0E2	2.4E3	1.8E2	3.4E2
FQE	2.3E0	3.5E2	1.7E1	2.3E0
$R(\lambda)$	2.4E0	1.3E1	2.0E1	2.4E0
$Q^{\hat{\pi}}(\lambda)$	6.7E0	4.2E2	1.5E1	6.7E0
TREE	2.4E0	1.5E1	2.2E1	2.4E0
IH	8.2E1	-	-	-

		IPS	
	STANDARD	PER-DECISION	
IS	1.4E0	4.9E2	
WIS	4.2E1	3.5E1	
NAIVE	1.1E1	-	

Table 133. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	5.7E0	9.4E1	1.8E2	1.6E1
Q-REG	7.4E1	5.8E1	5.4E1	7.4E1
MRDR	6.2E1	3.2E2	9.0E1	9.9E1
FQE	5.1E0	1.5E1	2.7E1	5.1E0
$R(\lambda)$	5.7E0	8.7E0	2.7E1	5.7E0
$Q^{\hat{\pi}}(\lambda)$	1.9E1	3.1E2	2.0E1	1.9E1
TREE	5.5E0	9.9E0	2.8E1	5.5E0
IH	2.6E1	-	-	-

		IPS	
	STANDARD PER-DECIS		
IS	1.0E0	1.2E2	
WIS	4.7E1	3.7E1	
NAIVE	7.3E0	-	

Table 134. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM	Hybrid)
	DIRECT	DR	WDR	MAGIC
AM	4.6E0	4.4E3	9.8E1	4.6E0
Q-Reg	1.1E2	1.4E3	1.3E3	1.1E2
MRDR	7.1E1	1.2E3	1.8E3	2.3E2
FQE	3.4E0	8.8E1	6.7E0	3.4E0
$R(\lambda)$	5.1E0	1.9E1	1.5E1	7.2E0
$Q^{\hat{\pi}}(\lambda)$	4.5E1	8.2E1	3.8E1	4.5E1
TREE	4.2E0	2.3E1	1.9E1	7.8E0
IH	1.3E1	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	5.2E1	1.4E2
WIS	4.1E1	2.4E1
NAIVE	6.8E0	-

Table 135. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic rewards. Sparse rewards.

DMHybrid DIRECT MAGIC DR WDR AM5.7E0 9.8E1 1.8E2 7.4E0 Q-REG 1.4E1 4.0E1 7.3E1 9.4E1 MRDR 1.9E1 4.2E0 2.5E2 2.4E2 FQE 4.6E0 6.3E0 2.8E1 4.6E0 $R(\lambda)$ 7.4E08.2E0 3.1E1 8.3E0 $Q^{\pi}(\lambda)$ 4.7E1 4.5E1 6.5E1 4.7E1 TREE 5.7E0 6.5E0 3.1E1 5.8E0 ΙH 5.2E0

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	6.9E0	
WIS	4.7E1	3.1E1	
NAIVE	4.0E0	-	

Table 136. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	1.7E0	1.6E3	1.2E2	1.6E0
Q-Reg	5.4E1	3.8E2	1.6E2	8.2E0
MRDR	9.3E1	4.8E1	6.0E2	6.0E2
FQE	1.9E0	1.1E2	1.4E1	2.0E0
$R(\lambda)$	4.2E0	8.9E0	1.1E1	9.0E0
$Q^{\hat{\pi}}(\lambda)$	4.3E0	3.8E1	1.5E1	4.3E0
TREE	3.4E0	3.0E1	7.7E0	6.6E0
IH	1.6E0	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	1.6E0	5.3E1
WIS	6.7E1	9.9E0
NAIVE	4.0E0	-

Table 137. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	1.5E0	1.3E5	6.1E1	1.2E1
Q-Reg	1.4E4	1.8E5	2.1E4	1.4E4
MRDR	8.9E3	3.8E5	3.2E4	8.9E3
FQE	1.1E0	2.4E4	1.2E1	1.2E0
$R(\lambda)$	5.1E0	8.1E0	1.7E1	5.1E0
$Q^{\pi}(\lambda)$	6.5E0	2.2E4	1.2E1	6.5E0
TREE	4.4E0	8.7E1	1.6E1	4.5E0
IH	8.1E-1	-	-	-

		IPS	
	STANDARD PER-DECISIO		
IS	5.7E0	1.7E4	
WIS	8.1E1	1.7E1	
NAIVE	4.7E0	-	

Table 138. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	1.1E0	4.6E1	6.9E1	5.1E0
Q-Reg	9.1E0	1.6E2	6.4E1	1.3E1
MRDR	1.3E1	4.5E1	3.1E1	1.5E1
FQE	9.8E-1	7.9E0	8.9E0	9.6E-1
$R(\lambda)$	2.6E0	8.7E0	1.0E1	2.6E0
$Q^{\pi}(\lambda)$	2.1E0	4.3E0	9.3E0	2.0E0
TREE	2.2E0	9.8E0	9.9E0	2.2E0
IH	9.8E-1	-	-	-

	IPS	
	STANDARD PER-DECIS	
IS	7.9E-1	8.5E0
WIS	1.9E1	9.9E0
NAIVE	3.7E0	-

Table 139. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.0E0	9.9E-1	8.5E0	1.0E0
Q-Reg	9.9E-1	1.3E0	1.8E3	4.0E1
MRDR	9.9E-1	1.3E0	1.4E3	1.3E1
FQE	1.0E0	1.0E0	2.5E0	1.0E0
$R(\lambda)$	1.0E0	9.9E-1	6.2E0	1.0E0
$Q^{\hat{\pi}}(\lambda)$	1.1E0	2.2E0	9.9E-1	9.7E-1
TRÈE	1.0E0	9.9E-1	6.2E0	1.0E0
IH	8.0E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	9.9E-1	9.9E-1	
WIS	6.2E0	6.2E0	
NAIVE	5.0E0	-	

Table 140. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.5E0	1.2E2	4.0E1	1.5E0	
Q-Reg	1.8E0	9.1E0	8.0E1	1.8E0	
MRDR	1.8E0	3.7E1	2.1E2	2.0E2	
FQE	1.3E0	1.6E0	6.4E0	1.3E0	
$R(\lambda)$	1.0E0	2.8E0	1.1E1	3.7E0	
$Q^{\hat{\pi}}(\lambda)$	2.8E0	2.6E3	2.6E0	2.7E0	
TRÈE	1.0E0	2.8E0	1.1E1	3.7E0	
IH	1.3E0	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	2.8E0	2.8E0	
WIS	1.1E1	1.1E1	
NAIVE	5.2E0	-	

Table 141. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.8E0	4.9E0	5.5E1	1.7E0	
Q-Reg	9.8E-1	9.0E-1	6.9E1	9.8E-1	
MRDR	9.6E-1	1.1E0	1.5E2	1.5E2	
FQE	1.5E0	1.4E0	5.5E0	1.5E0	
$R(\lambda)$	1.0E0	9.8E-1	6.4E0	1.1E0	
$Q^{\pi}(\lambda)$	8.5E0	1.1E1	9.8E0	7.0E0	
TREE	1.0E0	9.8E-1	6.4E0	1.1E0	
IH	1.0E0	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	9.8E-1	9.8E-1	
WIS	6.4E0	6.4E0	
NAIVE	4.5E0	-	

Table 142. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.3E0	8.9E2	8.3E1	1.8E0	
Q-Reg	1.0E0	1.0E1	8.1E4	3.3E0	
MRDR	1.0E0	1.3E1	4.2E5	4.2E5	
FQE	2.1E0	2.1E0	6.4E0	2.2E0	
$R(\lambda)$	1.0E0	1.0E0	1.1E1	2.5E0	
$Q^{\hat{\pi}}(\lambda)$	1.0E1	1.3E1	6.3E0	9.2E0	
TREE	1.0E0	1.0E0	1.1E1	2.5E0	
IH	1.0E0	-	-	-	

	IPS		
	STANDARD	Per-Decision	
IS	1.0E0	1.0E0	
WIS	1.1E1	1.1E1	
NAIVE	4.3E0	-	

Table 143. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic environment. Sparse rewards.

DM Hybrid DIRECT DR **WDR** MAGIC 4.4E-1 2.3E0 3.2E3 4.9E1 AM4.9E0 Q-REG 5.0E0 4.9E21.6E1 MRDR 5.2E1 4.6E1 1.1E3 1.1E3 FQE 4.1E-1 6.4E-14.8E0 4.2E-1 $R(\lambda)$ 1.0E0 9.0E0 2.8E0 4.3E0 $Q^{\hat{\pi}}(\hat{\lambda})$ 3.8E0 8.9E0 4.0E0 4.8E1 TREE 1.0E0 4.3E0 9.0E0 2.7E0 IΗ 6.8E-1

	IPS			
	STANDARD PER-DECISION			
IS	4.3E0	4.3E0		
WIS	9.0E0	9.0E0		
NAIVE	3.9E0	-		

Table 144. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.1E-1	1.4E3	2.2E1	1.9E-1	
Q-Reg	1.2E0	6.1E2	3.5E2	4.9E0	
MRDR	1.3E1	5.1E1	4.3E2	4.3E2	
FQE	9.2E-2	2.1E0	1.8E0	9.8E-2	
$R(\lambda)$	1.0E0	1.3E0	3.9E0	3.5E0	
$Q^{\hat{\pi}}(\lambda)$	1.3E0	3.1E1	2.7E0	1.5E0	
TREE	1.0E0	1.3E0	3.9E0	3.5E0	
IH	1.3E-1	-	-	_	

	IPS		
	STANDARD	PER-DECISION	
IS	1.3E0	1.3E0	
WIS	3.9E0	3.9E0	
NAIVE	4.1E0	-	

Table 145. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	8.1E-2	2.1E1	3.8E1	4.7E0
Q-Reg	1.6E0	3.5E1	4.0E1	2.4E0
MRDR	1.4E0	4.0E0	2.3E2	2.3E2
FQE	5.3E-2	9.1E0	6.5E0	5.1E-2
$R(\lambda)$	1.0E0	1.7E0	6.8E0	1.0E0
$Q^{\hat{\pi}}(\lambda)$	1.1E0	1.0E1	6.2E0	9.9E-1
TREE	1.0E0	1.7E0	6.8E0	1.0E0
IH	1.9E-1	-	-	-

	IPS				
	STANDARD PER-DECISION				
IS	1.7E0	1.7E0			
WIS	6.8E0	6.8E0			
NAIVE	3.9E0	-			

Table 146. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.6E-2	5.0E1	8.6E0	6.2E-2	
Q-Reg	1.1E0	2.8E1	1.8E1	4.8E0	
MRDR	8.7E-1	1.7E0	1.9E2	1.6E2	
FQE	3.4E-2	3.9E-1	2.8E0	3.4E-2	
$R(\lambda)$	1.0E0	1.1E0	4.9E0	1.2E0	
$Q^{\pi}(\lambda)$	6.6E-1	2.1E0	3.5E0	6.3E-1	
TREE	1.0E0	1.1E0	4.9E0	1.2E0	
IH	9.3E-2	-	-	-	

	IPS			
	STANDARD PER-DECISIO			
IS	1.1E0	1.1E0		
WIS	4.9E0	4.9E0		
NAIVE	4.0E0	-		

Table 147. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.6E1	8.4E3	3.2E2	1.7E1	
Q-Reg	9.9E3	6.2E5	4.0E4	9.9E3	
MRDR	7.8E3	1.4E6	1.3E4	7.9E3	
FQE	1.2E1	2.2E3	7.8E1	1.2E1	
$R(\lambda)$	2.2E1	9.1E3	9.7E1	2.5E1	
$Q^{\hat{\pi}}(\hat{\lambda})$	3.3E1	2.7E3	7.6E1	3.2E1	
TREE	2.0E1	1.3E4	8.8E1	2.0E1	
IH	3.2E2	-	-	-	

	IPS			
	STANDARD	Per-Decision		
IS	8.9E2	1.8E4		
WIS	1.3E2	6.6E1		
NAIVE	1.8E1	-		

Table 148. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.2E1	8.1E2	3.3E2	2.1E1	
Q-REG	3.0E3	6.1E3	5.2E3	5.3E2	
MRDR	1.9E3	3.3E4	1.6E4	8.5E3	
FQE	2.2E1	4.5E1	3.8E1	2.2E1	
$R(\lambda)$	2.5E1	3.6E1	1.7E2	2.5E1	
$Q^{\pi}(\lambda)$	2.3E2	3.3E2	2.1E2	2.3E2	
TREE	2.6E1	2.9E1	1.7E2	2.6E1	
IH	1.0E2	-	-	-	

	IPS			
	STANDARD	PER-DECISION		
IS	1.1E0	3.0E3		
WIS	2.4E2	1.9E2		
NAIVE	2.5E1	-		

Table 149. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	5.5E1	3.4E4	4.4E2	4.8E1	
Q-Reg	1.4E4	9.1E4	8.4E4	1.4E4	
MRDR	1.0E5	3.1E5	4.0E5	1.5E5	
FQE	5.2E1	2.0E4	1.5E2	5.2E1	
$R(\lambda)$	4.3E1	7.0E1	7.9E1	4.3E1	
$Q^{\hat{\pi}}(\lambda)$	1.5E2	1.5E4	1.6E2	1.5E2	
TREE	3.6E1	2.7E2	9.3E1	3.7E1	
IH	3.8E1	-	-	-	

	IPS				
	STANDARD PER-DECISION				
IS	1.0E0	2.1E4			
WIS	9.3E1	1.3E2			
NAIVE	1.8E1	-			

Table 150. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.7E1	2.8E3	6.0E2	2.8E1	
Q-Reg	3.1E2	6.9E1	4.1E2	3.1E2	
MRDR	3.5E2	7.6E2	3.1E3	3.4E2	
FQE	2.8E1	3.5E2	1.4E2	2.8E1	
$R(\lambda)$	4.0E1	1.4E2	1.1E2	3.9E1	
$Q^{\hat{\pi}}(\lambda)$	1.7E2	7.0E2	2.4E2	1.7E2	
TREE	3.3E1	1.5E2	1.2E2	3.2E1	
IH	3.2E1	-	-	-	

	IPS			
	STANDARD PER-DECISI			
IS	1.0E0	3.2E2		
WIS	2.2E2	1.3E2		
NAIVE	9.5E0	_		

Table 151. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.3E1	8.9E2	8.5E2	1.4E1	
Q-Reg	3.6E1	3.8E4	9.3E3	2.4E1	
MRDR	5.7E1	4.3E2	7.7E2	5.6E2	
FQE	1.1E1	6.0E1	1.1E2	1.1E1	
$R(\lambda)$	9.9E0	3.2E1	9.8E1	1.2E1	
$Q^{\hat{\pi}}(\lambda)$	4.2E1	2.4E2	1.5E2	4.4E1	
TREE	7.8E0	3.0E1	9.8E1	1.0E1	
IH	1.6E1	-	-	-	

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	1.0E0	3.1E1			
WIS	8.3E1	1.1E2			
NAIVE	7.8E0	-			

Table 152. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	8.9E0	5.1E2	5.2E2	8.4E0
Q-Reg	2.5E1	1.8E3	1.9E4	3.2E1
MRDR	1.2E2	5.4E1	2.5E3	2.4E3
FQE	8.3E0	3.4E1	8.0E1	8.3E0
$R(\lambda)$	1.6E1	3.6E1	8.5E1	1.7E1
$Q^{\pi}(\lambda)$	3.7E1	7.5E1	1.0E2	4.2E1
TREE	1.1E1	3.1E1	8.8E1	1.1E1
IH	8.7E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	9.9E1	2.4E1	
WIS	2.1E2	9.5E1	
NAIVE	5.2E0	-	

Table 153. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.0E0	2.3E2	4.1E2	4.9E0
Q-Reg	1.6E1	5.4E1	1.7E2	8.8E0
MRDR	2.7E1	4.2E1	3.0E2	1.4E2
FQE	2.4E0	1.6E1	1.0E1	5.3E0
$R(\lambda)$	3.4E0	2.0E1	1.7E1	3.2E0
$Q^{\hat{\pi}}(\hat{\lambda})$	1.3E1	2.9E1	6.4E0	1.2E1
TREE	4.3E0	2.0E1	1.7E1	4.0E0
IH	2.3E0	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	1.2E0	1.7E1		
WIS	4.2E1	1.6E1		
NAIVE	4.5E0	-		

Table 154. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	1.5E0	1.0E3	1.6E2	4.7E0
Q-Reg	3.2E2	9.7E3	2.3E3	3.2E2
MRDR	6.6E3	2.3E3	1.4E4	8.8E3
FQE	1.5E0	5.0E2	1.6E1	3.6E0
$R(\lambda)$	3.4E0	3.1E2	2.3E1	8.6E0
$Q^{\pi}(\lambda)$	9.1E0	5.0E2	1.5E1	7.6E0
TREE	2.1E0	3.7E2	2.1E1	6.5E0
IH	1.5E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	4.1E0	3.2E2	
WIS	4.5E1	2.0E1	
NAIVE	3.3E0	-	

Table 155. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.7E-2	4.9E-2	4.5E-2	4.1E-2
Q-REG	2.2E-1	9.6E-2	9.4E-2	1.8E-1
MRDR	3.6E-1	1.8E-1	1.6E-1	3.6E-1
FQE	3.3E-2	3.3E-2	3.3E-2	3.3E-2
$R(\lambda)$	3.3E-2	3.3E-2	3.3E-2	3.3E-2
$Q^{\pi}(\lambda)$	3.3E-2	3.3E-2	3.3E-2	3.3E-2
TREE	3.3E-2	3.4E-2	3.4E-2	3.3E-2
IH	1.2E-1	-	-	-

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	5.0E-1	2.7E-1			
WIS	3.6E-1	2.4E-1			
NAIVE	8.6E-1	-			

Table 156. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.0E-2	4.0E-2	3.7E-2	2.7E-2
Q-Reg	4.4E-2	5.5E-3	4.3E-3	7.1E-3
MRDR	7.9E-2	1.1E-2	7.6E-3	3.6E-2
FQE	4.2E-3	4.2E-3	4.2E-3	4.2E-3
$R(\lambda)$	4.2E-3	4.2E-3	4.2E-3	4.2E-3
$Q^{\hat{\pi}}(\lambda)$	4.2E-3	4.2E-3	4.2E-3	4.2E-3
TREE	4.2E-3	4.2E-3	4.2E-3	4.2E-3
IH	2.5E-2	-	-	-

	IPS			
	STANDARD PER-DECISIO			
IS	2.0E-1	7.1E-2		
WIS	4.4E-2	2.6E-2		
NAIVE	5.6E-1	-		

Table 157. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.0E-3	1.5E-2	1.4E-2	2.7E-3
Q-Reg	1.5E-2	1.2E-3	7.2E-4	4.2E-3
MRDR	6.7E-2	3.7E-3	2.0E-3	1.8E-2
FQE	4.2E-7	4.2E-7	4.2E-7	4.2E-7
$R(\lambda)$	4.2E-7	4.2E-7	4.2E-7	4.2E-7
$Q^{\hat{\pi}}(\lambda)$	4.2E-7	4.2E-7	4.2E-7	4.2E-7
TREE	1.6E-6	5.3E-7	5.4E-7	1.6E-6
IH	1.1E-2	-	-	-

	IPS	
	STANDARD	Per-Decision
IS	8.4E-2	3.8E-2
WIS	2.1E-2	1.2E-2
NAIVE	5.1E-1	-

Table 158. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.1E-3	9.5E-3	9.6E-3	2.3E-3
Q-Reg	2.0E-2	4.2E-4	1.6E-4	1.3E-3
MRDR	3.2E-2	1.6E-3	8.6E-4	2.1E-3
FQE	1.8E-5	1.8E-5	1.8E-5	1.8E-5
$R(\lambda)$	1.8E-5	1.8E-5	1.8E-5	1.8E-5
$Q^{\hat{\pi}}(\lambda)$	1.8E-5	1.8E-5	1.8E-5	1.8E-5
TRÈE	1.3E-5	1.6E-5	1.6E-5	1.3E-5
IH	9.4E-3	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	4.8E-2	2.5E-2		
WIS	1.4E-2	8.5E-3		
NAIVE	3.9E-1	-		

Table 159. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

DM Hybrid DIRECT DR WDR MAGIC AM2.3E-3 3.6E-3 3.8E-3 3.0E-3 Q-REG 6.4E-3 3.1E-4 2.3E-4 1.4E-3 MRDR 3.3E-22.0E-3 1.2E-3 1.5E-3 7.4E-87.4E-8 FQE 7.4E-8 7.4E-8 $R(\lambda)$ 7.3E-8 7.3E-8 7.2E-8 7.2E-8 $Q^{\pi}(\lambda)$ 7.6E-8 7.5E-8 7.5E-8 7.5E-8 1.0E-7 TREE 6.7E-86.8E-8 9.1E-8 ΙH 3.7E-3

	IPS		
	STANDARD PER-DECISION		
IS	2.2E-2	1.3E-2	
WIS	7.3E-3	4.5E-3	
NAIVE	4.7E-1	-	

Table 160. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.4E-4	3.0E-3	3.0E-3	1.8E-3
Q-Reg	3.5E-3	7.6E-6	6.5E-6	2.5E-5
MRDR	2.7E-2	1.7E-4	1.0E-4	2.7E-4
FQE	8.4E-7	8.4E-7	8.4E-7	8.4E-7
$R(\lambda)$	8.4E-7	8.4E-7	8.4E-7	8.4E-7
$Q^{\hat{\pi}}(\lambda)$	8.4E-7	8.4E-7	8.4E-7	8.5E-7
TREE	1.2E-7	7.9E-7	8.0E-7	1.3E-7
IH	4.6E-3	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.2E-2	9.2E-3	
WIS	2.9E-3	3.3E-3	
NAIVE	4.6E-1	-	

Table 161. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.7E-4	2.5E-3	2.4E-3	1.6E-3
Q-Reg	3.3E-3	2.9E-5	2.0E-5	1.1E-4
MRDR	1.7E-2	1.2E-4	8.9E-5	1.4E-4
FQE	1.4E-5	1.4E-5	1.4E-5	1.4E-5
$R(\lambda)$	1.4E-5	1.4E-5	1.4E-5	1.4E-5
$Q^{\hat{\pi}}(\lambda)$	1.4E-5	1.4E-5	1.4E-5	1.4E-5
TREE	1.0E-5	1.5E-5	1.5E-5	1.1E-5
IH	1.7E-3	-	-	-

	IPS	
	STANDARD	Per-Decision
IS	1.0E-2	5.8E-3
WIS	2.6E-3	1.7E-3
NAIVE	4.2E-1	-

Table 162. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.0E-4	5.5E-4	5.5E-4	2.2E-4
Q-Reg	9.3E-4	1.6E-5	1.6E-5	1.7E-5
MRDR	1.6E-2	1.8E-4	2.0E-4	2.4E-4
FQE	1.6E-5	1.6E-5	1.6E-5	1.6E-5
$R(\lambda)$	1.6E-5	1.6E-5	1.6E-5	1.6E-5
$Q^{\pi}(\lambda)$	1.6E-5	1.6E-5	1.6E-5	1.6E-5
TREE	1.2E-5	1.6E-5	1.6E-5	1.2E-5
IH	5.1E-4	-	-	-

		IPS	
	STANDARD	PER-DECISION	
IS	2.4E-3	1.8E-3	
WIS	8.6E-4	6.1E-4	
NAIVE	4.3E-1	-	

Table 163. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.9E-1	1.7E-1	1.7E-1	1.8E-1
Q-Reg	3.1E-1	2.7E-1	2.6E-1	2.8E-1
MRDR	3.2E-1	2.1E-1	2.3E-1	3.2E-1
FQE	1.9E-1	2.5E-1	2.4E-1	1.9E-1
$R(\lambda)$	2.3E-1	2.5E-1	2.4E-1	2.3E-1
$Q^{\hat{\pi}}(\lambda)$	2.2E-1	2.2E-1	2.2E-1	2.2E-1
TRÈE	2.4E-1	2.6E-1	2.6E-1	2.4E-1
IH	3.0E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	5.8E-1	
WIS	5.1E-1	3.7E-1	
NAIVE	5.6E-1	-	

Table 164. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.0E-1	2.0E-1	1.9E-1	1.0E-1
Q-Reg	2.2E-1	1.3E-1	1.2E-1	2.2E-1
MRDR	2.3E-1	1.5E-1	1.3E-1	2.3E-1
FQE	8.1E-2	9.0E-2	8.8E-2	8.1E-2
$R(\lambda)$	9.6E-2	9.9E-2	9.8E-2	9.6E-2
$Q^{\hat{\pi}}(\lambda)$	1.0E-1	1.2E-1	1.1E-1	1.0E-1
TRÈE	9.6E-2	9.4E-2	9.4E-2	9.6E-2
IH	1.3E-1	-	-	-

	IPS		
	STANDARD PER-DECISIO		
IS	4.3E-1	2.4E-1	
WIS	1.8E-1	1.5E-1	
NAIVE	5.9E-1	-	

Table 165. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.5E-2	4.9E-2	4.7E-2	4.3E-2
Q-Reg	3.1E-2	3.0E-2	3.0E-2	3.1E-2
MRDR	5.6E-2	3.1E-2	3.1E-2	4.9E-2
FQE	2.3E-2	2.6E-2	2.6E-2	2.3E-2
$R(\lambda)$	2.7E-2	2.9E-2	2.8E-2	2.7E-2
$Q^{\hat{\pi}}(\lambda)$	2.5E-2	3.0E-2	3.0E-2	2.5E-2
TREE	2.8E-2	2.7E-2	2.7E-2	2.8E-2
IH	2.0E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	2.7E-2	2.1E-2	
WIS	2.6E-2	2.1E-2	
NAIVE	4.6E-1	-	

Table 166. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.6E-2	3.4E-2	3.2E-2	1.5E-2
Q-Reg	2.9E-2	7.3E-3	7.3E-3	7.3E-3
MRDR	2.6E-2	6.4E-3	6.6E-3	2.3E-2
FQE	8.0E-3	8.8E-3	8.6E-3	7.9E-3
$R(\lambda)$	7.9E-3	7.8E-3	7.9E-3	7.9E-3
$Q^{\pi}(\lambda)$	8.1E-3	8.4E-3	8.3E-3	8.2E-3
TREE	8.1E-3	7.5E-3	7.7E-3	8.0E-3
IH	1.6E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	8.7E-2	3.3E-2	
WIS	1.9E-2	1.3E-2	
NAIVE	3.6E-1	-	

Table 167. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

DMHybrid DIRECT DR WDR MAGIC 2.8E-2 4.8E-2 5.0E-2 2.4E-2 AM 7.2E-3 Q-REG 1.8E-2 1.0E-2 1.0E-2 MRDR 4.5E-28.6E-3 1.1E-22.1E-21.0E-2 7.8E-3 1.0E-27.9E-3 FQE $R(\lambda)$ 1.0E-2 1.1E-21.1E-21.0E-2 $Q^{\pi}(\lambda)$ 1.0E-2 1.1E-2 1.1E-2 1.0E-2 1.1E-2 TREE 9.7E-3 1.1E-2 9.8E-3 ΙH 9.8E-3

	IPS			
	STANDARD PER-DECISION			
IS	5.4E-2	1.6E-2		
WIS	1.9E-2	1.0E-2		
NAIVE	4.7E-1	-		

Table 168. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.2E-3	1.9E-2	1.9E-2	8.0E-3
Q-Reg	1.2E-2	8.0E-3	7.8E-3	8.8E-3
MRDR	2.7E-2	7.7E-3	7.5E-3	1.0E-2
FQE	5.9E-3	7.3E-3	7.3E-3	5.9E-3
$R(\lambda)$	7.0E-3	7.3E-3	7.3E-3	7.0E-3
$Q^{\hat{\pi}}(\lambda)$	7.2E-3	7.3E-3	7.3E-3	7.2E-3
TRÈE	7.0E-3	7.3E-3	7.3E-3	7.0E-3
IH	7.7E-3	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	2.3E-2	1.5E-2	
WIS	9.3E-3	9.6E-3	
NAIVE	4.2E-1	-	

Table 169. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.9E-3	3.4E-3	3.3E-3	1.3E-3
Q-Reg	3.4E-3	9.3E-4	9.5E-4	2.5E-3
MRDR	2.5E-2	1.4E-3	1.4E-3	1.5E-3
FQE	6.7E-4	1.0E-3	1.0E-3	6.8E-4
$R(\lambda)$	8.5E-4	1.0E-3	1.0E-3	8.5E-4
$Q^{\pi}(\lambda)$	6.8E-4	1.0E-3	1.0E-3	6.9E-4
TREE	9.0E-4	1.0E-3	1.0E-3	9.1E-4
IH	3.1E-3	-	-	-

	IPS			
	STANDARD PER-DECIS			
IS	8.3E-3	5.4E-3		
WIS	2.9E-3	2.6E-3		
NAIVE	4.7E-1	-		

Table 170. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.1E-3	1.8E-3	1.8E-3	1.4E-3
Q-Reg	5.3E-4	1.0E-3	1.0E-3	4.2E-4
MRDR	1.9E-2	1.4E-3	1.5E-3	1.5E-3
FQE	9.5E-4	1.0E-3	1.0E-3	9.5E-4
$R(\lambda)$	1.0E-3	1.1E-3	1.1E-3	1.0E-3
$Q^{\pi}(\lambda)$	1.1E-3	1.0E-3	1.0E-3	1.1E-3
TREE	1.0E-3	1.1E-3	1.1E-3	1.0E-3
IH	8.7E-4	-	-	-

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	2.6E-3	1.6E-3			
WIS	1.1E-3	7.7E-4			
NAIVE	4.6E-1	-			

Table 171. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	7.7E-1	8.5E-1	9.0E-1	6.9E-1	
Q-REG	9.0E-1	1.2E0	1.1E0	1.0E0	
MRDR	6.6E-1	6.9E-1	7.3E-1	6.4E-1	
FQE	6.3E-1	9.1E-1	9.0E-1	6.3E-1	
$R(\lambda)$	9.5E-1	1.0E0	1.0E0	9.5E-1	
$Q^{\pi}(\lambda)$	8.6E-1	9.3E-1	9.3E-1	8.6E-1	
TREE	9.4E-1	9.9E-1	9.9E-1	9.4E-1	
IH	4.1E-1	-	-	_	

	IPS		
	STANDARD PER-DECISION		
IS	1.2E0	9.5E-1	
WIS	9.4E-1	8.5E-1	
NAIVE	1.2E0	-	

Table 172. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.1E-1	1.3E-1	1.6E-1	1.8E-1	
Q-Reg	2.6E-1	2.3E-1	2.2E-1	2.8E-1	
MRDR	1.2E-1	1.7E-1	1.6E-1	1.3E-1	
FQE	2.0E-1	1.4E-1	1.5E-1	2.0E-1	
$R(\lambda)$	1.7E-1	1.8E-1	1.8E-1	1.7E-1	
$Q^{\hat{\pi}}(\lambda)$	2.1E-1	1.8E-1	1.8E-1	2.1E-1	
TREE	1.5E-1	1.7E-1	1.7E-1	1.5E-1	
IH	1.9E-1	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	5.8E-1	3.0E-1	
WIS	2.9E-1	1.8E-1	
NAIVE	6.2E-1	-	

Table 173. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.1E-1	2.0E-1	2.0E-1	1.9E-1
Q-Reg	1.6E-1	1.0E-1	1.1E-1	1.5E-1
MRDR	1.6E-1	1.1E-1	1.1E-1	1.6E-1
FQE	1.3E-1	1.5E-1	1.4E-1	1.3E-1
$R(\lambda)$	1.2E-1	1.3E-1	1.2E-1	1.2E-1
$Q^{\hat{\pi}}(\lambda)$	1.1E-1	1.3E-1	1.2E-1	1.1E-1
TREE	1.4E-1	1.3E-1	1.3E-1	1.3E-1
IH	1.5E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	2.1E-1	1.6E-1	
WIS	1.8E-1	1.4E-1	
NAIVE	4.8E-1	-	

Table 174. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.3E-1	2.4E-1	2.4E-1	1.5E-1
Q-Reg	1.4E-1	9.1E-2	9.4E-2	1.0E-1
MRDR	1.2E-1	7.7E-2	8.0E-2	7.7E-2
FQE	6.2E-2	9.8E-2	9.7E-2	6.3E-2
$R(\lambda)$	8.8E-2	1.0E-1	1.0E-1	7.9E-2
$Q^{\hat{\pi}}(\lambda)$	9.6E-2	1.0E-1	1.0E-1	8.7E-2
TREE	7.9E-2	9.6E-2	9.6E-2	7.5E-2
IH	1.3E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	3.9E-1	2.1E-1	
WIS	1.9E-1	1.4E-1	
NAIVE	4.8E-1	-	

Table 175. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.7E-2	1.1E-1	1.2E-1	8.9E-2
Q-Reg	9.1E-2	9.0E-2	9.1E-2	8.0E-2
MRDR	1.3E-1	8.8E-2	9.3E-2	1.1E-1
FQE	7.8E-2	8.4E-2	8.6E-2	7.8E-2
$R(\lambda)$	9.3E-2	8.9E-2	9.0E-2	9.5E-2
$Q^{\hat{\pi}}(\lambda)$	1.0E-1	8.8E-2	8.9E-2	9.6E-2
TREE	1.0E-1	9.2E-2	9.2E-2	1.0E-1
IH	1.2E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.1E-1	9.7E-2		
WIS	9.8E-2	8.8E-2		
NAIVE	5.3E-1	-		

Table 176. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.0E-2	2.9E-2	3.0E-2	1.1E-2
Q-Reg	4.1E-2	4.1E-2	4.0E-2	3.7E-2
MRDR	8.6E-2	4.3E-2	4.3E-2	4.8E-2
FQE	2.4E-2	3.8E-2	3.8E-2	2.5E-2
$R(\lambda)$	3.7E-2	3.9E-2	3.9E-2	3.6E-2
$Q^{\hat{\pi}}(\lambda)$	3.9E-2	3.9E-2	4.0E-2	3.6E-2
TREE	3.6E-2	3.9E-2	3.9E-2	3.7E-2
IH	3.4E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	5.3E-2	4.6E-2	
WIS	4.3E-2	4.0E-2	
NAIVE	5.7E-1	-	

Table 177. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.2E-2	2.0E-2	2.0E-2	1.2E-2
Q-Reg	1.4E-2	1.1E-2	1.1E-2	9.6E-3
MRDR	6.6E-3	1.1E-2	1.1E-2	8.2E-3
FQE	5.3E-3	1.1E-2	1.1E-2	5.4E-3
$R(\lambda)$	9.0E-3	1.1E-2	1.1E-2	7.5E-3
$Q^{\pi}(\lambda)$	9.3E-3	1.1E-2	1.1E-2	7.2E-3
TREE	8.7E-3	1.1E-2	1.1E-2	7.9E-3
IH	8.6E-3	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.4E-2	1.6E-2	
WIS	1.2E-2	1.4E-2	
NAIVE	3.9E-1	-	

Table 178. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.2E-3	6.8E-3	6.9E-3	3.1E-3
Q-Reg	3.6E-3	1.8E-3	1.8E-3	3.3E-3
MRDR	2.1E-2	1.4E-3	1.2E-3	1.0E-2
FQE	1.7E-3	1.8E-3	1.8E-3	1.7E-3
$R(\lambda)$	1.8E-3	1.8E-3	1.8E-3	2.1E-3
$Q^{\hat{\pi}}(\lambda)$	2.5E-3	1.8E-3	1.8E-3	2.6E-3
TREE	1.8E-3	1.8E-3	1.8E-3	2.1E-3
IH	2.2E-3	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	4.8E-3	3.3E-3	
WIS	3.2E-3	2.7E-3	
NAIVE	4.6E-1	-	

Table 179. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

DM Hybrid DIRECT DR WDR MAGIC 5.9E-1 4.3E-1 4.3E-1 7.1E-1 AMQ-REG 3.9E-1 3.7E-1 3.6E-1 3.9E-1 MRDR 3.8E-1 2.8E-1 3.4E-1 3.8E-1 FQE 4.1E-1 3.0E-1 3.3E-1 4.1E-1 $R(\lambda)$ 3.2E-1 3.6E-1 3.2E-1 3.4E-1 $Q^{\hat{\pi}}(\hat{\lambda})$ 3.7E-1 3.4E-1 3.1E-1 3.7E-1 TREE 3.2E-1 3.5E-1 3.6E-1 3.2E-1 ΙH 7.9E-1

	IPS		
	STANDARD	PER-DECISION	
IS	1.4E0	9.7E-1	
WIS	6.9E-1	6.8E-1	
NAIVE	6.6E-1	-	

Table 180. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.2E0	1.7E0	1.9E0	1.4E0
Q-Reg	7.7E-1	9.8E-1	9.7E-1	7.7E-1
MRDR	6.6E-1	6.9E-1	7.4E-1	6.5E-1
FQE	9.8E-1	9.3E-1	1.0E0	9.8E-1
$R(\lambda)$	9.9E-1	1.0E0	1.0E0	9.9E-1
$Q^{\pi}(\lambda)$	9.6E-1	9.6E-1	1.0E0	9.6E-1
TREE	1.0E0	1.0E0	1.1E0	1.0E0
IH	7.1E-1	_	-	_

		IPS
	STANDARD	PER-DECISION
IS	7.6E-1	8.4E-1
WIS	7.7E-1	8.5E-1
NAIVE	1.6E0	-

Table 181. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.5E-1	4.6E-1	4.6E-1	2.6E-1
Q-Reg	3.4E-1	2.9E-1	3.0E-1	3.4E-1
MRDR	3.0E-1	2.4E-1	2.5E-1	3.0E-1
FQE	2.4E-1	3.0E-1	3.1E-1	2.4E-1
$R(\lambda)$	3.0E-1	3.1E-1	3.1E-1	3.0E-1
$Q^{\hat{\pi}}(\lambda)$	3.4E-1	3.3E-1	3.3E-1	3.4E-1
TREE	2.9E-1	3.0E-1	2.9E-1	2.9E-1
IH	2.7E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	3.8E-1	3.6E-1	
WIS	3.3E-1	3.1E-1	
NAIVE	9.0E-1	-	

Table 182. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.0E-1	1.1E-1	1.2E-1	9.8E-2
Q-Reg	6.8E-2	1.0E-1	1.1E-1	8.4E-2
MRDR	1.1E-1	8.4E-2	8.8E-2	1.0E-1
FQE	8.8E-2	9.1E-2	9.3E-2	8.8E-2
$R(\lambda)$	8.9E-2	1.1E-1	1.0E-1	8.7E-2
$Q^{\pi}(\lambda)$	8.8E-2	1.1E-1	1.1E-1	8.5E-2
TREE	8.3E-2	1.0E-1	9.7E-2	8.4E-2
IH	7.5E-2	-	-	-

		IPS
	STANDARD	Per-Decision
IS	5.9E-2	6.0E-2
WIS	8.1E-2	7.5E-2
NAIVE	5.3E-1	-

Table 183. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

DM Hybrid DIRECT DR WDR MAGIC 3.9E-2 8.8E-2 9.3E-2 6.4E-2 AM 7.0E-2 Q-REG 9.3E-2 7.4E-2 7.4E-2 MRDR 1.3E-1 8.1E-2 7.9E-2 1.3E-1 FQE 4.0E-2 6.3E-26.8E-2 4.0E-2 $R(\lambda)$ 6.1E-27.2E-2 4.7E-27.1E-2 $Q^{\hat{\pi}}(\hat{\lambda})$ 6.8E-2 7.4E-24.3E-2 7.2E-2 TREE 5.9E-2 5.3E-2 7.0E-2 7.0E-2ΙH 3.3E-2

	IPS		
	STANDARD	PER-DECISION	
IS	1.8E-1	9.7E-2	
WIS	1.6E-1	8.9E-2	
NAIVE	7.0E-1	-	

Table 184. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.9E-2	1.7E-1	1.6E-1	9.5E-2
Q-Reg	1.3E-1	7.6E-2	7.7E-2	1.1E-1
MRDR	1.3E-1	7.5E-2	7.0E-2	9.6E-2
FQE	4.9E-2	8.0E-2	7.9E-2	5.0E-2
$R(\lambda)$	7.0E-2	7.8E-2	7.8E-2	6.1E-2
$Q^{\hat{\pi}}(\lambda)$	7.1E-2	7.9E-2	7.8E-2	6.2E-2
TREE	7.0E-2	7.8E-2	7.8E-2	6.3E-2
IH	7.2E-2	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	2.0E-1	1.5E-1
WIS	1.4E-1	1.2E-1
NAIVE	4.9E-1	-

Table 185. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.7E-2	2.3E-2	2.3E-2	2.8E-2
Q-Reg	4.1E-2	3.5E-2	3.5E-2	3.4E-2
MRDR	5.6E-2	3.9E-2	3.9E-2	5.3E-2
FQE	2.4E-2	3.5E-2	3.5E-2	2.5E-2
$R(\lambda)$	3.3E-2	3.5E-2	3.5E-2	2.8E-2
$Q^{\hat{\pi}}(\lambda)$	3.3E-2	3.4E-2	3.5E-2	2.8E-2
TREE	3.2E-2	3.5E-2	3.5E-2	2.9E-2
IH	2.8E-2	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	4.9E-2	4.2E-2	
WIS	4.1E-2	3.8E-2	
NAIVE	5.1E-1	-	

Table 186. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.0E-2	2.2E-2	2.3E-2	1.7E-2
Q-Reg	1.5E-2	1.5E-2	1.5E-2	1.5E-2
MRDR	4.1E-2	1.6E-2	1.7E-2	2.6E-2
FQE	1.3E-2	1.5E-2	1.5E-2	1.3E-2
$R(\lambda)$	1.5E-2	1.5E-2	1.5E-2	1.4E-2
$Q^{\hat{\pi}}(\lambda)$	1.5E-2	1.5E-2	1.5E-2	1.4E-2
TREE	1.5E-2	1.5E-2	1.5E-2	1.5E-2
IH	1.4E-2	-	-	-

	IPS	
	STANDARD	Per-Decision
IS	2.5E-2	1.6E-2
WIS	2.2E-2	1.5E-2
NAIVE	5.0E-1	-

Table 187. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

DM Hybrid DIRECT DR WDR MAGIC 3.2E-1 AM8.6E-2 3.4E-1 2.6E-1 Q-REG 2.6E-1 9.6E-2 6.5E-2 3.9E-1 MRDR 2.7E-1 1.8E-1 1.4E-1 2.9E-1 1.9E-2 1.9E-2 1.9E-2 FQE 1.2E-1 $R(\lambda)$ 1.9E-2 1.9E-2 1.9E-2 1.2E-1 $Q^{\pi}(\lambda)$ 1.9E-2 1.9E-2 1.9E-2 1.2E-1 1.9E-2 TREE 1.9E-2 1.9E-2 1.2E-1 ΙH 1.7E-1

	IPS		
	STANDARD	PER-DECISION	
IS	6.0E-1	6.0E-1	
WIS	2.4E-1	2.4E-1	
NAIVE	3.6E-1	-	

Table 188. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.6E-2	1.1E-1	1.1E-1	2.7E-1
Q-Reg	8.3E-2	1.3E-2	1.3E-2	2.3E-1
MRDR	1.4E-1	3.1E-2	2.4E-2	3.1E-1
FQE	5.6E-4	5.6E-4	5.6E-4	2.0E-1
$R(\lambda)$	5.6E-4	5.6E-4	5.6E-4	2.0E-1
$Q^{\hat{\pi}}(\lambda)$	5.6E-4	5.6E-4	5.6E-4	2.0E-1
TREE	5.9E-4	5.3E-4	5.3E-4	2.0E-1
IH	1.4E-1	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	6.6E-2	6.6E-2
WIS	5.1E-2	5.1E-2
NAIVE	6.2E-1	-

Table 189. Graph, relative MSE. $T = 4, N = 32, \pi_b(a = 0) = 0.6, \pi_e(a = 0) = 0.8$. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.1E-2	4.8E-2	4.7E-2	2.7E-2
Q-Reg	1.2E-1	2.1E-2	9.0E-3	1.5E-2
MRDR	9.8E-2	2.2E-2	1.4E-2	4.7E-2
FQE	4.1E-4	4.1E-4	4.1E-4	4.1E-4
$R(\lambda)$	4.2E-4	4.1E-4	4.2E-4	4.2E-4
$Q^{\pi}(\lambda)$	4.1E-4	4.1E-4	4.1E-4	4.1E-4
TREE	4.3E-4	4.2E-4	4.2E-4	4.3E-4
IH	1.0E-1	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	1.1E-1	1.1E-1
WIS	7.3E-2	7.3E-2
NAIVE	5.4E-1	-

Table 190. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
			TITBRID	
	DIRECT	DR	WDR	MAGIC
AM	3.3E-2	3.9E-2	3.8E-2	2.5E-1
Q-Reg	8.1E-2	2.0E-3	7.1E-4	2.0E-1
MRDR	8.5E-2	1.5E-2	1.0E-2	2.1E-1
FQE	1.1E-9	1.1E-9	1.1E-9	2.0E-1
$R(\lambda)$	8.2E-9	5.3E-9	6.6E-9	2.0E-1
$Q^{\hat{\pi}}(\lambda)$	7.8E-10	1.1E-9	1.1E-9	2.0E-1
TREE	6.0E-6	5.9E-7	1.1E-6	2.0E-1
IH	1.2E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	7.4E-2	7.4E-2	
WIS	5.1E-2	5.1E-2	
NAIVE	6.0E-1	-	

Table 191. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

DM Hybrid DIRECT DR WDR MAGIC 2.2E-2 AM1.3E-2 1.2E-2 1.2E-2 Q-REG 3.0E-2 3.4E-4 2.1E-4 5.9E-4 2.4E-3MRDR 2.9E-21.1E-3 9.5E-3 6.9E-7 6.9E-7 FQE 6.9E-7 6.9E-7 $R(\lambda)$ 6.7E-7 6.8E-7 6.7E-7 6.7E-7 $Q^{\pi}(\lambda)$ 7.0E-7 6.9E-7 6.9E-7 7.0E-7 TREE 2.6E-69.0E-7 9.8E-7 2.4E-6 ΙH 1.2E-1

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	2.9E-2	2.9E-2		
WIS	1.9E-2	1.9E-2		
NAIVE	5.4E-1	-		

Table 192. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.7E-3	6.8E-3	6.9E-3	7.2E-3
Q-Reg	2.2E-2	1.0E-4	7.4E-5	9.7E-5
MRDR	2.3E-2	4.8E-4	5.5E-4	3.3E-3
FQE	4.9E-7	4.9E-7	4.9E-7	4.9E-7
$R(\lambda)$	5.1E-7	5.1E-7	5.1E-7	5.2E-7
$Q^{\hat{\pi}}(\lambda)$	4.9E-7	4.9E-7	4.9E-7	4.9E-7
TRÈE	9.3E-6	7.4E-7	7.6E-7	8.8E-6
IH	1.4E-2	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	2.4E-2	2.4E-2		
WIS	1.2E-2	1.2E-2		
NAIVE	4.6E-1	-		

Table 193. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.0E-3	5.5E-3	5.8E-3	4.0E-3
Q-Reg	2.5E-3	1.9E-5	2.0E-5	5.9E-4
MRDR	6.1E-3	5.4E-4	7.9E-4	2.0E-3
FQE	3.2E-5	3.2E-5	3.2E-5	3.2E-5
$R(\lambda)$	3.2E-5	3.2E-5	3.2E-5	3.2E-5
$Q^{\pi}(\lambda)$	3.2E-5	3.2E-5	3.2E-5	3.2E-5
TREE	1.1E-5	3.2E-5	3.2E-5	1.2E-5
IH	3.4E-3	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	2.8E-3	2.8E-3	
WIS	1.2E-3	1.2E-3	
NAIVE	4.6E-1	-	

Table 194. Graph, relative MSE. $T = 4, N = 1024, \pi_b(a = 0) = 0.6, \pi_e(a = 0) = 0.8$. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.1E-3	1.7E-3	1.6E-3	1.4E-3
Q-Reg	3.0E-3	6.3E-5	6.4E-5	5.8E-4
MRDR	4.6E-3	6.0E-4	8.6E-4	2.0E-3
FQE	6.3E-5	6.3E-5	6.3E-5	6.3E-5
$R(\lambda)$	6.3E-5	6.3E-5	6.3E-5	6.3E-5
$Q^{\pi}(\lambda)$	6.3E-5	6.3E-5	6.3E-5	6.3E-5
TREE	1.1E-4	6.2E-5	6.2E-5	1.0E-4
IH	2.2E-3	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	3.6E-3	3.6E-3	
WIS	1.5E-3	1.5E-3	
NAIVE	4.5E-1	-	

Table 195. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	3.3E0	3.9E0	3.1E0	2.8E0
Q-Reg	3.0E0	3.8E0	4.1E0	3.0E0
MRDR	1.5E0	2.5E0	3.0E0	1.5E0
FQE	2.8E0	3.8E0	3.8E0	2.8E0
$R(\lambda)$	4.3E0	4.6E0	4.6E0	4.3E0
$Q^{\hat{\pi}}(\lambda)$	4.6E0	4.8E0	4.8E0	4.6E0
TRÈE	3.7E0	4.1E0	4.1E0	3.7E0
IH	3.4E0	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	3.8E0	2.0E0	
WIS	3.5E0	2.4E0	
NAIVE	2.7E0	-	

Table 196. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.9E-1	2.2E0	1.9E0	4.7E-1
Q-Reg	3.8E-1	4.0E-1	3.7E-1	3.8E-1
MRDR	2.0E-1	1.8E-1	2.0E-1	2.0E-1
FQE	1.7E-1	1.9E-1	1.9E-1	1.7E-1
$R(\lambda)$	2.3E-1	2.3E-1	2.3E-1	2.3E-1
$Q^{\hat{\pi}}(\lambda)$	3.1E-1	2.6E-1	2.7E-1	3.1E-1
TRÈE	2.3E-1	2.2E-1	2.0E-1	2.3E-1
IH	4.1E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	5.3E-1	4.1E-1	
WIS	3.4E-1	2.5E-1	
NAIVE	5.1E-1	-	

Table 197. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.1E-1	6.6E-1	5.9E-1	4.2E-1
Q-Reg	5.2E-1	4.2E-1	3.8E-1	5.2E-1
MRDR	3.5E-1	4.1E-1	3.8E-1	3.4E-1
FQE	3.8E-1	3.2E-1	2.9E-1	3.8E-1
$R(\lambda)$	3.3E-1	3.2E-1	3.1E-1	3.3E-1
$Q^{\pi}(\lambda)$	3.2E-1	3.1E-1	3.0E-1	3.2E-1
TREE	3.7E-1	3.3E-1	3.2E-1	3.7E-1
IH	7.4E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	4.8E-1	4.2E-1	
WIS	4.7E-1	4.8E-1	
NAIVE	8.0E-1	-	

Table 198. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.7E-1	6.6E-1	6.9E-1	6.3E-1
Q-Reg	3.0E-1	2.5E-1	2.5E-1	3.1E-1
MRDR	1.3E-1	1.5E-1	1.5E-1	1.7E-1
FQE	3.5E-1	2.6E-1	2.5E-1	3.4E-1
$R(\lambda)$	2.6E-1	2.5E-1	2.4E-1	2.6E-1
$Q^{\hat{\pi}}(\lambda)$	2.8E-1	2.5E-1	2.4E-1	2.8E-1
TREE	2.7E-1	2.5E-1	2.5E-1	2.7E-1
IH	2.5E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	2.4E-1	1.8E-1	
WIS	3.4E-1	2.4E-1	
NAIVE	4.1E-1	-	

Table 199. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.4E-1	5.8E-1	5.6E-1	4.1E-1
Q-Reg	1.3E-1	9.3E-2	9.1E-2	1.1E-1
MRDR	9.6E-2	1.0E-1	1.0E-1	9.6E-2
FQE	5.4E-2	8.3E-2	8.6E-2	5.5E-2
$R(\lambda)$	7.1E-2	8.3E-2	8.5E-2	7.2E-2
$Q^{\hat{\pi}}(\lambda)$	7.2E-2	8.2E-2	8.5E-2	7.3E-2
TREE	6.7E-2	8.2E-2	8.4E-2	6.7E-2
IH	2.1E-1	-	-	-

		IPS		
		STANDARD	PER-DECISION	
IS		1.9E-1	1.2E-1	
WI	S	1.8E-1	1.1E-1	
NA	IVE	5.7E-1	-	

Table 200. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.5E-1	2.8E-1	2.8E-1	2.2E-1
Q-Reg	1.3E-1	1.4E-1	1.4E-1	1.4E-1
MRDR	1.6E-1	1.4E-1	1.4E-1	1.0E-1
FQE	1.2E-1	1.4E-1	1.4E-1	1.2E-1
$R(\lambda)$	1.4E-1	1.4E-1	1.4E-1	1.4E-1
$Q^{\hat{\pi}}(\lambda)$	1.5E-1	1.4E-1	1.4E-1	1.5E-1
TREE	1.4E-1	1.4E-1	1.4E-1	1.4E-1
IH	1.1E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	8.3E-2	1.4E-1	
WIS	7.7E-2	1.3E-1	
NAIVE	3.8E-1	-	

Table 201. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.0E-2	9.7E-2	9.6E-2	8.1E-2
Q-Reg	1.8E-2	1.9E-2	1.9E-2	1.9E-2
MRDR	1.4E-2	1.7E-2	1.7E-2	2.5E-2
FQE	1.6E-2	1.9E-2	1.9E-2	1.6E-2
$R(\lambda)$	1.8E-2	1.8E-2	1.9E-2	1.8E-2
$Q^{\pi}(\lambda)$	1.8E-2	1.9E-2	1.9E-2	1.8E-2
TREE	1.8E-2	1.8E-2	1.9E-2	1.8E-2
IH	1.7E-2	-	-	-

	IPS		
	STANDARD	Per-Decision	
IS	4.3E-2	1.8E-2	
WIS	5.0E-2	2.0E-2	
NAIVE	4.7E-1	-	

Table 202. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.7E-2	4.1E-2	4.1E-2	3.3E-2
Q-Reg	8.7E-3	1.0E-2	1.0E-2	9.4E-3
MRDR	1.0E-2	1.1E-2	1.2E-2	9.2E-3
FQE	9.3E-3	9.9E-3	1.0E-2	9.2E-3
$R(\lambda)$	9.8E-3	1.0E-2	1.0E-2	9.8E-3
$Q^{\pi}(\lambda)$	1.0E-2	1.0E-2	1.0E-2	1.0E-2
TREE	9.7E-3	9.9E-3	1.0E-2	9.7E-3
IH	9.7E-3	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	1.2E-2	7.8E-3		
WIS	1.3E-2	8.1E-3		
NAIVE	4.7E-1	-		

Table 203. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.5E0	3.5E0	3.2E0	2.3E0	
Q-Reg	3.3E0	2.1E0	2.2E0	2.3E0	
MRDR	1.7E0	1.6E0	1.8E0	1.7E0	
FQE	1.6E0	2.1E0	2.2E0	1.7E0	
$R(\lambda)$	2.3E0	2.6E0	2.6E0	2.4E0	
$Q^{\pi}(\lambda)$	2.5E0	2.7E0	2.7E0	2.5E0	
TREE	2.0E0	2.5E0	2.5E0	2.2E0	
IH	6.0E-1	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	4.2E0	4.2E0	
WIS	2.6E0	2.6E0	
NAIVE	2.1E0	-	

Table 204. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.7E0	2.4E0	2.3E0	2.2E0
Q-Reg	1.9E0	1.8E0	1.9E0	2.0E0
MRDR	1.4E0	9.2E-1	9.8E-1	1.3E0
FQE	1.8E0	1.9E0	1.8E0	2.0E0
$R(\lambda)$	2.1E0	2.2E0	2.1E0	2.2E0
$Q^{\hat{\pi}}(\lambda)$	2.1E0	2.3E0	2.2E0	2.2E0
TREE	2.0E0	2.0E0	2.0E0	2.2E0
IH	1.0E0	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	2.2E0	2.2E0	
WIS	1.6E0	1.6E0	
NAIVE	1.5E0	-	

Table 205. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.5E-1	1.3E0	1.5E0	5.7E-1
Q-Reg	7.7E-1	5.3E-1	5.2E-1	4.8E-1
MRDR	4.7E-1	5.4E-1	5.4E-1	3.3E-1
FQE	3.4E-1	3.9E-1	4.1E-1	3.3E-1
$R(\lambda)$	3.8E-1	4.4E-1	4.5E-1	3.5E-1
$Q^{\pi}(\lambda)$	4.9E-1	4.8E-1	5.0E-1	4.2E-1
TREE	3.4E-1	4.1E-1	4.2E-1	3.2E-1
IH	2.0E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	7.8E-1	7.8E-1	
WIS	5.8E-1	5.8E-1	
NAIVE	2.5E-1	-	

Table 206. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.7E-1	3.2E-1	3.3E-1	2.7E-1
Q-Reg	1.4E-1	2.1E-1	2.1E-1	2.1E-1
MRDR	9.0E-2	1.8E-1	2.1E-1	2.0E-1
FQE	4.2E-2	1.7E-1	1.8E-1	1.4E-1
$R(\lambda)$	1.2E-1	1.8E-1	1.9E-1	1.7E-1
$Q^{\pi}(\lambda)$	1.4E-1	1.8E-1	1.9E-1	1.8E-1
TREE	1.4E-1	1.8E-1	1.9E-1	1.9E-1
IH	6.5E-2	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	1.7E-1	1.7E-1		
WIS	1.6E-1	1.6E-1		
NAIVE	6.1E-1	-		

Table 207. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

DMHybrid DIRECT DR WDR MAGIC 2.0E-1 2.3E-1 2.3E-1 1.7E-1 AMQ-REG 1.1E-1 9.2E-2 9.1E-2 1.1E-1 MRDR 1.5E-1 1.1E-1 1.2E-1 1.2E-1 9.6E-2 9.6E-29.6E-2 FQE 9.5E-2 $R(\lambda)$ 1.0E-1 9.6E-2 9.6E-2 1.2E-1 $Q^{\pi}(\lambda)$ 1.1E-1 9.9E-2 9.8E-2 1.2E-1 TREE 1.2E-1 9.5E-2 9.5E-2 1.3E-1 ΙH 7.5E-2

	IPS		
	STANDARD PER-DECISIO		
IS	1.1E-1	1.1E-1	
WIS	1.0E-1	1.0E-1	
NAIVE	5.2E-1	-	

Table 208. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.3E-2	5.4E-2	5.3E-2	8.5E-2
Q-Reg	4.7E-2	3.1E-2	3.2E-2	5.4E-2
MRDR	5.4E-2	3.2E-2	3.3E-2	6.3E-2
FQE	3.5E-2	3.0E-2	3.1E-2	3.4E-2
$R(\lambda)$	3.4E-2	3.1E-2	3.0E-2	3.8E-2
$Q^{\hat{\pi}}(\lambda)$	3.2E-2	3.0E-2	3.1E-2	3.4E-2
TREE	3.7E-2	3.0E-2	3.0E-2	4.0E-2
IH	6.7E-2	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	4.5E-2	4.5E-2
WIS	4.6E-2	4.6E-2
NAIVE	6.2E-1	-

Table 209. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.4E-2	3.2E-2	3.2E-2	5.4E-2
Q-Reg	2.8E-2	1.7E-2	1.7E-2	2.6E-2
MRDR	2.5E-2	2.7E-2	2.9E-2	2.3E-2
FQE	1.4E-2	1.6E-2	1.7E-2	1.4E-2
$R(\lambda)$	1.6E-2	1.7E-2	1.7E-2	1.6E-2
$Q^{\hat{\pi}}(\hat{\lambda})$	1.6E-2	1.7E-2	1.7E-2	1.7E-2
TRÈE	1.6E-2	1.7E-2	1.7E-2	1.5E-2
IH	2.3E-2	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	2.7E-2	2.7E-2
WIS	2.4E-2	2.4E-2
NAIVE	5.2E-1	-

Table 210. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.0E-3	2.8E-2	2.8E-2	2.1E-2
Q-Reg	8.4E-3	8.3E-3	8.5E-3	7.9E-3
MRDR	1.3E-2	1.1E-2	1.2E-2	5.9E-3
FQE	9.3E-3	8.5E-3	8.6E-3	9.2E-3
$R(\lambda)$	7.9E-3	8.5E-3	8.6E-3	6.6E-3
$Q^{\hat{\pi}}(\lambda)$	8.5E-3	8.5E-3	8.6E-3	6.1E-3
TREE	7.7E-3	8.5E-3	8.6E-3	7.2E-3
IH	7.8E-3	-	-	-

		IPS	
	STANDARD PER-DECISION		
IS	9.0E-3	9.0E-3	
WIS	9.1E-3	9.1E-3	
NAIVE	5.1E-1	-	

Table 211. Graph, relative MSE. $T=4, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic environment. Stochastic rewards. Sparse rewards.

DM Hybrid DIRECT DR WDR MAGIC AM 2.4E1 2.5E1 2.4E1 1.7E1 Q-REG 1.7E1 2.1E1 2.0E1 1.7E1 MRDR 9.7E0 1.8E1 1.8E1 9.7E0 FQE 1.7E1 1.8E1 1.7E1 1.8E1 $R(\lambda)$ 1.8E1 1.9E1 1.8E1 1.8E1 $Q^{\pi}(\lambda)$ 1.7E1 1.8E1 1.8E1 1.7E1 TREE 1.9E1 1.9E1 1.8E1 1.9E1 ΙH 1.8E1

		IPS		
	STANDARD	STANDARD PER-DECISION		
IS	6.4E0	1.9E1		
WIS	6.9E0	1.9E1		
NAIVE	1.3E1	-		

Table 212. Graph, relative MSE. $T=4, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	5.6E0	9.8E0	1.1E1	5.8E0
Q-Reg	4.7E0	3.3E0	3.3E0	4.7E0
MRDR	3.1E0	3.0E0	3.0E0	3.1E0
FQE	2.1E0	3.7E0	3.6E0	2.1E0
$R(\lambda)$	2.7E0	3.4E0	3.3E0	2.7E0
$Q^{\hat{\pi}}(\lambda)$	3.6E0	4.0E0	3.8E0	3.6E0
TRÈE	2.4E0	3.2E0	3.0E0	2.4E0
IH	3.1E0	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	5.8E0	5.3E0
WIS	4.8E0	4.9E0
NAIVE	2.6E0	-

Table 213. Graph, relative MSE. $T=4, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	3.4E0	7.3E0	6.8E0	5.6E0
Q-Reg	1.9E0	1.5E0	1.5E0	1.9E0
MRDR	1.2E0	1.5E0	1.5E0	1.2E0
FQE	1.1E0	1.3E0	1.3E0	1.2E0
$R(\lambda)$	1.3E0	1.4E0	1.4E0	1.3E0
$Q^{\hat{\pi}}(\lambda)$	1.4E0	1.4E0	1.4E0	1.4E0
TREE	1.4E0	1.4E0	1.4E0	1.4E0
IH	7.4E-1	-	-	-

		IPS		
	STANDARD	PER-DECISION		
IS	1.7E0	1.7E0		
WIS	1.5E0	1.3E0		
NAIVE	1.2E0	-		

Table 214. Graph, relative MSE. $T=4, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	9.7E-1	1.8E0	1.7E0	1.1E0
Q-Reg	1.6E0	1.9E0	1.9E0	1.1E0
MRDR	8.0E-1	1.8E0	1.9E0	7.7E-1
FQE	5.3E-1	1.9E0	1.9E0	5.6E-1
$R(\lambda)$	1.5E0	1.9E0	2.0E0	1.0E0
$Q^{\pi}(\lambda)$	1.4E0	1.9E0	1.9E0	1.0E0
TREE	1.4E0	2.0E0	2.0E0	1.2E0
IH	4.9E-1	-	-	-

		IPS		
	STANDARD	Per-Decision		
IS	1.6E0	1.7E0		
WIS	1.8E0	1.8E0		
NAIVE	2.0E-1	-		

Table 215. Graph, relative MSE. $T=4, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.1E-1	1.2E0	1.2E0	7.7E-1
Q-Reg	3.5E-1	4.5E-1	4.4E-1	2.9E-1
MRDR	2.5E-1	3.9E-1	4.0E-1	2.7E-1
FQE	2.1E-1	4.8E-1	4.6E-1	2.2E-1
$R(\lambda)$	3.8E-1	4.7E-1	4.5E-1	3.6E-1
$Q^{\hat{\pi}}(\lambda)$	2.8E-1	4.6E-1	4.4E-1	3.2E-1
TREE	4.5E-1	4.7E-1	4.6E-1	4.1E-1
IH	1.3E-1	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	4.8E-1	3.3E-1		
WIS	5.1E-1	3.6E-1		
NAIVE	5.5E-1	-		

Table 216. Graph, relative MSE. $T=4, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.9E-1	5.4E-1	4.9E-1	2.2E-1
Q-Reg	2.8E-1	2.7E-1	2.6E-1	2.2E-1
MRDR	2.0E-1	2.8E-1	2.8E-1	1.9E-1
FQE	2.0E-1	3.0E-1	2.8E-1	2.0E-1
$R(\lambda)$	2.3E-1	2.8E-1	2.8E-1	1.9E-1
$Q^{\hat{\pi}}(\lambda)$	2.6E-1	2.9E-1	2.8E-1	2.0E-1
TREE	2.2E-1	2.8E-1	2.8E-1	2.0E-1
IH	1.8E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	3.3E-1	2.9E-1	
WIS	3.2E-1	2.8E-1	
NAIVE	6.2E-1	-	

Table 217. Graph, relative MSE. $T=4, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.7E-1	3.9E-1	4.0E-1	4.5E-1
Q-Reg	2.5E-1	2.2E-1	2.2E-1	2.2E-1
MRDR	2.1E-1	2.3E-1	2.3E-1	1.9E-1
FQE	1.4E-1	2.1E-1	2.1E-1	1.4E-1
$R(\lambda)$	1.9E-1	2.1E-1	2.1E-1	1.8E-1
$Q^{\hat{\pi}}(\hat{\lambda})$	2.1E-1	2.1E-1	2.1E-1	1.7E-1
TREE	2.0E-1	2.1E-1	2.1E-1	1.9E-1
IH	1.6E-1	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	2.8E-1	2.3E-1		
WIS	2.7E-1	2.3E-1		
NAIVE	3.8E-1	-		

Table 218. Graph, relative MSE. $T=4, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.2E-1	1.1E-1	1.2E-1	7.1E-2
Q-Reg	6.4E-2	6.6E-2	6.6E-2	6.0E-2
MRDR	5.1E-2	6.6E-2	6.7E-2	4.2E-2
FQE	3.6E-2	6.5E-2	6.6E-2	3.6E-2
$R(\lambda)$	6.1E-2	6.6E-2	6.6E-2	6.2E-2
$Q^{\pi}(\lambda)$	6.3E-2	6.7E-2	6.7E-2	6.0E-2
TREE	6.6E-2	6.5E-2	6.5E-2	6.8E-2
IH	2.6E-2	-	-	-

		IPS		
	STANDARD	PER-DECISION		
IS	9.3E-2	6.6E-2		
WIS	8.9E-2	6.2E-2		
NAIVE	4.8E-1	-		

Table 219. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.1E-1	2.2E-1	1.8E-1	1.0E-1
Q-Reg	4.6E-1	2.0E-1	2.3E-1	2.8E-1
MRDR	3.4E-1	1.3E0	8.6E-1	8.1E-1
FQE	1.1E-1	1.1E-1	1.1E-1	1.1E-1
$R(\lambda)$	1.1E-1	1.1E-1	1.1E-1	1.1E-1
$Q^{\hat{\pi}}(\lambda)$	1.1E-1	1.1E-1	1.1E-1	1.1E-1
TRÈE	1.8E-1	1.4E-1	1.2E-1	1.7E-1
IH	3.1E-2	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.1E0	3.0E-1	
WIS	1.3E-1	9.6E-2	
NAIVE	6.2E-1	-	

Table 220. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.3E-2	5.4E-2	5.8E-2	9.8E-3
Q-Reg	1.7E-1	1.7E-1	1.3E-1	1.0E-1
MRDR	1.7E-1	8.9E-1	4.9E-1	2.9E-1
FQE	6.1E-3	6.1E-3	6.1E-3	6.1E-3
$R(\lambda)$	6.0E-3	6.0E-3	6.0E-3	6.0E-3
$Q^{\pi}(\lambda)$	6.1E-3	6.1E-3	6.1E-3	6.1E-3
TREE	7.9E-2	2.8E-2	1.9E-2	4.6E-2
IH	8.8E-3	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	4.7E-1	1.9E-1		
WIS	7.5E-2	2.4E-2		
NAIVE	4.6E-1	-		

Table 221. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.0E-3	4.1E-2	2.7E-2	1.1E-3
Q-Reg	7.7E-2	8.8E-3	2.6E-3	4.8E-2
MRDR	7.4E-2	2.4E-1	1.4E-1	1.1E-1
FQE	5.0E-6	5.0E-6	5.0E-6	5.0E-6
$R(\lambda)$	7.0E-6	7.0E-6	7.0E-6	7.0E-6
$Q^{\hat{\pi}}(\lambda)$	5.0E-6	5.0E-6	5.0E-6	5.0E-6
TREE	5.6E-2	9.1E-3	6.5E-3	1.4E-2
IH	4.1E-3	-	-	-

	IPS		
	STANDARD PER-DECISIO		
IS	4.2E-1	1.1E-1	
WIS	3.1E-2	6.7E-3	
NAIVE	4.1E-1	-	

Table 222. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.3E-3	2.1E-2	2.3E-2	6.9E-4
Q-Reg	3.9E-2	5.4E-3	5.3E-3	1.9E-2
MRDR	5.3E-2	1.5E-1	1.2E-1	5.3E-2
FQE	5.5E-8	5.5E-8	5.5E-8	5.5E-8
$R(\lambda)$	1.9E-7	1.6E-7	1.5E-7	1.9E-7
$Q^{\hat{\pi}}(\lambda)$	5.6E-8	5.5E-8	5.5E-8	5.6E-8
TREE	5.9E-2	7.3E-3	3.6E-3	3.6E-3
IH	1.7E-3	-	-	-

		IPS		
	STANDARD	STANDARD PER-DECISION		
IS	1.4E-1	4.7E-2		
WIS	1.3E-2	4.3E-3		
NAIVE	4.7E-1	-		

Table 223. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.6E-4	7.1E-3	6.6E-3	2.6E-4
Q-Reg	1.6E-2	3.6E-3	2.3E-3	2.7E-3
MRDR	2.0E-2	1.8E-2	1.5E-2	2.8E-2
FQE	3.0E-6	3.0E-6	3.0E-6	3.0E-6
$R(\lambda)$	3.0E-6	3.0E-6	3.0E-6	3.0E-6
$Q^{\hat{\pi}}(\lambda)$	3.0E-6	3.0E-6	3.0E-6	3.0E-6
TREE	4.8E-2	2.6E-3	1.2E-3	1.2E-3
IH	6.7E-4	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	5.7E-2	2.0E-2	
WIS	6.0E-3	1.8E-3	
NAIVE	4.6E-1	-	

Table 224. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.8E-4	1.5E-3	1.6E-3	6.5E-4
Q-Reg	1.8E-2	2.3E-3	9.4E-4	9.9E-4
MRDR	1.4E-2	3.7E-2	2.5E-2	1.7E-2
FQE	1.0E-6	1.0E-6	1.0E-6	1.0E-6
$R(\lambda)$	1.0E-6	1.0E-6	1.0E-6	1.0E-6
$Q^{\hat{\pi}}(\lambda)$	1.0E-6	1.0E-6	1.0E-6	1.0E-6
TRÈE	3.0E-3	2.6E-4	2.4E-4	2.4E-4
IH	3.4E-4	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	2.6E-1	2.0E-2		
WIS	1.6E-2	1.8E-3		
NAIVE	4.5E-1	-		

Table 225. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.0E-5	3.3E-3	3.3E-3	1.5E-3
Q-Reg	4.5E-3	7.7E-5	3.3E-5	4.5E-5
MRDR	5.0E-3	2.6E-3	2.7E-3	5.7E-3
FQE	7.7E-7	7.7E-7	7.7E-7	7.7E-7
$R(\lambda)$	8.1E-7	8.0E-7	7.9E-7	8.1E-7
$Q^{\pi}(\lambda)$	7.6E-7	7.7E-7	7.7E-7	7.6E-7
TREE	2.2E-3	1.2E-4	1.1E-4	1.1E-4
IH	1.7E-4	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	4.0E-2	4.7E-3		
WIS	4.0E-3	6.0E-4		
NAIVE	4.5E-1	-		

Table 226. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.8E-5	4.7E-4	4.8E-4	7.2E-5
Q-Reg	2.1E-3	1.4E-5	1.0E-5	2.9E-5
MRDR	2.9E-3	5.0E-4	4.6E-4	9.6E-4
FQE	2.7E-7	2.7E-7	2.7E-7	2.7E-7
$R(\lambda)$	2.9E-7	2.8E-7	2.8E-7	2.9E-7
$Q^{\hat{\pi}}(\lambda)$	2.7E-7	2.7E-7	2.7E-7	2.7E-7
TREE	1.8E-3	7.7E-5	5.8E-5	5.8E-5
IH	9.6E-5	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.7E-2	2.1E-3	
WIS	1.8E-3	2.9E-4	
NAIVE	4.4E-1	-	

Table 227. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.5E-1	5.0E-1	4.3E-1	2.1E-1
Q-REG	4.3E-1	3.1E-1	2.6E-1	4.5E-1
MRDR	2.8E-1	1.1E0	4.5E-1	3.0E-1
FQE	1.3E-1	1.4E-1	1.3E-1	1.3E-1
$R(\lambda)$	1.4E-1	1.5E-1	1.4E-1	1.4E-1
$Q^{\hat{\pi}}(\lambda)$	1.5E-1	1.6E-1	1.5E-1	1.5E-1
TREE	2.1E-1	1.2E-1	7.0E-2	2.1E-1
IH	4.2E-2	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	4.9E-1	3.5E-1	
WIS	1.1E-1	6.0E-2	
NAIVE	4.4E-1	-	

Table 228. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.8E-2	3.2E-1	1.6E-1	7.6E-2
Q-Reg	3.4E-1	4.3E-1	1.1E-1	3.4E-1
MRDR	2.4E-1	2.3E0	6.4E-1	2.4E-1
FQE	1.4E-2	3.6E-2	1.0E-2	1.4E-2
$R(\lambda)$	1.6E-2	1.4E-2	1.6E-2	1.6E-2
$Q^{\hat{\pi}}(\lambda)$	2.8E-2	2.3E-2	1.9E-2	2.8E-2
TRÈE	6.7E-2	3.2E-2	1.7E-2	3.9E-2
IH	1.5E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.6E1	4.5E-1	
WIS	1.5E-1	3.2E-2	
NAIVE	4.3E-1	-	

Table 229. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.0E-2	8.5E-2	7.5E-2	1.1E-2
Q-Reg	7.3E-2	4.8E-2	4.6E-2	6.4E-2
MRDR	1.1E-1	2.0E-1	1.8E-1	1.3E-1
FQE	9.8E-3	2.6E-2	2.8E-2	9.8E-3
$R(\lambda)$	2.9E-2	3.5E-2	3.6E-2	2.9E-2
$Q^{\hat{\pi}}(\lambda)$	2.9E-2	4.4E-2	4.1E-2	2.9E-2
TREE	1.1E-1	4.7E-2	3.3E-2	4.4E-2
IH	8.4E-3	-	-	-

	IPS		
	STANDARD PER-DECISI		
IS	3.0E-1	6.2E-2	
WIS	5.4E-2	3.6E-2	
NAIVE	4.4E-1	-	

Table 230. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.1E-3	3.3E-2	2.1E-2	4.3E-3
Q-Reg	4.7E-2	1.4E-2	1.1E-2	3.3E-2
MRDR	5.8E-2	4.7E-2	4.1E-2	7.2E-2
FQE	4.1E-3	1.3E-2	1.2E-2	4.2E-3
$R(\lambda)$	5.8E-3	1.1E-2	1.1E-2	5.9E-3
$Q^{\hat{\pi}}(\lambda)$	6.7E-3	1.3E-2	1.2E-2	6.7E-3
TREE	5.5E-2	2.2E-2	1.9E-2	3.6E-2
IH	7.8E-3	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.7E-1	6.4E-2	
WIS	6.9E-2	2.4E-2	
NAIVE	4.1E-1	-	

Table 231. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

•	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.6E-3	2.6E-2	2.7E-2	6.3E-3
Q-Reg	1.5E-2	5.1E-3	5.4E-3	6.8E-3
MRDR	2.0E-2	2.8E-2	2.5E-2	2.5E-2
FQE	2.1E-3	4.9E-3	4.9E-3	2.1E-3
$R(\lambda)$	3.5E-3	4.3E-3	4.5E-3	3.5E-3
$Q^{\pi}(\lambda)$	4.7E-3	5.0E-3	4.8E-3	4.6E-3
TREE	4.7E-2	4.9E-3	5.1E-3	8.9E-3
IH	3.3E-3	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	7.6E-2	2.3E-2	
WIS	9.2E-3	4.5E-3	
NAIVE	4.5E-1	-	

Table 232. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.9E-3	1.2E-2	1.3E-2	4.5E-3
Q-Reg	3.0E-2	3.1E-3	3.4E-3	6.3E-3
MRDR	2.7E-2	1.4E-2	1.2E-2	1.4E-2
FQE	1.1E-3	4.3E-3	4.1E-3	1.1E-3
$R(\lambda)$	2.2E-3	3.7E-3	3.8E-3	2.3E-3
$Q^{\hat{\pi}}(\lambda)$	1.2E-3	3.8E-3	3.9E-3	1.3E-3
TRÈE	5.5E-3	4.2E-3	4.0E-3	5.4E-3
IH	2.1E-3	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.3E-1	3.5E-2	
WIS	1.8E-2	7.3E-3	
NAIVE	4.3E-1	-	

Table 233. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.6E-4	5.6E-3	5.2E-3	3.6E-3
Q-Reg	7.0E-3	8.2E-4	8.1E-4	9.0E-4
MRDR	4.1E-3	4.5E-3	4.3E-3	6.7E-3
FQE	5.5E-4	9.0E-4	8.6E-4	5.6E-4
$R(\lambda)$	6.5E-4	8.1E-4	8.0E-4	6.6E-4
$Q^{\pi}(\lambda)$	7.3E-4	8.8E-4	8.3E-4	7.1E-4
TREE	3.0E-3	7.8E-4	8.4E-4	1.8E-3
IH	7.5E-4	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	3.0E-2	8.1E-3	
WIS	5.7E-3	1.3E-3	
NAIVE	4.5E-1	-	

Table 234. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.9E-4	3.4E-3	3.7E-3	4.2E-4
Q-Reg	1.8E-3	9.9E-4	1.0E-3	1.6E-3
MRDR	1.8E-3	1.1E-3	1.2E-3	1.4E-3
FQE	2.9E-4	9.4E-4	9.8E-4	3.1E-4
$R(\lambda)$	6.8E-4	9.7E-4	9.9E-4	6.9E-4
$Q^{\pi}(\lambda)$	5.4E-4	9.7E-4	1.0E-3	5.5E-4
TREE	1.9E-3	8.3E-4	8.9E-4	3.1E-4
IH	1.7E-4	-	-	-

		IPS		
	STANDARD	PER-DECISION		
IS	1.0E-2	1.1E-3		
WIS	3.7E-3	6.1E-4		
NAIVE	4.4E-1	-		

Table 235. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.5E-1	7.2E-1	4.7E-1	1.5E-1
Q-Reg	5.4E-1	5.1E-1	4.9E-1	5.4E-1
MRDR	4.1E-1	1.2E0	6.6E-1	4.2E-1
FQE	1.7E-1	3.8E-1	2.2E-1	1.7E-1
$R(\lambda)$	1.7E-1	1.7E-1	1.7E-1	1.7E-1
$Q^{\hat{\pi}}(\lambda)$	1.9E-1	2.6E-1	2.0E-1	1.9E-1
TRÈE	2.4E-1	2.1E-1	2.1E-1	2.3E-1
IH	1.5E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	8.9E-1	4.6E-1	
WIS	5.1E-1	2.9E-1	
NAIVE	5.8E-1	-	

Table 236. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

•	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.2E-2	6.0E-1	4.5E-1	3.4E-2
Q-Reg	1.8E-1	1.5E-1	2.2E-1	1.8E-1
MRDR	1.3E-1	1.2E0	1.1E0	2.3E-1
FQE	2.3E-2	8.7E-2	7.4E-2	2.3E-2
$R(\lambda)$	3.2E-2	3.9E-2	4.3E-2	3.2E-2
$Q^{\hat{\pi}}(\lambda)$	6.0E-2	5.8E-2	6.2E-2	5.9E-2
TRÈE	1.0E-1	9.5E-2	8.1E-2	8.7E-2
IH	3.2E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	4.8E-1	1.6E-1	
WIS	1.4E-1	8.3E-2	
NAIVE	3.2E-1	-	

Table 237. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.1E-2	9.7E-2	1.4E-1	3.5E-2
Q-Reg	1.4E-1	1.0E-1	1.0E-1	1.4E-1
MRDR	1.4E-1	4.5E-1	3.6E-1	1.7E-1
FQE	2.0E-2	7.2E-2	5.7E-2	2.0E-2
$R(\lambda)$	4.0E-2	3.8E-2	3.9E-2	4.1E-2
$Q^{\hat{\pi}}(\lambda)$	3.7E-2	6.4E-2	4.3E-2	3.7E-2
TREE	5.9E-2	5.1E-2	5.2E-2	6.2E-2
IH	1.2E-2	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	6.5E-1	1.4E-1		
WIS	7.6E-2	7.5E-2		
NAIVE	4.2E-1	-		

Table 238. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.5E-2	9.3E-2	1.2E-1	2.1E-2	
Q-REG	6.1E-2	4.5E-2	4.5E-2	6.4E-2	
MRDR	7.1E-2	7.7E-2	6.9E-2	9.1E-2	
FQE	1.3E-2	3.3E-2	4.1E-2	1.3E-2	
$R(\lambda)$	2.9E-2	3.4E-2	3.9E-2	3.1E-2	
$Q^{\hat{\pi}}(\lambda)$	3.0E-2	2.9E-2	3.8E-2	3.0E-2	
TREE	6.5E-2	3.6E-2	5.1E-2	6.3E-2	
IH	1.0E-2	-	-	-	

	IPS				
	STANDARD PER-DECISIO				
IS	1.5E-1	5.7E-2			
WIS	7.1E-2	4.5E-2			
NAIVE	4.5E-1	-			

Table 239. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic environment. Dense rewards.

DMHybrid DIRECT DR WDR MAGIC 4.5E-2 3.7E-2 1.0E-2 1.3E-2 AMQ-REG 2.9E-2 1.1E-2 1.2E-21.9E-2 MRDR 4.1E-2 2.9E-2 2.5E-2 5.2E-2 FQE 6.5E-39.9E-3 9.6E-3 6.7E-3 $R(\lambda)$ 6.7E-39.3E-3 9.6E-3 7.3E-3 $Q^{\dot{\pi}}(\dot{\lambda})$ 8.7E-3 8.5E-3 9.2E-3 5.9E-3 TREE 5.4E-2 1.2E-2 1.1E-2 3.8E-2 ΙH 8.9E-3

	IPS			
	STANDARD PER-DECISION			
IS	1.6E-1	2.8E-2		
WIS	6.2E-2	1.2E-2		
NAIVE	4.5E-1	-		

Table 240. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM	Hybrid		
	DIRECT	DR	WDR	MAGIC
AM	4.0E-3	4.6E-2	5.3E-2	3.6E-2
Q-Reg	1.3E-2	6.0E-3	6.2E-3	8.9E-3
MRDR	2.8E-2	1.7E-2	1.8E-2	3.0E-2
FQE	2.5E-3	7.1E-3	7.4E-3	2.5E-3
$R(\lambda)$	7.0E-3	7.1E-3	7.0E-3	7.0E-3
$Q^{\hat{\pi}}(\lambda)$	3.8E-3	7.7E-3	7.9E-3	3.8E-3
TREE	3.4E-2	1.1E-2	1.1E-2	1.9E-2
IH	3.7E-3	-	-	-

	IPS			
	STANDARD PER-DECI			
IS	5.2E-2	1.5E-2		
WIS	1.7E-2	1.2E-2		
NAIVE	4.8E-1	-		

Table 241. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic environment. Dense rewards

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	8.8E-4	1.2E-2	1.2E-2	1.7E-3	
Q-Reg	4.5E-3	5.9E-3	5.9E-3	5.7E-3	
MRDR	4.1E-3	6.1E-3	6.2E-3	8.0E-3	
FQE	1.1E-3	6.8E-3	6.4E-3	1.2E-3	
$R(\lambda)$	3.3E-3	5.9E-3	5.7E-3	3.3E-3	
$Q^{\hat{\pi}}(\lambda)$	3.7E-3	7.1E-3	6.8E-3	3.8E-3	
TREE	1.2E-2	5.7E-3	5.7E-3	1.2E-2	
IH	9.3E-4	-	-	-	

	IPS			
	STANDARD PER-DECISI			
IS	2.6E-2	5.3E-3		
WIS	1.3E-2	6.3E-3		
NAIVE	4.3E-1	-		

Table 242. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.1E-3	8.3E-3	8.7E-3	2.1E-3	
Q-Reg	3.8E-3	1.6E-3	1.6E-3	2.1E-3	
MRDR	3.1E-3	1.6E-3	1.5E-3	3.0E-3	
FQE	7.7E-4	1.7E-3	1.7E-3	7.5E-4	
$R(\lambda)$	1.6E-3	1.8E-3	1.8E-3	1.6E-3	
$Q^{\pi}(\lambda)$	9.0E-4	1.4E-3	1.5E-3	1.0E-3	
TREE	2.8E-3	1.9E-3	2.0E-3	9.0E-4	
IH	9.4E-4	-	-	-	

	IPS			
	STANDARD PER-DECIS			
IS	1.9E-2	5.0E-3		
WIS	5.4E-3	1.7E-3		
NAIVE	4.3E-1	-		

Table 243. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.4E-1	1.3E0	1.2E0	4.2E-1	
Q-Reg	6.3E-1	4.7E-1	3.6E-1	6.3E-1	
MRDR	4.6E-1	2.6E0	7.8E-1	4.6E-1	
FQE	3.6E-1	4.3E-1	4.6E-1	3.6E-1	
$R(\lambda)$	4.5E-1	4.7E-1	4.5E-1	4.5E-1	
$Q^{\hat{\pi}}(\lambda)$	5.8E-1	6.4E-1	6.1E-1	5.8E-1	
TREE	4.7E-1	3.9E-1	3.6E-1	4.7E-1	
IH	2.3E-1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	7.5E-1		
WIS	6.5E-1	4.4E-1		
NAIVE	5.2E-1	-		

Table 244. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.3E-1	8.6E-1	6.6E-1	1.7E-1	
Q-Reg	1.7E-1	4.6E-1	5.4E-1	2.8E-1	
MRDR	2.3E-1	1.1E0	1.3E0	2.3E-1	
FQE	1.5E-1	5.2E-1	3.2E-1	1.5E-1	
$R(\lambda)$	2.5E-1	2.8E-1	2.6E-1	2.5E-1	
$Q^{\pi}(\lambda)$	3.0E-1	5.6E-1	4.7E-1	3.0E-1	
TREE	3.1E-1	2.3E-1	2.1E-1	3.1E-1	
IH	9.8E-2	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	1.5E0	1.8E-1	
WIS	5.1E-1	1.5E-1	
NAIVE	6.2E-1	-	

Table 245. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.8E-2	7.0E-1	3.6E-1	4.9E-2
Q-Reg	3.3E-1	8.7E-2	9.3E-2	2.5E-1
MRDR	2.5E-1	3.5E-1	3.2E-1	3.6E-1
FQE	3.0E-2	2.5E-1	1.6E-1	3.0E-2
$R(\lambda)$	7.0E-2	1.1E-1	9.1E-2	7.0E-2
$Q^{\hat{\pi}}(\lambda)$	7.4E-2	3.0E-1	2.0E-1	7.4E-2
TREE	1.3E-1	8.5E-2	7.5E-2	1.3E-1
IH	5.5E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.7E0	3.5E-1	
WIS	5.4E-1	1.4E-1	
NAIVE	4.2E-1	-	

Table 246. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.6E-2	2.9E-1	2.8E-1	3.4E-2
Q-Reg	6.5E-2	3.7E-2	3.3E-2	4.4E-2
MRDR	7.7E-2	1.6E-1	1.2E-1	8.8E-2
FQE	1.6E-2	2.7E-2	2.1E-2	1.6E-2
$R(\lambda)$	2.0E-2	2.6E-2	2.4E-2	2.0E-2
$Q^{\hat{\pi}}(\lambda)$	2.3E-2	3.8E-2	2.9E-2	2.0E-2
TREE	8.8E-2	2.4E-2	2.5E-2	8.6E-2
IH	1.4E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	6.6E-1	6.9E-2	
WIS	1.3E-1	2.2E-2	
NAIVE	4.8E-1	-	

Table 247. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic environment. Stochastic rewards. Dense rewards.

DM HybridDIRECT DR WDR MAGIC 5.4E-3 6.0E-3 1.2E-1 1.2E-1 AM 1.0E-2 Q-REG 1.2E-1 1.2E-2 2.6E-2 MRDR 5.7E-2 2.0E-1 1.2E-1 5.5E-2 FQE 2.6E-3 5.2E-2 3.1E-2 2.8E-3 $R(\lambda)$ 1.2E-2 2.6E-2 2.0E-21.2E-2 $Q^{\hat{\pi}}(\hat{\lambda})$ 1.1E-2 2.6E-2 4.4E-21.1E-2 TREE 4.5E-2 3.6E-22.8E-23.6E-2ΙH 4.1E-3

	IPS		
	STANDARD	PER-DECISION	
IS	3.0E-1	1.1E-1	
WIS	4.4E-2	4.1E-2	
NAIVE	4.5E-1	-	

Table 248. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.6E-3	4.6E-2	4.5E-2	6.2E-3
Q-Reg	8.3E-3	1.4E-2	1.5E-2	1.1E-2
MRDR	1.3E-2	2.5E-2	2.3E-2	9.2E-3
FQE	3.8E-3	1.1E-2	1.3E-2	3.5E-3
$R(\lambda)$	9.0E-3	1.3E-2	1.3E-2	9.5E-3
$Q^{\hat{\pi}}(\hat{\lambda})$	6.3E-3	1.4E-2	1.6E-2	6.4E-3
TREE	3.1E-2	1.5E-2	1.5E-2	1.9E-2
IH	3.5E-3	_	_	_

	IPS		
	STANDARD	PER-DECISION	
IS	1.4E-1	1.1E-2	
WIS	6.6E-2	1.2E-2	
NAIVE	4.4E-1	-	

Table 249. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.3E-3	3.0E-2	3.2E-2	9.4E-3
Q-Reg	1.4E-2	1.4E-2	1.3E-2	1.1E-2
MRDR	8.3E-3	6.6E-3	6.3E-3	9.8E-3
FQE	5.0E-3	1.2E-2	1.2E-2	5.0E-3
$R(\lambda)$	8.8E-3	1.3E-2	1.3E-2	8.1E-3
$Q^{\hat{\pi}}(\lambda)$	1.3E-2	1.3E-2	1.3E-2	1.0E-2
TREE	9.4E-3	1.3E-2	1.3E-2	1.0E-2
IH	4.2E-3	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E-1	1.5E-2	
WIS	2.1E-2	1.2E-2	
NAIVE	4.2E-1	-	

Table 250. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.3E-3	1.1E-2	1.1E-2	3.8E-3
Q-Reg	6.0E-3	5.0E-3	5.0E-3	5.6E-3
MRDR	5.2E-3	3.4E-3	3.8E-3	4.9E-3
FQE	1.1E-3	5.3E-3	5.6E-3	1.1E-3
$R(\lambda)$	2.7E-3	5.1E-3	5.2E-3	2.4E-3
$Q^{\pi}(\lambda)$	2.0E-3	5.0E-3	5.3E-3	1.5E-3
TREE	1.3E-2	5.0E-3	5.4E-3	8.8E-3
IH	9.8E-4	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	9.2E-3	5.4E-3	
WIS	7.3E-3	5.4E-3	
NAIVE	4.6E-1	-	

Table 251. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.5E-1	1.3E0	3.6E0	5.7E-1
Q-Reg	1.7E0	3.2E0	5.5E0	1.7E0
MRDR	8.7E-1	6.3E0	5.8E0	9.5E-1
FQE	4.3E-1	4.3E-1	4.3E-1	4.8E-1
$R(\lambda)$	4.6E-1	4.5E-1	4.4E-1	5.1E-1
$Q^{\hat{\pi}}(\lambda)$	4.3E-1	4.3E-1	4.3E-1	4.8E-1
TRÈE	9.8E-1	7.1E-1	6.6E-1	8.5E-1
IH	6.9E-1	-	-	-

	IPS		
	STANDARD	Per-Decision	
IS	1.0E0	1.0E0	
WIS	7.0E-1	7.0E-1	
NAIVE	2.6E-1	-	

Table 252. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.2E-1	2.0E0	1.3E0	2.2E-1
Q-Reg	5.6E-1	1.8E0	4.8E-1	7.8E-1
MRDR	4.4E-1	2.2E0	1.9E0	4.2E-1
FQE	4.6E-2	4.6E-2	4.6E-2	1.4E-1
$R(\lambda)$	4.6E-2	4.6E-2	4.5E-2	1.4E-1
$Q^{\hat{\pi}}(\lambda)$	4.6E-2	4.6E-2	4.6E-2	1.4E-1
TREE	9.9E-1	5.7E-1	1.6E-1	8.7E-1
IH	6.2E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	6.2E-1	6.2E-1	
WIS	1.7E-1	1.7E-1	
NAIVE	3.5E-1	-	

Table 253. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.4E-2	5.0E-1	3.4E-1	1.2E-1
Q-Reg	1.6E0	5.3E-1	1.7E-1	6.3E-1
MRDR	1.1E0	8.7E0	3.6E0	1.8E0
FQE	3.6E-3	3.6E-3	3.6E-3	1.0E-1
$R(\lambda)$	4.6E-3	4.5E-3	4.5E-3	1.0E-1
$Q^{\pi}(\lambda)$	3.6E-3	3.6E-3	3.6E-3	1.0E-1
TREE	9.1E-1	8.5E-1	1.0E-1	5.3E-1
IH	6.1E-1	-	-	-

	IPS			
	STANDARD PER-DECISIO			
IS	1.9E0	1.9E0		
WIS	1.6E-1	1.6E-1		
NAIVE	3.3E-1	-		

Table 254. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.3E-2	3.2E-1	3.4E-1	1.3E-1
Q-Reg	1.8E-1	1.1E-1	9.6E-2	2.1E-1
MRDR	1.8E-1	7.2E-1	8.0E-1	4.9E-1
FQE	2.0E-6	2.0E-6	2.0E-6	1.0E-1
$R(\lambda)$	4.4E-5	3.2E-5	4.2E-5	1.0E-1
$Q^{\hat{\pi}}(\lambda)$	2.0E-6	2.0E-6	2.0E-6	1.0E-1
TREE	9.7E-1	1.8E-1	9.3E-2	6.1E-1
IH	6.2E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.8E-1	1.8E-1		
WIS	9.1E-2	9.1E-2		
NAIVE	4.7E-1	-		

Table 255. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

DMHybrid DIRECT DR WDR MAGIC 8.8E-3 6.9E-1 6.1E-1 3.1E-1 AMQ-REG 1.7E-1 4.7E-2 2.4E-21.1E-1 MRDR 1.5E-1 4.0E-1 2.9E-1 3.1E-1 2.2E-7 2.2E-7 FQE 2.2E-7 2.2E-7 $R(\lambda)$ 1.2E-5 9.7E-6 1.0E-5 1.2E-5 $Q^{\pi}(\lambda)$ 2.3E-7 2.2E-7 2.2E-7 2.3E-7 TREE 9.4E-1 1.6E-1 7.0E-21.9E-1 ΙH 6.1E-1

	IPS			
	STANDARD PER-DECISION			
IS	2.2E-1	2.2E-1		
WIS	7.6E-2	7.6E-2		
NAIVE	5.7E-1	-		

Table 256. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.5E-3	1.7E-1	1.6E-1	9.2E-2
Q-Reg	2.2E-1	4.6E-3	2.1E-3	2.1E-2
MRDR	1.7E-1	1.5E-1	1.2E-1	2.2E-1
FQE	9.0E-6	9.0E-6	9.0E-6	9.0E-6
$R(\lambda)$	2.3E-5	2.3E-5	2.6E-5	2.3E-5
$Q^{\hat{\pi}}(\lambda)$	9.0E-6	9.0E-6	9.0E-6	9.0E-6
TREE	9.5E-1	2.0E-1	1.3E-1	1.6E-1
IH	1.3E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	2.3E-1	2.3E-1	
WIS	1.3E-1	1.3E-1	
NAIVE	4.9E-1	-	

Table 257. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.2E-3	6.4E-2	5.8E-2	3.4E-3
Q-Reg	4.8E-2	2.4E-3	1.7E-3	4.4E-3
MRDR	3.3E-2	5.0E-2	6.0E-2	6.9E-2
FQE	2.2E-5	2.2E-5	2.2E-5	2.2E-5
$R(\lambda)$	2.3E-5	2.0E-5	2.0E-5	2.3E-5
$Q^{\hat{\pi}}(\lambda)$	2.2E-5	2.2E-5	2.2E-5	2.2E-5
TREE	9.6E-1	4.5E-2	2.2E-2	2.2E-2
IH	3.8E-3	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	5.3E-2	5.3E-2	
WIS	2.5E-2	2.5E-2	
NAIVE	4.3E-1	-	

Table 258. Graph, relative MSE. $T = 16, N = 1024, \pi_b(a = 0) = 0.6, \pi_e(a = 0) = 0.8$. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.3E-3	3.0E-2	2.6E-2	8.3E-4
Q-Reg	2.4E-2	5.7E-4	3.2E-4	1.1E-3
MRDR	2.0E-2	1.6E-2	1.1E-2	1.1E-2
FQE	2.5E-5	2.5E-5	2.5E-5	2.5E-5
$R(\lambda)$	2.5E-5	2.7E-5	2.8E-5	2.5E-5
$Q^{\pi}(\lambda)$	2.5E-5	2.5E-5	2.5E-5	2.5E-5
TREE	1.0E0	2.5E-2	1.4E-2	1.4E-2
IH	2.0E-3	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	2.5E-2	2.5E-2	
WIS	1.4E-2	1.4E-2	
NAIVE	4.5E-1	-	

Table 259. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	8.4E0	9.8E1	4.9E1	8.9E0
Q-Reg	1.5E1	2.1E1	1.5E1	1.5E1
MRDR	6.6E0	1.2E1	8.7E0	6.6E0
FQE	5.8E0	1.1E1	1.1E1	5.8E0
$R(\lambda)$	8.3E0	8.3E0	8.6E0	8.3E0
$Q^{\pi}(\lambda)$	1.0E1	1.1E1	9.3E0	1.0E1
TREE	7.8E0	9.3E0	1.0E1	7.8E0
IH	7.8E0	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.1E1	1.5E1	
WIS	2.0E1	1.0E1	
NAIVE	6.0E0	-	

Table 260. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	6.4E0	5.9E1	4.9E1	1.3E1
Q-Reg	2.8E0	3.8E0	3.2E0	2.8E0
MRDR	1.9E0	1.9E0	1.1E0	1.9E0
FQE	5.9E0	1.1E1	9.5E0	5.9E0
$R(\lambda)$	6.3E0	5.5E0	6.1E0	6.3E0
$Q^{\pi}(\lambda)$	7.1E0	8.0E0	7.9E0	7.1E0
TREE	3.8E0	3.7E0	4.7E0	4.1E0
IH	2.9E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	4.1E0	3.2E0	
WIS	9.0E0	3.2E0	
NAIVE	4.5E0	-	

Table 261. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	5.9E0	2.1E1	2.1E1	6.5E0
Q-REG	1.4E0	7.3E0	4.9E0	1.4E0
MRDR	1.0E0	8.8E0	4.6E0	1.1E0
FQE	5.2E0	7.4E0	3.7E0	5.2E0
$R(\lambda)$	3.3E0	4.1E0	3.3E0	3.3E0
$Q^{\hat{\pi}}(\lambda)$	3.8E0	9.6E0	3.6E0	3.8E0
TREE	4.0E0	2.1E0	2.4E0	4.0E0
IH	4.6E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	6.8E0	1.5E0	
WIS	9.5E0	1.9E0	
NAIVE	2.3E0	-	

Table 262. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	1.8E0	1.3E1	1.2E1	3.7E0
Q-REG	2.3E0	2.6E0	2.5E0	2.2E0
MRDR	1.7E0	2.9E0	2.8E0	1.6E0
FQE	1.1E0	2.8E0	2.5E0	1.1E0
$R(\lambda)$	1.2E0	2.0E0	2.0E0	1.2E0
$Q^{\hat{\pi}}(\lambda)$	1.7E0	2.6E0	2.6E0	1.7E0
TREE	1.4E0	2.1E0	2.0E0	1.4E0
IH	1.9E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	7.0E0	2.3E0	
WIS	7.7E0	2.0E0	
NAIVE	2.0E0	-	

Table 263. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic rewards. Sparse rewards.

DMHybrid DIRECT DR WDR MAGIC 9.5E-1 1.9E0 1.5E0 1.2E0 AMQ-REG 9.8E-1 1.4E0 1.3E0 9.2E-1 MRDR 1.1E0 2.8E0 1.5E0 7.3E-1 FQE 9.1E-1 1.5E0 1.1E0 8.8E-1 $R(\lambda)$ 9.9E-1 1.2E0 1.0E0 9.9E-1 1.3E0 $Q^{\pi}(\lambda)$ 7.5E-1 1.1E0 7.4E-1 TREE 1.2E0 1.1E0 1.0E0 9.4E-1 ΙH 1.5E0

	IPS		
	STANDARD PER-DECISION		
IS	3.7E0	1.1E0	
WIS	4.0E0	8.9E-1	
NAIVE	1.2E0	-	

Table 264. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.6E-1	4.4E0	4.2E0	2.1E0
Q-Reg	1.1E0	1.4E0	1.3E0	1.3E0
MRDR	9.4E-1	7.8E-1	7.8E-1	8.2E-1
FQE	3.8E-1	1.3E0	1.1E0	3.8E-1
$R(\lambda)$	6.6E-1	1.1E0	1.1E0	6.7E-1
$Q^{\pi}(\lambda)$	8.0E-1	1.4E0	1.2E0	7.9E-1
TREE	1.0E0	9.8E-1	9.4E-1	9.4E-1
IH	3.6E-1	-	-	-

		IPS		
	STANDARD	STANDARD PER-DECISION		
IS	3.1E0	1.1E0		
WIS	2.5E0	1.0E0		
NAIVE	7.1E-1	-		

Table 265. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.7E-1	1.2E0	1.2E0	2.2E-1
Q-Reg	3.2E-1	2.5E-1	2.4E-1	3.4E-1
MRDR	3.9E-1	2.8E-1	2.7E-1	3.8E-1
FQE	1.0E-1	2.2E-1	2.3E-1	9.9E-2
$R(\lambda)$	2.1E-1	2.3E-1	2.4E-1	2.1E-1
$Q^{\hat{\pi}}(\lambda)$	1.9E-1	2.2E-1	2.3E-1	1.8E-1
TREE	9.8E-1	3.0E-1	2.7E-1	4.3E-1
IH	1.1E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.4E0	3.3E-1	
WIS	1.2E0	2.9E-1	
NAIVE	6.3E-1	-	

Table 266. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.2E-1	9.3E-1	9.2E-1	5.7E-2	
Q-REG	4.0E-1	4.2E-1	4.1E-1	4.3E-1	
MRDR	2.9E-1	3.2E-1	3.3E-1	3.6E-1	
FQE	1.0E-1	4.0E-1	4.0E-1	1.6E-1	
$R(\lambda)$	1.7E-1	4.0E-1	4.0E-1	1.8E-1	
$Q^{\hat{\pi}}(\lambda)$	1.5E-1	4.0E-1	4.0E-1	1.5E-1	
TREE	7.1E-1	3.7E-1	4.1E-1	2.5E-1	
IH	1.0E-1	-	-	-	

		IPS		
	STANDARD	PER-DECISION		
IS	8.7E-1	3.8E-1		
WIS	9.6E-1	4.2E-1		
NAIVE	3.5E-1	-		

Table 267. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.0E0	1.4E1	1.5E1	1.5E0	
Q-Reg	1.9E1	5.1E1	4.1E1	3.0E0	
MRDR	3.0E0	2.7E1	2.0E1	8.9E-1	
FQE	1.1E0	6.4E0	4.8E0	1.3E0	
$R(\lambda)$	1.2E0	1.2E0	1.3E0	1.2E0	
$Q^{\hat{\pi}}(\lambda)$	2.0E0	4.3E0	3.0E0	1.4E0	
TREE	1.0E0	4.0E0	3.8E0	1.3E0	
IH	6.5E-1	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	8.7E0	8.7E0	
WIS	5.2E0	5.2E0	
NAIVE	1.7E0	-	

Table 268. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	8.3E-1	5.5E0	3.7E0	1.3E0
Q-Reg	8.0E0	3.6E1	1.9E1	7.8E0
MRDR	6.4E0	1.4E2	5.2E1	6.4E0
FQE	3.8E-1	1.9E0	1.8E0	3.6E-1
$R(\lambda)$	1.1E0	1.4E0	1.4E0	1.0E0
$Q^{\hat{\pi}}(\lambda)$	1.4E0	1.1E0	1.4E0	1.1E0
TREE	8.9E-1	4.9E0	2.4E0	8.5E-1
IH	6.8E-1	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	1.2E1	1.2E1		
WIS	2.9E0	2.9E0		
NAIVE	1.3E0	-		

Table 269. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	5.8E-1	1.5E1	1.0E1	3.8E-1	
Q-Reg	2.1E0	3.5E0	2.7E0	2.8E0	
MRDR	1.6E0	4.5E0	3.1E0	3.4E0	
FQE	4.4E-1	1.9E0	2.0E0	4.5E-1	
$R(\lambda)$	1.4E0	1.7E0	1.7E0	1.4E0	
$Q^{\pi}(\lambda)$	9.7E-1	1.9E0	1.6E0	1.0E0	
TREE	9.4E-1	1.9E0	2.7E0	9.4E-1	
IH	7.4E-1	-	-	-	

	IPS		
	STANDARD PER-DECISIO		
IS	2.1E0	2.1E0	
WIS	2.8E0	2.8E0	
NAIVE	1.7E0	-	

Table 270. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.2E-1	3.3E0	2.5E0	2.3E-1
Q-Reg	1.0E0	3.8E-1	4.0E-1	9.3E-1
MRDR	7.3E-1	2.2E0	2.0E0	8.4E-1
FQE	7.8E-2	5.6E-1	5.6E-1	7.9E-2
$R(\lambda)$	2.2E-1	3.7E-1	3.6E-1	2.3E-1
$Q^{\pi}(\lambda)$	3.8E-1	3.6E-1	3.4E-1	2.9E-1
TREE	1.0E0	9.9E-1	8.7E-1	7.8E-1
IH	5.1E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	1.0E0	
WIS	8.7E-1	8.7E-1	
NAIVE	5.9E-1	-	

Table 271. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.3E-1	4.4E0	4.4E0	1.4E0
Q-Reg	9.4E-1	7.4E-1	6.7E-1	9.4E-1
MRDR	7.9E-1	3.9E0	1.6E0	6.7E-1
FQE	6.3E-2	1.2E0	8.2E-1	6.7E-2
$R(\lambda)$	3.6E-1	7.2E-1	6.3E-1	3.7E-1
$\mathbf{Q}^{\dot{\pi}}(\lambda)$	2.5E-1	1.1E0	8.5E-1	2.2E-1
TREE	9.6E-1	8.0E-1	6.8E-1	8.1E-1
IH	5.3E-1	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	9.0E-1	9.0E-1		
WIS	7.3E-1	7.3E-1		
NAIVE	6.2E-1	-		

Table 272. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.3E-2	9.9E-1	9.3E-1	1.1E-1
Q-Reg	4.5E-1	3.4E-1	3.2E-1	2.9E-1
MRDR	5.2E-1	5.6E-1	5.1E-1	6.8E-1
FQE	3.6E-2	3.9E-1	3.2E-1	3.3E-2
$R(\lambda)$	1.6E-1	3.0E-1	2.9E-1	1.7E-1
$Q^{\pi}(\lambda)$	1.3E-1	3.7E-1	3.2E-1	1.4E-1
TREE	1.0E0	4.5E-1	4.0E-1	7.1E-1
IH	3.5E-2	_	_	_

	IPS		
	STANDARD PER-DECISION		
IS	4.5E-1	4.5E-1	
WIS	4.0E-1	4.0E-1	
NAIVE	4.0E-1	-	

Table 273. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.2E-2	4.7E-1	3.9E-1	5.3E-2
Q-Reg	2.4E-1	1.4E-1	1.3E-1	1.5E-1
MRDR	1.8E-1	4.8E-1	4.2E-1	2.5E-1
FQE	3.5E-2	1.8E-1	1.5E-1	3.6E-2
$R(\lambda)$	8.3E-2	1.5E-1	1.4E-1	8.4E-2
$Q^{\hat{\pi}}(\lambda)$	1.0E-1	1.8E-1	1.5E-1	8.0E-2
TREE	9.9E-1	2.5E-1	2.2E-1	3.7E-1
IH	3.4E-2	-	-	-

	IPS			
	STANDARD PER-DECIS			
IS	2.5E-1	2.5E-1		
WIS	2.2E-1	2.2E-1		
NAIVE	4.8E-1	-		

Table 274. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	6.8E-3	1.3E-1	1.2E-1	1.3E-2	
Q-Reg	1.2E-1	7.4E-2	7.1E-2	1.1E-1	
MRDR	1.2E-1	6.5E-2	6.6E-2	8.8E-2	
FQE	1.1E-2	6.6E-2	6.4E-2	1.2E-2	
$R(\lambda)$	3.2E-2	6.8E-2	6.8E-2	3.3E-2	
$Q^{\pi}(\lambda)$	2.9E-2	7.1E-2	6.8E-2	2.6E-2	
TREE	1.0E0	1.2E-1	8.1E-2	8.1E-2	
IH	1.7E-2	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	1.2E-1	1.2E-1	
WIS	8.1E-2	8.1E-2	
NAIVE	4.4E-1	-	

Table 275. Graph, relative MSE. $T=16, N=8, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

DM Hybrid DIRECT DR WDR MAGIC AM 4.8E1 2.2E2 1.2E2 5.4E1 Q-REG 1.3E2 3.2E2 2.2E2 1.3E2 MRDR 4.4E2 2.0E2 6.7E1 6.2E1 FQE 2.6E1 5.2E1 2.6E1 1.1E2 $R(\lambda)$ 4.2E15.2E1 4.0E1 4.2E1 $Q^{\pi}(\lambda)$ 3.8E1 3.8E1 8.1E1 3.6E1 TREE 3.3E1 7.7E1 6.3E1 4.0E1 ΙH 5.8E1

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	1.3E2	1.3E2			
WIS	1.5E2	8.6E1			
NAIVE	2.8E1	-			

Table 276. Graph, relative MSE. $T=16, N=16, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	2.2E1	1.0E2	1.0E2	2.2E1
Q-Reg	5.3E1	4.9E1	4.8E1	5.6E1
MRDR	2.6E1	7.5E1	4.5E1	2.0E1
FQE	2.2E1	4.0E1	4.9E1	2.1E1
$R(\lambda)$	3.0E1	3.5E1	3.6E1	3.0E1
$Q^{\hat{\pi}}(\lambda)$	6.0E1	7.3E1	7.4E1	6.0E1
TREE	1.4E1	1.5E1	2.5E1	1.4E1
IH	1.7E1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.1E2	4.3E1	
WIS	9.0E1	4.3E1	
NAIVE	1.5E1	-	

Table 277. Graph, relative MSE. $T=16, N=32, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.2E0	1.5E2	1.1E2	4.8E0
Q-Reg	6.5E1	5.4E1	3.7E1	6.5E1
MRDR	4.3E1	3.2E2	1.2E2	4.3E1
FQE	3.9E0	3.8E1	2.2E1	3.9E0
$R(\lambda)$	9.8E0	1.5E1	1.3E1	9.8E0
$Q^{\hat{\pi}}(\lambda)$	1.4E1	3.1E1	1.8E1	1.4E1
TREE	2.3E1	2.6E1	1.8E1	2.2E1
IH	3.8E0	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	5.5E1	6.3E1	
WIS	3.6E1	2.7E1	
NAIVE	5.0E0	-	

Table 278. Graph, relative MSE. $T=16, N=64, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	6.6E0	3.5E1	3.5E1	6.5E0
Q-Reg	5.7E0	1.4E1	1.5E1	6.0E0
MRDR	5.9E0	5.2E0	5.1E0	4.8E0
FQE	5.3E0	8.5E0	9.7E0	5.2E0
$R(\lambda)$	7.9E0	8.3E0	8.5E0	7.8E0
$Q^{\hat{\pi}}(\lambda)$	7.1E0	8.7E0	9.4E0	6.5E0
TREE	3.0E0	5.8E0	7.5E0	3.0E0
IH	3.8E0	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.3E1	6.6E0	
WIS	2.0E1	8.6E0	
NAIVE	5.9E0	-	

Table 279. Graph, relative MSE. $T=16, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	6.1E0	2.7E1	3.0E1	3.8E0	
Q-Reg	2.5E0	6.2E0	5.2E0	2.4E0	
MRDR	3.9E0	6.3E0	5.8E0	1.6E0	
FQE	2.7E0	3.5E0	3.3E0	2.6E0	
$R(\lambda)$	3.1E0	3.7E0	3.5E0	2.9E0	
$Q^{\hat{\pi}}(\hat{\lambda})$	5.1E0	5.0E0	4.4E0	5.1E0	
TREE	3.1E0	2.4E0	2.3E0	2.9E0	
IH	3.2E0	-	-	-	

		IPS			
	STANDARD	STANDARD PER-DECISION			
IS	3.7E0	2.7E0			
WIS	4.7E0	2.6E0			
NAIVE	1.9E0	-			

Table 280. Graph, relative MSE. $T=16, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.5E0	7.6E0	8.0E0	1.6E0	
Q-Reg	2.1E0	3.0E0	3.0E0	1.7E0	
MRDR	1.9E0	1.7E0	1.8E0	1.7E0	
FQE	4.5E-1	2.2E0	2.5E0	4.8E-1	
$R(\lambda)$	1.8E0	2.7E0	2.8E0	1.8E0	
$Q^{\pi}(\lambda)$	1.4E0	2.3E0	2.5E0	1.3E0	
TREE	1.2E0	3.0E0	3.1E0	1.1E0	
IH	4.5E-1	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	3.4E0	2.2E0	
WIS	4.4E0	2.6E0	
NAIVE	2.4E-1	-	

Table 281. Graph, relative MSE. $T=16, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	8.5E-1	1.0E1	1.0E1	5.1E0	
Q-Reg	1.4E0	1.2E0	1.3E0	1.3E0	
MRDR	7.9E-1	1.5E0	1.4E0	7.6E-1	
FQE	4.4E-1	1.1E0	1.1E0	4.3E-1	
$R(\lambda)$	4.3E-1	1.1E0	1.1E0	4.4E-1	
$Q^{\hat{\pi}}(\hat{\lambda})$	5.0E-1	1.1E0	1.1E0	4.6E-1	
TREE	1.2E0	1.0E0	1.1E0	7.9E-1	
IH	4.2E-1	-	-	-	

		IPS			
	STANDARD	STANDARD PER-DECISION			
IS	6.1E0	1.3E0			
WIS	5.5E0	1.4E0			
NAIVE	9.9E-1	-			

Table 282. Graph, relative MSE. $T=16, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.5E-1	9.0E-1	8.9E-1	6.2E-1
Q-Reg	5.9E-1	5.3E-1	5.6E-1	5.7E-1
MRDR	6.7E-1	5.7E-1	6.2E-1	6.5E-1
FQE	3.6E-1	5.9E-1	5.9E-1	3.4E-1
$R(\lambda)$	5.9E-1	6.0E-1	5.8E-1	5.9E-1
$Q^{\pi}(\lambda)$	4.2E-1	5.8E-1	5.6E-1	4.3E-1
TREE	1.5E0	6.0E-1	5.8E-1	1.4E0
IH	3.2E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	5.8E0	5.9E-1	
WIS	6.2E0	5.6E-1	
NAIVE	8.9E-1	-	

H.2. Detailed Results for Graph-POMDP

Table 283. Graph-POMDP, relative MSE. $T=2,N=256,H=2,\pi_b(a=0)=0.2,\pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.4E-1	1.2E-1	3.7E-2	3.7E-2
Q-Reg	1.5E-1	1.1E-2	3.5E-3	1.4E-2
MRDR	1.4E0	1.2E-2	3.3E-2	3.3E-2
FQE	8.7E-1	2.3E-2	2.7E-3	5.7E-2
$R(\lambda)$	5.1E-1	1.0E-2	2.4E-3	3.2E-2
$Q^{\hat{\pi}}(\lambda)$	5.3E-2	9.2E-3	2.8E-3	2.9E-2
TREE	3.8E-1	8.2E-3	2.4E-3	2.2E-2
IH	8.4E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.3E-1	4.5E-2	
WIS	1.5E-2	6.0E-3	
NAIVE	3.8E0	-	

Table 284. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.6E-1	3.8E-2	1.3E-2	1.3E-2
Q-Reg	1.6E-1	1.4E-3	1.4E-3	8.7E-3
MRDR	1.5E0	1.4E-3	2.8E-2	2.8E-2
FQE	9.5E-1	9.7E-3	1.3E-3	1.3E-3
$R(\lambda)$	5.6E-1	3.3E-3	9.9E-4	6.1E-3
$Q^{\hat{\pi}}(\hat{\lambda})$	5.5E-2	5.9E-3	1.7E-3	9.7E-4
TREE	4.2E-1	2.3E-3	1.0E-3	2.4E-2
IH	9.4E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	6.3E-2	2.0E-2		
WIS	1.5E-2	4.3E-3		
NAIVE	4.0E0	-		

Table 285. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	9.7E-1	6.7E-3	2.4E-3	2.4E-3	
Q-Reg	1.7E-1	3.1E-3	1.2E-3	4.3E-3	
MRDR	1.5E0	3.0E-3	5.5E-3	5.5E-3	
FQE	1.0E0	7.8E-3	1.6E-3	1.6E-3	
$R(\lambda)$	6.3E-1	5.0E-3	1.4E-3	1.4E-3	
$Q^{\hat{\pi}}(\lambda)$	9.0E-2	3.3E-3	1.3E-3	1.3E-3	
TREE	4.9E-1	4.5E-3	1.3E-3	1.4E-3	
IH	1.0E0	-	-	-	

	IPS				
	STANDARD PER-DECISION				
IS	1.6E-2	1.0E-2			
WIS	3.3E-3	2.5E-3			
NAIVE	4.0E0	-			

Table 286. Graph-POMDP, relative MSE. $T=2,N=256,H=2,\pi_b(a=0)=0.2,\pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.1E0	1.9E-1	1.5E-1	4.5E-1	
Q-Reg	2.8E-1	5.7E-2	4.4E-2	3.4E-2	
MRDR	1.7E0	5.4E-2	1.2E-1	1.2E-1	
FQE	1.0E0	8.1E-2	4.8E-2	2.9E-1	
$R(\lambda)$	6.7E-1	6.4E-2	4.4E-2	1.5E-1	
$Q^{\hat{\pi}}(\lambda)$	1.0E-1	5.7E-2	4.3E-2	7.4E-2	
TREE	5.3E-1	6.2E-2	4.4E-2	1.1E-1	
IH	1.0E0	-	-	-	

	IPS			
	STANDARD PER-DECISIO			
IS	9.8E-2	6.0E-2		
WIS	9.7E-2	5.3E-2		
NAIVE	4.0E0	-		

Table 287. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.1E0	4.7E-2	3.0E-2	6.2E-2
Q-Reg	2.0E-1	1.4E-2	1.1E-2	6.8E-3
MRDR	1.5E0	1.6E-2	3.9E-2	3.9E-2
FQE	1.2E0	1.4E-2	8.1E-3	8.1E-3
$R(\lambda)$	7.3E-1	9.8E-3	8.5E-3	6.3E-2
$Q^{\hat{\pi}}(\lambda)$	1.3E-1	1.2E-2	9.1E-3	2.1E-2
TREE	5.9E-1	9.5E-3	8.5E-3	6.1E-2
IH	1.2E0	-	-	-

	IPS				
	STANDARD PER-DECISIO				
IS	5.5E-2	1.7E-2			
WIS	2.7E-2	8.6E-3			
NAIVE	4.1E0	-			

Table 288. Graph-POMDP, relative MSE. $T=2,N=1024,H=2,\pi_b(a=0)=0.2,\pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.6E-1	4.2E-2	2.6E-2	2.6E-2
Q-Reg	1.7E-1	1.1E-2	1.1E-2	2.1E-2
MRDR	1.5E0	1.0E-2	1.4E-2	1.4E-2
FQE	1.0E0	2.1E-2	1.2E-2	1.2E-2
$R(\lambda)$	6.1E-1	1.6E-2	1.1E-2	1.3E-2
$Q^{\hat{\pi}}(\lambda)$	6.4E-2	1.0E-2	1.1E-2	2.4E-2
TREE	4.7E-1	1.4E-2	1.1E-2	2.1E-2
IH	1.0E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	5.0E-2	2.3E-2	
WIS	3.3E-2	1.4E-2	
NAIVE	4.0E0	-	

Table 289. Graph-POMDP, relative MSE. $T=2,N=256,H=2,\pi_b(a=0)=0.2,\pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.6E0	6.5E-1	5.5E-1	1.5E0
Q-Reg	5.7E-1	3.1E-1	2.9E-1	3.5E-1
MRDR	2.3E0	3.0E-1	2.1E-1	2.1E-1
FQE	1.7E0	4.2E-1	3.3E-1	8.5E-1
$R(\lambda)$	1.2E0	3.7E-1	3.1E-1	6.8E-1
$Q^{\hat{\pi}}(\lambda)$	4.0E-1	2.9E-1	3.0E-1	3.1E-1
TRÈE	1.1E0	3.5E-1	3.1E-1	6.3E-1
IH	1.8E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	8.2E-1	4.1E-1		
WIS	7.3E-1	3.6E-1		
NAIVE	4.6E0	-		

Table 290. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.1E0	1.2E-1	1.1E-1	4.8E-1
Q-Reg	2.7E-1	5.5E-2	5.7E-2	6.1E-2
MRDR	1.8E0	5.6E-2	6.8E-2	6.8E-2
FQE	1.2E0	5.7E-2	5.6E-2	3.6E-1
$R(\lambda)$	7.8E-1	5.4E-2	5.6E-2	1.8E-1
$Q^{\hat{\pi}}(\lambda)$	1.2E-1	5.4E-2	5.6E-2	3.6E-2
TREE	6.3E-1	5.3E-2	5.6E-2	1.2E-1
IH	1.2E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.5E-1	5.9E-2	
WIS	1.4E-1	5.6E-2	
NAIVE	4.3E0	-	

Table 291. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.0E0	1.4E-1	6.8E-2	6.8E-2	
Q-Reg	2.8E-1	6.7E-2	6.4E-2	7.1E-2	
MRDR	1.7E0	6.6E-2	7.6E-2	7.6E-2	
FQE	1.1E0	1.1E-1	7.3E-2	2.3E-1	
$R(\lambda)$	6.9E-1	8.6E-2	6.7E-2	1.4E-1	
$Q^{\hat{\pi}}(\lambda)$	9.0E-2	6.1E-2	6.4E-2	9.5E-2	
TREE	5.4E-1	8.1E-2	6.7E-2	9.0E-2	
IH	1.1E0	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	3.0E-1	1.2E-1	
WIS	2.3E-1	8.3E-2	
NAIVE	4.0E0	-	

Table 292. Graph-POMDP, relative MSE. $T=2,N=256, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	7.4E-1	2.8E-1	2.6E-1	8.4E-1	
Q-Reg	1.8E-1	1.7E-1	1.6E-1	1.5E-1	
MRDR	1.4E0	1.7E-1	2.6E-1	2.6E-1	
FQE	6.5E-1	1.8E-1	1.5E-1	1.2E-1	
$R(\lambda)$	3.5E-1	1.8E-1	1.5E-1	9.7E-2	
$Q^{\hat{\pi}}(\lambda)$	1.2E-1	2.3E-1	1.7E-1	4.9E-1	
TREE	2.4E-1	1.8E-1	1.5E-1	1.3E-1	
IH	6.8E-1	-	-	-	

	IPS		
	STANDARD	Per-Decision	
IS	1.9E-1	1.8E-1	
WIS	2.4E-1	1.4E-1	
NAIVE	3.7E0	-	

Table 293. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.2E0	4.5E-1	4.7E-1	6.5E-1	
Q-Reg	2.2E-1	7.6E-2	7.0E-2	8.5E-2	
MRDR	1.7E0	7.9E-2	1.0E-1	1.0E-1	
FQE	1.1E0	6.3E-2	6.4E-2	2.8E-1	
$R(\lambda)$	6.8E-1	6.7E-2	6.7E-2	2.1E-1	
$Q^{\hat{\pi}}(\lambda)$	1.1E-1	7.8E-2	6.8E-2	1.8E-1	
TREE	5.5E-1	6.8E-2	6.7E-2	1.8E-1	
IH	1.1E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E-1	6.6E-2		
WIS	9.4E-2	6.0E-2		
NAIVE	4.3E0	-		

Table 294. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.0E0	1.3E-1	1.5E-1	4.7E-1	
Q-Reg	1.9E-1	7.4E-2	6.9E-2	7.2E-2	
MRDR	1.5E0	7.2E-2	9.2E-2	9.2E-2	
FQE	1.0E0	6.5E-2	6.3E-2	2.6E-1	
$R(\lambda)$	6.3E-1	6.7E-2	6.5E-2	1.8E-1	
$Q^{\hat{\pi}}(\lambda)$	1.4E-1	8.0E-2	6.8E-2	1.0E-1	
TREE	5.0E-1	6.9E-2	6.5E-2	1.5E-1	
IH	1.0E0	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E-1	5.4E-2	
WIS	9.0E-2	5.9E-2	
NAIVE	4.0E0	-	

Table 295. Graph-POMDP, relative MSE. $T=2,N=256,H=2,\pi_b(a=0)=0.2,\pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.9E0	3.0E-1	7.8E-2	7.8E-2
Q-Reg	1.2E-1	1.5E-2	5.8E-3	3.1E-2
MRDR	1.0E0	5.3E-3	3.2E-2	3.2E-2
FQE	3.8E0	2.7E-1	2.6E-2	2.6E-2
$R(\lambda)$	1.8E0	1.4E-1	1.8E-2	1.8E-2
$Q^{\hat{\pi}}(\lambda)$	2.6E-1	2.5E-2	8.4E-3	1.1E-1
TREE	1.8E0	1.4E-1	1.9E-2	1.9E-2
IH	3.8E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.2E-1	1.2E-1	
WIS	2.6E-2	2.6E-2	
NAIVE	3.9E0	-	

Table 296. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.0E0	5.6E-2	4.2E-2	4.2E-2
Q-Reg	1.1E-1	6.1E-3	5.0E-3	3.7E-2
MRDR	1.1E0	5.1E-3	8.5E-3	8.5E-3
FQE	4.0E0	3.7E-2	7.4E-3	7.4E-3
$R(\lambda)$	1.9E0	2.1E-2	6.5E-3	6.5E-3
$Q^{\hat{\pi}}(\lambda)$	3.8E-1	8.4E-3	5.5E-3	2.6E-2
TRÈE	1.9E0	2.2E-2	6.6E-3	6.6E-3
IH	4.1E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.7E-2	1.7E-2	
WIS	7.4E-3	7.4E-3	
NAIVE	4.1E0	-	

Table 297. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.0E0	5.5E-2	1.7E-2	1.7E-2
Q-Reg	1.0E-1	2.4E-3	9.6E-4	1.2E-2
MRDR	1.1E0	1.3E-3	1.2E-2	1.2E-2
FQE	3.9E0	7.6E-2	5.7E-3	5.7E-3
$R(\lambda)$	1.8E0	3.7E-2	3.8E-3	3.8E-3
$Q^{\pi}(\lambda)$	2.7E-1	7.3E-3	1.6E-3	1.6E-3
TREE	1.8E0	3.8E-2	4.1E-3	4.1E-3
IH	3.9E0	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	3.3E-2	3.3E-2
WIS	5.8E-3	5.8E-3
NAIVE	3.9E0	-

Table 298. Graph-POMDP, relative MSE. $T=2,N=256,H=2,\pi_b(a=0)=0.2,\pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.0E0	4.9E-1	3.5E-1	7.6E-1
Q-Reg	3.4E-1	9.8E-2	8.9E-2	2.6E-1
MRDR	1.6E0	9.5E-2	1.6E-1	1.6E-1
FQE	4.2E0	3.0E-1	9.9E-2	3.3E-1
$R(\lambda)$	2.1E0	1.8E-1	9.1E-2	6.2E-1
$Q^{\pi}(\lambda)$	3.1E-1	8.9E-2	8.8E-2	2.2E-1
TREE	2.1E0	1.8E-1	9.1E-2	6.3E-1
IH	4.2E0	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	8.2E-2	1.3E-1
WIS	8.0E-2	9.5E-2
NAIVE	3.8E0	-

Table 299. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

DMHybrid DIRECT DR WDR MAGIC 3.3E-1 2.2E-1 AM4.7E0 2.2E-1 Q-REG 2.3E-1 9.5E-2 9.9E-2 3.2E-1MRDR 1.4E-1 1.5E0 1.0E-1 1.4E-1 FQE 4.7E0 1.0E-1 9.0E-2 9.0E-2 $R(\lambda)$ 2.2E0 8.7E-2 9.3E-2 9.3E-2 $Q^{\pi}(\lambda)$ 2.7E-1 9.2E-2 9.7E-22.4E-1 9.2E-2 TREE 2.2E0 9.2E-2 8.7E-2 ΙH 4.5E0

	IPS		
	STANDARD	PER-DECISION	
IS	9.1E-2	8.7E-2	
WIS	1.1E-1	9.0E-2	
NAIVE	4.5E0	-	

Table 300. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.0E0	2.4E-1	1.8E-1	1.8E-1
Q-Reg	1.8E-1	5.5E-2	5.2E-2	1.3E-1
MRDR	1.2E0	5.1E-2	5.4E-2	5.4E-2
FQE	4.0E0	1.2E-1	5.9E-2	5.9E-2
$R(\lambda)$	1.9E0	8.8E-2	5.6E-2	5.6E-2
$Q^{\hat{\pi}}(\hat{\lambda})$	2.7E-1	6.2E-2	5.3E-2	1.4E-1
TREE	1.9E0	8.9E-2	5.7E-2	5.7E-2
IH	4.0E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E-1	8.4E-2	
WIS	7.7E-2	5.8E-2	
NAIVE	3.9E0	-	

Table 301. Graph-POMDP, relative MSE. $T=2,N=256, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.8E0	1.2E0	6.4E-1	2.0E0
Q-Reg	3.7E-1	2.7E-1	2.6E-1	3.4E-1
MRDR	1.1E0	2.5E-1	2.8E-1	2.8E-1
FQE	3.8E0	8.6E-1	3.3E-1	3.3E-1
$R(\lambda)$	1.7E0	6.0E-1	3.1E-1	3.4E-1
$Q^{\hat{\pi}}(\lambda)$	2.5E-1	3.9E-1	2.9E-1	2.5E-1
TREE	1.7E0	6.1E-1	3.2E-1	3.4E-1
IH	3.7E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	6.6E-1	6.6E-1		
WIS	3.4E-1	3.4E-1		
NAIVE	3.9E0	-		

Table 302. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.1E0	4.3E-1	3.5E-1	1.3E0
Q-Reg	2.2E-1	8.7E-2	9.7E-2	2.1E-1
MRDR	1.3E0	1.1E-1	2.0E-1	2.0E-1
FQE	4.1E0	1.6E-1	8.0E-2	8.0E-2
$R(\lambda)$	1.9E0	1.0E-1	8.2E-2	3.8E-1
$Q^{\hat{\pi}}(\lambda)$	3.1E-1	7.9E-2	9.1E-2	2.9E-1
TREE	1.9E0	1.0E-1	8.2E-2	5.3E-1
IH	4.1E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	9.1E-2	9.1E-2		
WIS	7.9E-2	7.9E-2		
NAIVE	3.9E0	-		

Table 303. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

DMHybrid DIRECT DR WDR MAGIC AM4.2E0 2.6E-1 2.4E-1 2.4E-1 Q-REG 1.5E-1 1.3E-1 1.3E-1 1.5E-1 MRDR 1.5E-1 1.0E0 1.3E-1 1.5E-1 FQE 4.0E0 1.6E-1 1.3E-1 1.3E-1 $R(\lambda)$ 1.9E0 1.4E-1 1.3E-1 1.3E-1 $Q^{\pi}(\lambda)$ 3.1E-1 1.2E-1 1.3E-1 2.8E-1 TREE 1.9E0 1.3E-1 1.4E-1 1.3E-1 ΙH 4.1E0

	IPS			
	STANDARD PER-DECISION			
IS	1.3E-1	1.3E-1		
WIS	1.3E-1	1.3E-1		
NAIVE	3.9E0	-		

Table 304. Graph-POMDP, relative MSE. $T=2,N=256, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.1E0	1.9E0	1.7E0	2.9E0	
Q-Reg	8.9E-1	1.2E0	1.4E0	1.5E0	
MRDR	1.1E0	1.4E0	1.7E0	1.6E0	
FQE	3.5E0	1.3E0	1.3E0	2.4E0	
$R(\lambda)$	1.7E0	1.2E0	1.3E0	2.3E0	
$Q^{\hat{\pi}}(\lambda)$	8.2E-1	1.2E0	1.4E0	1.5E0	
TREE	1.8E0	1.2E0	1.3E0	1.5E0	
IH	3.6E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	2.6E0	1.3E0		
WIS	2.8E0	1.3E0		
NAIVE	4.0E0	-		

Table 305. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.5E0	1.3E0	1.1E0	3.6E0
Q-Reg	1.1E-1	1.7E-1	1.9E-1	1.9E-1
MRDR	9.0E-1	1.9E-1	2.8E-1	4.1E-1
FQE	4.3E0	2.6E-1	1.7E-1	1.7E-1
$R(\lambda)$	1.8E0	1.9E-1	1.7E-1	8.7E-1
$Q^{\hat{\pi}}(\hat{\lambda})$	2.6E-1	1.6E-1	1.8E-1	3.4E-1
TREE	1.9E0	1.9E-1	1.7E-1	8.7E-1
IH	4.3E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	5.0E-1	2.4E-1		
WIS	4.0E-1	1.7E-1		
NAIVE	4.4E0	-		

Table 306. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.5E0	5.7E-1	5.1E-1	1.7E0
Q-Reg	4.2E-1	2.7E-1	2.7E-1	4.3E-1
MRDR	1.7E0	2.6E-1	2.4E-1	3.1E-1
FQE	4.4E0	3.4E-1	2.8E-1	2.8E-1
$R(\lambda)$	2.2E0	3.1E-1	2.8E-1	1.0E0
$Q^{\hat{\pi}}(\hat{\lambda})$	4.7E-1	2.8E-1	2.7E-1	5.2E-1
TREE	2.2E0	3.1E-1	2.8E-1	1.0E0
IH	4.5E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	5.6E-1	3.2E-1		
WIS	5.1E-1	2.8E-1		
NAIVE	4.0E0	-		

Table 307. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.0E-1	2.4E1	1.0E-1	9.6E-2	
Q-Reg	1.7E0	1.1E1	1.6E-1	1.8E0	
MRDR	1.6E0	1.3E1	5.6E1	5.6E1	
FQE	1.4E-2	1.9E0	6.3E-3	4.9E-3	
$R(\lambda)$	2.7E-2	4.9E0	1.3E-1	1.3E-1	
$Q^{\pi}(\lambda)$	7.1E-3	3.2E0	7.5E-3	2.5E-3	
TREE	8.4E-3	1.0E1	1.9E-1	1.9E-1	
IH	1.0E-2	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	9.1E-1	4.3E0	
WIS	8.8E-1	2.8E-1	
NAIVE	4.0E0	-	

Table 308. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.1E-1	1.3E0	1.1E-1	1.1E-1
Q-Reg	1.5E0	3.7E-1	9.0E-1	1.4E0
MRDR	5.4E-1	1.1E0	5.3E0	5.3E0
FQE	1.9E-2	1.1E0	2.0E-2	8.7E-3
$R(\lambda)$	3.1E-2	1.8E0	1.3E-1	1.3E-1
$Q^{\pi}(\lambda)$	9.1E-3	8.0E-1	1.2E-2	1.7E-3
TREE	6.3E-3	1.8E0	1.4E-1	1.4E-1
IH	1.4E-2	_	_	_

	IPS		
	STANDARD	PER-DECISION	
IS	3.9E0	2.0E0	
WIS	9.4E-1	2.4E-1	
NAIVE	4.0E0	-	

Table 309. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.1E-1	1.6E-1	8.1E-2	9.2E-2
Q-Reg	2.6E-1	3.7E-1	6.6E-2	1.5E-1
MRDR	2.0E-1	2.6E-1	1.2E0	1.2E0
FQE	1.4E-2	2.9E-2	9.0E-3	4.1E-3
$R(\lambda)$	2.0E-2	1.3E-1	9.4E-2	9.8E-2
$Q^{\hat{\pi}}(\lambda)$	5.8E-3	2.6E-2	7.4E-3	6.2E-4
TREE	2.0E-3	1.6E-1	1.0E-1	1.0E-1
IH	1.1E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.1E0	2.1E-1	
WIS	7.0E-1	1.6E-1	
NAIVE	4.0E0	-	

Table 310. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.3E-1	2.2E-1	1.6E-1	1.0E-1
Q-Reg	4.5E-1	1.4E-1	1.3E-1	3.9E-1
MRDR	4.2E-1	3.5E-1	2.9E0	2.9E0
FQE	2.1E-2	2.8E-2	6.1E-2	1.1E-2
$R(\lambda)$	5.6E-2	1.5E-1	2.0E-1	2.0E-1
$Q^{\pi}(\lambda)$	2.0E-2	1.8E-2	5.1E-2	1.2E-2
TREE	2.4E-2	2.0E-1	2.2E-1	2.2E-1
IH	1.7E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	9.2E-1	3.8E-1	
WIS	1.2E0	3.4E-1	
NAIVE	4.0E0	-	

Table 311. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.1E-1	1.9E-1	2.0E-1	1.5E-1
Q-Reg	4.2E-1	1.2E-1	2.0E-1	3.3E-1
MRDR	3.6E-1	3.3E-1	2.1E0	2.0E0
FQE	1.8E-2	3.5E-2	7.0E-2	7.4E-3
$R(\lambda)$	2.8E-2	1.5E-1	1.9E-1	2.0E-1
$Q^{\hat{\pi}}(\lambda)$	8.5E-3	2.7E-2	6.7E-2	3.0E-3
TREE	1.3E-2	2.2E-1	2.0E-1	2.0E-1
IH	1.5E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	9.9E-1	3.4E-1	
WIS	1.1E0	2.9E-1	
NAIVE	4.0E0	-	

Table 312. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.2E-1	6.2E0	3.8E-1	1.7E-1
Q-Reg	5.5E-1	3.2E0	1.3E-1	2.4E-1
MRDR	8.5E-1	3.5E-1	3.1E0	3.1E0
FQE	1.5E-2	9.2E-1	5.0E-2	7.0E-3
$R(\lambda)$	5.6E-2	1.8E0	2.1E-1	2.1E-1
$Q^{\hat{\pi}}(\lambda)$	1.1E-2	7.0E-1	4.3E-2	4.2E-3
TRÈE	1.5E-2	2.3E0	2.7E-1	2.7E-1
IH	1.1E-2	-	-	_

	IPS			
	STANDARD PER-DECISION			
IS	9.6E-1	5.7E-1		
WIS	1.0E0	3.1E-1		
NAIVE	4.0E0	-		

Table 313. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.5E-1	1.9E0	4.0E-1	1.9E-1
Q-Reg	3.3E-1	2.1E-1	5.7E-1	2.4E-1
MRDR	3.1E-1	4.6E-1	7.9E0	7.9E0
FQE	4.1E-2	2.9E-1	1.6E-1	2.7E-2
$R(\lambda)$	4.9E-2	4.7E-1	3.5E-1	9.8E-2
$Q^{\hat{\pi}}(\lambda)$	6.0E-2	2.2E-1	1.6E-1	4.8E-2
TREE	2.7E-2	7.1E-1	3.7E-1	9.0E-2
IH	3.8E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	5.6E-1		
WIS	2.1E0	4.5E-1		
NAIVE	4.1E0	-		

Table 314. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.3E-1	1.5E-1	6.2E-1	3.5E-1	
Q-Reg	3.1E-1	1.8E-1	1.1E-1	2.5E-1	
MRDR	3.7E-1	4.8E-1	2.3E0	2.4E0	
FQE	2.7E-2	1.4E-1	1.7E-1	2.3E-2	
$R(\lambda)$	6.0E-2	1.4E-1	4.1E-1	1.8E-1	
$Q^{\pi}(\lambda)$	4.1E-2	2.0E-1	1.7E-1	3.0E-2	
TREE	3.7E-2	1.8E-1	4.2E-1	1.9E-1	
IH	2.1E-2	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	2.7E-1		
WIS	1.5E0	5.2E-1		
NAIVE	4.0E0	-		

Table 315. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.2E-1	4.7E-1	3.4E-1	1.8E-1	
Q-Reg	4.6E-1	1.1E-1	1.5E-1	3.9E-1	
MRDR	7.4E-1	3.3E-1	5.8E0	5.7E0	
FQE	1.9E-2	5.2E-2	1.1E-1	8.6E-3	
$R(\lambda)$	2.7E-2	1.2E-1	1.9E-1	2.0E-1	
$Q^{\hat{\pi}}(\lambda)$	2.2E-2	6.0E-2	9.6E-2	1.0E-2	
TREE	9.0E-3	2.3E-1	1.9E-1	1.9E-1	
IH	1.3E-2	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	3.3E-1		
WIS	1.0E0	2.8E-1		
NAIVE	4.0E0	-		

Table 316. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.9E-1	1.6E0	4.9E-1	3.4E-1	
Q-Reg	6.5E-1	2.7E-1	3.4E-1	4.8E-1	
MRDR	9.1E-1	7.5E-1	2.3E0	2.3E0	
FQE	1.3E-2	2.7E-1	2.1E-1	9.4E-3	
$R(\lambda)$	3.6E-2	2.9E-1	2.8E-1	1.8E-1	
$Q^{\pi}(\lambda)$	3.9E-2	2.6E-1	2.2E-1	3.8E-2	
TREE	3.2E-2	3.3E-1	3.2E-1	2.5E-1	
IH	1.3E-2	-	-	-	

	IPS		
	STANDARD	Per-Decision	
IS	9.8E-1	8.4E-1	
WIS	1.1E0	4.2E-1	
NAIVE	4.0E0	-	

Table 317. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.1E-1	1.5E0	5.9E-1	3.1E-1
Q-Reg	4.6E-1	2.4E-1	6.2E-1	7.0E-1
MRDR	3.6E-1	7.8E-1	8.9E0	9.0E0
FQE	3.7E-2	3.0E-1	2.5E-1	5.8E-2
$R(\lambda)$	4.4E-2	4.4E-1	4.3E-1	1.8E-1
$Q^{\pi}(\lambda)$	3.0E-2	2.5E-1	2.4E-1	4.8E-2
TREE	3.9E-2	5.1E-1	4.5E-1	2.3E-1
IH	3.2E-2	-	-	-

	IPS			
	STANDARD PER-DECISIO			
IS	9.0E-1	5.9E-1		
WIS	1.2E0	5.6E-1		
NAIVE	4.0E0	-		

Table 318. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.2E-1	3.1E-1	1.2E0	3.1E-1
Q-Reg	4.1E-1	3.4E-1	4.6E-1	3.9E-1
MRDR	7.7E-1	3.9E-1	2.6E0	2.7E0
FQE	2.0E-2	4.8E-2	3.1E-1	3.4E-2
$R(\lambda)$	4.7E-2	1.2E-1	4.4E-1	8.2E-2
$Q^{\hat{\pi}}(\lambda)$	3.6E-2	3.3E-2	2.8E-1	5.0E-2
TREE	2.2E-2	1.4E-1	4.4E-1	2.1E-1
IH	1.6E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	3.2E-1		
WIS	1.7E0	5.3E-1		
NAIVE	4.0E0	-		

Table 319. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.5E0	5.6E0	5.8E0	1.6E0	
Q-Reg	8.9E-1	6.1E-1	1.7E0	1.2E0	
MRDR	7.5E-1	1.5E0	1.2E2	1.2E2	
FQE	4.3E0	3.9E0	1.7E0	4.1E0	
$R(\lambda)$	1.0E0	9.2E-1	1.6E0	1.4E0	
$Q^{\hat{\pi}}(\lambda)$	1.6E-1	9.3E-2	1.2E0	1.6E-1	
TREE	1.0E0	9.2E-1	1.6E0	1.4E0	
IH	4.3E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	9.2E-1	9.2E-1		
WIS	1.6E0	1.6E0		
NAIVE	4.1E0	-		

Table 320. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.5E0	1.7E0	5.4E0	1.6E0
Q-Reg	9.9E-1	4.1E0	1.1E0	9.1E-1
MRDR	9.6E-1	1.2E0	4.9E1	4.9E1
FQE	4.1E0	3.6E0	1.8E0	3.8E0
$R(\lambda)$	1.0E0	1.1E0	1.8E0	1.4E0
$Q^{\hat{\pi}}(\lambda)$	9.6E-2	2.3E-1	9.7E-1	8.7E-2
TREE	1.0E0	1.1E0	1.8E0	1.4E0
IH	4.0E0	_	_	-

	IPS			
	STANDARD PER-DECISION			
IS	1.1E0	1.1E0		
WIS	1.8E0	1.8E0		
NAIVE	4.0E0	-		

Table 321. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.5E0	3.2E0	8.7E0	1.7E0	
Q-Reg	9.8E-1	7.0E-1	2.0E0	9.2E-1	
MRDR	1.0E0	1.2E0	1.7E1	1.6E1	
FQE	4.1E0	3.8E0	2.6E0	3.2E0	
$R(\lambda)$	1.0E0	9.6E-1	2.6E0	1.8E0	
$Q^{\pi}(\lambda)$	3.8E-2	3.8E-2	1.2E0	3.2E-2	
TREE	1.0E0	9.6E-1	2.6E0	1.8E0	
IH	4.0E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	9.6E-1	9.6E-1		
WIS	2.6E0	2.6E0		
NAIVE	4.1E0	-		

Table 322. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	9.3E-1	8.7E1	3.8E1	1.7E0	
Q-Reg	8.0E1	4.4E1	2.2E1	3.6E1	
MRDR	4.0E2	5.3E0	1.8E2	2.1E2	
FQE	2.4E0	4.8E1	1.1E1	2.4E0	
$R(\lambda)$	2.0E0	3.3E1	1.2E1	2.1E0	
$Q^{\hat{\pi}}(\lambda)$	2.9E0	3.7E1	1.1E1	2.5E0	
TREE	2.0E0	4.1E1	1.1E1	2.2E0	
IH	2.2E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	8.2E-1	3.5E1		
WIS	4.5E1	1.1E1		
NAIVE	3.3E0	-		

Table 323. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	1.2E0	1.2E2	6.4E1	2.0E0
Q-Reg	6.2E1	1.6E1	6.4E1	6.2E1
MRDR	2.6E1	1.5E1	1.2E2	1.1E2
FQE	4.8E0	1.1E2	5.0E1	3.8E0
$R(\lambda)$	2.8E0	1.0E2	5.4E1	2.4E0
$Q^{\pi}(\lambda)$	3.2E0	9.6E1	4.9E1	3.4E0
TREE	2.2E0	1.0E2	5.0E1	1.6E0
IH	5.1E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	2.0E1	1.1E2		
WIS	4.5E1	5.0E1		
NAIVE	4.0E0	-		

Table 324. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	1.9E0	5.9E2	2.2E1	1.8E0
Q-Reg	5.9E1	3.0E1	4.1E1	8.0E1
MRDR	4.5E1	2.1E1	2.7E1	2.9E1
FQE	4.9E0	9.5E1	3.5E1	4.8E0
$R(\lambda)$	2.3E0	7.9E1	3.5E1	2.6E0
$Q^{\hat{\pi}}(\lambda)$	1.4E0	7.0E1	3.8E1	1.4E0
TREE	2.7E0	8.3E1	3.5E1	2.8E0
IH	4.9E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.2E2	8.1E1		
WIS	3.1E1	3.5E1		
NAIVE	3.9E0	_		

Table 325. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.5E0	7.2E2	2.1E1	6.9E0	
Q-Reg	1.0E0	7.9E-1	4.6E0	2.6E0	
MRDR	2.3E0	2.0E0	3.1E2	3.1E2	
FQE	3.8E0	3.5E0	6.4E0	3.7E0	
$R(\lambda)$	1.0E0	1.0E0	6.1E0	2.0E0	
$Q^{\pi}(\lambda)$	2.0E0	2.7E0	3.2E0	1.9E0	
TREE	1.0E0	1.0E0	6.1E0	2.0E0	
IH	3.6E0	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	1.0E0	
WIS	6.1E0	6.1E0	
NAIVE	4.0E0	-	

Table 326. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.5E0	1.9E0	1.4E1	2.4E0	
Q-Reg	1.1E0	1.3E0	6.8E0	1.1E0	
MRDR	1.4E0	1.1E0	2.3E1	1.7E1	
FQE	4.4E0	4.2E0	7.5E0	4.3E0	
$R(\lambda)$	1.0E0	1.0E0	7.5E0	1.0E0	
$Q^{\hat{\pi}}(\hat{\lambda})$	9.0E-1	5.1E-1	6.4E0	8.7E-1	
TREE	1.0E0	1.0E0	7.5E0	1.0E0	
IH	4.3E0	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	1.0E0	
WIS	7.5E0	7.5E0	
NAIVE	3.9E0	-	

Table 327. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.4E0	2.4E1	1.4E1	1.5E0	
Q-Reg	1.2E0	1.4E0	4.9E0	4.5E0	
MRDR	1.3E0	2.0E0	6.9E1	6.8E1	
FQE	4.1E0	4.1E0	4.7E0	3.9E0	
$R(\lambda)$	1.0E0	1.2E0	4.8E0	1.1E0	
$Q^{\hat{\pi}}(\lambda)$	3.0E-1	3.1E-1	3.4E0	2.9E-1	
TREE	1.0E0	1.2E0	4.8E0	1.1E0	
IH	3.9E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.2E0	1.2E0		
WIS	4.8E0	4.8E0		
NAIVE	3.9E0	-		

Table 328. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8$. Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.1E0	2.2E3	2.5E2	5.5E0	
Q-Reg	4.7E1	6.3E1	5.0E1	5.9E1	
MRDR	5.4E1	3.7E1	1.5E2	1.4E2	
FQE	1.2E1	9.4E1	7.6E1	1.0E1	
$R(\lambda)$	2.6E1	1.2E2	8.2E1	2.5E1	
$Q^{\hat{\pi}}(\hat{\lambda})$	1.2E1	6.8E1	7.0E1	1.1E1	
TREE	2.3E1	1.1E2	8.4E1	1.8E1	
IH	1.2E1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	9.9E-1	6.0E1		
WIS	6.5E1	7.3E1		
NAIVE	7.5E0	-		

Table 329. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

DM	Hybrid)
DIRECT	DR	WDR	MAGIC
2.3E0	8.8E2	2.6E2	3.6E0
1.8E2	9.7E1	6.5E1	1.4E2
5.9E2	7.4E1	2.2E2	4.0E2
9.0E0	1.7E2	6.7E1	1.1E1
1.2E1	1.6E2	7.1E1	1.4E1
1.4E1	2.2E2	7.1E1	1.5E1
1.3E1	1.7E2	6.8E1	1.4E1
8.8E0	-	-	-
	DIRECT 2.3E0 1.8E2 5.9E2 9.0E0 1.2E1 1.4E1 1.3E1	DIRECT DR 2.3E0 8.8E2 1.8E2 9.7E1 5.9E2 7.4E1 9.0E0 1.7E2 1.2E1 1.6E2 1.4E1 2.2E2 1.3E1 1.7E2	DIRECT DR WDR 2.3E0 8.8E2 2.6E2 1.8E2 9.7E1 6.5E1 5.9E2 7.4E1 2.2E2 9.0E0 1.7E2 6.7E1 1.2E1 1.6E2 7.1E1 1.4E1 2.2E2 7.1E1 1.3E1 1.7E2 6.8E1

	IPS			
	STANDARD PER-DECISION			
IS	1.7E0	2.0E2		
WIS	5.0E1	6.6E1		
NAIVE	5.9E0	-		

Table 330. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.9E0	1.7E2	1.9E2	4.1E0
Q-Reg	1.9E2	3.1E2	2.0E2	1.8E2
MRDR	1.6E3	7.5E0	9.4E1	9.9E1
FQE	4.8E0	5.6E1	5.6E1	5.0E0
$R(\lambda)$	2.8E0	4.6E1	5.4E1	2.8E0
$Q^{\hat{\pi}}(\lambda)$	2.7E0	4.4E1	4.7E1	3.0E0
TREE	2.6E0	4.6E1	5.4E1	2.8E0
IH	4.9E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.2E0	4.7E1		
WIS	5.0E1	5.4E1		
NAIVE	5.0E0	-		

Table 331. Graph-POMDP, relative MSE. $T=2,N=256,H=2,\pi_b(a=0)=0.6,\pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.1E-1	3.6E-3	2.8E-3	2.8E-3
Q-Reg	1.1E-1	2.9E-3	2.0E-3	3.2E-3
MRDR	4.7E-1	2.5E-3	1.4E-3	1.9E-3
FQE	1.2E-1	5.4E-3	4.1E-3	3.2E-3
$R(\lambda)$	1.7E-1	3.0E-3	1.8E-3	7.2E-3
$Q^{\hat{\pi}}(\lambda)$	5.7E-2	2.7E-3	2.0E-3	2.0E-3
TRÈE	2.1E-1	2.9E-3	1.7E-3	6.9E-3
IH	1.2E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.7E-2	1.4E-2		
WIS	9.0E-3	7.5E-3		
NAIVE	4.6E-1	-		

Table 332. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.1E-1	2.9E-3	2.0E-3	2.0E-3
Q-Reg	1.0E-1	1.7E-3	1.3E-3	1.3E-3
MRDR	4.8E-1	1.9E-3	2.0E-3	2.0E-3
FQE	1.1E-1	1.8E-3	1.3E-3	1.3E-3
$R(\lambda)$	1.6E-1	1.8E-3	1.4E-3	6.9E-3
$Q^{\pi}(\lambda)$	5.0E-2	1.9E-3	1.4E-3	1.4E-3
TREE	2.0E-1	1.8E-3	1.4E-3	4.4E-3
IH	1.1E-1	_	-	_

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	6.5E-3	4.5E-3			
WIS	2.9E-3	2.2E-3			
NAIVE	4.5E-1	-			

Table 333. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.2E-1	2.2E-3	1.8E-3	1.8E-3
Q-Reg	1.1E-1	4.8E-4	3.7E-4	3.7E-4
MRDR	5.1E-1	5.1E-4	6.4E-4	6.4E-4
FQE	1.3E-1	7.5E-4	5.7E-4	5.7E-4
$R(\lambda)$	1.8E-1	4.7E-4	3.6E-4	1.3E-3
$Q^{\pi}(\lambda)$	6.4E-2	4.8E-4	3.6E-4	3.6E-4
TREE	2.2E-1	4.7E-4	3.6E-4	1.6E-3
IH	1.3E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	3.4E-3	2.5E-3		
WIS	1.6E-3	1.2E-3		
NAIVE	4.9E-1	-		

Table 334. Graph-POMDP, relative MSE. $T=2,N=256,H=2,\pi_b(a=0)=0.6,\pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	7.4E-2	2.5E-2	2.3E-2	2.6E-2	
Q-Reg	8.4E-2	8.2E-3	7.2E-3	3.2E-2	
MRDR	4.0E-1	8.0E-3	7.3E-3	8.7E-3	
FQE	7.8E-2	9.9E-3	8.7E-3	2.2E-2	
$R(\lambda)$	1.3E-1	8.3E-3	7.2E-3	3.1E-2	
$Q^{\pi}(\lambda)$	3.9E-2	8.3E-3	7.3E-3	7.3E-3	
TREE	1.6E-1	8.2E-3	7.0E-3	2.1E-2	
IH	7.9E-2	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.8E-2	1.9E-2		
WIS	1.1E-2	1.2E-2		
NAIVE	3.6E-1	-		

Table 335. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.5E-1	1.8E-2	1.7E-2	2.0E-2
Q-Reg	1.2E-1	5.7E-3	5.5E-3	4.3E-3
MRDR	4.9E-1	5.6E-3	5.4E-3	1.0E-2
FQE	1.3E-1	6.4E-3	6.2E-3	4.4E-3
$R(\lambda)$	1.8E-1	5.6E-3	5.4E-3	1.1E-2
$Q^{\hat{\pi}}(\lambda)$	6.5E-2	5.5E-3	5.4E-3	5.4E-3
TRÈE	2.2E-1	5.6E-3	5.3E-3	1.2E-2
IH	1.3E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.1E-2	9.9E-3	
WIS	7.6E-3	7.7E-3	
NAIVE	4.7E-1	-	

Table 336. Graph-POMDP, relative MSE. $T=2,N=1024,H=2,\pi_b(a=0)=0.6,\pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.0E-1	4.7E-3	4.4E-3	4.4E-3
Q-Reg	1.1E-1	2.1E-3	1.9E-3	1.9E-3
MRDR	5.0E-1	2.1E-3	2.2E-3	2.2E-3
FQE	1.2E-1	2.2E-3	1.9E-3	1.9E-3
$R(\lambda)$	1.8E-1	2.1E-3	1.9E-3	4.7E-3
$Q^{\hat{\pi}}(\hat{\lambda})$	6.5E-2	2.2E-3	1.9E-3	1.9E-3
TRÈE	2.2E-1	2.1E-3	1.9E-3	2.3E-3
IH	1.2E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	2.7E-3	2.7E-3		
WIS	2.1E-3	2.1E-3		
NAIVE	4.8E-1	-		

Table 337. Graph-POMDP, relative MSE. $T=2, N=256, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.9E-1	5.8E-2	5.8E-2	2.4E-1
Q-Reg	1.7E-1	4.9E-2	4.9E-2	1.5E-1
MRDR	5.6E-1	4.8E-2	4.9E-2	1.7E-1
FQE	2.2E-1	4.8E-2	4.8E-2	1.6E-1
$R(\lambda)$	2.6E-1	4.8E-2	4.8E-2	1.5E-1
$Q^{\hat{\pi}}(\lambda)$	1.3E-1	4.8E-2	4.8E-2	1.6E-1
TRÈE	3.1E-1	4.8E-2	4.8E-2	1.5E-1
IH	2.1E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	5.5E-2	4.9E-2	
WIS	5.4E-2	4.9E-2	
NAIVE	5.8E-1	-	

Table 338. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Harman	
	DM	Hybrid		
	DIRECT	DR	WDR	MAGIC
AM	1.4E-1	2.7E-2	2.5E-2	5.6E-2
Q-Reg	1.2E-1	2.5E-2	2.3E-2	6.1E-2
MRDR	4.8E-1	2.4E-2	2.1E-2	2.9E-2
FQE	1.3E-1	2.8E-2	2.6E-2	4.8E-2
$R(\lambda)$	1.8E-1	2.5E-2	2.3E-2	4.2E-2
$Q^{\pi}(\lambda)$	6.7E-2	2.4E-2	2.3E-2	5.9E-2
TREE	2.3E-1	2.5E-2	2.2E-2	4.4E-2
IH	1.3E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	4.5E-2	3.6E-2	
WIS	3.8E-2	3.0E-2	
NAIVE	4.7E-1	-	

Table 339. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.6E-1	2.3E-2	2.1E-2	4.0E-2
Q-Reg	1.4E-1	2.0E-2	1.9E-2	2.1E-2
MRDR	5.1E-1	2.0E-2	1.8E-2	3.4E-2
FQE	1.6E-1	2.2E-2	2.1E-2	2.1E-2
$R(\lambda)$	2.1E-1	2.0E-2	1.9E-2	3.2E-2
$Q^{\hat{\pi}}(\lambda)$	8.7E-2	2.0E-2	1.9E-2	1.6E-2
TREE	2.6E-1	2.0E-2	1.9E-2	3.5E-2
IH	1.5E-1	-	-	-

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	3.5E-2	2.8E-2		
WIS	2.9E-2 2.3E- 2			
NAIVE	4.7E-1	-		

Table 340. Graph-POMDP, relative MSE. $T=2,N=256, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.7E-1	7.7E-2	7.4E-2	2.4E-1	
Q-Reg	1.9E-1	6.1E-2	6.0E-2	7.9E-2	
MRDR	5.6E-1	5.9E-2	5.5E-2	1.6E-1	
FQE	2.5E-1	6.8E-2	6.6E-2	1.3E-1	
$R(\lambda)$	2.6E-1	6.1E-2	6.0E-2	9.2E-2	
$Q^{\pi}(\lambda)$	1.3E-1	6.0E-2	6.0E-2	9.5E-2	
TREE	3.2E-1	6.2E-2	5.9E-2	1.1E-1	
IH	2.5E-1	-	-	-	

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	7.7E-2	7.9E-2		
WIS	6.8E-2	7.3E-2		
NAIVE	5.8E-1	-		

Table 341. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.8E-1	5.5E-2	5.6E-2	1.4E-1
Q-Reg	1.4E-1	3.5E-2	3.6E-2	5.7E-2
MRDR	5.5E-1	3.6E-2	3.8E-2	9.5E-2
FQE	1.8E-1	3.3E-2	3.3E-2	5.8E-2
$R(\lambda)$	2.3E-1	3.5E-2	3.6E-2	4.8E-2
$Q^{\pi}(\lambda)$	1.0E-1	3.6E-2	3.6E-2	5.3E-2
TREE	2.8E-1	3.5E-2	3.6E-2	5.2E-2
IH	1.8E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	3.2E-2	3.0E-2	
WIS	3.4E-2	3.2E-2	
NAIVE	5.6E-1	-	

Table 342. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.3E-1	4.4E-2	4.4E-2	1.0E-1
Q-Reg	1.5E-1	1.8E-2	1.8E-2	2.8E-2
MRDR	5.3E-1	1.8E-2	1.7E-2	7.9E-2
FQE	1.7E-1	1.9E-2	1.9E-2	2.6E-2
$R(\lambda)$	2.2E-1	1.8E-2	1.8E-2	1.9E-2
$Q^{\pi}(\lambda)$	9.6E-2	1.8E-2	1.8E-2	2.5E-2
TREE	2.6E-1	1.8E-2	1.8E-2	2.2E-2
IH	1.6E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	2.2E-2	2.2E-2	
WIS	2.0E-2	2.1E-2	
NAIVE	5.2E-1	-	

Table 343. Graph-POMDP, relative MSE. $T=2,N=256,H=2,\pi_b(a=0)=0.6,\pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.4E-1	2.4E-2	1.9E-2	1.9E-2
Q-Reg	2.9E-2	2.9E-3	3.0E-3	5.6E-3
MRDR	3.0E-1	3.3E-3	4.3E-3	4.3E-3
FQE	4.5E-1	8.0E-3	4.7E-3	4.7E-3
$R(\lambda)$	1.6E-1	2.7E-3	2.8E-3	2.8E-3
$Q^{\hat{\pi}}(\lambda)$	3.4E-2	2.5E-3	2.7E-3	3.3E-3
TREE	2.7E-1	3.1E-3	2.8E-3	2.8E-3
IH	4.5E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E-2	1.0E-2	
WIS	4.6E-3	4.6E-3	
NAIVE	4.3E-1	-	

Table 344. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.9E-1	7.9E-3	7.8E-3	7.8E-3	
Q-Reg	3.4E-2	9.8E-4	8.9E-4	1.2E-3	
MRDR	3.7E-1	8.9E-4	1.1E-3	1.1E-3	
FQE	5.1E-1	4.9E-3	2.8E-3	2.8E-3	
$R(\lambda)$	1.8E-1	1.3E-3	9.2E-4	9.2E-4	
$Q^{\hat{\pi}}(\lambda)$	4.1E-2	1.0E-3	9.0E-4	9.0E-4	
TREE	2.8E-1	1.7E-3	1.0E-3	1.0E-3	
IH	5.1E-1	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	5.7E-3	5.7E-3	
WIS	2.8E-3	2.8E-3	
NAIVE	5.2E-1	-	

Table 345. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.4E-1	2.2E-3	1.9E-3	1.9E-3
Q-Reg	2.4E-2	4.4E-4	5.1E-4	5.1E-4
MRDR	3.1E-1	5.0E-4	8.5E-4	8.5E-4
FQE	4.4E-1	7.9E-4	4.7E-4	4.7E-4
$R(\lambda)$	1.4E-1	3.8E-4	4.7E-4	4.7E-4
$Q^{\pi}(\lambda)$	2.4E-2	4.6E-4	5.2E-4	5.2E-4
TREE	2.5E-1	3.6E-4	4.1E-4	4.1E-4
IH	4.4E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.1E-3	1.1E-3	
WIS	4.6E-4	4.6E-4	
NAIVE	4.4E-1	-	

Table 346. Graph-POMDP, relative MSE. $T=2,N=256,H=2,\pi_b(a=0)=0.6,\pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.0E-1	6.3E-2	6.2E-2	1.3E-1
Q-Reg	4.4E-2	1.6E-2	1.6E-2	5.0E-2
MRDR	3.6E-1	1.6E-2	1.7E-2	2.4E-2
FQE	5.3E-1	1.9E-2	1.7E-2	1.7E-2
$R(\lambda)$	1.9E-1	1.5E-2	1.6E-2	5.1E-2
$Q^{\hat{\pi}}(\lambda)$	5.3E-2	1.5E-2	1.6E-2	4.2E-2
TREE	2.9E-1	1.5E-2	1.5E-2	2.7E-2
IH	5.3E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.8E-2	1.8E-2	
WIS	1.6E-2	1.6E-2	
NAIVE	5.1E-1	-	

Table 347. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.7E-1	2.7E-2	2.6E-2	6.4E-2
Q-Reg	4.7E-2	1.7E-2	1.6E-2	5.7E-2
MRDR	3.5E-1	1.6E-2	1.6E-2	1.6E-2
FQE	5.0E-1	1.9E-2	1.8E-2	1.8E-2
$R(\lambda)$	2.0E-1	1.7E-2	1.6E-2	2.7E-2
$Q^{\hat{\pi}}(\hat{\lambda})$	5.0E-2	1.6E-2	1.6E-2	5.2E-2
TRÈE	3.2E-1	1.7E-2	1.6E-2	1.6E-2
IH	5.0E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.9E-2	2.0E-2	
WIS	1.8E-2	1.8E-2	
NAIVE	4.8E-1	-	

Table 348. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.0E-1	2.2E-2	1.9E-2	1.9E-2
Q-Reg	3.8E-2	1.1E-2	1.0E-2	3.5E-2
MRDR	3.5E-1	1.0E-2	9.1E-3	9.1E-3
FQE	4.9E-1	1.7E-2	1.4E-2	1.4E-2
$R(\lambda)$	1.7E-1	1.2E-2	1.1E-2	1.1E-2
$Q^{\hat{\pi}}(\hat{\lambda})$	4.0E-2	1.1E-2	1.1E-2	3.3E-2
TREE	2.8E-1	1.3E-2	1.1E-2	1.1E-2
IH	4.9E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.8E-2	1.8E-2	
WIS	1.5E-2	1.4E-2	
NAIVE	5.0E-1	-	

Table 349. Graph-POMDP, relative MSE. $T=2,N=256, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.0E-1	1.4E-1	1.4E-1	3.8E-1
Q-Reg	6.8E-2	3.7E-2	3.7E-2	6.7E-2
MRDR	3.6E-1	3.7E-2	3.8E-2	1.0E-1
FQE	5.3E-1	3.5E-2	3.5E-2	7.1E-2
$R(\lambda)$	2.2E-1	3.6E-2	3.7E-2	1.7E-1
$Q^{\hat{\pi}}(\hat{\lambda})$	7.1E-2	3.6E-2	3.7E-2	6.8E-2
TRÈE	3.5E-1	3.6E-2	3.7E-2	1.1E-1
IH	5.3E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	3.7E-2	3.7E-2	
WIS	3.6E-2	3.6E-2	
NAIVE	4.8E-1	-	

Table 350. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.8E-1	9.3E-2	8.9E-2	2.1E-1
Q-Reg	6.2E-2	2.0E-2	2.0E-2	6.6E-2
MRDR	3.9E-1	2.0E-2	2.0E-2	2.0E-2
FQE	5.5E-1	2.1E-2	2.0E-2	2.0E-2
$R(\lambda)$	2.2E-1	2.0E-2	2.0E-2	3.3E-2
$Q^{\hat{\pi}}(\lambda)$	5.9E-2	2.0E-2	2.0E-2	6.3E-2
TREE	3.5E-1	2.0E-2	2.0E-2	2.0E-2
IH	5.5E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	2.1E-2	2.1E-2	
WIS	2.0E-2	2.0E-2	
NAIVE	5.1E-1	-	

Table 351. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.1E-1	2.2E-2	2.1E-2	4.2E-2
Q-Reg	1.7E-2	7.4E-3	7.3E-3	1.8E-2
MRDR	2.7E-1	7.2E-3	6.8E-3	6.8E-3
FQE	4.0E-1	9.5E-3	8.6E-3	8.6E-3
$R(\lambda)$	1.2E-1	7.7E-3	7.4E-3	7.4E-3
$Q^{\hat{\pi}}(\lambda)$	2.1E-2	7.5E-3	7.4E-3	2.2E-2
TREE	2.2E-1	8.0E-3	7.5E-3	7.5E-3
IH	4.1E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	9.9E-3	9.9E-3	
WIS	8.6E-3	8.6E-3	
NAIVE	3.9E-1	-	

Table 352. Graph-POMDP, relative MSE. $T=2,N=256, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.8E-1	1.2E-1	1.1E-1	2.4E-1
Q-Reg	9.1E-2	1.1E-1	1.1E-1	1.3E-1
MRDR	3.1E-1	1.1E-1	1.0E-1	2.1E-1
FQE	3.8E-1	1.3E-1	1.2E-1	2.9E-1
$R(\lambda)$	1.8E-1	1.1E-1	1.1E-1	1.8E-1
$Q^{\hat{\pi}}(\hat{\lambda})$	8.4E-2	1.1E-1	1.1E-1	1.3E-1
TREE	2.6E-1	1.2E-1	1.1E-1	2.5E-1
IH	3.8E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.6E-1	1.4E-1	
WIS	1.4E-1	1.2E-1	
NAIVE	4.1E-1	-	

Table 353. Graph-POMDP, relative MSE. $T=2, N=512, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.1E-1	3.1E-2	3.0E-2	3.5E-1
Q-Reg	3.7E-2	1.5E-2	1.5E-2	9.5E-2
MRDR	3.3E-1	1.5E-2	1.6E-2	1.2E-1
FQE	5.7E-1	1.7E-2	1.5E-2	9.4E-2
$R(\lambda)$	1.7E-1	1.5E-2	1.5E-2	1.5E-1
$Q^{\hat{\pi}}(\hat{\lambda})$	3.5E-2	1.5E-2	1.5E-2	8.7E-2
TREE	3.1E-1	1.5E-2	1.5E-2	1.6E-1
IH	5.7E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	2.2E-2	1.7E-2	
WIS	1.8E-2	1.5E-2	
NAIVE	5.2E-1	-	

Table 354. Graph-POMDP, relative MSE. $T=2, N=1024, H=2, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.0E-1	1.0E-1	1.0E-1	2.1E-1
Q-Reg	7.0E-2	5.7E-2	5.7E-2	1.2E-1
MRDR	3.5E-1	5.8E-2	5.9E-2	1.2E-1
FQE	5.5E-1	5.6E-2	5.6E-2	5.6E-2
$R(\lambda)$	1.9E-1	5.6E-2	5.7E-2	2.0E-1
$Q^{\hat{\pi}}(\hat{\lambda})$	7.0E-2	5.7E-2	5.7E-2	1.2E-1
TREE	3.2E-1	5.6E-2	5.7E-2	1.1E-1
IH	5.5E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	5.3E-2	5.7E-2	
WIS	5.3E-2	5.6E-2	
NAIVE	5.2E-1	-	

Table 355. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.9E-1	5.4E-3	2.1E-3	2.1E-3
Q-Reg	1.5E-2	4.4E-4	1.2E-4	1.4E-4
MRDR	1.5E-2	5.2E-3	3.6E-3	9.4E-3
FQE	3.0E-3	2.2E-4	1.7E-4	1.2E-4
$R(\lambda)$	7.3E-4	2.4E-4	1.7E-4	4.9E-4
$Q^{\hat{\pi}}(\lambda)$	6.5E-3	2.9E-4	1.6E-4	2.6E-4
TREE	4.8E-3	3.2E-3	4.8E-4	5.2E-4
IH	8.7E-4	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	2.2E-1	2.2E-2	
WIS	1.3E-2	2.0E-3	
NAIVE	4.3E-1	-	

Table 356. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.8E-1	1.4E-3	7.6E-4	7.6E-4	
Q-Reg	9.3E-3	1.2E-4	3.5E-5	8.2E-5	
MRDR	9.5E-3	6.3E-4	4.5E-4	3.1E-3	
FQE	2.6E-3	1.1E-4	5.7E-5	3.6E-5	
$R(\lambda)$	6.3E-4	1.2E-4	3.0E-5	7.5E-5	
$Q^{\pi}(\lambda)$	6.0E-3	9.9E-5	3.0E-5	1.0E-4	
TREE	5.7E-3	8.0E-4	2.0E-4	7.0E-4	
IH	1.3E-3	_	-	_	

	IPS		
	STANDARD PER-DECISION		
IS	4.7E-2	2.9E-3	
WIS	4.8E-3	4.1E-4	
NAIVE	4.4E-1	-	

Table 357. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Dense rewards

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.9E-1	9.8E-4	8.9E-4	8.9E-4
Q-Reg	1.2E-2	1.2E-4	1.1E-4	1.4E-4
MRDR	1.2E-2	7.0E-4	7.1E-4	3.0E-3
FQE	2.8E-3	1.4E-4	1.3E-4	6.3E-5
$R(\lambda)$	6.5E-4	1.7E-4	1.1E-4	1.6E-4
$Q^{\pi}(\lambda)$	6.3E-3	1.3E-4	1.1E-4	1.7E-4
TREE	6.1E-3	4.3E-4	2.3E-4	5.9E-4
IH	1.5E-3	-	-	-

	IPS		
	STANDARD PER-DECISIO		
IS	2.0E-2	3.3E-3	
WIS	2.1E-3	4.9E-4	
NAIVE	4.5E-1	-	

Table 358. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.9E-1	1.8E-2	9.5E-3	9.5E-3
Q-Reg	2.1E-2	4.2E-3	3.1E-3	5.6E-3
MRDR	1.6E-2	6.6E-3	6.3E-3	1.2E-2
FQE	3.0E-3	2.7E-3	2.4E-3	1.0E-3
$R(\lambda)$	2.8E-3	2.7E-3	2.5E-3	2.3E-3
$Q^{\hat{\pi}}(\lambda)$	7.0E-3	2.8E-3	2.6E-3	2.1E-3
TRÈE	6.5E-3	1.8E-3	2.1E-3	2.4E-3
IH	1.8E-3	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.1E-1	3.1E-2		
WIS	7.3E-3	2.5E-3		
NAIVE	4.4E-1	-		

Table 359. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

DM Hybrid DIRECT DR WDR MAGIC 3.9E-1 8.8E-3 5.4E-3 5.4E-3 AM Q-REG 4.7E-3 2.0E-3 2.1E-3 1.5E-3 MRDR 4.9E-3 1.7E-3 1.8E-3 2.9E-3 FQE 2.1E-3 2.5E-3 2.3E-3 1.6E-3 2.4E-3 2.3E-3 $R(\lambda)$ 1.1E-3 2.3E-3 $Q^{\hat{\pi}}(\lambda)$ 5.4E-32.3E-3 2.1E-3 1.4E-3 TREE 2.5E-3 3.7E-32.6E-38.8E-4 ΙH 5.9E-4

	IPS		
	STANDARD	PER-DECISION	
IS	9.5E-2	1.3E-2	
WIS	6.9E-3	3.0E-3	
NAIVE	4.2E-1	-	

Table 360. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.9E-1	4.1E-3	3.8E-3	3.8E-3
Q-Reg	5.7E-3	7.9E-4	8.6E-4	1.2E-3
MRDR	5.7E-3	1.4E-3	1.5E-3	1.1E-3
FQE	3.4E-3	7.7E-4	8.4E-4	3.0E-4
$R(\lambda)$	1.1E-3	7.7E-4	8.4E-4	6.2E-4
$Q^{\hat{\pi}}(\lambda)$	7.8E-3	7.8E-4	8.5E-4	1.7E-3
TRÈE	6.3E-3	9.5E-4	9.5E-4	1.3E-3
IH	1.5E-3	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.1E-2	2.3E-3		
WIS	2.9E-3	8.8E-4		
NAIVE	4.4E-1	-		

Table 361. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.7E-1	4.2E-3	6.1E-3	6.1E-3
Q-Reg	1.9E-2	6.6E-3	6.1E-3	3.0E-3
MRDR	2.0E-2	1.3E-2	9.6E-3	8.7E-3
FQE	3.6E-3	7.4E-3	6.7E-3	3.2E-3
$R(\lambda)$	4.5E-3	7.2E-3	6.4E-3	6.1E-3
$Q^{\hat{\pi}}(\lambda)$	1.2E-2	6.9E-3	6.5E-3	6.1E-3
TREE	6.1E-3	5.0E-3	6.7E-3	4.9E-3
IH	3.8E-3	-	-	-

	IPS		
	STANDARD PER-DECISI		
IS	6.3E-2	9.8E-3	
WIS	1.7E-2	4.4E-3	
NAIVE	4.5E-1	-	

Table 362. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.9E-1	1.2E-2	9.3E-3	9.3E-3
Q-Reg	1.0E-2	2.7E-3	2.2E-3	6.5E-3
MRDR	1.0E-2	3.6E-3	3.3E-3	6.3E-3
FQE	4.0E-3	2.9E-3	2.3E-3	1.1E-3
$R(\lambda)$	2.7E-3	2.8E-3	2.3E-3	2.4E-3
$Q^{\hat{\pi}}(\lambda)$	6.9E-3	2.9E-3	2.3E-3	1.9E-3
TREE	7.7E-3	2.6E-3	1.8E-3	3.0E-3
IH	1.6E-3	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	3.4E-2	1.1E-2	
WIS	8.3E-3	2.7E-3	
NAIVE	4.4E-1	-	

Table 363. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.0E-1	4.9E-3	5.2E-3	5.2E-3
Q-Reg	7.4E-3	1.8E-3	2.0E-3	3.2E-3
MRDR	9.1E-3	2.0E-3	2.1E-3	4.2E-3
FQE	4.9E-3	2.3E-3	2.1E-3	6.1E-4
$R(\lambda)$	1.5E-3	2.0E-3	2.0E-3	1.6E-3
$Q^{\hat{\pi}}(\lambda)$	6.9E-3	2.0E-3	2.0E-3	1.3E-3
TREE	5.5E-3	2.8E-3	2.1E-3	1.9E-3
IH	2.4E-3	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	2.8E-2	7.7E-3		
WIS	7.4E-3	2.4E-3		
NAIVE	4.6E-1	-		

Table 364. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

•	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.2E-1	5.4E-2	4.4E-2	4.4E-2
Q-Reg	3.2E-2	1.8E-2	1.8E-2	2.8E-2
MRDR	3.0E-2	1.9E-2	1.7E-2	2.5E-2
FQE	7.2E-3	1.8E-2	1.9E-2	4.1E-3
$R(\lambda)$	1.2E-2	1.9E-2	1.9E-2	1.2E-2
$Q^{\hat{\pi}}(\lambda)$	1.7E-2	1.8E-2	1.9E-2	1.1E-2
TREE	1.2E-2	1.9E-2	1.8E-2	8.5E-3
IH	6.9E-3	-	-	-

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	1.5E-1	3.7E-2			
WIS	8.6E-2	1.9E-2			
NAIVE	4.7E-1	-			

Table 365. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.7E-1	2.3E-2	1.5E-2	1.5E-2
Q-Reg	2.5E-2	8.4E-3	8.5E-3	1.4E-2
MRDR	2.2E-2	6.6E-3	7.7E-3	1.5E-2
FQE	4.5E-3	9.8E-3	8.5E-3	3.0E-3
$R(\lambda)$	6.2E-3	9.0E-3	8.7E-3	6.0E-3
$Q^{\pi}(\lambda)$	9.4E-3	9.6E-3	8.6E-3	4.6E-3
TREE	8.5E-3	1.0E-2	8.0E-3	6.0E-3
IH	2.8E-3	-	-	-

	IPS			
	STANDARD PER-DECISIO			
IS	9.1E-2	1.9E-2		
WIS	3.2E-2	8.8E-3		
NAIVE	4.5E-1	-		

Table 366. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Dense rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.3E-1	1.2E-2	1.1E-2	1.1E-2
Q-Reg	1.8E-2	1.2E-2	1.2E-2	1.3E-2
MRDR	1.7E-2	1.2E-2	1.2E-2	1.1E-2
FQE	7.5E-3	1.3E-2	1.2E-2	7.0E-3
$R(\lambda)$	7.6E-3	1.2E-2	1.2E-2	6.2E-3
$Q^{\hat{\pi}}(\lambda)$	1.2E-2	1.2E-2	1.2E-2	6.5E-3
TREE	1.5E-2	1.3E-2	1.2E-2	9.3E-3
IH	4.9E-3	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.4E-2	1.4E-2		
WIS	1.4E-2	1.2E-2		
NAIVE	4.6E-1	-		

Table 367. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.6E-1	6.2E-2	3.7E-2	2.5E-1
Q-Reg	1.2E-1	3.4E-2	3.2E-2	6.4E-2
MRDR	1.1E-1	1.9E-1	1.2E-1	1.2E-1
FQE	5.2E-1	1.3E-1	5.1E-2	1.7E-1
$R(\lambda)$	2.9E-2	3.5E-2	3.3E-2	3.1E-2
$Q^{\hat{\pi}}(\lambda)$	1.8E-2	3.5E-2	3.3E-2	1.9E-2
TREE	1.0E0	1.6E-1	5.0E-2	5.0E-2
IH	5.1E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.6E-1	1.6E-1		
WIS	5.0E-2	5.0E-2		
NAIVE	4.9E-1	-		

Table 368. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.4E-1	7.0E-2	6.6E-2	1.6E-1
Q-Reg	1.1E-2	9.5E-3	9.4E-3	1.1E-2
MRDR	1.0E-2	3.6E-2	2.2E-2	6.7E-3
FQE	3.9E-1	8.8E-3	3.9E-3	3.8E-2
$R(\lambda)$	7.6E-3	5.5E-3	6.6E-3	9.2E-3
$Q^{\hat{\pi}}(\lambda)$	1.3E-2	7.5E-3	8.5E-3	1.1E-2
TRÈE	1.0E0	2.1E-2	3.9E-3	3.9E-3
IH	4.0E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	2.1E-2	2.1E-2		
WIS	3.9E-3	3.9E-3		
NAIVE	4.2E-1	-		

Table 369. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.5E-1	4.5E-2	2.2E-2	2.2E-2
Q-Reg	3.7E-2	3.6E-3	3.8E-3	1.1E-2
MRDR	4.4E-2	1.7E-2	1.3E-2	2.6E-2
FQE	4.4E-1	2.4E-2	8.4E-3	8.4E-3
$R(\lambda)$	4.1E-3	5.1E-3	4.7E-3	5.1E-3
$Q^{\hat{\pi}}(\lambda)$	4.3E-3	5.6E-3	4.3E-3	2.9E-3
TREE	1.0E0	3.6E-2	8.3E-3	8.3E-3
IH	4.5E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	3.6E-2	3.6E-2		
WIS	8.3E-3	8.3E-3		
NAIVE	4.3E-1	-		

Table 370. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.1E0	3.2E0	3.0E0	7.4E-1	
Q-Reg	6.0E-1	4.6E-1	4.4E-1	6.4E-1	
MRDR	4.0E-1	4.8E-1	4.2E-1	4.8E-1	
FQE	4.1E-1	7.2E-1	5.5E-1	4.5E-1	
$R(\lambda)$	2.3E-1	6.4E-1	5.4E-1	3.1E-1	
$Q^{\hat{\pi}}(\lambda)$	1.6E-1	6.6E-1	5.0E-1	1.6E-1	
TREE	1.5E0	7.3E-1	5.5E-1	1.0E0	
IH	4.0E-1	-	-	-	

		IPS		
	STANDARD	PER-DECISION		
IS	2.6E0	7.9E-1		
WIS	2.6E0	5.6E-1		
NAIVE	5.5E-1	-		

Table 371. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.3E0	6.4E-1	5.9E-1	1.1E0
Q-REG	2.6E-1	2.6E-1	2.6E-1	2.4E-1
MRDR	2.2E-1	2.6E-1	2.8E-1	2.0E-1
FQE	6.7E-1	2.9E-1	2.8E-1	6.1E-1
$R(\lambda)$	1.5E-1	2.6E-1	2.6E-1	1.5E-1
$Q^{\hat{\pi}}(\lambda)$	1.4E-1	2.6E-1	2.6E-1	1.5E-1
TREE	1.1E0	3.1E-1	2.8E-1	1.0E0
IH	6.8E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	6.6E-1	2.9E-1		
WIS	7.7E-1	2.8E-1		
NAIVE	5.6E-1	-		

Table 372. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic rewards. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.3E-1	5.5E-1	5.8E-1	6.2E-1
Q-Reg	3.4E-1	3.6E-1	3.6E-1	3.2E-1
MRDR	2.9E-1	3.6E-1	3.5E-1	2.7E-1
FQE	4.8E-1	3.9E-1	3.8E-1	4.1E-1
$R(\lambda)$	2.1E-1	3.6E-1	3.6E-1	2.1E-1
$Q^{\hat{\pi}}(\lambda)$	1.4E-1	3.8E-1	3.6E-1	1.3E-1
TREE	1.1E0	3.8E-1	3.8E-1	6.3E-1
IH	4.8E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.3E0	3.8E-1		
WIS	1.3E0	3.8E-1		
NAIVE	4.5E-1	-		

Table 373. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.5E-1	8.5E-1	8.5E-1	8.1E-1
Q-Reg	1.5E-1	2.0E-1	2.5E-1	1.5E-1
MRDR	1.5E-1	2.0E-1	2.8E-1	2.9E-1
FQE	5.4E-1	1.7E-1	2.1E-1	4.7E-1
$R(\lambda)$	1.4E-1	2.3E-1	2.5E-1	1.4E-1
$Q^{\hat{\pi}}(\lambda)$	8.4E-2	2.0E-1	2.5E-1	8.6E-2
TREE	1.0E0	1.7E-1	2.0E-1	4.9E-1
IH	5.5E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.7E-1	1.7E-1	
WIS	2.0E-1	2.0E-1	
NAIVE	5.7E-1	-	

Table 374. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	8.5E-1	4.8E-1	4.6E-1	5.6E-1	
Q-Reg	1.8E-1	1.3E-1	1.3E-1	1.6E-1	
MRDR	1.6E-1	1.4E-1	1.4E-1	1.7E-1	
FQE	4.2E-1	1.9E-1	1.5E-1	2.6E-1	
$R(\lambda)$	3.6E-2	1.4E-1	1.3E-1	3.4E-2	
$Q^{\hat{\pi}}(\hat{\lambda})$	3.2E-2	1.5E-1	1.3E-1	2.7E-2	
TREE	1.0E0	2.1E-1	1.5E-1	1.5E-1	
IH	4.3E-1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	2.1E-1	2.1E-1		
WIS	1.5E-1	1.5E-1		
NAIVE	4.6E-1	-		

Table 375. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Sparse rewards.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.6E-1	3.1E-1	3.1E-1	5.2E-1
Q-Reg	1.2E-1	7.7E-2	7.5E-2	6.4E-2
MRDR	9.3E-2	6.4E-2	7.0E-2	8.1E-2
FQE	4.6E-1	1.3E-1	9.2E-2	2.2E-1
$R(\lambda)$	3.0E-2	8.2E-2	7.6E-2	3.0E-2
$Q^{\hat{\pi}}(\lambda)$	2.6E-2	8.5E-2	7.5E-2	2.1E-2
TREE	1.0E0	1.4E-1	9.3E-2	9.3E-2
IH	4.7E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.4E-1	1.4E-1		
WIS	9.3E-2	9.3E-2		
NAIVE	4.9E-1	-		

Table 376. Graph-POMDP, relative MSE. $T=16, N=256, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8$. Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	9.1E-1	6.1E0	6.5E0	1.0E0
Q-Reg	4.1E0	4.2E0	4.1E0	3.5E0
MRDR	3.7E0	3.4E0	3.4E0	3.0E0
FQE	9.8E-1	3.7E0	3.9E0	6.2E-1
$R(\lambda)$	3.3E0	4.1E0	4.1E0	2.7E0
$Q^{\hat{\pi}}(\hat{\lambda})$	2.8E0	4.4E0	4.3E0	2.3E0
TREE	3.6E0	3.8E0	4.0E0	2.9E0
IH	1.1E0	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.2E1	3.7E0	
WIS	1.0E1	4.0E0	
NAIVE	1.1E0	-	

Table 377. Graph-POMDP, relative MSE. $T=16, N=512, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM	Hybrid)
	DIRECT	DR	WDR	MAGIC
AM	1.2E0	4.7E0	4.5E0	2.0E0
Q-Reg	9.4E-1	1.3E0	1.3E0	1.0E0
MRDR	1.0E0	1.6E0	1.6E0	1.1E0
FQE	1.3E0	1.2E0	1.2E0	1.3E0
$R(\lambda)$	1.0E0	1.2E0	1.3E0	1.1E0
$Q^{\hat{\pi}}(\lambda)$	9.0E-1	1.2E0	1.2E0	9.0E-1
TREE	2.7E0	1.1E0	1.2E0	2.7E0
IH	1.3E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	2.2E0	1.2E0		
WIS	2.3E0	1.2E0		
NAIVE	1.1E0	-		

Table 378. Graph-POMDP, relative MSE. $T=16, N=1024, H=6, \pi_b(a=0)=0.6, \pi_e(a=0)=0.8.$ Stochastic environment. Stochastic rewards. Sparse rewards.

	DM		TT	
	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.8E-1	1.2E0	1.1E0	1.1E0
Q-Reg	5.6E-1	5.1E-1	5.0E-1	4.7E-1
MRDR	5.9E-1	5.9E-1	5.8E-1	4.8E-1
FQE	5.7E-1	5.6E-1	5.3E-1	5.0E-1
$R(\lambda)$	2.7E-1	5.1E-1	5.0E-1	2.2E-1
$Q^{\hat{\pi}}(\lambda)$	1.3E-1	4.6E-1	4.6E-1	1.5E-1
TREE	9.9E-1	5.9E-1	5.4E-1	9.6E-1
IH	5.7E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	7.3E-1	5.8E-1		
WIS	7.8E-1	5.4E-1		
NAIVE	6.7E-1	-		

H.3. Detailed Results for Graph Mountain Car (Graph-MC)

Table 379. Graph-MC, relative MSE. $T=250, N=128, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	5.4E-1	5.1E-1	5.0E0	4.2E0	
Q-Reg	1.5E2	1.3E1	3.5E3	1.5E2	
MRDR	9.7E2	1.3E1	7.6E4	2.7E4	
FQE	4.0E-1	4.0E-1	1.8E-1	1.5E-1	
$R(\lambda)$	4.4E-1	9.4E0	1.7E1	1.7E1	
$Q^{\hat{\pi}}(\lambda)$	1.9E130	1.9E129	1.8E131	1.0E0	
TREE	4.4E-1	9.4E0	1.7E1	1.7E1	
IH	2.0E1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	9.4E0		
WIS	2.0E1	2.0E1		
NAIVE	2.0E1	-		

Table 380. Graph-MC, relative MSE. $T=250, N=256, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.8E-1	5.1E-1	5.4E0	4.8E0	
Q-Reg	3.3E-1	3.7E-1	4.3E0	1.1E0	
MRDR	1.8E-1	4.3E-1	1.1E4	1.1E4	
FQE	3.7E-1	3.7E-1	1.8E-1	1.3E-1	
$R(\lambda)$	3.7E-1	3.7E-1	1.6E1	1.6E1	
$Q^{\hat{\pi}}(\lambda)$	9.1E118	7.5E117	1.5E119	1.0E0	
TREE	3.9E-1	3.7E-1	1.6E1	1.5E1	
IH	2.1E1	-	-	-	

		IPS		
	STANDARD	PER-DECISION		
IS	1.0E0	3.3E-1		
WIS	2.1E1	2.1E1		
NAIVE	2.1E1	-		

Table 381. Graph-MC, relative MSE. $T=250, N=512, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$

	DM	DM Hybrid		
	DIRECT	DR	WDR	MAGIC
AM	3.9E-1	3.9E-1	5.8E0	5.3E0
Q-REG	2.7E-1	1.1E0	2.7E0	8.8E-1
MRDR	1.4E-1	1.6E0	9.1E3	9.1E3
FQE	2.9E-1	2.9E-1	1.3E-1	8.8E-2
$R(\lambda)$	3.0E-1	3.0E-1	1.7E1	1.7E1
$Q^{\hat{\pi}}(\hat{\lambda})$	4.9E146	9.2E145	1.0E147	1.0E0
TREE	3.1E-1	2.9E-1	1.8E1	1.7E1
IH	2.1E1	-	-	-

	IPS		
	STANDARD	Per-Decision	
IS	1.0E0	4.4E-1	
WIS	2.1E1	2.1E1	
NAIVE	2.1E1	-	

Table 382. Graph-MC, relative MSE. $T=250, N=1024, \pi_b(a=0)=0.2, \pi_e(a=0)=0.8.$

	DM Hybrid			
	DIRECT	DR	WDR	MAGIC
AM	3.6E-1	3.9E-1	5.7E0	5.2E0
Q-REG	2.8E-1	2.4E-1	3.8E0	2.0E0
MRDR	1.7E-1	2.8E-1	2.2E4	2.2E4
FQE	2.6E-1	2.6E-1	1.1E-1	7.7E-2
$R(\lambda)$	2.6E-1	2.6E-1	1.6E1	1.5E1
$Q^{\pi}(\lambda)$	9.6E121	1.5E122	7.2E124	1.0E0
TREE	2.8E-1	2.6E-1	1.6E1	1.6E1
IH	2.0E1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	2.7E-1		
WIS	2.0E1	2.0E1		
NAIVE	2.0E1	-		

Table 383. Graph-MC, relative MSE. $T=250, N=128, \pi_b(a=0)=0.5, \pi_e(a=0)=0.5.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.5E-1	1.7E-3	1.7E-3	3.7E-3	
Q-Reg	1.7E-2	1.5E-4	1.5E-4	1.3E-4	
MRDR	3.0E-2	5.0E-3	5.0E-3	6.1E-3	
FQE	3.5E-2	6.0E-5	6.0E-5	6.5E-5	
$R(\lambda)$	1.7E-1	7.0E-4	7.0E-4	1.5E-3	
$Q^{\hat{\pi}}(\hat{\lambda})$	1.6E-1	7.1E-4	7.1E-4	1.6E-3	
TREE	7.9E-1	3.0E-4	3.0E-4	5.4E-4	
IH	2.2E-4	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	2.2E-4	2.2E-4	
WIS	2.2E-4	2.2E-4	
NAIVE	2.2E-4	-	

Table 384. Graph-MC, relative MSE. $T=250, N=256, \pi_b(a=0)=0.5, \pi_e(a=0)=0.5.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.4E-1	2.8E-4	2.8E-4	6.7E-4	
Q-Reg	1.5E-2	2.6E-5	2.6E-5	2.3E-5	
MRDR	2.6E-2	1.5E-3	1.5E-3	1.4E-3	
FQE	3.4E-2	9.0E-6	9.0E-6	1.1E-5	
$R(\lambda)$	1.7E-1	2.2E-4	2.2E-4	3.0E-4	
$Q^{\hat{\pi}}(\lambda)$	1.6E-1	2.2E-4	2.2E-4	3.0E-4	
TREE	7.9E-1	1.1E-4	1.1E-4	1.5E-4	
IH	1.0E-4	-	-	-	

	IPS		
	STANDARD PER-DECISIO		
IS	1.0E-4	1.0E-4	
WIS	1.0E-4	1.0E-4	
NAIVE	1.0E-4	-	

Table 385. Graph-MC, relative MSE. $T=250, N=512, \pi_b(a=0)=0.5, \pi_e(a=0)=0.5.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.4E-1	1.9E-4	1.9E-4	2.7E-4	
Q-Reg	1.7E-2	3.0E-6	3.0E-6	6.0E-6	
MRDR	2.9E-2	2.7E-4	2.7E-4	2.6E-4	
FQE	3.4E-2	4.0E-6	4.0E-6	9.0E-6	
$R(\lambda)$	1.6E-1	7.5E-5	7.5E-5	6.6E-5	
$Q^{\hat{\pi}}(\lambda)$	1.5E-1	7.4E-5	7.4E-5	7.4E-5	
TREE	7.9E-1	6.6E-5	6.6E-5	8.7E-5	
IH	5.9E-5	-	-	-	

	IPS			
	STANDARD	STANDARD PER-DECISION		
IS	5.9E-5	5.9E-5		
WIS	5.9E-5	5.9E-5		
NAIVE	5.9E-5	-		

Table 386. Graph-MC, relative MSE. $T=250, N=1024, \pi_b(a=0)=0.5, \pi_e(a=0)=0.5.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.4E-1	4.7E-5	4.7E-5	5.3E-5	
Q-Reg	1.7E-2	5.0E-6	5.0E-6	6.0E-6	
MRDR	2.9E-2	5.5E-5	5.5E-5	6.5E-5	
FQE	3.6E-2	3.0E-6	3.0E-6	2.0E-6	
$R(\lambda)$	1.6E-1	3.5E-5	3.5E-5	4.1E-5	
$Q^{\hat{\pi}}(\lambda)$	1.5E-1	3.4E-5	3.4E-5	4.1E-5	
TREE	7.9E-1	2.4E-5	2.4E-5	2.4E-5	
IH	3.2E-5	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	3.2E-5	3.2E-5	
WIS	3.2E-5	3.2E-5	
NAIVE	3.2E-5	-	

Table 387. Graph-MC, relative MSE. $T=250, N=128, \pi_b(a=0)=0.5, \pi_e(a=0)=0.6.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	7.6E-2	9.9E-3	7.1E-3	7.1E-3	
Q-Reg	7.3E-3	4.1E-4	7.1E-4	1.1E-3	
MRDR	5.2E-3	3.4E-2	4.8E-2	5.6E-2	
FQE	4.0E-3	4.7E-5	2.7E-4	2.7E-4	
$R(\lambda)$	1.1E-1	3.0E-3	2.4E-3	2.4E-3	
$Q^{\hat{\pi}}(\lambda)$	1.4E-3	3.1E-4	9.8E-5	9.8E-5	
TREE	5.8E-1	9.5E-3	1.5E-2	1.5E-2	
IH	4.0E-1	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	4.6E-2	1.1E-2	
WIS	1.4E-2	1.8E-2	
NAIVE	1.6E0	-	

Table 388. Graph-MC, relative MSE. $T=250, N=256, \pi_b(a=0)=0.5, \pi_e(a=0)=0.6.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	5.4E-2	4.8E-3	5.5E-3	5.5E-3	
Q-Reg	3.3E-3	1.8E-4	2.6E-4	4.1E-4	
MRDR	2.6E-3	4.3E-3	4.6E-3	6.7E-3	
FQE	2.8E-3	2.3E-4	2.4E-4	2.4E-4	
$R(\lambda)$	8.7E-2	6.2E-4	4.9E-4	4.9E-4	
$Q^{\hat{\pi}}(\lambda)$	7.0E-6	2.0E-4	2.0E-4	2.0E-4	
TREE	5.7E-1	2.6E-3	1.5E-3	1.5E-3	
IH	2.9E-1	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	1.8E-2	3.6E-3	
WIS	2.7E-3	2.2E-3	
NAIVE	1.6E0	-	

Table 389. Graph-MC, relative MSE. $T=250, N=512, \pi_b(a=0)=0.5, \pi_e(a=0)=0.6.$

	DM	Hybrid		
	DIRECT	DR	WDR	MAGIC
AM	5.2E-2	4.3E-3	3.6E-3	3.5E-3
Q-Reg	5.9E-3	3.4E-5	2.4E-5	2.5E-5
MRDR	4.5E-3	4.4E-3	4.8E-3	7.4E-3
FQE	4.3E-3	8.0E-6	1.4E-5	1.4E-5
$R(\lambda)$	7.9E-2	4.9E-4	1.5E-4	1.5E-4
$Q^{\hat{\pi}}(\hat{\lambda})$	2.1E-4	6.0E-6	6.0E-6	6.0E-6
TREE	5.7E-1	2.4E-3	1.1E-3	1.1E-3
IH	3.4E-1	-	-	-

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	1.7E-2	3.0E-3			
WIS	2.9E-3	1.8E-3			
NAIVE	1.5E0	-			

Table 390. Graph-MC, relative MSE. $T=250, N=1024, \pi_b(a=0)=0.5, \pi_e(a=0)=0.6.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.6E-2	7.3E-4	5.2E-4	5.5E-4	
Q-Reg	3.3E-3	5.3E-6	7.0E-6	8.0E-6	
MRDR	1.8E-3	1.3E-3	1.4E-3	2.1E-3	
FQE	4.3E-3	2.6E-6	6.0E-6	6.0E-6	
$R(\lambda)$	7.7E-2	6.6E-5	2.3E-4	2.4E-4	
$Q^{\pi}(\lambda)$	1.6E-4	9.4E-7	2.0E-6	2.0E-6	
TREE	5.7E-1	4.0E-4	7.1E-4	7.2E-4	
IH	2.9E-1	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	1.2E-2	6.8E-4	
WIS	1.2E-3	9.2E-4	
NAIVE	1.5E0	-	

Table 391. Graph-MC, relative MSE. $T=250, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.5.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.2E-1	4.3E-1	1.7E-1	3.1E-1	
Q-Reg	2.0E0	1.2E3	5.0E0	1.4E-1	
MRDR	2.2E0	4.6E0	1.4E1	1.5E1	
FQE	2.9E-2	1.4E0	2.9E-2	2.9E-2	
$R(\lambda)$	3.2E-1	6.0E-1	6.7E-2	1.2E-1	
$Q^{\hat{\pi}}(\lambda)$	1.0E-1	5.3E-1	6.4E-2	9.8E-2	
TREE	7.2E-1	1.5E0	3.4E-2	1.6E-1	
IH	3.6E-2	-	-	-	

-	IPS		
	STANDARD	PER-DECISION	
IS	2.5E2	1.5E0	
WIS	1.5E-1	3.1E-2	
NAIVE	3.1E-1	-	

Table 392. Graph-MC, relative MSE. $T=250, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.5.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.8E-1	1.1E-1	9.0E-2	1.8E-1	
Q-Reg	8.4E-2	2.4E-1	9.5E-2	5.0E-2	
MRDR	8.3E-2	6.2E-1	1.0E0	9.9E-1	
FQE	6.2E-3	4.8E-3	3.2E-3	6.5E-3	
$R(\lambda)$	3.0E-1	1.1E-1	2.6E-2	8.6E-2	
$Q^{\hat{\pi}}(\lambda)$	1.4E-2	2.5E-2	5.5E-3	9.3E-3	
TREE	7.1E-1	8.2E-2	1.3E-2	1.1E-1	
IH	3.3E-2	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	6.1E-1	9.1E-2	
WIS	7.4E-2	9.5E-3	
NAIVE	3.0E-1	-	

Table 393. Graph-MC, relative MSE. $T=250, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.5.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.6E-1	1.0E-1	6.6E-2	1.2E-1	
Q-Reg	6.1E-2	2.2E0	2.1E-1	6.4E-2	
MRDR	4.6E-2	4.8E-1	1.2E0	1.2E0	
FQE	5.6E-3	4.8E-3	3.8E-3	5.8E-3	
$R(\lambda)$	2.7E-1	6.4E-2	2.0E-2	6.8E-2	
$Q^{\hat{\pi}}(\lambda)$	2.1E-2	2.2E-2	7.3E-3	1.5E-2	
TREE	7.2E-1	5.7E-2	1.1E-2	7.7E-2	
IH	1.4E-2	-	-	-	

		IPS			
	STANDARD	STANDARD PER-DECISION			
IS	6.9E-1	5.7E-2			
WIS	9.1E-2	8.8E-3			
NAIVE	3.0E-1	-			

Table 394. Graph-MC, relative MSE. $T=250, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.5.$

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.4E-1	7.4E-2	3.7E-2	7.3E-2
Q-Reg	4.2E-2	2.0E-2	3.0E-2	3.0E-2
MRDR	3.2E-2	2.6E-1	5.1E-1	3.4E-1
FQE	5.4E-3	2.3E-2	2.7E-3	5.7E-3
$R(\lambda)$	2.7E-1	2.4E-2	1.2E-2	6.3E-2
$Q^{\hat{\pi}}(\lambda)$	1.7E-2	1.5E-2	2.7E-3	1.1E-2
TREE	7.2E-1	4.0E-2	6.1E-3	5.4E-2
IH	1.6E-3	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	3.3E0	4.0E-2	
WIS	5.7E-2	3.1E-3	
NAIVE	3.1E-1	-	

Table 395. Graph-MC, relative MSE. $T=250, N=128, \pi_b(a=0)=0.6, \pi_e(a=0)=0.6.$

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.8E-2	7.7E-3	7.7E-3	1.8E-2
Q-Reg	2.2E-3	3.5E-4	3.5E-4	1.7E-4
MRDR	2.3E-3	1.8E-2	1.8E-2	8.0E-3
FQE	6.3E-5	3.7E-5	3.7E-5	3.6E-5
$R(\lambda)$	3.9E-4	1.4E-4	1.4E-4	2.9E-4
$Q^{\hat{\pi}}(\hat{\lambda})$	3.6E-4	1.3E-4	1.3E-4	2.7E-4
TREE	4.7E-1	1.4E-3	1.4E-3	3.5E-3
IH	1.8E-3	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.8E-3	1.8E-3	
WIS	1.8E-3	1.8E-3	
NAIVE	1.8E-3	-	

Table 396. Graph-MC, relative MSE. $T=250, N=256, \pi_b(a=0)=0.6, \pi_e(a=0)=0.6.$

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.3E-2	1.1E-3	1.1E-3	3.7E-3
Q-Reg	9.7E-4	8.9E-5	8.9E-5	8.2E-5
MRDR	1.0E-3	3.1E-3	3.1E-3	2.1E-3
FQE	4.8E-5	3.7E-5	3.7E-5	3.6E-5
$R(\lambda)$	1.4E-4	8.4E-5	8.4E-5	9.7E-5
$Q^{\hat{\pi}}(\lambda)$	1.3E-4	8.3E-5	8.3E-5	9.4E-5
TREE	4.7E-1	2.8E-4	2.8E-4	9.3E-4
IH	3.6E-4	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	3.6E-4	3.6E-4		
WIS	3.6E-4	3.6E-4		
NAIVE	3.6E-4	-		

Table 397. Graph-MC, relative MSE. $T=250, N=512, \pi_b(a=0)=0.6, \pi_e(a=0)=0.6.$

	DM	Hybrid		
	DIRECT	DR	WDR	MAGIC
AM	3.0E-2	1.2E-3	1.2E-3	1.3E-3
Q-Reg	6.6E-4	1.7E-5	1.7E-5	1.7E-5
MRDR	7.6E-4	6.6E-4	6.6E-4	6.1E-4
FQE	1.3E-5	1.8E-5	1.8E-5	1.8E-5
$R(\lambda)$	1.0E-5	1.5E-5	1.5E-5	1.3E-5
$Q^{\hat{\pi}}(\hat{\lambda})$	1.1E-5	1.5E-5	1.5E-5	1.3E-5
TREE	4.6E-1	3.4E-4	3.4E-4	3.8E-4
IH	3.4E-4	-	-	-

	IPS				
	STANDARD	STANDARD PER-DECISION			
IS	3.4E-4	3.4E-4			
WIS	3.4E-4	3.4E-4			
NAIVE	3.4E-4	-			

Table 398. Graph-MC, relative MSE. $T=250, N=1024, \pi_b(a=0)=0.6, \pi_e(a=0)=0.6.$

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.0E-2	6.0E-4	6.0E-4	7.1E-4
Q-Reg	4.3E-4	2.0E-6	2.0E-6	2.0E-6
MRDR	4.6E-4	2.1E-4	2.1E-4	2.0E-4
FQE	2.0E-6	2.0E-6	2.0E-6	2.0E-6
$R(\lambda)$	6.0E-6	3.0E-6	3.0E-6	3.0E-6
$Q^{\hat{\pi}}(\lambda)$	5.0E-6	3.0E-6	3.0E-6	3.0E-6
TREE	4.5E-1	4.9E-4	4.9E-4	4.7E-4
IH	4.0E-4	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	4.0E-4	4.0E-4	
WIS	4.0E-4	4.0E-4	
NAIVE	4.0E-4	-	

Table 399. Graph-MC, relative MSE. $T=250, N=128, \pi_b(a=0)=0.8, \pi_e(a=0)=0.2.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	9.0E-1	9.0E-1	9.1E-1	9.0E-1	
Q-Reg	8.2E-1	9.0E-1	8.5E-1	8.1E-1	
MRDR	7.4E-1	9.2E-1	8.8E0	8.7E0	
FQE	8.7E-1	8.7E-1	8.7E-1	8.7E-1	
$R(\lambda)$	8.7E-1	8.7E-1	5.8E-1	6.1E-1	
$Q^{\hat{\pi}}(\hat{\lambda})$	1.1E125	1.0E123	9.3E124	1.0E0	
TREE	8.8E-1	8.7E-1	5.8E-1	6.1E-1	
IH	4.5E-1	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	8.2E-1	
WIS	6.7E-1	4.9E-1	
NAIVE	6.7E-1	-	

Table 400. Graph-MC, relative MSE. $T=250, N=256, \pi_b(a=0)=0.8, \pi_e(a=0)=0.2.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	8.8E-1	8.9E-1	9.0E-1	8.8E-1	
Q-Reg	7.2E-1	1.0E0	1.2E0	6.0E-1	
MRDR	7.0E-1	1.2E0	9.2E0	8.5E0	
FQE	8.3E-1	8.3E-1	8.3E-1	8.3E-1	
$R(\lambda)$	8.4E-1	8.0E-1	5.5E-1	5.8E-1	
$Q^{\hat{\pi}}(\lambda)$	5.4E107	1.6E110	2.9E108	1.0E0	
TREE	8.5E-1	7.9E-1	5.4E-1	5.7E-1	
IH	3.3E-1	-	_	-	

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	6.9E-1	
WIS	6.7E-1	4.6E-1	
NAIVE	6.7E-1	-	

Table 401. Graph-MC, relative MSE. $T=250, N=512, \pi_b(a=0)=0.8, \pi_e(a=0)=0.2.$

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	8.7E-1	8.8E-1	9.0E-1	8.7E-1	
Q-Reg	8.4E-1	8.3E-1	7.1E-1	7.8E-1	
MRDR	7.5E-1	8.4E-1	1.5E1	1.5E1	
FQE	8.5E-1	8.5E-1	8.5E-1	8.5E-1	
$R(\lambda)$	8.5E-1	8.5E-1	5.1E-1	5.4E-1	
$Q^{\hat{\pi}}(\hat{\lambda})$	3.4E114	3.3E112	3.9E114	1.0E0	
TREE	8.5E-1	8.5E-1	5.1E-1	5.5E-1	
IH	4.3E-1	-	-	-	

		IPS		
	STANDARD	STANDARD PER-DECISION		
IS	1.0E0	8.4E-1		
WIS	6.7E-1	4.3E-1		
NAIVE	6.7E-1	-		

Table 402. Graph-MC, relative MSE. $T=250, N=1024, \pi_b(a=0)=0.8, \pi_e(a=0)=0.2.$

	DM		Hybrid	
	DM		HIIDKID	
	DIRECT	DR	WDR	MAGIC
AM	8.5E-1	8.1E-1	8.2E-1	8.5E-1
Q-Reg	7.5E-1	8.3E-1	9.6E-1	9.5E-1
MRDR	6.9E-1	1.1E0	7.6E0	7.6E0
FQE	8.2E-1	8.2E-1	8.2E-1	8.2E-1
$R(\lambda)$	8.2E-1	8.1E-1	5.4E-1	5.6E-1
$Q^{\pi}(\lambda)$	2.4E112	1.5E110	6.8E112	1.0E0
TREE	8.3E-1	7.9E-1	5.2E-1	5.4E-1
IH	3.5E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	7.3E-1	
WIS	6.7E-1	4.1E-1	
NAIVE	6.7E-1	-	

H.4. Detailed Results for Mountain Car (MC)

Table 403. MC, relative MSE. Model Type: linear. $T=250, N=128, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	7.0E-2	1.3E-2	7.6E-3	1.0E-2	
Q-Reg	2.2E-1	2.6E-3	7.9E-4	2.4E-3	
MRDR	9.2E-1	8.2E-3	5.9E-4	5.9E-4	
FQE	5.7E-1	7.5E-3	1.2E-3	1.1E-3	
$R(\lambda)$	1.7E-1	1.8E-3	3.8E-4	3.9E-3	
$Q^{\hat{\pi}}(\hat{\lambda})$	1.5E-1	1.9E-3	3.9E-4	5.1E-3	
TREE	1.7E-1	1.8E-3	3.7E-4	4.3E-3	
IH	3.2E-2	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	4.4E-2	9.4E-3		
WIS	5.2E-4	6.0E-4		
NAIVE	3.2E-2	-		

Table 404. MC, relative MSE. Model Type: linear. $T=250, N=256, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	7.5E-2	7.0E-3	4.7E-4	1.0E-3	
Q-REG	2.0E-1	2.5E-3	8.6E-4	6.5E-4	
MRDR	9.2E-1	2.9E-3	2.6E-4	1.4E-4	
FQE	5.8E-1	5.3E-3	1.7E-3	1.7E-3	
$R(\lambda)$	1.7E-1	2.2E-3	5.6E-4	2.2E-3	
$Q^{\hat{\pi}}(\lambda)$	1.5E-1	2.3E-3	6.1E-4	1.5E-3	
TREE	1.7E-1	2.2E-3	5.8E-4	1.6E-3	
IH	3.2E-2	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	2.7E-2	3.3E-3	
WIS	4.7E-4	3.4E-4	
NAIVE	3.2E-2	-	

Table 405. MC, relative MSE. Model Type: NN. $T=250, N=128, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.0E-1	9.7E-3	9.5E-3	1.5E-2
Q-Reg	2.0E-1	2.7E-3	8.9E-4	3.2E-3
MRDR	8.8E-1	6.7E-3	8.4E-4	5.2E-4
FQE	1.2E-2	9.5E-4	7.6E-4	7.6E-4
$R(\lambda)$	8.9E-3	6.4E-3	1.7E-3	5.2E-3
$Q^{\hat{\pi}}(\lambda)$	1.4E-1	5.7E-3	1.6E-3	2.8E-3
TREE	8.7E-2	5.7E-3	1.5E-3	6.6E-3
IH	3.3E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	4.0E-2	7.4E-3		
WIS	8.2E-4	8.1E-4		
NAIVE	3.3E-2	-		

Table 406. MC, relative MSE. Model Type: NN. $T=250, N=256, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.4E-2	2.1E-3	2.0E-3	1.1E-2
Q-Reg	1.1E-1	1.1E-3	4.2E-4	4.5E-4
MRDR	5.9E-1	2.5E-3	6.1E-4	1.2E-3
FQE	8.6E-3	2.2E-4	1.8E-4	1.9E-4
$R(\lambda)$	1.6E-1	1.5E-3	1.2E-3	2.5E-3
$Q^{\pi}(\lambda)$	3.4E-2	1.2E-3	6.9E-4	2.6E-3
TREE	1.4E-2	1.9E-3	1.1E-3	1.5E-3
IH	3.1E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	2.6E-2	5.2E-3		
WIS	8.1E-4	5.0E-4		
NAIVE	3.1E-2	-		

Table 407. MC, relative MSE. Model Type: linear. $T=250, N=128, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=1.00$ -Greedy(DDQN).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	5.2E-1	4.1E-1	2.4E0	6.5E-1	
Q-REG	9.2E-1	8.6E-1	2.0E0	5.5E-1	
MRDR	9.9E-1	8.8E-1	4.3E-1	4.4E-1	
FQE	6.0E-1	5.2E-1	2.3E-1	2.7E-1	
$R(\lambda)$	7.3E-1	6.4E-1	4.6E-1	4.8E-1	
$Q^{\hat{\pi}}(\lambda)$	6.4E-1	5.9E-1	4.6E-1	4.5E-1	
TREE	7.3E-1	6.4E-1	4.7E-1	4.8E-1	
IH	4.1E-1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	8.9E-1		
WIS	4.9E-1	4.3E-1		
NAIVE	4.1E-1	-		

Table 408. MC, relative MSE. Model Type: linear. $T=250, N=256, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=1.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.2E-1	6.3E-1	2.2E0	2.1E-1
Q-Reg	8.8E-1	7.6E-1	5.8E-1	6.5E-1
MRDR	9.9E-1	8.6E-1	4.0E-1	4.2E-1
FQE	6.0E-1	4.9E-1	2.4E-1	3.0E-1
$R(\lambda)$	7.3E-1	6.2E-1	4.0E-1	4.5E-1
$Q^{\hat{\pi}}(\lambda)$	6.4E-1	5.4E-1	3.5E-1	4.4E-1
TREE	7.3E-1	6.2E-1	4.1E-1	4.5E-1
IH	4.0E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	8.6E-1		
WIS	4.7E-1	4.1E-1		
NAIVE	4.1E-1	-		

Table 409. MC, relative MSE. Model Type: NN. $T=250, N=128, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=1.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.5E-1	8.5E-1	7.1E-1	7.7E-1
Q-Reg	9.2E-1	7.5E-1	4.3E-1	4.8E-1
MRDR	1.0E0	8.5E-1	4.2E-1	4.5E-1
FQE	6.1E-2	4.9E-2	6.8E-2	4.8E-2
$R(\lambda)$	3.5E-1	3.0E-1	2.0E-1	1.6E-1
$Q^{\hat{\pi}}(\lambda)$	1.8E0	1.4E0	3.1E0	3.8E0
TREE	2.4E-1	2.4E-1	2.1E-1	1.5E-1
IH	4.1E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	8.4E-1		
WIS	4.7E-1	4.3E-1		
NAIVE	4.1E-1	-		

Table 410. MC, relative MSE. Model Type: NN. $T=250, N=256, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=1.00$ -Greedy(DDQN).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	9.5E-1	8.3E-1	5.7E-1	7.4E-1	
Q-Reg	9.2E-1	7.8E-1	4.4E-1	4.9E-1	
MRDR	1.0E0	8.7E-1	4.2E-1	4.4E-1	
FQE	4.3E-2	3.3E-2	1.6E-2	3.2E-2	
$R(\lambda)$	5.7E-1	4.6E-1	2.0E-1	2.2E-1	
$Q^{\hat{\pi}}(\lambda)$	2.1E1	2.0E1	1.8E1	2.7E1	
TREE	4.5E-1	3.4E-1	1.3E-1	1.7E-1	
IH	4.0E-1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	8.7E-1		
WIS	4.8E-1	4.2E-1		
NAIVE	4.1E-1	-		

Table 411. MC, relative MSE. Model Type: linear. $T=250, N=128, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.4E-1	3.2E-1	3.0E0	2.7E0	
Q-REG	6.6E-1	5.1E2	7.9E26	6.7E1	
MRDR	9.4E-1	9.7E0	5.0E0	4.9E0	
FQE	2.2E-1	1.8E0	2.9E0	2.7E0	
$R(\lambda)$	1.6E-2	3.4E0	1.9E0	1.6E0	
$Q^{\hat{\pi}}(\lambda)$	9.1E-2	3.6E0	2.0E0	1.8E0	
TREE	1.5E-2	3.4E0	1.9E0	1.7E0	
IH	5.2E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	1.1E0		
WIS	-	-		
NAIVE	5.2E0	-		

Table 412. MC, relative MSE. Model Type: linear. $T=250, N=256, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.2E-1	1.7E-1	3.6E0	3.4E0
Q-Reg	7.0E-1	2.5E23	1.5E27	1.9E1
MRDR	9.6E-1	6.7E-1	5.1E0	5.0E0
FQE	2.1E-1	2.3E-1	3.0E0	2.8E0
$R(\lambda)$	1.6E-2	8.0E-3	1.8E0	1.7E0
$Q^{\hat{\pi}}(\hat{\lambda})$	8.2E-2	2.2E-1	1.6E0	1.6E0
TREE	1.5E-2	8.4E-3	1.7E0	1.7E0
IH	5.1E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	6.7E-1	
WIS	-	-	
NAIVE	5.1E0	-	

Table 413. MC, relative MSE. Model Type: NN. $T=250, N=128, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	8.6E-1	7.3E-1	1.7E1	4.4E0	
Q-Reg	8.4E-1	5.5E-1	1.1E1	1.1E1	
MRDR	9.4E-1	5.0E-1	7.5E0	7.3E0	
FQE	1.2E-1	3.9E-1	2.3E0	2.2E0	
$R(\lambda)$	1.5E0	1.7E0	4.9E0	4.7E0	
$Q^{\pi}(\lambda)$	5.4E0	5.6E0	4.5E0	4.4E0	
TREE	2.0E0	2.3E0	4.8E0	4.6E0	
IH	5.2E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	5.4E-1		
WIS	-	-		
NAIVE	5.2E0	-		

Table 414. MC, relative MSE. Model Type: NN. $T=250, N=256, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.6E-1	4.7E-1	5.1E0	4.8E0
Q-Reg	7.7E-1	4.8E-1	1.1E1	1.1E1
MRDR	9.0E-1	4.7E-1	7.0E0	6.9E0
FQE	2.7E-2	1.3E-1	3.2E-1	2.9E-1
$R(\lambda)$	1.0E0	1.2E0	4.3E0	4.3E0
$Q^{\pi}(\lambda)$	5.6E0	5.6E0	4.4E0	4.5E0
TREE	2.1E0	2.5E0	4.7E0	4.7E0
IH	5.1E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	5.3E-1		
WIS	-	-		
NAIVE	5.1E0	-		

Table 415. MC, relative MSE. Model Type: linear. $T=250, N=128, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.10$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.3E-2	3.2E-1	2.9E0	1.0E0
Q-Reg	6.3E-1	4.6E-1	4.7E0	1.6E0
MRDR	8.5E-1	6.0E-1	3.5E-1	3.5E-1
FQE	2.7E-1	1.4E-1	7.1E-1	6.1E-1
$R(\lambda)$	6.5E-2	1.8E-2	3.5E-1	2.9E-1
$Q^{\hat{\pi}}(\lambda)$	9.4E-3	1.1E-1	4.9E-1	3.9E-1
TREE	6.3E-2	1.8E-2	3.7E-1	3.1E-1
IH	3.0E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	5.4E-1		
WIS	2.8E-1	3.6E-1		
NAIVE	3.0E0	-		

Table 416. MC, relative MSE. Model Type: linear. $T=250, N=256, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.10$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.4E-2	1.1E0	4.1E0	7.6E-1
Q-Reg	6.6E-1	4.2E-1	7.1E-1	7.4E-1
MRDR	7.4E-1	4.6E-1	5.2E-1	3.5E-1
FQE	2.6E-1	1.7E-1	7.0E-1	3.9E-1
$R(\lambda)$	6.2E-2	1.6E-2	2.9E-1	2.4E-1
$Q^{\hat{\pi}}(\hat{\lambda})$	1.4E-2	1.0E-1	4.8E-1	4.7E-1
TREE	5.9E-2	1.6E-2	3.3E-1	2.6E-1
IH	3.1E0	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	5.7E-1	
WIS	2.3E-2	2.2E-1	
NAIVE	3.1E0	-	

Table 417. MC, relative MSE. Model Type: NN. $T=250, N=128, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.10$ -Greedy(DDQN).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	8.7E-1	6.1E-1	7.7E-1	7.5E-1	
Q-Reg	8.1E-1	5.7E-1	3.7E-1	4.1E-1	
MRDR	9.6E-1	6.9E-1	6.5E-1	6.6E-1	
FQE	1.7E-2	3.9E-3	3.9E-1	2.0E-1	
$R(\lambda)$	8.1E-1	9.0E-1	1.5E0	1.4E0	
$Q^{\pi}(\lambda)$	3.7E0	3.8E0	3.2E0	3.2E0	
TREE	8.9E-1	9.8E-1	1.6E0	1.4E0	
IH	3.1E0	-	-	-	

	IPS	
	STANDARD	PER-DECISION
IS	1.0E0	7.1E-1
WIS	2.0E-1	3.9E-1
NAIVE	3.1E0	-

Table 418. MC, relative MSE. Model Type: NN. $T=250, N=256, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.10$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.8E-1	4.0E-1	5.8E-1	4.0E-1
Q-Reg	7.6E-1	4.1E-1	1.9E-1	2.1E-1
MRDR	9.7E-1	5.3E-1	3.1E-1	3.2E-1
FQE	5.0E-3	1.4E-1	1.3E-1	1.5E-2
$R(\lambda)$	1.1E0	1.2E0	2.1E0	1.6E0
$Q^{\pi}(\lambda)$	2.7E0	2.8E0	2.4E0	2.4E0
TREE	8.8E-1	1.0E0	1.6E0	1.3E0
IH	3.1E0	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	1.0E0	5.4E-1
WIS	8.7E-2	2.5E-1
NAIVE	3.1E0	-

Table 419. MC, relative MSE. Model Type: linear. $T=250, N=128, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.0E-2	1.3E-2	7.6E-3	1.0E-2
Q-Reg	2.2E-1	2.6E-3	7.9E-4	2.4E-3
MRDR	9.2E-1	8.2E-3	5.9E-4	5.9E-4
FQE	5.7E-1	7.5E-3	1.2E-3	1.1E-3
$R(\lambda)$	1.7E-1	1.8E-3	3.8E-4	3.9E-3
$Q^{\hat{\pi}}(\lambda)$	1.5E-1	1.9E-3	3.9E-4	5.1E-3
TREE	1.7E-1	1.8E-3	3.7E-4	4.3E-3
IH	3.2E-2	-	-	-

	IPS		
	STANDARD	Per-Decision	
IS	4.4E-2	9.4E-3	
WIS	5.2E-4	6.0E-4	
NAIVE	3.2E-2	-	

Table 420. MC, relative MSE. Model Type: linear. $T=250, N=256, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.5E-2	7.0E-3	4.7E-4	1.0E-3
Q-Reg	2.0E-1	2.5E-3	8.6E-4	6.5E-4
MRDR	9.2E-1	2.9E-3	2.6E-4	1.4E-4
FQE	5.8E-1	5.3E-3	1.7E-3	1.7E-3
$R(\lambda)$	1.7E-1	2.2E-3	5.6E-4	2.2E-3
$Q^{\hat{\pi}}(\hat{\lambda})$	1.5E-1	2.3E-3	6.1E-4	1.5E-3
TREE	1.7E-1	2.2E-3	5.8E-4	1.6E-3
IH	3.2E-2	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	2.7E-2	3.3E-3	
WIS	4.7E-4	3.4E-4	
NAIVE	3.2E-2	-	

Table 421. MC, relative MSE. Model Type: NN. $T=250, N=128, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.0E-1	9.7E-3	9.5E-3	1.5E-2
Q-Reg	2.0E-1	2.7E-3	8.9E-4	3.2E-3
MRDR	8.8E-1	6.7E-3	8.4E-4	5.2E-4
FQE	1.2E-2	9.5E-4	7.6E-4	7.6E-4
$R(\lambda)$	8.9E-3	6.4E-3	1.7E-3	5.2E-3
$Q^{\hat{\pi}}(\lambda)$	1.4E-1	5.7E-3	1.6E-3	2.8E-3
TREE	8.7E-2	5.7E-3	1.5E-3	6.6E-3
IH	3.3E-2	-	-	-

		IPS		
	STANDARD	PER-DECISION		
IS	4.0E-2	7.4E-3		
WIS	8.2E-4	8.1E-4		
NAIVE	3.3E-2	-		

Table 422. MC, relative MSE. Model Type: NN. $T=250, N=256, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.4E-2	2.1E-3	2.0E-3	1.1E-2
Q-Reg	1.1E-1	1.1E-3	4.2E-4	4.5E-4
MRDR	5.9E-1	2.5E-3	6.1E-4	1.2E-3
FQE	8.6E-3	2.2E-4	1.8E-4	1.9E-4
$R(\lambda)$	1.6E-1	1.5E-3	1.2E-3	2.5E-3
$Q^{\pi}(\lambda)$	3.4E-2	1.2E-3	6.9E-4	2.6E-3
TREE	1.4E-2	1.9E-3	1.1E-3	1.5E-3
IH	3.1E-2	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	2.6E-2	5.2E-3	
WIS	8.1E-4	5.0E-4	
NAIVE	3.1E-2	-	

Table 423. MC, relative MSE. Model Type: linear. $T=250, N=128, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=1.00$ -Greedy(DDQN).

DMHybrid WDR DIRECT DR MAGIC AM5.2E-1 4.1E-1 2.4E0 6.5E-1 Q-REG 9.2E-1 8.6E-1 2.0E05.5E-1 MRDR9.9E-1 4.3E-1 8.8E-1 4.4E-1 FQE 6.0E-1 5.2E-1 2.3E-1 2.7E-1 $\overset{\mathsf{R}(\lambda)}{\mathsf{Q}^{\pi}(\lambda)}$ 7.3E-1 6.4E-14.6E-1 4.8E-1 6.4E-1 5.9E-1 4.6E-1 4.5E-1 T_{REE} 7.3E-1 4.7E-16.4E-1 4.8E-1 ΙH 4.1E-1

	IPS		
	STANDARD	Per-Decision	
IS	1.0E0	8.9E-1	
WIS	4.9E-1	4.3E-1	
NAIVE	4.1E-1	-	

Table 424. MC, relative MSE. Model Type: linear. $T=250, N=256, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=1.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.2E-1	6.3E-1	2.2E0	2.1E-1
Q-Reg	8.8E-1	7.6E-1	5.8E-1	6.5E-1
MRDR	9.9E-1	8.6E-1	4.0E-1	4.2E-1
FQE	6.0E-1	4.9E-1	2.4E-1	3.0E-1
$R(\lambda)$	7.3E-1	6.2E-1	4.0E-1	4.5E-1
$Q^{\hat{\pi}}(\hat{\lambda})$	6.4E-1	5.4E-1	3.5E-1	4.4E-1
TREE	7.3E-1	6.2E-1	4.1E-1	4.5E-1
IH	4.0E-1	-	-	-

	IPS		
	STANDARD PER-DECISIO		
IS	1.0E0	8.6E-1	
WIS	4.7E-1	4.1E-1	
NAIVE	4.1E-1	-	

Table 425. MC, relative MSE. Model Type: NN. T = 250, N = 128, $\pi_b = 0.10$ -Greedy(DDQN), $\pi_e = 1.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.5E-1	8.5E-1	7.1E-1	7.7E-1
Q-REG	9.2E-1	7.5E-1	4.3E-1	4.8E-1
MRDR	1.0E0	8.5E-1	4.2E-1	4.5E-1
FQE	6.1E-2	4.9E-2	6.8E-2	4.8E-2
$R(\lambda)$	3.5E-1	3.0E-1	2.0E-1	1.6E-1
$Q^{\hat{\pi}}(\lambda)$	1.8E0	1.4E0	3.1E0	3.8E0
TREE	2.4E-1	2.4E-1	2.1E-1	1.5E-1
IH	4.1E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	8.4E-1	
WIS	4.7E-1	4.3E-1	
NAIVE	4.1E-1	-	

Table 426. MC, relative MSE. Model Type: NN. $T=250, N=256, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=1.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.5E-1	8.3E-1	5.7E-1	7.4E-1
Q-Reg	9.2E-1	7.8E-1	4.4E-1	4.9E-1
MRDR	1.0E0	8.7E-1	4.2E-1	4.4E-1
FQE	4.3E-2	3.3E-2	1.6E-2	3.2E-2
$R(\lambda)$	5.7E-1	4.6E-1	2.0E-1	2.2E-1
$Q^{\pi}(\lambda)$	2.1E1	2.0E1	1.8E1	2.7E1
TREE	4.5E-1	3.4E-1	1.3E-1	1.7E-1
IH	4.0E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	8.7E-1	
WIS	4.8E-1	4.2E-1	
NAIVE	4.1E-1	-	

Table 427. MC, relative MSE. Model Type: linear. $T=250, N=128, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

DMHybrid DIRECT DR WDR MAGIC 3.2E-1 AM1.4E-1 3.0E0 2.7E0 Q-REG 6.6E-15.1E2 7.9E26 6.7E1 5.0E0 MRDR 9.4E-19.7E0 4.9E0 FQE 2.2E-11.8E0 2.9E0 2.7E0 $R(\lambda)$ 1.6E-2 3.4E0 1.9E0 1.6E0 $Q^{\pi}(\lambda)$ 9.1E-2 3.6E0 2.0E0 1.8E0 TREE 1.5E-2 1.9E0 1.7E0 3.4E0 ΙH 5.2E0

	IPS			
	STANDARD PER-DECISIO			
IS	1.0E0	1.1E0		
WIS	-	-		
NAIVE	5.2E0	-		

Table 428. MC, relative MSE. Model Type: linear. $T=250, N=256, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.2E-1	1.7E-1	3.6E0	3.4E0
Q-Reg	7.0E-1	2.5E23	1.5E27	1.9E1
MRDR	9.6E-1	6.7E-1	5.1E0	5.0E0
FQE	2.1E-1	2.3E-1	3.0E0	2.8E0
$R(\lambda)$	1.6E-2	8.0E-3	1.8E0	1.7E0
$Q^{\hat{\pi}}(\hat{\lambda})$	8.2E-2	2.2E-1	1.6E0	1.6E0
TREE	1.5E-2	8.4E-3	1.7E0	1.7E0
IH	5.1E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	6.7E-1	
WIS	-	-	
NAIVE	5.1E0	-	

Table 429. MC, relative MSE. Model Type: NN. $T=250, N=128, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid)
	DIRECT	DR	WDR	MAGIC
AM	8.6E-1	7.3E-1	1.7E1	4.4E0
Q-Reg	8.4E-1	5.5E-1	1.1E1	1.1E1
MRDR	9.4E-1	5.0E-1	7.5E0	7.3E0
FQE	1.2E-1	3.9E-1	2.3E0	2.2E0
$R(\lambda)$	1.5E0	1.7E0	4.9E0	4.7E0
$Q^{\pi}(\lambda)$	5.4E0	5.6E0	4.5E0	4.4E0
TREE	2.0E0	2.3E0	4.8E0	4.6E0
IH	5.2E0	-	-	-

	IPS		
	STANDARD	Per-Decision	
IS	1.0E0	5.4E-1	
WIS	-	-	
NAIVE	5.2E0	-	

Table 430. MC, relative MSE. Model Type: NN. $T=250, N=256, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.6E-1	4.7E-1	5.1E0	4.8E0
Q-Reg	7.7E-1	4.8E-1	1.1E1	1.1E1
MRDR	9.0E-1	4.7E-1	7.0E0	6.9E0
FQE	2.7E-2	1.3E-1	3.2E-1	2.9E-1
$R(\lambda)$	1.0E0	1.2E0	4.3E0	4.3E0
$Q^{\pi}(\lambda)$	5.6E0	5.6E0	4.4E0	4.5E0
TREE	2.1E0	2.5E0	4.7E0	4.7E0
IH	5.1E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	5.3E-1	
WIS	-	-	
NAIVE	5.1E0	-	

Table 431. MC, relative MSE. Model Type: linear. $T=250, N=128, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.10$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.3E-2	3.2E-1	2.9E0	1.0E0
Q-Reg	6.3E-1	4.6E-1	4.7E0	1.6E0
MRDR	8.5E-1	6.0E-1	3.5E-1	3.5E-1
FQE	2.7E-1	1.4E-1	7.1E-1	6.1E-1
$R(\lambda)$	6.5E-2	1.8E-2	3.5E-1	2.9E-1
$Q^{\hat{\pi}}(\lambda)$	9.4E-3	1.1E-1	4.9E-1	3.9E-1
TREE	6.3E-2	1.8E-2	3.7E-1	3.1E-1
IH	3.0E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	5.4E-1		
WIS	2.8E-1	3.6E-1		
NAIVE	3.0E0	-		

Table 432. MC, relative MSE. Model Type: linear. $T=250, N=256, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.10$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.4E-2	1.1E0	4.1E0	7.6E-1
Q-Reg	6.6E-1	4.2E-1	7.1E-1	7.4E-1
MRDR	7.4E-1	4.6E-1	5.2E-1	3.5E-1
FQE	2.6E-1	1.7E-1	7.0E-1	3.9E-1
$R(\lambda)$	6.2E-2	1.6E-2	2.9E-1	2.4E-1
$Q^{\hat{\pi}}(\hat{\lambda})$	1.4E-2	1.0E-1	4.8E-1	4.7E-1
TREE	5.9E-2	1.6E-2	3.3E-1	2.6E-1
IH	3.1E0	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	5.7E-1	
WIS	2.3E-2	2.2E-1	
NAIVE	3.1E0	-	

Table 433. MC, relative MSE. Model Type: NN. $T=250, N=128, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.10$ -Greedy(DDQN).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	8.7E-1	6.1E-1	7.7E-1	7.5E-1	
Q-Reg	8.1E-1	5.7E-1	3.7E-1	4.1E-1	
MRDR	9.6E-1	6.9E-1	6.5E-1	6.6E-1	
FQE	1.7E-2	3.9E-3	3.9E-1	2.0E-1	
$R(\lambda)$	8.1E-1	9.0E-1	1.5E0	1.4E0	
$Q^{\hat{\pi}}(\lambda)$	3.7E0	3.8E0	3.2E0	3.2E0	
TREE	8.9E-1	9.8E-1	1.6E0	1.4E0	
IH	3.1E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	7.1E-1		
WIS	2.0E-1	3.9E-1		
NAIVE	3.1E0	-		

Table 434. MC, relative MSE. Model Type: NN. $T=250, N=256, \pi_b=1.00$ -Greedy(DDQN), $\pi_e=0.10$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	8.8E-1	4.0E-1	5.8E-1	4.0E-1
Q-Reg	7.6E-1	4.1E-1	1.9E-1	2.1E-1
MRDR	9.7E-1	5.3E-1	3.1E-1	3.2E-1
FQE	5.0E-3	1.4E-1	1.3E-1	1.5E-2
$R(\lambda)$	1.1E0	1.2E0	2.1E0	1.6E0
$Q^{\hat{\pi}}(\lambda)$	2.7E0	2.8E0	2.4E0	2.4E0
TREE	8.8E-1	1.0E0	1.6E0	1.3E0
IH	3.1E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	5.4E-1		
WIS	8.7E-2	2.5E-1		
NAIVE	3.1E0	-		

H.5. Detailed Results for Pixel-Based Mountain Car (Pix-MC)

Table 435. Pixel MC, relative MSE. Model Type: conv. $T=150, N=512, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.6E5	2.1E5	2.0E5	1.1E5	
Q-REG	6.8E-3	8.8E-3	9.1E-3	9.5E-3	
MRDR	4.7E-3	3.0E-2	4.1E-2	1.8E-2	
FQE	3.2E-3	1.1E-3	1.8E-3	9.8E-4	
$R(\lambda)$	-	-	-	-	
$Q^{\pi}(\lambda)$	-	-	-	-	
TREE	-	-	-	-	
IH	4.4E-3	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	4.4E-1	5.5E-3		
WIS	5.2E-3	3.8E-4		
NAIVE	1.0E-5	-		

Table 436. Pixel MC, relative MSE. Model Type: conv. $T=150, N=512, \pi_b=0.25$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.4E5	8.2E4	7.7E5	5.7E5
Q-Reg	1.0E-1	3.6E-3	6.9E-3	1.0E-2
MRDR	1.3E-1	9.3E-3	8.6E-3	4.4E-3
FQE	2.6E-3	7.1E-4	6.4E-4	1.7E-4
$R(\lambda)$	-	-	-	-
$Q^{\pi}(\lambda)$	-	-	-	-
TREE	-	-	-	-
IH	1.3E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	1.6E-2	
WIS	-	-	
NAIVE	7.5E-5	-	

Table 437. Pixel MC, relative MSE. Model Type: conv. $T=150, N=512, \pi_b=0.25$ -Greedy(DDQN), $\pi_e=0.10$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.1E2	1.2E3	2.6E3	9.5E2
Q-Reg	3.8E-3	3.8E-2	3.1E-2	2.2E-2
MRDR	3.6E-2	4.5E-3	4.2E-3	2.6E-3
FQE	1.5E-3	8.0E-4	8.9E-4	7.3E-4
$R(\lambda)$	-	-	-	-
$Q^{\pi}(\lambda)$	-	-	-	-
TREE	-	-	-	-
IH	3.8E-3	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	2.3E-1	4.1E-3
WIS	3.4E-4	7.8E-5
NAIVE	3.1E-5	-

H.6. Detailed Results for Gridworld

Table 438. Gridworld, relative MSE. $T=25, N=64, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.2E-2	4.3E-2	4.4E-2	4.0E-2
Q-Reg	1.0E-1	4.7E-2	4.2E-2	4.7E-2
MRDR	1.6E-1	4.5E-2	3.1E-2	2.9E-2
FQE	3.7E-2	3.6E-2	3.6E-2	3.7E-2
$R(\lambda)$	1.4E0	7.1E-2	2.6E-2	2.0E-2
$Q^{\hat{\pi}}(\lambda)$	2.3E0	7.5E-2	6.4E-2	2.9E-2
TREE	1.1E0	6.6E-2	7.4E-3	5.7E-3
IH	2.1E-2	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	4.7E-2	6.4E-2
WIS	1.6E-2	6.6E-3
NAIVE	9.6E-2	-

Table 439. Gridworld, relative MSE. $T=25, N=128, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.1E-3	1.9E-3	2.1E-3	1.7E-3
Q-Reg	3.9E-2	1.5E-2	1.3E-2	1.7E-2
MRDR	7.8E-2	1.1E-2	6.9E-3	8.7E-3
FQE	1.2E-2	1.1E-2	1.0E-2	1.1E-2
$R(\lambda)$	1.4E0	1.7E-2	1.2E-2	1.3E-2
$Q^{\hat{\pi}}(\lambda)$	1.6E0	1.5E-2	8.8E-3	7.1E-3
TREE	9.1E-1	1.7E-2	2.6E-3	3.0E-3
IH	1.1E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.3E-2	1.9E-2	
WIS	4.1E-3	1.1E-3	
NAIVE	9.6E-2	-	

Table 440. Gridworld, relative MSE. $T=25, N=256, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.0E-3	1.5E-3	1.6E-3	1.2E-3
Q-Reg	6.6E-3	3.4E-3	3.3E-3	4.1E-3
MRDR	2.9E-2	3.0E-3	2.4E-3	3.5E-3
FQE	3.5E-3	2.3E-3	2.2E-3	2.6E-3
$R(\lambda)$	1.7E-1	5.7E-3	4.4E-3	6.4E-3
$Q^{\hat{\pi}}(\lambda)$	2.3E-1	4.0E-3	2.9E-3	2.7E-3
TREE	4.3E-1	4.8E-3	1.2E-3	1.9E-3
IH	4.8E-3	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	4.8E-3	6.3E-3
WIS	1.0E-3	5.1E-4
NAIVE	1.1E-1	-

Table 441. Gridworld, relative MSE. $T=25, N=512, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.9E-4	1.2E-3	1.2E-3	1.0E-3
Q-Reg	5.0E-3	4.1E-4	4.0E-4	5.9E-4
MRDR	2.8E-2	4.1E-4	3.8E-4	6.1E-4
FQE	5.6E-4	2.4E-4	2.3E-4	3.1E-4
$R(\lambda)$	2.8E-3	4.4E-4	4.1E-4	1.3E-3
$Q^{\pi}(\lambda)$	2.4E-4	2.4E-4	2.4E-4	2.5E-4
TREE	3.4E-1	4.2E-4	2.4E-4	3.9E-4
IH	4.1E-3	-	-	-

	IPS	
	STANDARD	PER-DECISION
IS	9.9E-4	1.3E-3
WIS	4.0E-4	3.7E-4
NAIVE	9.0E-2	-

Table 442. Gridworld, relative MSE. $T=25, N=1024, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

DM Hybrid DIRECT DR WDR MAGIC 3.6E-4 1.7E-4 1.7E-4 1.1E-4 AMQ-REG 2.0E-3 3.2E-4 3.2E-45.4E-4 MRDR 2.3E-2 4.0E-43.9E-47.9E-4 FQE 8.2E-4 3.3E-43.3E-4 4.7E-4 $R(\lambda)$ 3.0E-3 4.6E-4 4.6E-4 1.5E-3 $Q^{\pi}(\lambda)$ 3.6E-4 3.4E-4 3.6E-4 3.4E-4 TREE 3.4E-18.2E-43.9E-4 3.6E-4 1.8E-3 ΙH

	IPS		
	STANDARD	PER-DECISION	
IS	7.4E-4	8.4E-4	
WIS	3.1E-4	3.4E-4	
NAIVE	9.1E-2	-	

Table 443. Gridworld, relative MSE. $T=25, N=64, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.0E-1	1.4E-1	1.2E-1	1.8E-1
Q-Reg	2.4E-1	1.1E-1	1.5E-1	3.8E-2
MRDR	2.9E-1	8.6E-2	6.5E-2	3.1E-2
FQE	3.7E-2	7.8E-3	2.5E-3	6.6E-3
$R(\lambda)$	7.7E-1	1.5E-1	2.8E-2	5.0E-2
$Q^{\hat{\pi}}(\lambda)$	5.7E-2	8.2E-3	5.2E-3	5.9E-3
TREE	1.0E0	1.7E-1	2.9E-2	1.6E-1
IH	4.8E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.2E-1	1.8E-1		
WIS	1.1E-2	3.4E-2		
NAIVE	1.2E0	-		

Table 444. Gridworld, relative MSE. $T=25, N=128, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.4E-2	1.2E-2	1.7E-2	1.2E-2
Q-Reg	2.0E-1	3.4E-2	2.7E-2	3.2E-2
MRDR	2.1E-1	5.8E-2	2.0E-2	2.1E-2
FQE	2.8E-2	2.1E-3	9.8E-4	1.5E-3
$R(\lambda)$	5.1E-1	6.1E-2	8.3E-3	1.1E-2
$Q^{\hat{\pi}}(\lambda)$	1.4E-3	1.3E-3	1.1E-3	1.5E-3
TREE	8.9E-1	1.5E-1	1.3E-2	3.8E-2
IH	3.5E-2	-	-	-

	IPS			
	STANDARD PER-DECISIO			
IS	2.6E-1	2.9E-1		
WIS	1.8E-2	3.7E-2		
NAIVE	1.2E0	-		

Table 445. Gridworld, relative MSE. $T=25, N=256, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.8E-2	2.1E-2	2.0E-2	1.8E-2
Q-Reg	7.9E-2	3.5E-3	2.9E-3	4.6E-3
MRDR	8.1E-2	1.1E-2	6.7E-3	4.1E-3
FQE	2.1E-2	4.8E-4	2.7E-4	4.9E-4
$R(\lambda)$	4.7E-1	1.5E-2	3.4E-3	4.4E-3
$Q^{\hat{\pi}}(\lambda)$	4.2E-4	3.1E-4	2.9E-4	4.9E-4
TREE	8.7E-1	4.2E-2	4.9E-3	8.9E-3
IH	2.8E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	6.6E-2	8.5E-2	
WIS	3.8E-3	1.1E-2	
NAIVE	1.3E0	-	

Table 446. Gridworld, relative MSE. $T=25, N=512, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

DMHybrid DIRECT DR WDR MAGIC 1.1E-2 7.2E-3 6.9E-3 5.6E-3 AMQ-REG 2.0E-2 6.2E-4 6.2E-49.3E-4 MRDR 9.3E-2 1.0E-3 3.5E-46.8E-4 FQE 1.8E-2 1.1E-4 4.6E-5 1.4E-4 $\overset{}{\overset{}{Q}^{\pi}(\lambda)}$ 4.6E-1 6.7E-31.8E-3 3.1E-3 2.9E-5 5.7E-5 3.6E-5 9.1E-5 TREE 8.8E-1 6.1E-3 1.3E-2 3.5E-3 ΙH 2.7E-2

	IPS			
	STANDARD PER-DECISION			
IS	1.2E-2	1.7E-2		
WIS	1.1E-3	4.2E-3		
NAIVE	1.2E0	-		

Table 447. Gridworld, relative MSE. $T=25, N=1024, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.4E-2	3.3E-3	3.3E-3	2.7E-3
Q-Reg	6.6E-3	1.5E-4	1.5E-4	2.3E-4
MRDR	6.0E-2	6.6E-4	2.7E-4	4.9E-4
FQE	1.8E-2	5.9E-5	5.4E-5	9.0E-5
$R(\lambda)$	4.8E-1	2.6E-3	2.2E-4	5.8E-4
$Q^{\hat{\pi}}(\lambda)$	2.7E-5	3.3E-5	3.5E-5	5.7E-5
TREE	9.0E-1	6.1E-3	5.6E-4	2.0E-3
IH	2.8E-2	-	-	-

	IPS			
	STANDARD PER-DECISI			
IS	7.8E-3	9.6E-3		
WIS	5.3E-4	1.1E-3		
NAIVE	1.2E0	-		

Table 448. Gridworld, relative MSE. $T=25, N=64, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.0E-1	3.6E-1	3.2E-1	5.3E-1
Q-Reg	2.7E0	2.7E0	2.4E0	1.3E0
MRDR	1.3E0	8.1E0	2.7E0	1.1E0
FQE	1.2E-1	1.1E-1	1.9E-2	1.4E-2
$R(\lambda)$	1.2E0	1.2E0	2.4E-1	1.1E0
$Q^{\hat{\pi}}(\lambda)$	1.8E-2	2.3E-2	1.1E-2	1.5E-2
TREE	1.2E0	1.5E0	3.4E-1	1.3E0
IH	1.0E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.1E0	1.6E0		
WIS	2.9E-1	4.5E-1		
NAIVE	3.9E0	-		

Table 449. Gridworld, relative MSE. $T=25, N=128, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.9E-1	2.9E-1	2.5E-1	2.6E-1
Q-Reg	5.6E-1	1.8E0	1.0E0	3.1E-1
MRDR	5.0E-1	3.7E0	1.7E0	4.6E-1
FQE	1.1E-1	2.0E-2	1.3E-2	2.9E-3
$R(\lambda)$	1.2E0	7.4E-1	9.3E-2	6.9E-1
$Q^{\pi}(\lambda)$	2.9E-3	3.4E-3	2.7E-3	2.9E-3
TREE	1.1E0	9.9E-1	1.3E-1	1.0E0
IH	8.7E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	5.5E-1	1.4E0	
WIS	2.9E-1	1.6E-1	
NAIVE	4.0E0	-	

Table 450. Gridworld, relative MSE. $T=25, N=256, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

DMHybridDIRECT DR WDR MAGIC 2.5E-1 2.1E-1 1.9E-1 1.3E-1 AMQ-REG 3.0E-1 1.6E-1 2.5E-1 1.3E-1 MRDR 4.8E-1 3.9E-1 2.6E-1 3.6E-1 FQE 1.2E-1 5.1E-3 3.0E-32.4E-3 $\overset{}{\overset{}{Q}^{\pi}(\lambda)}$ 1.2E0 4.0E-1 9.6E-2 5.2E-1 1.4E-3 7.4E-4 1.4E-3 1.1E-3 TREE 1.2E0 7.2E-1 4.7E-1 1.2E-1 8.0E-2 ΙH

	IPS		
	STANDARD	PER-DECISION	
IS	3.0E-1	5.2E-1	
WIS	5.1E-2	1.3E-1	
NAIVE	4.3E0	-	

Table 451. Gridworld, relative MSE. $T=25, N=512, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.2E-1	1.1E-1	9.3E-2	7.2E-2
Q-Reg	9.5E-1	9.2E-1	2.8E-1	3.6E-1
MRDR	6.0E-1	4.9E0	1.5E0	2.3E0
FQE	1.3E-1	1.3E-2	2.5E-3	2.2E-3
$R(\lambda)$	1.2E0	1.2E0	1.2E-1	2.3E-1
$Q^{\hat{\pi}}(\lambda)$	1.8E-3	1.6E-3	4.6E-4	6.2E-4
TREE	1.1E0	1.5E0	1.5E-1	3.3E-1
IH	7.0E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.3E0	1.7E0	
WIS	3.9E-2	1.6E-1	
NAIVE	4.3E0	-	

Table 452. Gridworld, relative MSE. $T=25, N=1024, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.5E-1	6.4E-2	5.2E-2	4.3E-2
Q-Reg	6.2E-2	1.5E-2	2.1E-2	5.7E-2
MRDR	1.7E-1	7.2E-2	3.9E-2	3.1E-1
FQE	1.3E-1	2.8E-3	8.5E-4	6.1E-4
$R(\lambda)$	1.3E0	2.4E-1	3.3E-2	1.1E-1
$Q^{\hat{\pi}}(\hat{\lambda})$	3.0E-4	3.0E-4	9.8E-5	1.2E-4
TREE	1.2E0	2.8E-1	4.0E-2	1.4E-1
IH	6.7E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS WIS	1.9E-1 8.5E-3	3.0E-1 4.2E-2	
NAIVE	4.2E0	4.2E-2	

Table 453. Gridworld, relative MSE. $T=25, N=64, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.1E0	1.4E0	1.3E0	1.1E0
Q-Reg	1.1E0	3.2E0	2.1E1	8.9E-1
MRDR	1.1E0	8.4E-1	1.1E0	1.6E0
FQE	3.6E-1	2.0E-1	1.2E-1	2.4E-1
$R(\lambda)$	1.2E0	1.2E0	1.3E0	2.3E0
$Q^{\pi}(\lambda)$	2.1E-1	2.2E-1	1.6E-1	2.2E-1
TREE	1.2E0	1.3E0	1.4E0	2.6E0
IH	2.1E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	9.4E-1	1.4E0	
WIS	1.8E0	1.6E0	
NAIVE	8.3E0	-	

Table 454. Gridworld, relative MSE. $T=25, N=128, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

DM Hybrid DIRECT DR WDR MAGIC 7.2E-1 2.3E0 1.2E0 AM1.0E0 Q-REG 4.3E0 3.9E1 2.5E1 4.0E0 MRDR 1.2E0 7.2E1 1.6E1 4.3E0 FQE 3.1E-1 9.5E-2 8.6E-2 5.9E-2 $\overset{\mathsf{R}(\lambda)}{\mathsf{Q}^{\pi}(\lambda)}$ 1.3E0 7.8E0 1.4E0 2.4E0 7.5E-2 8.0E-2 6.9E-2 7.3E-2TREE 2.7E0 1.2E0 1.1E1 1.6E0 7.9E-2 ΙH

	IPS		
	STANDARD	PER-DECISION	
IS	8.7E0	1.2E1	
WIS	1.2E0	1.7E0	
NAIVE	7.6E0	-	

Table 455. Gridworld, relative MSE. $T=25, N=256, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	9.4E-1	1.1E0	1.9E0	8.7E-1
Q-Reg	2.9E0	1.3E1	1.2E1	1.1E0
MRDR	8.0E0	3.9E1	2.9E1	5.3E0
FQE	2.7E-1	1.8E-1	3.8E-2	2.1E-2
$R(\lambda)$	1.3E0	3.8E0	5.4E-1	1.7E0
$Q^{\pi}(\lambda)$	3.0E-2	4.1E-2	2.7E-2	2.8E-2
TREE	1.2E0	4.6E0	5.7E-1	1.8E0
IH	1.4E-1	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	3.8E0	4.8E0	
WIS	4.8E-1	5.7E-1	
NAIVE	8.1E0	-	

Table 456. Gridworld, relative MSE. $T=25, N=512, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.0E-1	3.5E0	5.8E-1	3.7E-1
Q-Reg	9.8E0	1.3E1	2.7E0	3.9E1
MRDR	4.3E0	1.4E2	1.5E1	2.2E1
FQE	2.9E-1	5.0E-1	1.5E-2	1.5E-2
$R(\lambda)$	1.3E0	3.8E1	3.6E-1	1.4E0
$Q^{\hat{\pi}}(\lambda)$	1.9E-2	1.7E-1	8.8E-3	1.6E-2
TREE	1.2E0	4.2E1	3.8E-1	1.5E0
IH	9.5E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	3.0E1	4.3E1		
WIS	1.4E-1	4.0E-1		
NAIVE	7.9E0	-		

Table 457. Gridworld, relative MSE. $T=25, N=1024, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.7E-1	4.1E-1	1.3E-1	1.9E-1
Q-Reg	3.1E0	1.2E0	1.0E0	4.6E-1
MRDR	1.0E0	1.3E0	1.5E0	1.1E0
FQE	2.8E-1	4.7E-2	1.7E-2	9.6E-3
$R(\lambda)$	1.3E0	1.1E0	7.8E-1	1.5E0
$Q^{\pi}(\lambda)$	8.1E-3	8.7E-3	3.7E-3	3.2E-3
TREE	1.2E0	1.2E0	8.3E-1	1.5E0
IH	1.1E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	7.0E-1	1.2E0		
WIS	2.1E-1	8.5E-1		
NAIVE	7.9E0	-		

Table 458. Gridworld, relative MSE. $T=25, N=64, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.1E0	1.1E0	1.6E0	1.2E0	
Q-Reg	1.5E0	2.4E0	3.8E0	2.0E0	
MRDR	1.2E0	2.3E0	3.7E0	4.0E0	
FQE	1.2E0	1.2E0	1.2E0	1.2E0	
$R(\lambda)$	1.1E0	1.2E0	3.6E0	4.3E0	
$Q^{\hat{\pi}}(\lambda)$	9.9E0	1.3E1	1.0E1	8.6E0	
TREE	1.1E0	1.2E0	3.8E0	4.4E0	
IH	1.3E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	1.7E0		
WIS	8.4E0	3.9E0		
NAIVE	1.1E1	-		

Table 459. Gridworld, relative MSE. $T=25, N=128, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	1.2E0	1.1E0	1.2E0	1.2E0
Q-Reg	1.3E0	2.0E0	9.4E0	1.8E0
MRDR	1.2E0	1.1E0	4.1E0	3.7E0
FQE	8.8E-1	8.6E-1	7.4E-1	8.3E-1
$R(\lambda)$	1.2E0	1.3E0	5.1E0	4.8E0
$Q^{\pi}(\lambda)$	8.9E-1	7.6E-1	5.6E-1	7.3E-1
TREE	1.1E0	1.3E0	5.4E0	5.1E0
IH	6.9E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	1.4E0		
WIS	5.6E0	5.6E0		
NAIVE	1.1E1	-		

Table 460. Gridworld, relative MSE. $T=25, N=256, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.2E0	1.2E0	2.4E0	1.2E0	
Q-Reg	1.7E0	3.2E0	1.1E1	1.8E0	
MRDR	1.7E0	2.8E0	1.2E0	2.5E0	
FQE	4.4E-1	3.4E-1	2.1E-1	2.9E-1	
$R(\lambda)$	1.2E0	1.5E0	4.1E0	4.1E0	
$Q^{\hat{\pi}}(\lambda)$	3.7E-1	2.9E0	3.5E-1	3.3E-1	
TREE	1.1E0	1.6E0	4.4E0	4.2E0	
IH	7.2E-1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	2.0E0		
WIS	4.6E0	4.6E0		
NAIVE	1.1E1	-		

Table 461. Gridworld, relative MSE. $T=25, N=512, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.2E0	4.3E0	1.4E0	1.3E0	
Q-Reg	3.8E0	5.7E2	3.4E1	3.0E0	
MRDR	1.6E0	1.5E3	1.0E2	2.8E0	
FQE	3.9E-1	8.3E0	9.1E-2	1.4E-1	
$R(\lambda)$	1.2E0	2.2E0	4.1E0	4.0E0	
$Q^{\pi}(\lambda)$	9.6E-2	4.3E-1	7.1E-2	8.0E-2	
TREE	1.2E0	4.2E0	4.3E0	4.2E0	
IH	2.8E-1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	9.7E-1	5.5E0		
WIS	2.3E0	4.4E0		
NAIVE	1.1E1	-		

Table 462. Gridworld, relative MSE. $T=25, N=1024, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.).

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.1E0	1.1E0	2.1E0	9.2E-1	
Q-Reg	1.5E0	8.0E1	8.4E1	2.2E1	
MRDR	1.0E0	8.5E0	3.1E1	1.9E1	
FQE	3.5E-1	1.6E-1	5.1E-2	7.5E-2	
$R(\lambda)$	1.2E0	6.2E0	1.6E0	2.7E0	
$Q^{\hat{\pi}}(\lambda)$	3.6E-2	3.0E-2	1.9E-2	3.6E-2	
TREE	1.1E0	9.4E0	1.6E0	2.7E0	
IH	2.4E-1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	8.9E0	1.0E1		
WIS	1.7E0	1.6E0		
NAIVE	1.1E1	-		

H.7. Detailed Results for Pixel Gridworld

Table 463. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.3E0	2.3E1	2.3E1	7.3E-1	
Q-Reg	1.1E-1	4.3E-3	4.1E-3	4.5E-3	
MRDR	1.5E-1	1.5E-2	9.7E-3	1.4E-2	
FQE	1.8E-2	1.9E-3	1.8E-3	3.6E-3	
$R(\lambda)$	1.3E-3	8.3E-4	8.1E-4	6.9E-4	
$Q^{\hat{\pi}}(\lambda)$	2.0E-3	2.0E-3	2.0E-3	2.1E-3	
TREE	2.9E-3	1.6E-3	1.6E-3	2.0E-3	
IH	2.1E-3	-	-	-	

	IPS				
	STANDARD PER-DECISIO				
IS	1.1E-2	1.4E-2			
WIS	1.3E-3	3.9E-3			
NAIVE	9.6E-2	-			

Table 464. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.4E2	8.0E1	7.7E1	6.0E1	
Q-Reg	4.7E-2	2.3E-3	2.1E-3	1.9E-3	
MRDR	1.9E-1	6.4E-3	4.6E-3	4.1E-3	
FQE	8.9E-3	9.6E-4	1.1E-3	2.0E-3	
$R(\lambda)$	1.2E-3	6.5E-4	6.3E-4	5.8E-4	
$Q^{\pi}(\lambda)$	3.0E-3	1.4E-3	1.4E-3	1.3E-3	
TREE	3.1E-3	6.0E-4	6.2E-4	1.1E-3	
IH	1.1E-3	-	-	-	

	IPS			
	STANDARD PER-DECISIO			
IS	5.0E-3	6.3E-3		
WIS	1.6E-3	2.2E-3		
NAIVE	9.6E-2	-		

Table 465. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.0E3	1.2E1	9.5E1	1.4E2	
Q-Reg	8.7E-3	5.5E-4	6.7E-4	7.5E-4	
MRDR	2.2E-1	3.0E-3	3.2E-3	2.4E-3	
FQE	2.9E-3	2.7E-4	2.7E-4	6.3E-4	
$R(\lambda)$	6.4E-4	2.3E-4	2.3E-4	2.4E-4	
$Q^{\pi}(\lambda)$	3.5E-3	2.6E-4	2.6E-4	5.0E-4	
TREE	7.5E-4	3.8E-4	3.9E-4	4.0E-4	
IH	1.6E-3	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	6.2E-3	8.9E-3		
WIS	7.3E-4	2.3E-3		
NAIVE	1.1E-1	-		

Table 466. Pixel-Gridworld, relative MSE. $T=25, N=512, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.4E2	6.8E0	6.8E0	5.4E0	
Q-Reg	5.2E-3	3.2E-4	2.8E-4	2.9E-4	
MRDR	1.2E-1	1.2E-3	6.1E-4	4.5E-4	
FQE	9.7E-4	1.7E-4	1.7E-4	2.3E-4	
$R(\lambda)$	2.0E-3	2.2E-4	2.0E-4	9.3E-4	
$Q^{\hat{\pi}}(\hat{\lambda})$	9.5E-4	2.0E-4	2.0E-4	5.9E-4	
TREE	1.9E-3	1.8E-4	2.0E-4	3.7E-4	
IH	2.6E-4	-	-	-	

	IPS			
	STANDARD PER-DECI			
IS	2.3E-3	3.2E-3		
WIS	1.9E-4	4.2E-4		
NAIVE	9.0E-2	-		

Table 467. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.3E0	7.7E2	1.4E0	2.7E-1	
Q-Reg	3.5E0	7.3E2	1.4E0	1.7E-1	
MRDR	4.0E0	7.9E2	1.5E0	8.3E-1	
FQE	1.6E-2	7.7E0	2.3E-2	4.7E-3	
$R(\lambda)$	3.2E-3	8.4E0	4.8E-3	1.9E-3	
$Q^{\pi}(\lambda)$	2.5E-3	1.6E1	3.0E-3	3.9E-3	
TREE	1.3E-3	9.8E0	7.8E-3	2.0E-3	
IH	3.2E-2	-	-	-	

	IPS				
	STANDARD PER-DECISION				
IS	2.3E2	3.9E2			
WIS	9.2E-2	2.0E-2			
NAIVE	7.4E-2	-			

Table 468. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.4E0	2.5E1	1.0E0	6.0E-1	
Q-Reg	6.9E-1	1.6E1	8.4E-2	5.7E-2	
MRDR	1.8E0	2.0E1	1.1E-1	8.7E-2	
FQE	1.9E-2	3.3E-1	1.6E-2	4.0E-3	
$R(\lambda)$	1.2E-3	2.7E-2	3.1E-3	1.0E-3	
$Q^{\hat{\pi}}(\lambda)$	2.3E-3	1.8E-1	6.5E-3	2.5E-3	
TREE	8.0E-4	5.1E-2	4.5E-3	1.1E-3	
IH	2.0E-2	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	1.7E1	2.6E1	
WIS	6.7E-2	3.1E-2	
NAIVE	9.2E-2	-	

Table 469. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.8E1	8.4E0	4.9E-1	7.1E-1
Q-Reg	1.6E-1	2.3E-2	1.6E-2	9.0E-3
MRDR	7.6E-1	1.0E-1	2.1E-2	3.4E-2
FQE	1.9E-3	2.5E-4	2.9E-4	5.0E-4
$R(\lambda)$	2.2E-3	1.6E-3	1.4E-3	2.4E-3
$Q^{\pi}(\lambda)$	2.3E-3	4.8E-3	2.3E-3	2.1E-3
TREE	4.8E-4	2.3E-3	1.3E-3	1.2E-3
IH	4.2E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	6.0E-2	1.5E-1	
WIS	9.2E-3	1.2E-2	
NAIVE	9.1E-2	-	

Table 470. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM	Hybrid		
	DIRECT	DR	WDR	MAGIC
AM	1.5E4	1.3E16	2.8E4	1.3E4
Q-Reg	1.6E2	3.8E12	1.4E3	8.8E1
MRDR	1.2E1	1.2E14	7.2E2	3.7E1
FQE	1.8E-1	6.1E12	1.6E0	5.0E-2
$R(\lambda)$	2.4E-2	7.2E12	6.8E-1	2.7E-2
$Q^{\hat{\pi}}(\lambda)$	3.7E-2	8.4E12	4.4E-1	3.0E-2
TREE	3.8E-2	7.9E12	5.3E-1	4.2E-2
IH	4.4E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	5.6E12	1.0E12		
WIS	5.9E0	5.8E0		
NAIVE	4.6E-1	-		

Table 471. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.2E2	7.4E17	3.7E3	2.1E2
Q-REG	1.4E2	1.1E14	9.9E2	1.3E2
MRDR	5.2E0	6.9E7	3.3E1	3.4E0
FQE	1.4E-1	2.4E13	2.5E0	1.7E-1
$R(\lambda)$	4.6E-2	6.7E12	8.0E-1	1.2E-1
$Q^{\pi}(\lambda)$	5.4E-2	3.6E12	5.8E-1	9.0E-2
TREE	4.7E-2	6.2E12	1.2E0	8.6E-2
IH	9.1E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	2.2E12	1.3E14	
WIS	2.5E0	3.9E0	
NAIVE	4.4E-1	-	

Table 472. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=0.20$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.7E1	8.7E4	7.9E1	1.2E1	
Q-Reg	1.5E1	4.6E6	5.8E0	5.0E0	
MRDR	4.4E1	1.0E6	1.4E1	1.0E0	
FQE	8.2E-2	9.0E6	9.3E-1	7.6E-2	
$R(\lambda)$	3.4E-2	9.6E6	8.0E-1	1.1E-1	
$Q^{\hat{\pi}}(\hat{\lambda})$	1.7E-2	7.7E6	1.1E0	5.8E-2	
TREE	1.7E-2	1.4E7	5.8E-1	6.4E-2	
IH	7.9E-1	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	2.0E8	1.2E6	
WIS	3.9E0	2.3E0	
NAIVE	3.6E-1	-	

Table 473. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.5E3	3.5E3	3.4E3	2.9E3
Q-Reg	1.6E-1	6.3E-2	1.7E-2	9.5E-2
MRDR	4.4E-1	1.3E-1	6.1E-2	3.0E-1
FQE	1.5E-1	6.2E-3	2.7E-3	1.3E-2
$R(\lambda)$	5.9E-3	3.1E-3	3.2E-3	4.7E-3
$Q^{\pi}(\lambda)$	8.3E-3	1.2E-3	1.4E-3	2.2E-3
TREE	4.2E-3	4.9E-3	5.2E-3	7.1E-3
IH	5.3E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.2E-1	1.7E-1	
WIS	8.2E-3	2.8E-2	
NAIVE	1.2E0	-	

Table 474. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.3E0	1.7E0	6.1E-1	1.7E0	
Q-Reg	1.6E-1	1.1E-2	9.1E-3	5.6E-2	
MRDR	5.8E-1	1.6E-1	5.8E-2	1.6E-1	
FQE	7.3E-2	5.5E-3	1.4E-3	1.5E-3	
$R(\lambda)$	2.9E-3	6.6E-4	6.3E-4	9.7E-4	
$Q^{\hat{\pi}}(\lambda)$	9.2E-3	1.2E-3	5.3E-4	1.1E-3	
TREE	4.1E-3	1.2E-3	4.7E-4	1.1E-3	
IH	5.0E-2	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	2.7E-1	3.2E-1	
WIS	1.2E-2	1.9E-2	
NAIVE	1.2E0	-	

Table 475. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.8E0	1.2E0	1.1E0	5.1E-1
Q-Reg	2.1E-1	6.0E-3	5.9E-3	1.1E-2
MRDR	6.0E-1	6.5E-3	4.9E-3	1.4E-2
FQE	1.6E-1	1.1E-3	6.7E-4	1.1E-3
$R(\lambda)$	3.2E-3	4.1E-4	3.8E-4	1.3E-3
$Q^{\hat{\pi}}(\lambda)$	8.8E-3	5.6E-4	2.4E-4	1.4E-3
TREE	2.7E-3	3.8E-4	1.8E-4	5.3E-4
IH	5.5E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	9.8E-2	1.2E-1		
WIS	2.9E-3	1.1E-2		
NAIVE	1.3E0	-		

Table 476. Pixel-Gridworld, relative MSE. $T=25, N=512, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.4E0	1.9E-1	2.0E-1	3.0E-1
Q-Reg	8.1E-2	7.2E-4	6.0E-4	1.0E-3
MRDR	2.5E-1	4.6E-3	3.9E-3	3.5E-3
FQE	3.3E-3	7.8E-5	6.9E-5	2.5E-4
$R(\lambda)$	2.8E-3	7.4E-5	9.3E-5	3.1E-4
$Q^{\hat{\pi}}(\hat{\lambda})$	1.8E-3	3.4E-4	2.2E-4	7.4E-4
TREE	1.6E-3	1.1E-4	8.2E-5	8.3E-5
IH	3.8E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.9E-2	2.6E-2	
WIS	1.6E-3	6.4E-3	
NAIVE	1.2E0	-	

Table 477. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.9E1	4.6E8	5.0E1	1.5E1
Q-Reg	2.1E0	1.1E14	5.1E-1	5.0E-1
MRDR	2.3E0	4.2E19	4.3E0	1.5E0
FQE	1.5E-1	5.6E15	7.8E-2	3.8E-2
$R(\lambda)$	2.0E-3	1.2E16	1.8E-3	1.1E-3
$Q^{\pi}(\lambda)$	6.6E-3	6.2E15	3.4E-2	1.3E-2
TREE	4.5E-3	2.7E15	3.6E-2	8.7E-3
IH	7.0E-2	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	3.0E19	1.9E8		
WIS	3.5E-1	2.1E-1		
NAIVE	1.1E0	-		

Table 478. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.8E1	1.1E3	1.6E1	9.8E0	
Q-Reg	1.7E0	1.9E2	8.2E-1	9.4E-2	
MRDR	1.9E0	4.2E2	4.6E0	8.6E-1	
FQE	5.8E-2	2.2E1	2.0E-2	6.6E-3	
$R(\lambda)$	1.1E-3	2.5E0	1.9E-2	1.8E-3	
$Q^{\hat{\pi}}(\hat{\lambda})$	3.2E-2	9.8E0	3.3E-3	2.4E-3	
TREE	1.6E-3	6.7E0	3.7E-3	1.7E-3	
IH	1.1E-1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.2E3	1.9E2		
WIS	7.3E-1	1.7E-1		
NAIVE	1.3E0	-		

Table 479. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	5.5E1	1.8E2	5.5E0	5.4E1	
Q-Reg	8.3E-1	2.2E1	3.1E-1	9.7E-2	
MRDR	8.5E-1	1.8E1	1.7E-1	1.7E-1	
FQE	3.1E-2	8.9E-2	3.1E-3	2.4E-3	
$R(\lambda)$	1.7E-3	4.1E-2	2.9E-3	3.0E-3	
$Q^{\hat{\pi}}(\lambda)$	2.1E-3	5.9E-2	3.0E-3	1.6E-3	
TREE	5.9E-3	6.5E-2	6.0E-3	5.2E-3	
IH	1.0E-1	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	9.1E0	1.1E1	
WIS	4.1E-2	5.1E-2	
NAIVE	1.2E0	-	

Table 480. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.2E2	3.3E10	3.3E3	6.6E2
Q-Reg	4.9E0	1.2E6	1.4E1	2.7E0
MRDR	1.3E0	6.1E8	8.2E0	1.5E0
FQE	5.2E-1	1.1E6	8.4E-1	2.8E-1
$R(\lambda)$	7.0E-2	1.3E9	5.6E-1	1.1E-1
$Q^{\hat{\pi}}(\lambda)$	1.6E-1	4.6E8	3.5E-1	9.2E-2
TREE	7.1E-2	5.3E8	1.3E0	8.5E-2
IH	1.4E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	2.2E11	1.9E5		
WIS	2.3E0	6.9E0		
NAIVE	4.3E0	-		

Table 481. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.3E3	1.4E6	1.4E3	3.7E3	
Q-Reg	4.1E1	2.6E3	2.1E0	6.4E0	
MRDR	2.1E0	1.2E9	2.3E0	1.1E0	
FQE	3.9E-1	2.1E5	4.1E-2	7.1E-2	
$R(\lambda)$	5.2E-2	6.0E7	8.5E-2	3.6E-2	
$Q^{\hat{\pi}}(\lambda)$	1.1E-1	1.8E7	9.5E-2	8.3E-2	
TREE	6.2E-2	1.9E7	6.9E-2	3.2E-2	
IH	1.3E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.9E10	3.1E2		
WIS	2.0E0	1.2E0		
NAIVE	4.0E0	-		

Table 482. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=0.40$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.3E3	4.9E2	4.9E2	6.2E2	
Q-Reg	2.5E0	1.8E-1	1.2E-2	2.1E-1	
MRDR	2.6E0	1.9E0	1.7E-1	5.1E-1	
FQE	2.3E-1	7.0E-2	5.0E-2	2.6E-2	
$R(\lambda)$	3.5E-2	9.0E-2	5.9E-2	1.9E-2	
$Q^{\hat{\pi}}(\lambda)$	6.5E-2	5.3E-2	6.6E-2	7.1E-2	
TREE	6.2E-2	7.8E-2	4.2E-2	2.7E-2	
IH	1.0E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.3E0	1.1E0		
WIS	3.1E-1	1.9E-1		
NAIVE	3.9E0	-		

Table 483. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	6.0E1	2.6E2	5.1E1	5.8E1	
Q-REG	1.9E0	3.2E0	6.4E-1	1.6E0	
MRDR	7.7E-1	2.5E0	1.2E0	1.0E0	
FQE	2.5E-1	4.1E-1	3.4E-2	6.9E-2	
$R(\lambda)$	7.2E-3	3.6E-2	2.7E-2	8.7E-3	
$Q^{\hat{\pi}}(\lambda)$	3.0E-2	4.0E-2	1.6E-2	9.5E-3	
TREE	9.0E-3	3.4E-1	2.7E-2	6.2E-3	
IH	4.7E-1	-	-	-	

	IPS			
	STANDARD	PER-DECISION		
IS	2.8E0	3.9E0		
WIS	2.0E-1	5.6E-1		
NAIVE	3.9E0	-		

Table 484. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.9E1	5.7E1	2.8E1	5.9E0	
Q-Reg	3.5E0	3.1E1	6.2E-1	8.6E-1	
MRDR	7.3E-1	4.1E0	2.7E-1	1.1E0	
FQE	9.3E-2	3.1E-2	7.4E-3	6.6E-3	
$R(\lambda)$	3.5E-3	5.3E-3	2.9E-3	1.8E-3	
$Q^{\hat{\pi}}(\lambda)$	2.7E-2	1.1E-2	4.1E-3	1.2E-2	
TREE	7.5E-3	5.1E-2	2.9E-3	6.5E-3	
IH	4.9E-1	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	2.3E0	5.1E0		
WIS	2.3E-1	1.7E-1		
NAIVE	4.0E0	-		

Table 485. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	5.4E1	1.5E1	1.9E1	5.2E1
Q-Reg	3.1E-1	1.5E-1	3.8E-2	1.8E-1
MRDR	1.1E0	6.3E-1	3.7E-1	3.8E-1
FQE	7.4E-2	5.9E-3	6.5E-3	6.3E-3
$R(\lambda)$	8.1E-3	5.5E-3	4.7E-3	2.3E-3
$Q^{\pi}(\lambda)$	4.1E-2	1.1E-2	2.1E-3	1.9E-2
TREE	3.2E-3	3.3E-3	2.4E-3	2.9E-3
IH	5.2E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	3.1E-1	5.1E-1		
WIS	6.3E-2	1.4E-1		
NAIVE	4.3E0	-		

Table 486. Pixel-Gridworld, relative MSE. $T=25, N=512, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.5E0	8.7E0	7.0E0	2.9E0
Q-Reg	1.5E0	1.5E0	1.5E-1	6.4E-1
MRDR	2.7E0	3.8E0	1.0E0	1.5E0
FQE	3.4E-2	3.2E-3	1.3E-3	1.2E-3
$R(\lambda)$	1.1E-2	1.6E-2	4.9E-3	6.3E-3
$Q^{\pi}(\lambda)$	4.6E-2	2.0E-2	2.0E-2	7.1E-3
TREE	8.3E-3	9.8E-3	5.0E-3	2.2E-3
IH	5.2E-1	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.2E0	1.7E0		
WIS	4.6E-2	1.5E-1		
NAIVE	4.3E0	-		

Table 487. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	9.8E0	1.6E20	9.4E1	1.1E1	
Q-Reg	1.0E1	7.3E10	1.3E1	2.3E1	
MRDR	2.2E0	2.9E20	2.8E0	3.5E0	
FQE	3.9E-1	1.5E17	6.8E-2	1.7E-1	
$R(\lambda)$	2.7E-2	2.5E18	1.2E-2	2.2E-2	
$Q^{\hat{\pi}}(\lambda)$	5.1E-2	6.0E18	3.5E-2	2.2E-2	
TREE	1.9E-2	2.5E18	2.5E-2	2.3E-2	
IH	4.8E-1	-	-	-	

	IPS				
	STANDARD PER-DECISION				
IS	1.5E21	2.4E5			
WIS	2.1E-1	9.3E-1			
NAIVE	4.5E0	-			

Table 488. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	1.4E1	5.5E14	8.8E1	1.2E1	
Q-REG	6.2E-1	3.0E15	1.3E-1	6.5E-1	
MRDR	9.1E-1	1.6E18	3.4E-1	5.6E-1	
FQE	1.4E-1	3.4E15	2.5E-2	3.9E-2	
$R(\lambda)$	9.0E-3	3.0E15	1.1E-2	3.3E-3	
$Q^{\hat{\pi}}(\lambda)$	2.2E-2	3.3E16	1.3E-2	2.2E-2	
TREE	1.2E-2	9.1E16	1.9E-2	7.9E-3	
IH	5.9E-1	-	-	_	

	IPS			
	STANDARD PER-DECISI			
IS	6.1E18	4.4E5		
WIS	1.7E-1	5.0E-1		
NAIVE	4.5E0	-		

Table 489. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.6E1	1.7E2	5.0E1	8.0E1	
Q-Reg	5.7E-1	6.4E-1	1.2E-1	4.9E-1	
MRDR	4.3E-1	1.6E1	5.3E0	3.5E-1	
FQE	4.1E-2	9.9E-2	3.3E-3	1.8E-3	
$R(\lambda)$	2.3E-2	9.6E-2	7.3E-3	3.6E-3	
$Q^{\pi}(\lambda)$	3.5E-2	1.0E-1	1.8E-2	1.1E-2	
TREE	7.9E-3	4.2E-2	2.6E-3	3.6E-3	
IH	5.7E-1	-	-	-	

	IPS			
	STANDARD PER-DECISIO			
IS	7.5E-1	2.3E0		
WIS	1.0E-1	1.0E-1		
NAIVE	4.4E0	-		

Table 490. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.0E2	4.8E16	3.0E2	2.2E2	
Q-Reg	9.9E0	8.3E11	2.0E1	1.6E1	
MRDR	1.5E0	6.5E18	4.0E1	2.7E1	
FQE	9.0E-1	2.3E13	5.2E-1	5.9E-1	
$R(\lambda)$	2.9E-1	1.4E16	3.1E-1	2.2E-1	
$Q^{\hat{\pi}}(\lambda)$	5.0E-1	1.3E16	2.8E-1	2.8E-1	
TREE	1.8E-1	1.3E15	3.5E-1	1.6E-1	
IH	2.9E0	-	-	-	

	IPS			
	STANDARD PER-DECISIO			
IS	2.6E18	2.2E5		
WIS	2.2E0	2.2E0		
NAIVE	1.2E1	-		

Table 491. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	7.4E1	2.9E12	1.1E4	9.7E1	
Q-Reg	1.6E0	7.0E7	1.6E0	2.4E0	
MRDR	3.8E0	6.6E14	4.1E0	4.3E0	
FQE	1.0E0	8.8E12	3.4E-1	4.5E-1	
$R(\lambda)$	4.0E-1	2.5E14	2.3E-1	1.7E-1	
$Q^{\pi}(\lambda)$	5.9E-1	8.7E13	3.2E-1	3.0E-1	
TREE	3.9E-1	9.4E13	3.1E-1	2.5E-1	
IH	3.1E0	-	-	-	

	IPS				
	STANDARD PER-DECISION				
IS	5.6E16	6.2E1			
WIS	2.4E0	2.3E0			
NAIVE	1.2E1	-			

Table 492. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=0.60$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	9.2E3	3.3E5	6.5E4	9.4E3	
Q-Reg	1.4E0	3.6E2	1.4E0	1.7E0	
MRDR	9.9E-1	2.3E11	1.4E0	1.5E0	
FQE	3.9E-1	4.7E8	7.3E-2	7.2E-2	
$R(\lambda)$	1.2E-1	7.3E10	1.1E-1	6.9E-2	
$Q^{\hat{\pi}}(\lambda)$	3.8E-1	3.0E10	3.7E-1	2.5E-1	
TREE	1.7E-1	2.9E10	1.2E-1	9.9E-2	
IH	3.2E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	2.2E13	3.1E0		
WIS	1.6E0	1.8E0		
NAIVE	1.2E1	-		

Table 493. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.6E0	8.2E1	2.3E1	4.1E0
Q-Reg	1.2E0	1.8E0	1.1E0	2.5E0
MRDR	1.2E0	8.8E-1	9.6E-1	2.1E0
FQE	2.8E0	1.5E0	4.1E-1	1.7E0
$R(\lambda)$	1.4E-1	1.7E-1	9.5E-2	7.1E-2
$Q^{\pi}(\lambda)$	6.2E0	4.3E0	7.5E-1	3.7E0
TREE	4.2E-1	3.7E-1	2.6E-1	3.7E-1
IH	1.6E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	1.5E0		
WIS	1.9E0	1.5E0		
NAIVE	8.3E0	-		

Table 494. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM	Hybrid		
	DIRECT	DR	WDR	MAGIC
AM	9.8E1	1.5E3	4.9E2	9.6E1
Q-Reg	1.2E0	3.3E0	1.5E0	2.3E0
MRDR	8.4E-1	9.7E0	3.0E1	1.3E0
FQE	4.7E-1	2.0E-1	9.9E-2	1.7E-1
$R(\lambda)$	7.3E-2	1.7E-1	7.9E-2	3.9E-2
$Q^{\pi}(\lambda)$	1.8E-1	3.9E-1	6.2E-2	3.9E-2
TREE	2.2E-2	7.2E-1	6.4E-2	4.9E-2
IH	1.4E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	3.0E0	4.4E0		
WIS	1.2E0	1.5E0		
NAIVE	7.6E0	-		

Table 495. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	7.4E0	1.2E2	1.4E2	5.2E0	
Q-Reg	1.4E0	3.1E0	3.0E0	2.0E0	
MRDR	9.4E-1	3.1E0	3.2E0	1.1E0	
FQE	3.6E-1	6.2E-2	4.6E-2	4.5E-2	
$R(\lambda)$	9.0E-2	1.9E-1	3.5E-2	2.8E-2	
$Q^{\hat{\pi}}(\lambda)$	8.3E-2	1.8E-1	7.4E-2	5.4E-2	
TREE	2.5E-1	1.1E-1	7.8E-2	1.3E-1	
IH	1.5E0	-	-	-	

	IPS		
	STANDARD	PER-DECISION	
IS	1.4E0	1.9E0	
WIS	4.7E-1	6.6E-1	
NAIVE	8.1E0	-	

Table 496. Pixel-Gridworld, relative MSE. $T=25, N=512, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.5E1	1.1E5	4.2E2	7.3E1	
Q-Reg	7.1E0	4.2E2	4.9E-1	1.0E1	
MRDR	2.2E0	5.1E1	1.6E0	9.1E0	
FQE	9.6E-2	1.3E-1	7.7E-3	1.9E-2	
$R(\lambda)$	1.3E-2	1.7E0	2.0E-2	2.3E-2	
$Q^{\hat{\pi}}(\lambda)$	1.6E-1	7.0E0	3.9E-2	6.3E-2	
TREE	3.3E-2	2.6E0	4.0E-2	2.6E-2	
IH	1.5E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	2.7E1	3.7E1		
WIS	1.4E-1	4.1E-1		
NAIVE	7.9E0	-		

Table 497. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	4.8E0	1.4E15	7.4E0	5.3E0
Q-Reg	1.2E0	4.7E11	2.7E0	3.9E0
MRDR	1.3E0	2.5E16	1.6E0	3.6E0
FQE	1.3E0	1.1E12	2.2E-1	7.5E-1
$R(\lambda)$	6.5E-2	4.7E14	8.6E-2	4.9E-2
$Q^{\pi}(\lambda)$	6.2E-1	1.1E16	4.7E-1	4.3E-1
TREE	2.7E-1	1.4E15	1.4E-1	1.8E-1
IH	1.3E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	2.6E17	7.6E1		
WIS	1.1E0	3.0E0		
NAIVE	8.3E0	-		

Table 498. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM	Hybrid		
	DIRECT	DR	WDR	MAGIC
AM	1.0E1	2.2E17	3.9E1	8.0E0
Q-Reg	1.6E0	4.1E14	3.6E0	2.9E0
MRDR	1.2E0	4.1E20	1.2E0	1.7E0
FQE	2.6E-1	2.4E18	4.5E-2	1.1E-1
$R(\lambda)$	6.0E-2	2.0E19	4.1E-2	4.0E-2
$Q^{\hat{\pi}}(\lambda)$	5.0E-1	4.0E19	1.7E-1	1.3E-1
TREE	8.5E-2	1.6E19	2.4E-2	4.6E-2
IH	1.4E0	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	2.1E20	3.3E2	
WIS	4.4E-1	2.1E0	
NAIVE	8.2E0	-	

Table 499. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	7.2E1	1.1E15	4.1E2	7.2E1	
Q-Reg	5.1E0	5.6E11	3.1E0	1.1E1	
MRDR	4.7E0	4.5E18	1.1E0	5.1E0	
FQE	7.8E-2	2.5E13	1.5E-2	2.7E-2	
$R(\lambda)$	5.7E-2	2.6E17	3.3E-2	2.1E-2	
$Q^{\pi}(\lambda)$	2.2E-1	2.3E17	1.7E-1	4.4E-2	
TREE	2.5E-1	2.4E17	1.5E-2	6.1E-2	
IH	1.4E0	-	-	-	

	IPS			
	STANDARD	PER-DECISION		
IS	6.1E18	2.8E1		
WIS	5.0E-1	1.5E0		
NAIVE	8.2E0	-		

Table 500. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	7.7E0	2.7E17	2.5E1	3.3E1
Q-Reg	5.4E0	5.8E14	2.0E1	2.1E1
MRDR	2.1E0	8.4E21	1.3E1	6.4E0
FQE	1.0E1	1.9E19	4.0E0	5.8E0
$R(\lambda)$	1.7E1	1.6E20	7.1E0	1.5E1
$Q^{\hat{\pi}}(\hat{\lambda})$	2.0E1	1.3E21	3.4E0	8.0E0
TREE	1.7E1	1.4E20	4.4E0	1.1E1
IH	5.9E0	-	-	-

	IPS			
	STANDARD PER-DECISIO			
IS	1.1E22	9.6E2		
WIS	7.2E0	5.2E0		
NAIVE	2.4E1	-		

Table 501. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	6.0E0	8.8E15	5.2E1	5.3E0
Q-Reg	1.9E0	3.6E8	2.9E0	4.4E0
MRDR	8.1E0	3.9E18	1.3E0	1.5E1
FQE	1.8E0	9.4E16	1.2E0	8.6E-1
$R(\lambda)$	5.2E-1	3.4E16	6.3E-1	4.8E-1
$Q^{\pi}(\lambda)$	4.5E-1	3.7E17	9.5E-1	6.6E-1
TREE	1.1E0	6.6E17	8.0E-1	5.4E-1
IH	5.0E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.6E19	7.4E1		
WIS	5.3E0	3.1E0		
NAIVE	2.3E1	-		

Table 502. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=0.80$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.8E1	1.3E13	8.8E2	5.2E1	
Q-Reg	2.3E0	1.9E9	4.1E0	5.1E0	
MRDR	1.4E0	5.5E14	3.8E0	3.7E0	
FQE	1.1E0	2.9E13	3.1E-1	1.5E-1	
$R(\lambda)$	1.2E0	7.5E13	3.6E-1	5.6E-1	
$Q^{\pi}(\lambda)$	7.8E0	3.0E14	1.3E0	1.8E0	
TREE	1.3E0	2.1E14	5.8E-1	8.0E-1	
IH	5.7E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.9E16	1.1E1		
WIS	8.9E0	3.6E0		
NAIVE	2.4E1	-		

Table 503. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.1E0	3.7E1	1.4E1	5.7E0	
Q-Reg	1.1E0	1.8E0	3.2E0	4.0E0	
MRDR	1.1E0	8.8E-1	3.6E0	4.5E0	
FQE	2.3E1	2.1E1	6.1E0	1.4E1	
$R(\lambda)$	9.1E2	2.0E3	7.6E4	9.1E2	
$Q^{\hat{\pi}}(\lambda)$	5.7E1	5.7E1	7.3E1	4.8E1	
TREE	2.0E1	2.1E2	4.2E2	6.9E1	
IH	2.0E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	1.7E0		
WIS	1.0E1	3.3E0		
NAIVE	1.1E1	-		

Table 504. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.2E0	2.1E1	3.8E1	7.6E0	
Q-Reg	1.3E0	1.6E0	5.9E0	6.1E0	
MRDR	1.1E0	9.1E-1	8.0E0	8.1E0	
FQE	1.1E1	1.1E1	2.8E0	6.5E0	
$R(\lambda)$	5.1E0	4.8E0	2.7E0	2.8E0	
$Q^{\hat{\pi}}(\hat{\lambda})$	3.1E1	2.7E1	9.8E0	1.2E1	
TREE	3.2E0	3.2E0	1.4E0	2.0E0	
IH	2.1E0	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	1.4E0	
WIS	5.6E0	5.7E0	
NAIVE	1.1E1	-	

Table 505. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	3.1E0	6.2E1	3.2E1	3.5E0
Q-Reg	1.2E0	2.0E0	3.8E0	4.3E0
MRDR	1.1E0	1.2E0	2.5E0	2.9E0
FQE	3.9E-1	3.7E-1	3.3E-1	2.9E-1
$R(\lambda)$	7.7E0	1.1E1	5.9E0	5.8E0
$Q^{\pi}(\lambda)$	3.9E1	3.8E1	9.5E0	1.5E1
TREE	5.8E-1	4.0E0	3.0E-1	3.1E-1
IH	2.1E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	1.7E0		
WIS	4.1E0	4.2E0		
NAIVE	1.1E1	-		

Table 506. Pixel-Gridworld, relative MSE. $T=25, N=512, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b known. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	4.9E0	1.1E3	3.7E2	6.3E0	
Q-Reg	1.9E0	2.4E0	3.5E0	4.4E0	
MRDR	9.7E-1	2.5E1	4.7E0	2.7E0	
FQE	5.7E-1	7.3E0	6.6E-2	1.3E-1	
$R(\lambda)$	5.4E-1	5.6E0	1.1E-1	2.2E-1	
$Q^{\hat{\pi}}(\hat{\lambda})$	3.6E1	3.2E1	6.4E0	1.2E1	
TREE	1.3E0	3.5E1	3.5E-1	6.5E-1	
IH	2.1E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	9.8E-1	8.7E0		
WIS	2.4E0	4.2E0		
NAIVE	1.1E1	-		

Table 507. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	2.1E0	7.8E0	9.7E0	4.8E0	
Q-Reg	1.2E0	1.6E0	6.5E0	4.8E0	
MRDR	1.1E0	9.6E-1	8.1E0	7.5E0	
FQE	2.8E1	2.3E1	1.2E1	1.3E1	
$R(\lambda)$	1.5E1	1.6E1	8.7E0	9.0E0	
$Q^{\hat{\pi}}(\lambda)$	5.2E1	3.4E1	1.4E1	1.9E1	
TREE	4.8E1	4.9E1	1.1E1	1.2E1	
IH	2.1E0	-	-	-	

	IPS			
	STANDARD PER-DECISION			
IS	1.0E0	1.4E0		
WIS	1.0E1	6.6E0		
NAIVE	1.0E1	-		

Table 508. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	5.3E0	2.5E19	4.1E1	6.5E0	
Q-Reg	1.1E0	1.4E16	4.5E0	3.9E0	
MRDR	8.8E-1	5.1E21	5.2E0	5.1E0	
FQE	2.4E1	2.2E19	8.0E0	9.4E0	
$R(\lambda)$	3.1E1	1.6E20	4.6E0	4.6E0	
$Q^{\hat{\pi}}(\lambda)$	2.7E1	1.7E20	8.7E0	1.2E1	
TREE	1.4E2	2.1E19	7.1E0	7.6E0	
IH	2.0E0	-	-	-	

	IPS			
	STANDARD PER-DECISIO			
IS	4.1E21	1.5E1		
WIS	8.1E0	4.2E0		
NAIVE	1.0E1	-		

Table 509. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	2.6E0	6.8E15	3.5E1	2.6E0
Q-Reg	1.2E0	2.0E8	3.8E0	4.5E0
MRDR	1.2E0	6.1E16	4.5E0	5.5E0
FQE	2.5E0	1.4E13	5.8E-1	1.1E0
$R(\lambda)$	1.4E0	3.0E16	4.2E-1	8.0E-1
$Q^{\pi}(\lambda)$	2.8E1	6.0E16	9.4E0	9.8E0
TREE	2.1E0	1.9E16	5.1E-1	7.7E-1
IH	2.0E0	-	-	-

	IPS			
	STANDARD PER-DECISION			
IS	8.8E15	1.4E0		
WIS	2.9E0	3.9E0		
NAIVE	1.0E1	-		

Table 510. Pixel-Gridworld, relative MSE. $T=25, N=64, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrid		
	DIRECT	DR	WDR	MAGIC	
AM	3.4E0	5.4E1	1.7E1	1.5E1	
Q-Reg	1.7E0	3.6E0	1.9E1	1.8E1	
MRDR	1.2E0	8.5E-1	3.0E1	2.7E1	
FQE	1.2E2	1.1E2	3.4E1	6.6E1	
$R(\lambda)$	1.1E3	2.5E3	1.4E4	5.1E3	
$Q^{\hat{\pi}}(\lambda)$	1.0E2	1.8E2	1.9E1	4.6E1	
TREE	2.8E1	4.5E1	3.2E1	1.7E1	
IH	7.8E0	-	-	-	

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	3.3E0	
WIS	2.8E1	1.7E1	
NAIVE	3.6E1	-	

Table 511. Pixel-Gridworld, relative MSE. $T=25, N=128, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	5.9E0	8.7E1	7.0E1	1.1E1
Q-Reg	1.5E0	2.1E0	8.5E0	9.3E0
MRDR	1.1E0	6.7E0	1.2E1	1.4E1
FQE	6.1E1	6.0E1	1.6E1	2.4E1
$R(\lambda)$	2.3E1	2.4E2	1.0E1	1.0E1
$Q^{\pi}(\lambda)$	3.9E1	8.4E1	1.6E1	2.5E1
TREE	1.8E1	2.3E1	9.1E0	1.1E1
IH	8.4E0	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	2.6E3	1.8E0	
WIS	3.6E1	1.2E1	
NAIVE	3.8E1	-	

Table 512. Pixel-Gridworld, relative MSE. $T=25, N=256, \pi_b=1.00$ -Greedy(V iter.), $\pi_e=0.10$ -Greedy(V iter.). Note: we use the same policy as in Gridworld. π_b unknown. Stochastic environment.

	DM		Hybrii)
	DIRECT	DR	WDR	MAGIC
AM	4.4E0	1.8E3	2.4E2	7.3E0
Q-Reg	1.7E0	3.4E0	6.7E0	8.5E0
MRDR	1.5E0	8.8E2	8.6E0	8.9E0
FQE	6.5E0	5.0E1	2.3E0	2.7E0
$R(\lambda)$	1.4E1	3.8E4	4.5E0	5.7E0
$Q^{\hat{\pi}}(\lambda)$	6.6E1	5.0E4	9.6E0	3.2E1
TREE	1.6E1	1.5E4	3.9E0	6.4E0
IH	7.8E0	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	3.7E5	2.1E0	
WIS	2.2E1	8.2E0	
NAIVE	3.5E1	-	

H.8. Detailed Results for Enduro

Table 513. Enduro, relative MSE. Model Type: conv. $T=500, N=512, \pi_b=0.10$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	-	-	-	-
Q-Reg	8.9E-1	8.3E-1	9.8E-1	2.9E-1
MRDR	9.3E-1	1.5E0	1.5E0	3.1E-1
FQE	3.2E-1	8.5E-2	1.4E-1	3.8E-2
$R(\lambda)$	-	-	-	-
$Q^{\pi}(\lambda)$	-	-	-	-
TREE	-	-	-	-
IH	1.0E-2	-	-	-

	IPS		
	STANDARD	PER-DECISION	
IS	1.0E0	8.5E-1	
WIS	8.9E-2	8.2E-2	
NAIVE	1.1E-2	-	

Table 514. Enduro, relative MSE. Model Type: conv. $T=500, N=512, \pi_b=0.25$ -Greedy(DDQN), $\pi_e=0.00$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	-	-	-	-
Q-Reg	1.0E0	1.1E0	4.7E0	6.6E0
MRDR	1.0E0	1.0E0	1.3E-1	1.5E-1
FQE	7.1E-1	7.1E-1	7.4E-2	8.9E-2
$R(\lambda)$	-	-	-	-
$Q^{\hat{\pi}}(\lambda)$	-	-	-	-
TREE	-	-	-	-
IH	1.5E-1	-	-	-

	IPS		
	STANDARD PER-DECISION		
IS	1.0E0	1.0E0	
WIS	3.1E-1	5.4E-2	
NAIVE	1.6E-1	-	

Table 515. Enduro, relative MSE. Model Type: conv. $T=500, N=512, \pi_b=0.25$ -Greedy(DDQN), $\pi_e=0.10$ -Greedy(DDQN).

	DM		Hybrid	
	DIRECT	DR	WDR	MAGIC
AM	_	_	_	_
Q-Reg	9.0E-1	6.8E-1	7.5E-1	5.1E-1
MRDR	1.0E0	4.1E0	2.2E1	3.4E-1
FQE	6.5E-1	3.5E-1	8.6E-2	4.5E-2
$R(\lambda)$	-	-	-	-
$Q^{\hat{\pi}}(\lambda)$	-	-	-	-
TREE	-	-	-	-
IH	9.5E-2	-	-	-

		IPS	
	STANDARD	PER-DECISION	
IS	1.0E0	8.8E-1	
WIS	1.5E-1	8.9E-2	
NAIVE	9.9E-2	-	